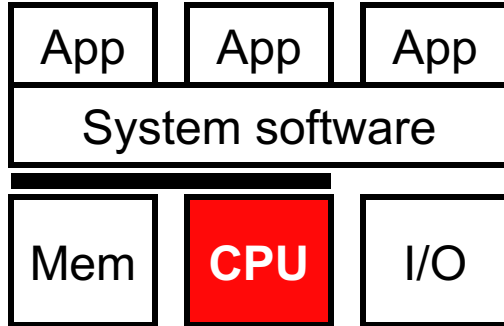# CIS 5710
# Computer Organization and Design

## Unit 4: Single-Cycle Datapath

Based on slides by Profs. Benedict Brown, C.J. Taylor, Amir Roth & Milo Martin
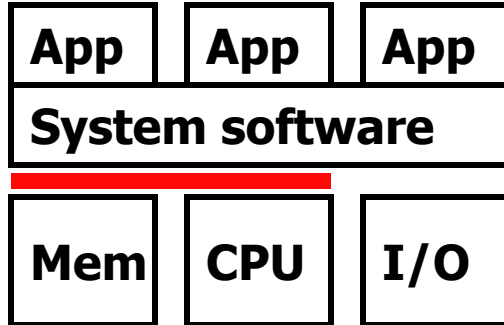
# This Unit: Single-Cycle Datapath

| App | App | App |
|-----|-----|-----|
| System software | | |

| Mem | **CPU** | I/O |
|-----|-----|-----|

- Overview of ISAs
- Datapath storage elements
- MIPS Datapath
- MIPS Control

# Readings

- P&H
  - Sections 4.1 – 4.4

# Recall from CIS 2400…

# 2400 Review: ISA

| | | |
|---|---|---|
| **App** | **App** | **App** |
| **System software** | | |

| | | |
|---|---|---|
| **Mem** | **CPU** | **I/O** |

- App/OS are software … execute on hardware
- HW/SW interface is **ISA (instruction set architecture)**
  - A **"contract"** between SW and HW
  - Encourages compatibility, allows SW/HW to evolve independently
  - **Functional definition** of HW storage locations & operations
    - Storage locations: registers, memory
    - Operations: add, multiply, branch, load, store, etc.
  - **Precise description** of how to invoke & access them
    - Instructions (bit-patterns hardware interprets as commands)

# 2400 Review: LC4 ISA

| App | App | App |
|-----|-----|-----|
| **System software** | | |

| Mem | CPU | I/O |
|-----|-----|-----|
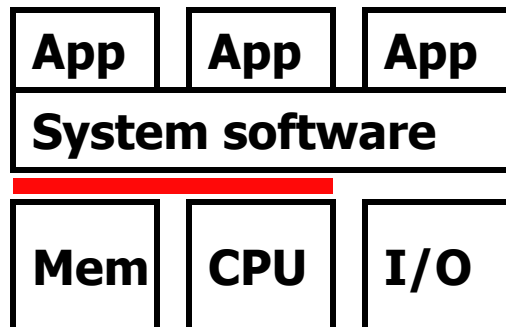
- **LC4**: a toy ISA you (may) know
  - 16-bit ISA (what does this mean?)
  - 16-bit insns
  - 8 registers (integer)
  - ~30 different insns
  - Simple OS support
  - 16b (not byte!) addressable
- **Assembly language**
  - Human-readable ISA representation
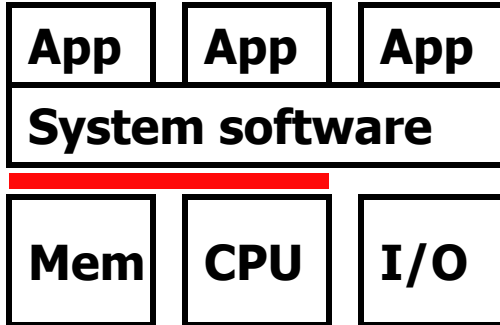
```
       .DATA
array .BLKW #100
sum    .FILL #0
    .CODE
    .FALIGN
array_sum
    CONST R5, #0
    LEA R1, array
    LEA R2, sum
array_sum_loop
    LDR R3, R1, #0
    LDR R4, R2, #0
    ADD R4, R3, R4
    STR R4, R2, #0
    ADD R1, R1, #1
    ADD R5, R5, #1
    CMPI R5, #100
    BRn array_sum_loop
```

# MIPS: A Real ISA

| App | App | App |
|-----|-----|-----|

**System software**
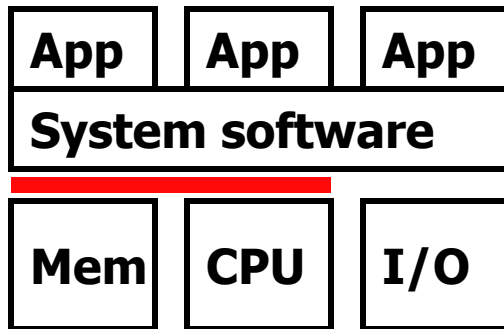
| Mem | CPU | I/O |
|-----|-----|-----|

- **MIPS**: example of real ISA
  - 32/64-bit operations
  - 32-bit insns
  - 64 registers
    - 32 integer, 32 floating point
  - ~100 different insns
  - Full OS support

**Example code is MIPS, but all ISAs are pretty similar**

```
        .data
array:  .space 100
sum:    .word 0
        .text


array_sum:
    li $5, 0
    la $1, array
    la $2, sum
array_sum_loop:
    lw $3, 0($1)
    lw $4, 0($2)
    add $4, $3, $4
    sw $4, 0($2)
    addi $1, $1, 1
    addi $5, $5, 1
    li $6, 100
    blt $5, $6, array_sum_loop
```

# 2400 Review: Assembly Language

| App | App | App |
|-----|-----|-----|

| System software |
|-----------------|

| Mem | CPU | I/O |
|-----|-----|-----|

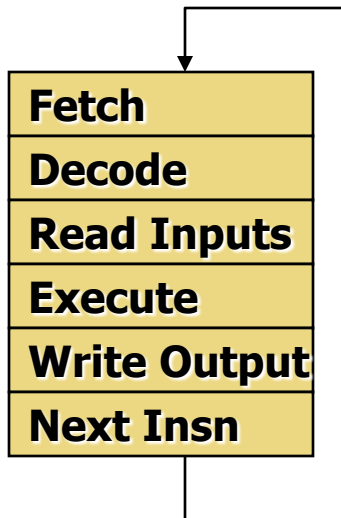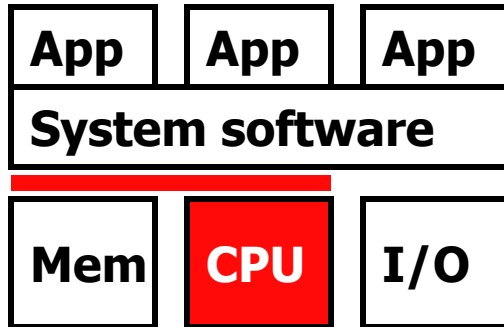| Machine code | Assembly code |
|:---:|:---|
| **x9A00** | CONST R5, #0 |
| **x9200** | CONST R1, array |
| **xD320** | HICONST R1, array |
| **x9464** | CONST R2, sum |
| **xD520** | HICONST R2, sum |
| **x6640** | LDR R3, R1, #0 |
| **x6880** | LDR R4, R2, #0 |
| **x18C4** | ADD R4, R3, R4 |
| **x7880** | STR R4, R2, #0 |
| **x1261** | ADD R1, R1, #1 |
| **x1BA1** | ADD R5, R5, #1 |
| **x2B64** | CMPI R5, #100 |
| **x03F8** | BRn array_sum_loop |

- **Assembly language**
  - Human-readable representation
- **Machine language**
  - Machine-readable representation
  - 1s and 0s (often displayed in "hex")
- **Assembler**
  - Translates assembly to machine

# 2400 Review: Insn Execution Model

| App | App | App |
|-----|-----|-----|

**System software**

| Mem | **CPU** | I/O |
|-----|---------|-----|

| Fetch |
|-------|
| Decode |
| Read Inputs |
| Execute |
| Write Output |
| Next Insn |

**Instruction → Insn**

- The computer is just finite state machine
  - **Registers** (few of them, but fast)
  - **Memory** (lots of memory, but slower)
  - **Program counter** (next insn to execute)
    - Sometimes called "instruction pointer"
- A computer executes **instructions**
  - **Fetches** next instruction from memory
  - **Decodes** it (figure out what it does)
  - **Reads** its **inputs** (registers & memory)
  - **Executes** it (adds, multiply, etc.)
  - **Write** its **outputs** (registers & memory)
  - **Next insn** (adjust the program counter)
- **Program is just "data in memory"**
  - Makes computers programmable ("universal")

# What is an ISA?

# What Is An ISA?

- **ISA (instruction set architecture)**
  - A well-defined hardware/software interface
  - The **"contract"** between software and hardware
    - **Functional definition** of storage locations & operations
      - Storage locations: registers, memory
      - Operations: add, multiply, branch, load, store, etc
    - **Precise description** of how to invoke & access them
- Not in the "contract": non-functional aspects
  - How operations are implemented
  - Which operations are fast and which are slow
  - Which operations take more power and which take less
- Instructions
  - Bit-patterns hardware interprets as commands
  - Instruction $\rightarrow$ Insn (instruction is too long to write in slides)

# LC4 vs Real ISAs

- ## LC4 has the basic features of a real-world ISAs
  - ± LC4 lacks some realism
    - Address size is only 16 bits
    - Only one data type (16-bit signed integer)
    - Little support for system software, none for multiprocessing (later)
    - always access 16b words in memory

- ## Many real-world ISAs to choose from:
  - Intel x86 (laptops, desktop, and servers)
  - MIPS (used throughout in book)
  - ARM (in all your mobile devices)
  - PowerPC (servers & game consoles)
  - SPARC (servers)
  - RISC-V (an open source ISA)
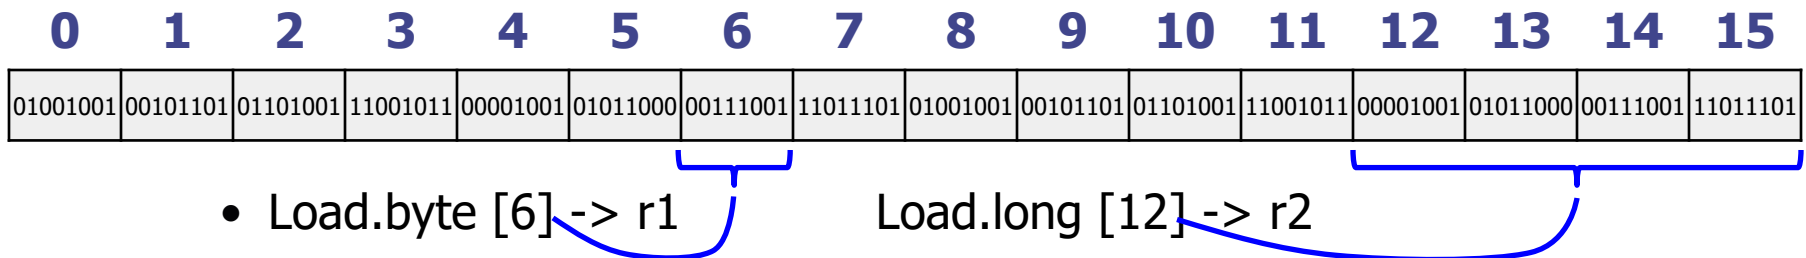  - Historical: IBM 370, VAX, Alpha, PA-RISC, 68k, …

# Some Key Attributes of ISAs

- Instruction encoding
  - Fixed length (16-bit for LC4, 32-bit for MIPS)
  - Variable length (x86 1 byte to 16 bytes, average of ~3 bytes)
  - Limited variability (ARM insns are 16b or 32b)
- Number and type of registers
  - LC-4 has 8 registers
  - MIPS has 32 "integer" registers and 32 "floating point" registers
  - ARM & x86 both have 16 "integer" regs and 16 "floating point" regs
- Address space
  - LC4: 16-bit addresses at 16-bit granularity (128KB total)
  - ARM: 32-bit addresses at 8-bit granularly (4GB total)
  - Modern x86 and ARM64: 64-bit addresses (16 exabytes!)
- Memory addressing modes
  - MIPS & LC4: address calculated by "reg+offset"
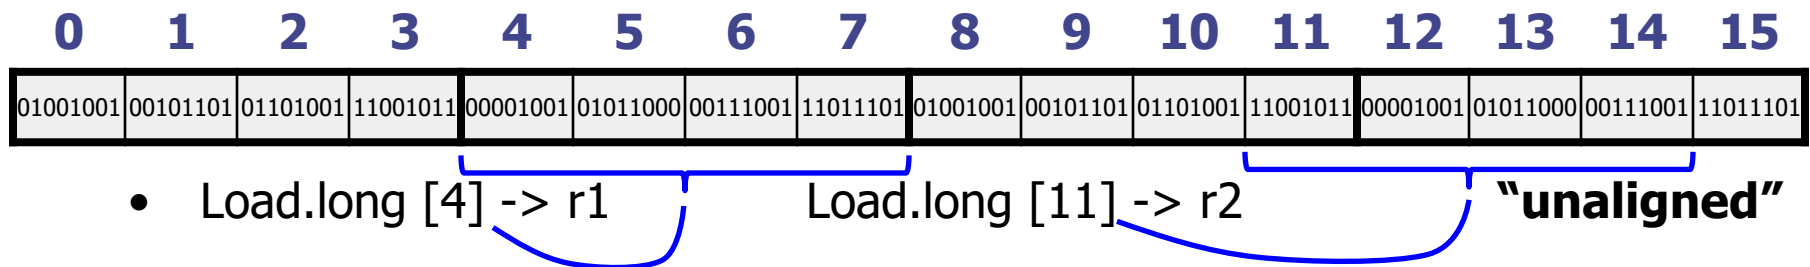  - x86 and others have much more complicated addressing modes

# Access Granularity & Alignment

- **Byte addressability**
  - An address points to a byte (8 bits) of data
  - The ISA's minimum granularity to read or write memory
  - ISAs also support wider load/stores
    - "Half" (2 bytes), "Longs" (4 bytes), "Quads" (8 bytes)

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|
| 01001001 | 00101101 | 01101001 | 11001011 | 00001001 | 01011000 | 00111001 | 11011101 | 01001001 | 00101101 | 01101001 | 11001011 | 00001001 | 01011000 | 00111001 | 11011101 |

- Load.byte [6] -> r1       Load.long [12] -> r2

However, physical memory systems operate on **even larger chunks**

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|
| 01001001 | 00101101 | 01101001 | 11001011 | 00001001 | 01011000 | 00111001 | 11011101 | 01001001 | 00101101 | 01101001 | 11001011 | 00001001 | 01011000 | 00111001 | 11011101 |

- Load.long [4] -> r1       Load.long [11] -> r2       **"unaligned"**

- **Access alignment**: if address % size != 0, then it is "unaligned"
  - A single unaligned access may require multiple physical memory accesses

# Handling Unaligned Accesses

- **Access alignment**: if address % size != 0, then it is "unaligned"
  - A single unaligned access may require multiple physical memory accesses

- How to handle such unaligned accesses?
  1. Disallow (unaligned operations are considered illegal)
     - MIPS, ARMv5 and earlier took this route
  2. Support in hardware? (allow such operations)
     - x86, ARMv6+ allow regular loads/stores to be unaligned
       - Unaligned access still slower, adds significant hardware complexity
  3. Trap to software routine?
     - Simpler hardware, but high penalty when unaligned
  4. In software (compiler can use regular instructions when possibly unaligned
     - Load, shift, load, shift, and  (slow, needs help from compiler)

# How big is this struct in C?

```
struct foo {
  char c;
  int i;
}
```

# Another Addressing Issue: Endian-ness

- **Endian-ness**: arrangement of bytes in a multi-byte number
  - Big-endian: sensible order (e.g., MIPS, PowerPC, ARM)
    - A 4-byte integer: "00000000 00000000 00000010 00000011" is 515
  - Little-endian: reverse order (e.g., x86)
    - A 4-byte integer: "00000011 00000010 00000000 00000000" is 515
  - Why little endian?

| 00000011 00000010 | 00000000 | 00000000 |
|---|---|---|
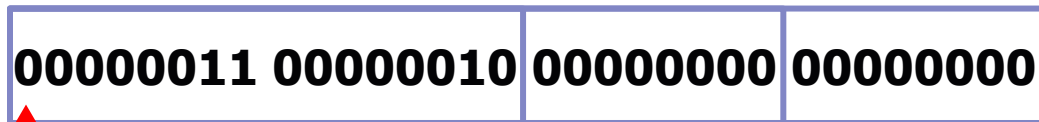
**starting address**

**integer casts are free on little-endian architectures**

# ISA Code Examples

# Array Sum Loop: LC4

```
    .DATA
array .BLKW #100
sum   .FILL #0
    .CODE
    .FALIGN
array_sum
    CONST R5, #0
    LEA R1, array
    LEA R2, sum
L1
    LDR R3, R1, #0
    LDR R4, R2, #0
    ADD R4, R3, R4
    STR R4, R2, #0
    ADD R1, R1, #1
    ADD R5, R5, #1
    CMPI R5, #100
    BRn L1
```

```
int array[100];
int sum;
void array_sum() {
    for (int i=0; i<100;i++)
    {
        sum += array[i];
    }
}
```

CIS 5710  |  Dr. Joe Devietti  |  ISAs & Single Cycle

# Array Sum Loop: LC4 ➜ MIPS

```
    .DATA                   .data                MIPS (right) similar to LC4
array .BLKW #100        array: .space 100
sum    .FILL #0         sum:    .word 0
    .CODE
    .FALIGN                 .text                Syntactic differences:
array_sum              array_sum:                register names begin with $
    CONST R5, #0           li $5, 0              immediates are un-prefixed
    LEA R1, array          la $1, array
    LEA R2, sum            la $2, sum
L1                     L1:                        Only simple addressing modes
    LDR R3, R1, #0         lw $3, 0($1)          syntax: displacement(reg)
    LDR R4, R2, #0         lw $4, 0($2)
    ADD R4, R3, R4        add $4, $3, $4
    STR R4, R2, #0        sw $4, 0($2)
    ADD R1, R1, #1       addi $1, $1, 1
    ADD R5, R5, #1       addi $5, $5, 1          Left-most register is generally
    CMPI R5, #100         li $6, 100             destination register
    BRn L1              blt $5, $6, L1
```

# Array Sum Loop: LC4 ➔ x86

```
    .DATA                        .LFE2                    x86 (right) is different
array .BLKW #100                 .comm array,400,32
sum    .FILL #0                  .comm sum,4,4
    .CODE                                                 Syntactic differences:
    .FALIGN                      .globl array_sum         register names begin with %
array_sum                    array_sum:                   immediates begin with $
    CONST R5, #0                 movl $0, -4(%rbp)
    LEA R1, array
    LEA R2, sum                                           %rbp is base (frame) pointer
L1                           .L1:
    LDR R3, R1, #0               movl -4(%rbp), %eax
    LDR R4, R2, #0               movl array(,%eax,4), %edx
    ADD R4, R3, R4               movl sum(%rip), %eax
    STR R4, R2, #0               addl %edx, %eax          Many addressing modes
    ADD R1, R1, #1               movl %eax, sum(%rip)
    ADD R5, R5, #1               addl $1, -4(%rbp)
    CMPI R5, #100               cmpl $99,-4(%rbp)
    BRn L1                       jle .L1
```

# x86 Operand Model

```
.LFE2
    .comm array,400,32
    .comm sum,4,4

    .globl array_sum
array_sum:
    movl $0, -4(%rbp)


.L1:
    movl -4(%rbp), %eax
    movl array(,%eax,4), %edx
    movl sum(%rip), %eax
    addl %edx, %eax
    movl %eax, sum(%rip)
    addl $1, -4(%rbp)
    cmpl $99,-4(%rbp)
    jle .L1
```

- x86 uses explicit accumulators
  - Both register and memory
  - Distinguished by addressing mode

Two operand insns
(right-most is typically source & destination)

Register accumulator: %eax = %eax + %edx

"L" insn suffix and "%e..." reg.
prefix mean "32-bit value"

Memory accumulator:
Memory[%rbp-4] = Memory[%rbp-4] + 1

## MOVUPS—Move Unaligned Packed Single-Precision Floating-Point Values

| Opcode/ Instruction | Op/ En | 64/32-bit Mode | CPUID Feature Flag | Description |
|---|---|---|---|---|
| 0F 10 /r<br>MOVUPS xmm1, xmm2/m128 | RM | V/V | SSE | Move packed single-precision floating-point values from xmm2/m128 to xmm1. |
| VEX.128.0F.WIG 10 /r<br>VMOVUPS xmm1, xmm2/m128 | RM | V/V | AVX | Move unaligned packed single-precision floating-point from xmm2/mem to xmm1. |
| VEX.256.0F.WIG 10 /r<br>VMOVUPS ymm1, ymm2/m256 | RM | V/V | AVX | Move unaligned packed single-precision floating-point from ymm2/mem to ymm1. |
| 0F 11 /r<br>MOVUPS xmm2/m128, xmm1 | MR | V/V | SSE | Move packed single-precision floating-point values from xmm1 to xmm2/m128. |
| VEX.128.0F.WIG 11 /r<br>VMOVUPS xmm2/m128, xmm1 | MR | V/V | AVX | Move unaligned packed single-precision floating-point from xmm1 to xmm2/mem. |
| VEX.256.0F.WIG 11 /r<br>VMOVUPS ymm2/m256, ymm1 | MR | V/V | AVX | Move unaligned packed single-precision floating-point from ymm1 to ymm2/mem. |

# Implementing an ISA

# Implementing an ISA
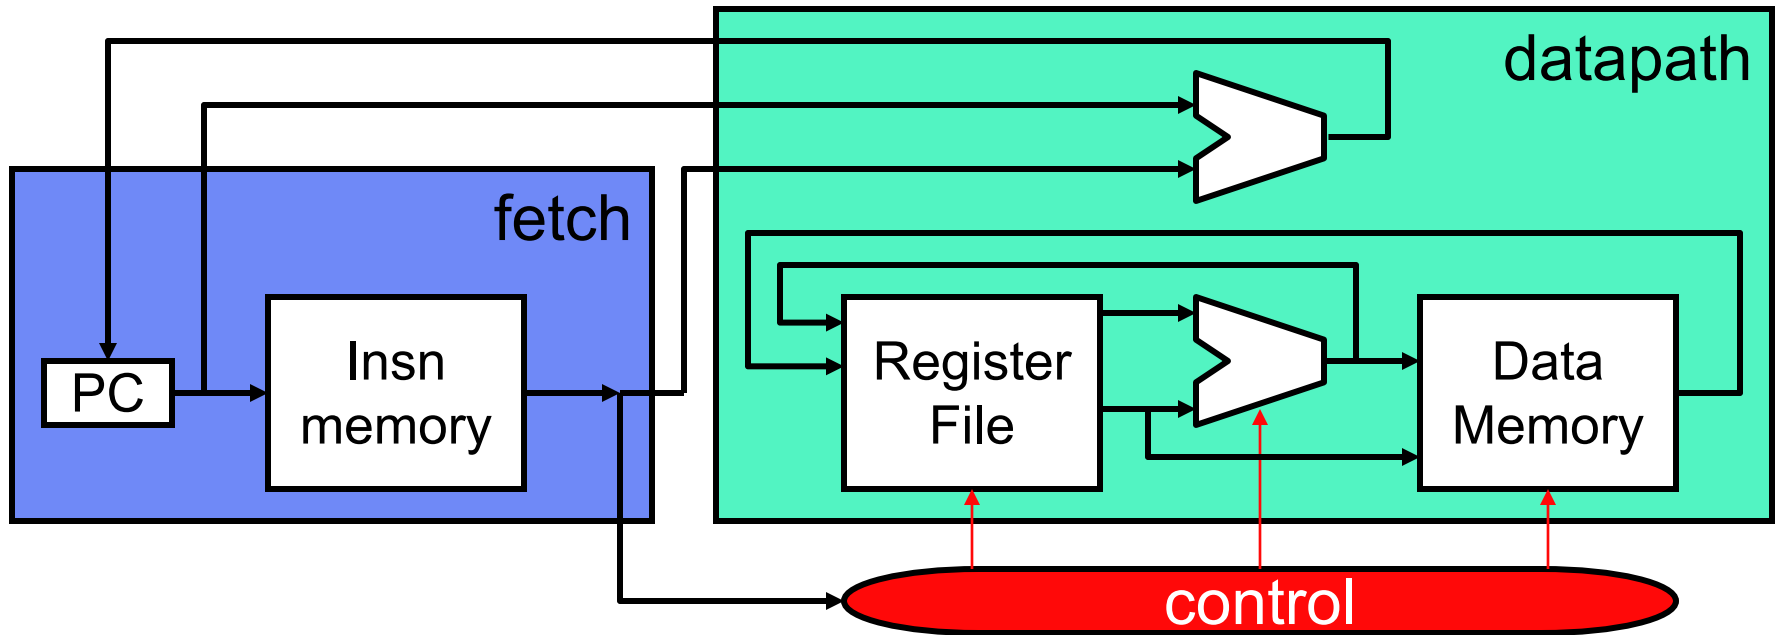


- **Datapath**: performs computation (registers, ALUs, etc.)
  - ISA specific: can implement every insn (single-cycle: in one pass!)
- **Control**: determines which computation is performed
  - Routes data through datapath (which regs, which ALU op)
- **Fetch**: get insn, translate opcode into control
- **Fetch** → **Decode** → **Execute** "cycle"

# Two Types of Components
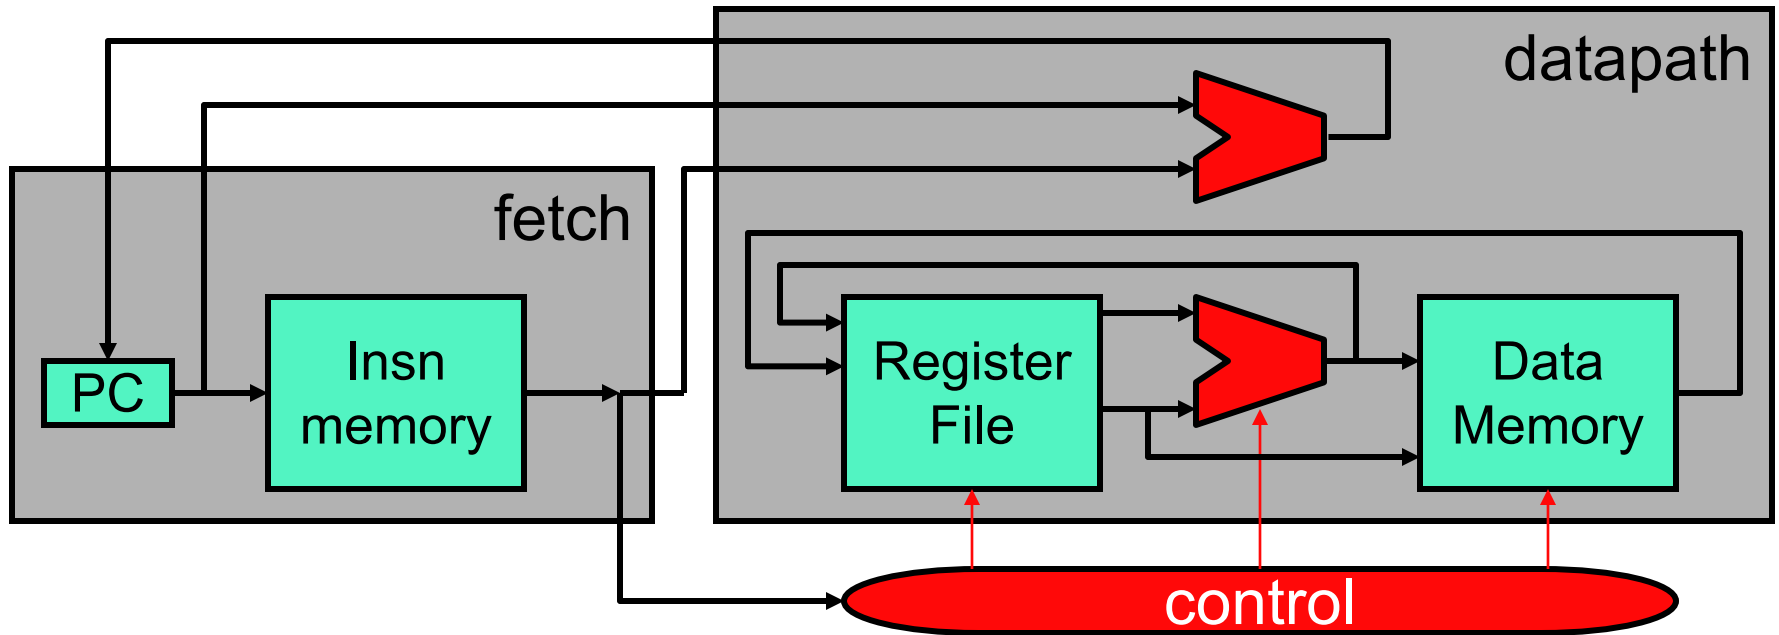


- **Purely combinational**: stateless computation
  - ALUs, muxes, control
  - Arbitrary Boolean functions
- **Combinational/sequential**: storage
  - PC, insn/data memories, register file
  - Internally contain some combinational components

# Example Datapath
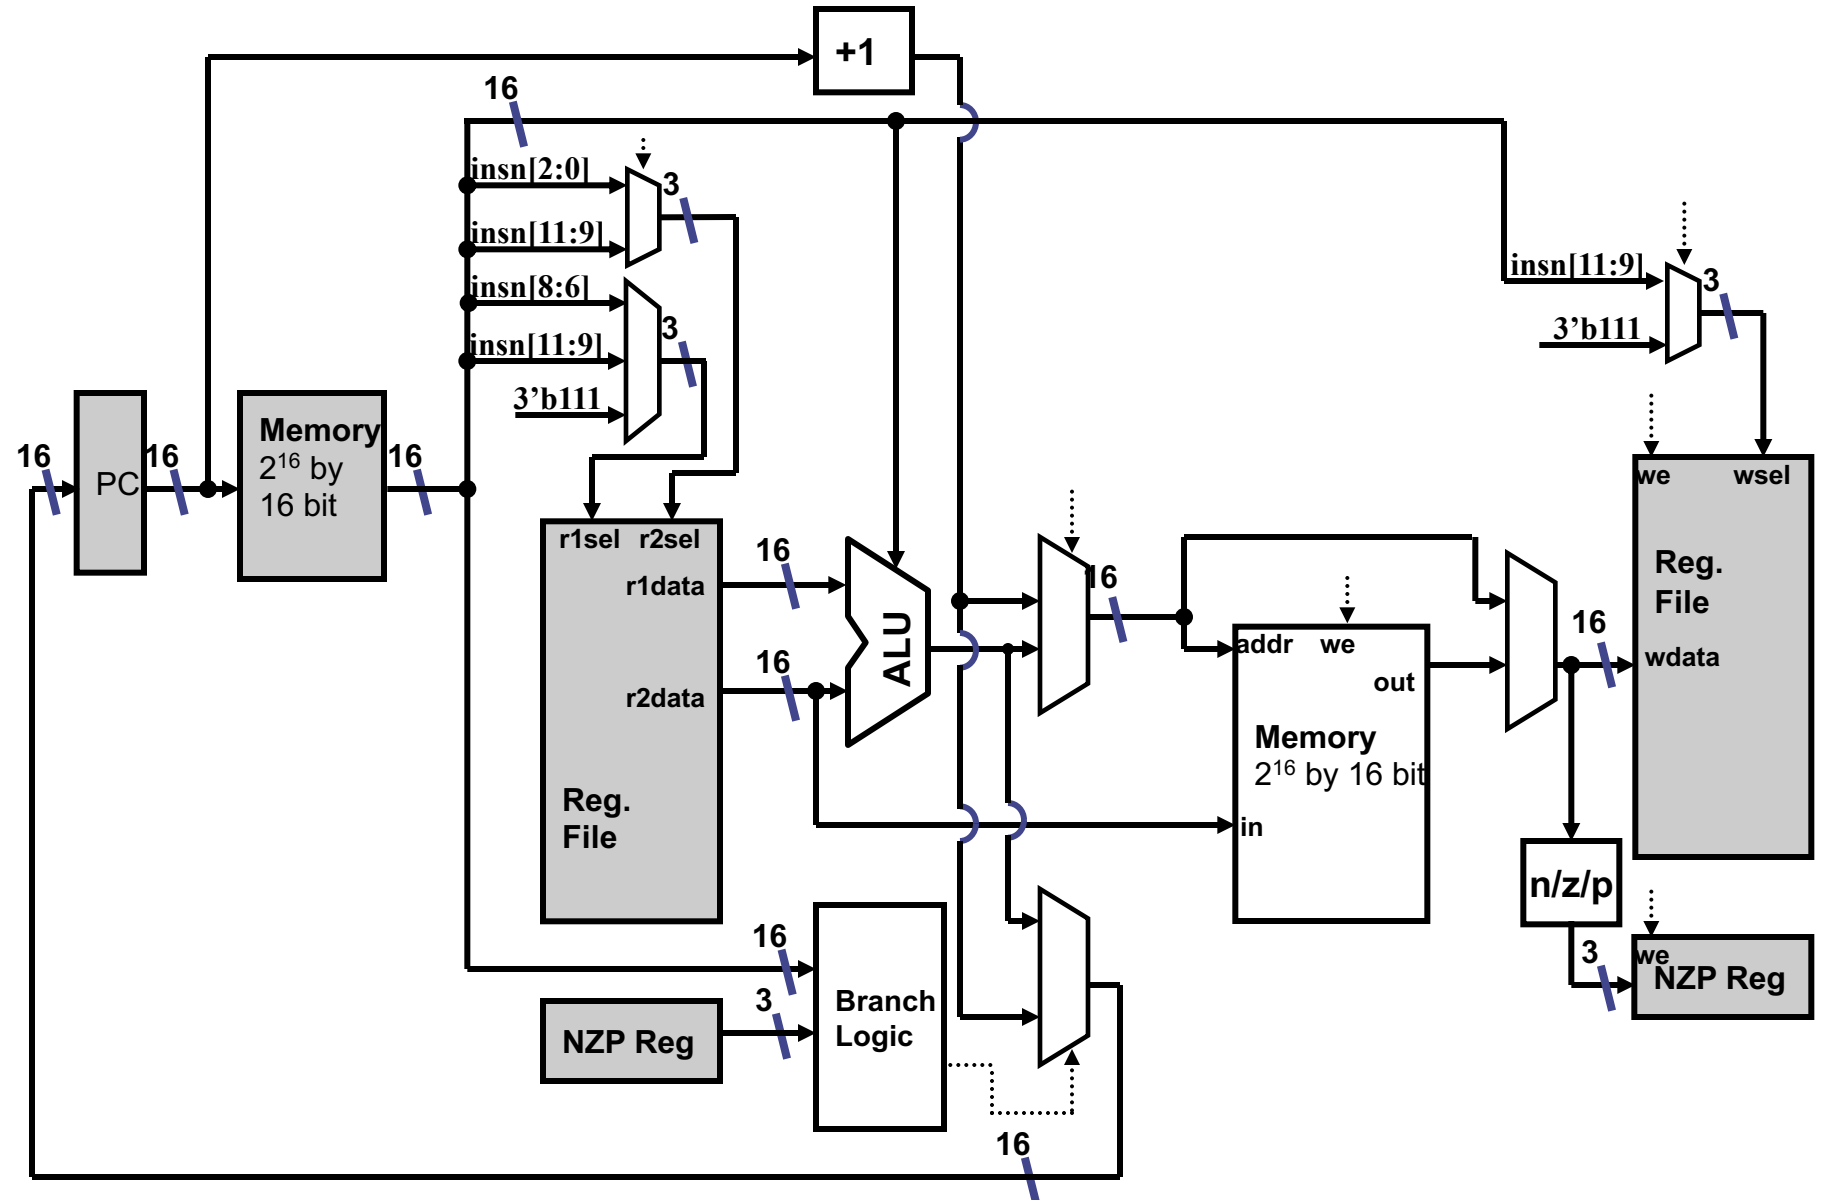
**LC4 Datapath**

# MIPS Datapath

# Unified vs Split Memory Architecture
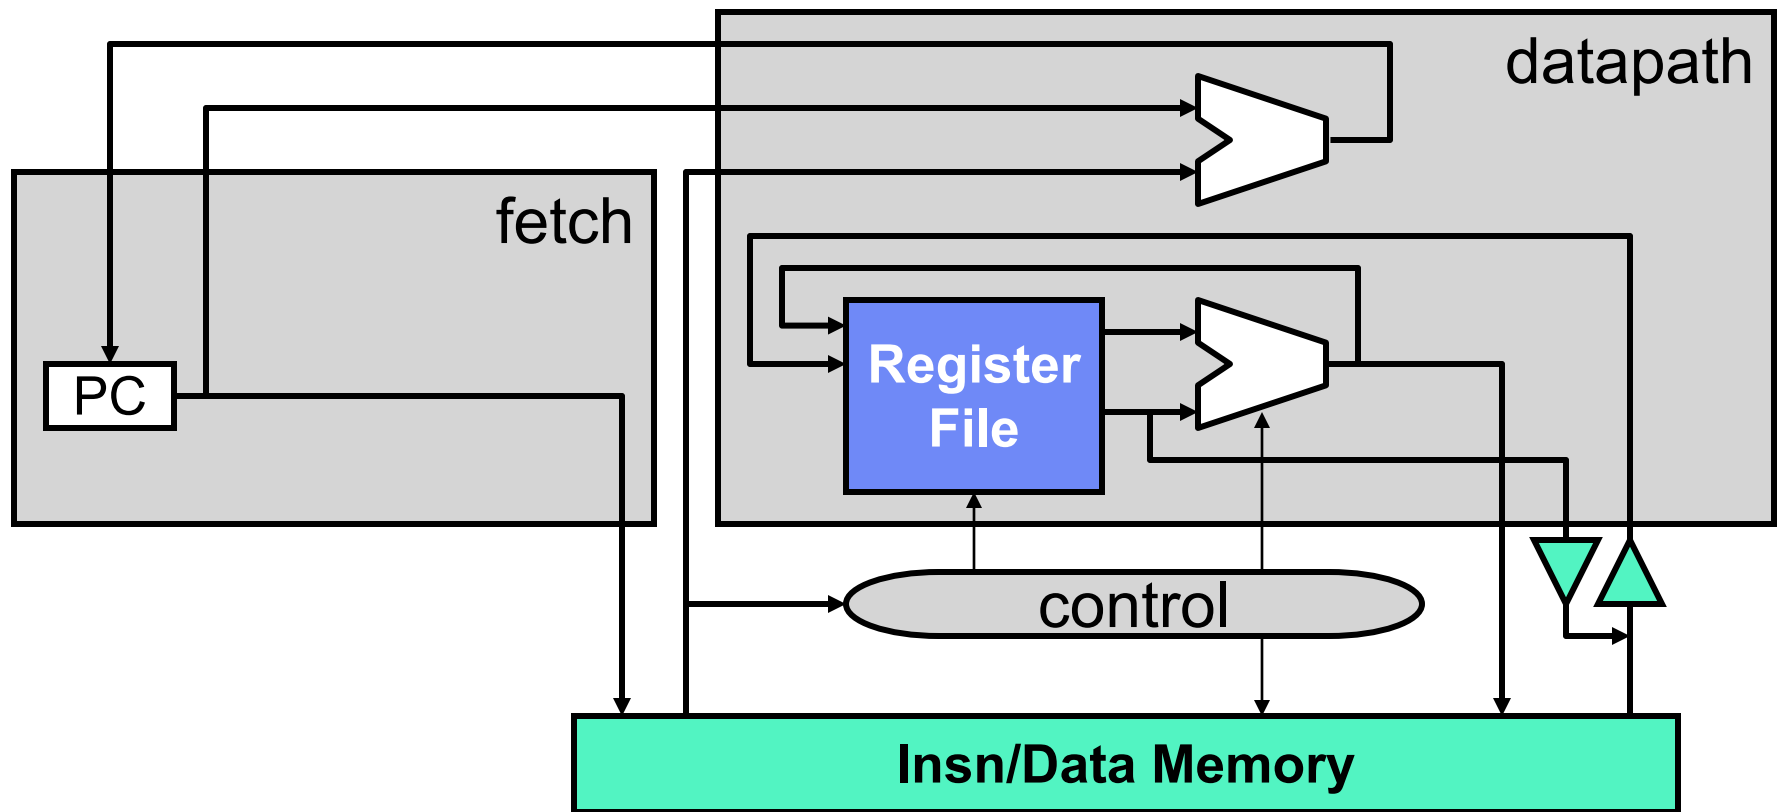


- **Unified architecture**: unified insn/data memory
- **"Harvard" architecture**: split insn/data memories

# Datapath for MIPS ISA

- MIPS: 32-bit instructions, registers are $0, $2... $31

- Consider only the following instructions

```
add  $1,$2,$3       $1 = $2 + $3        (add)
addi $1,$2,3        $1 = $2 + 3         (add immed)
lw $1,4($3)         $1 = Memory[4+$3]   (load)
sw $1,4($3)         Memory[4+$3] = $1   (store)
beq $1,$2,PC_relative_target   (branch equal)
j absolute_target          (unconditional jump)
```

- Why only these?
  - Most other instructions are the same from datapath viewpoint
  - The ones that aren't are left for you to figure out ☺

# MIPS Instruction layout

## MIPS machine language

| Name | Format | Example | | | | | | Comments |
|------|--------|------|------|------|------|------|------|----------|
| add | R | 0 | 18 | 19 | 17 | 0 | 32 | add $s1,$s2,$s3 |
| sub | R | 0 | 18 | 19 | 17 | 0 | 34 | sub $s1,$s2,$s3 |
| addi | I | 8 | 18 | 17 | 100 | | | addi $s1,$s2,100 |
| lw | I | 35 | 18 | 17 | 100 | | | lw $s1,100($s2) |
| sw | I | 43 | 18 | 17 | 100 | | | sw $s1,100($s2) |
| Field size | | 6 bits | 5 bits | 5 bits | 5 bits | 5 bits | 6 bits | All MIPS instructions are 32 bits long |
| R-format | R | op | rs | rt | rd | shamt | funct | Arithmetic instruction format |
| I-format | I | op | rs | rt | address | | | Data transfer format |

**FIGURE 2.6   MIPS architecture revealed through Section 2.5.** The two MIPS instruction formats so far are R and I. The first 16 bits are the same: both contain an *op* field, giving the base operation; an *rs* field, giving one of the sources; and the *rt* field, which specifies the other source operand, except for load word, where it specifies the destination register. R-format divides the last 16 bits into an *rd* field, specifying the destination register; the *shamt* field, which Section 2.6 explains; and the *funct* field, which specifies the specific operation of R-format instructions. I-format combines the last 16 bits into a single *address* field.

# Start With Fetch
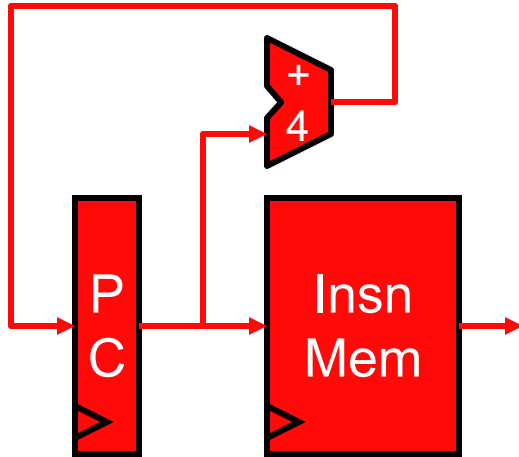


- PC and instruction memory (split insn/data architecture, for now)
- A +4 incrementer computes default next instruction PC
- How would Verilog for this look given insn memory as interface?

# First Instruction: **add**
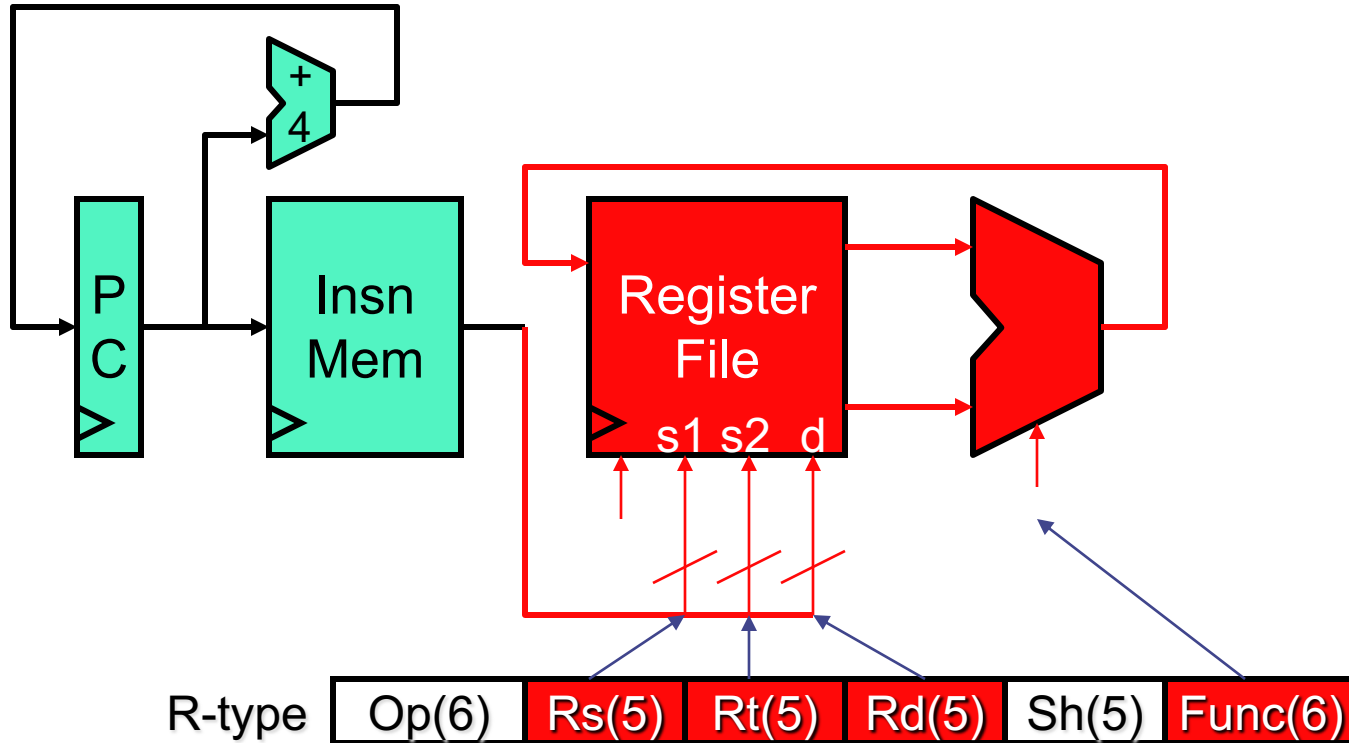


- Add register file
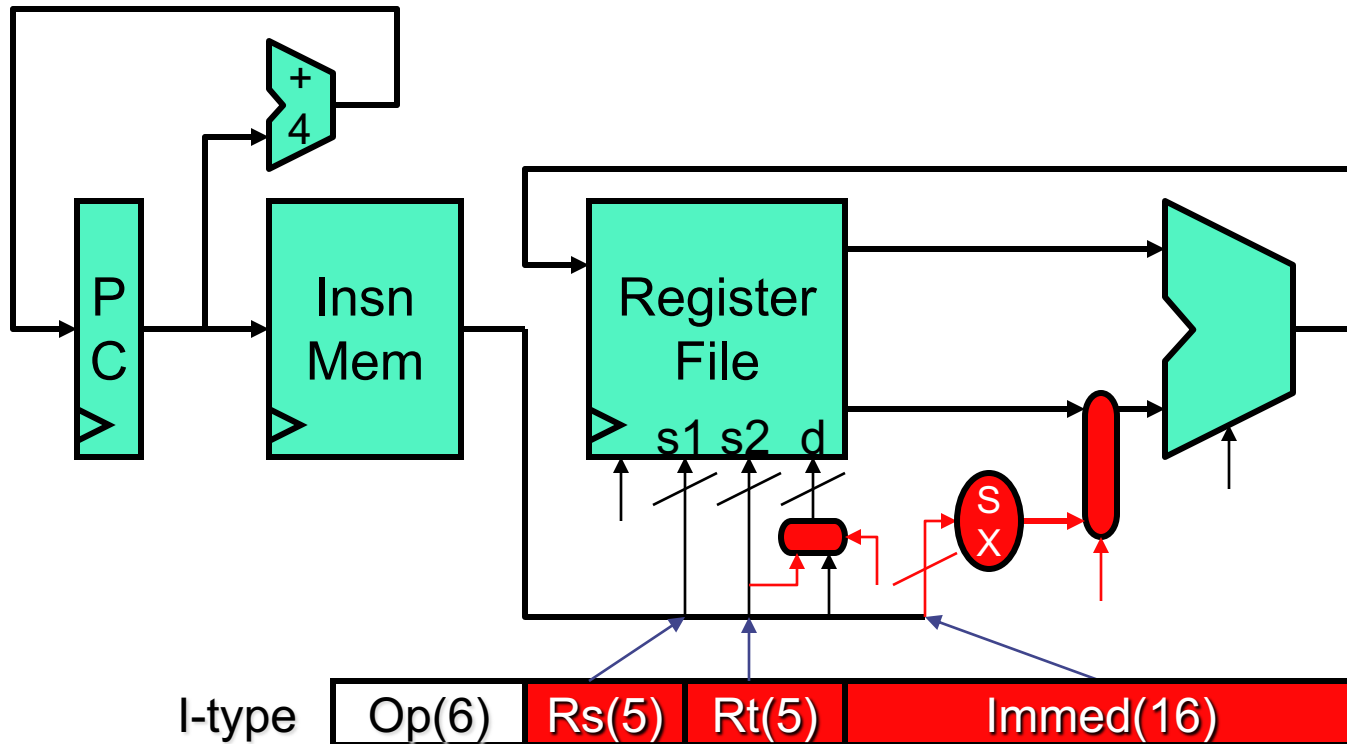- Add arithmetic/logical unit (ALU)

# Wire Select in Verilog

- How to pull out individual fields of an insn? **Wire select**

```verilog
wire [31:0] insn;
wire [5:0] op = insn[31:26];
wire [4:0] rs = insn[25:21];
wire [4:0] rt = insn[20:16];
wire [4:0] rd = insn[15:11];
wire [4:0] sh = insn[10:6];
wire [5:0] func = insn[5:0];
```

R-type | Op(6) | Rs(5) | Rt(5) | Rd(5) | Sh(5) | Func(6)

# Second Instruction: **addi**



- Destination register can now be either Rd or Rt
- Add sign extension unit and mux into second ALU input

# Verilog Wire Concatenation

- Recall two Verilog constructs
    - **Wire concatenation**: `{bus0, bus1, … , busn}`
    - **Wire repeat:** `{repeat_x_times{w0}}`

- How do you specify sign extension? **Wire concatenation**

```
wire [31:0] insn;

wire [15:0] imm16 = insn[15:0];

wire [31:0] sximm16 = {{16{imm16[15]}}, imm16};
```

| I-type | Op(6) | Rs(5) | Rt(5) | Immed(16) |
|--------|-------|-------|-------|-----------|

# Third Instruction: lw



I-type

| Op(6) | Rs(5) | Rt(5) | Immed(16) |
|-------|-------|-------|-----------|
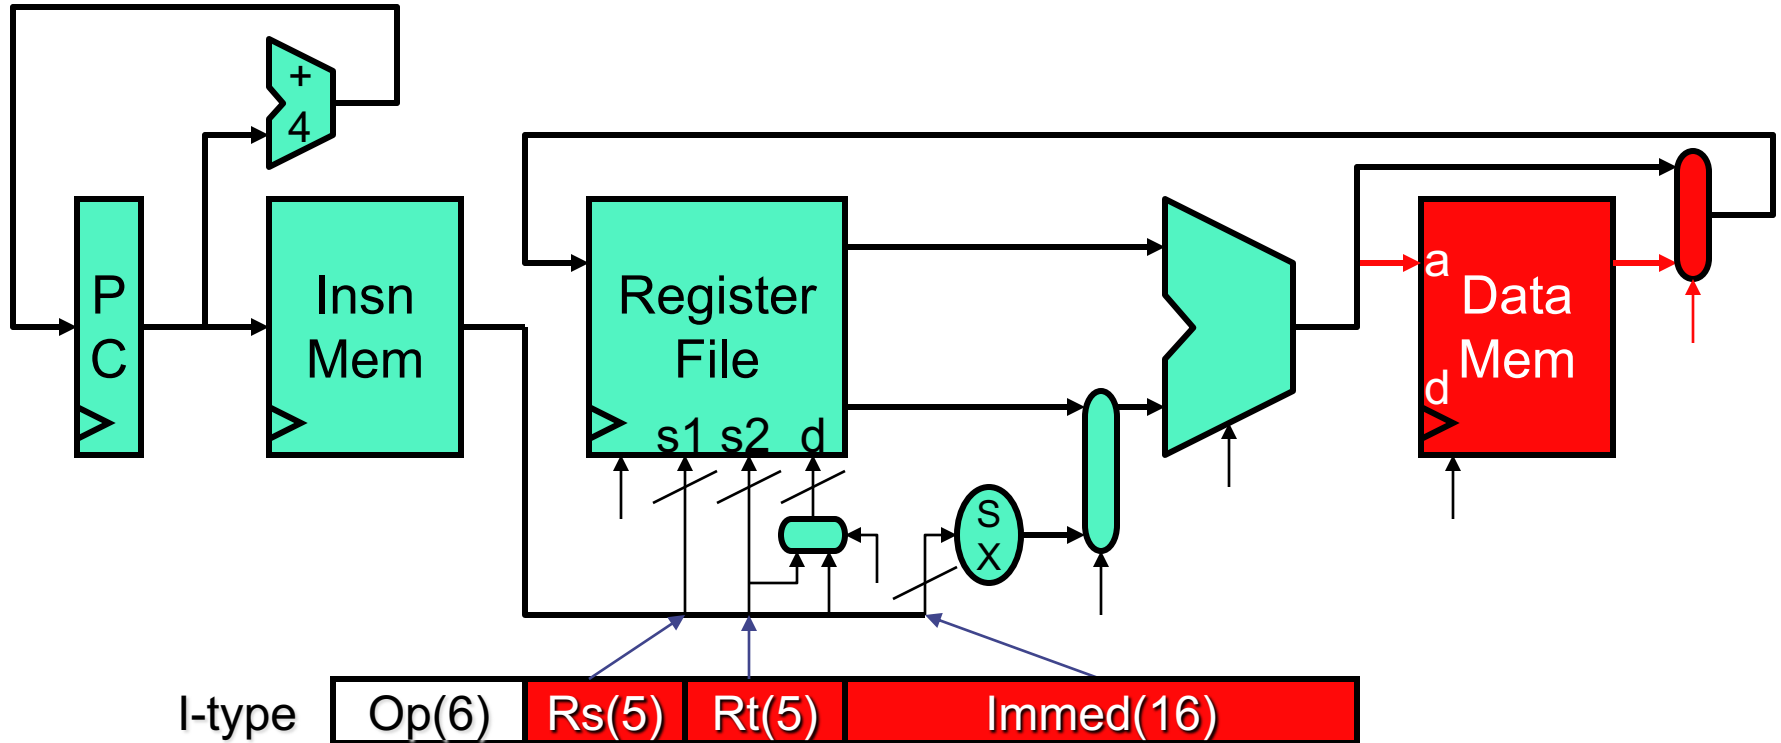
- Add data memory, address is ALU output
- Add register write data mux to select memory output or ALU output

# Fourth Instruction: **sw**



I-type

| Op(6) | Rs(5) | Rt(5) | Immed(16) |
|-------|-------|-------|-----------|
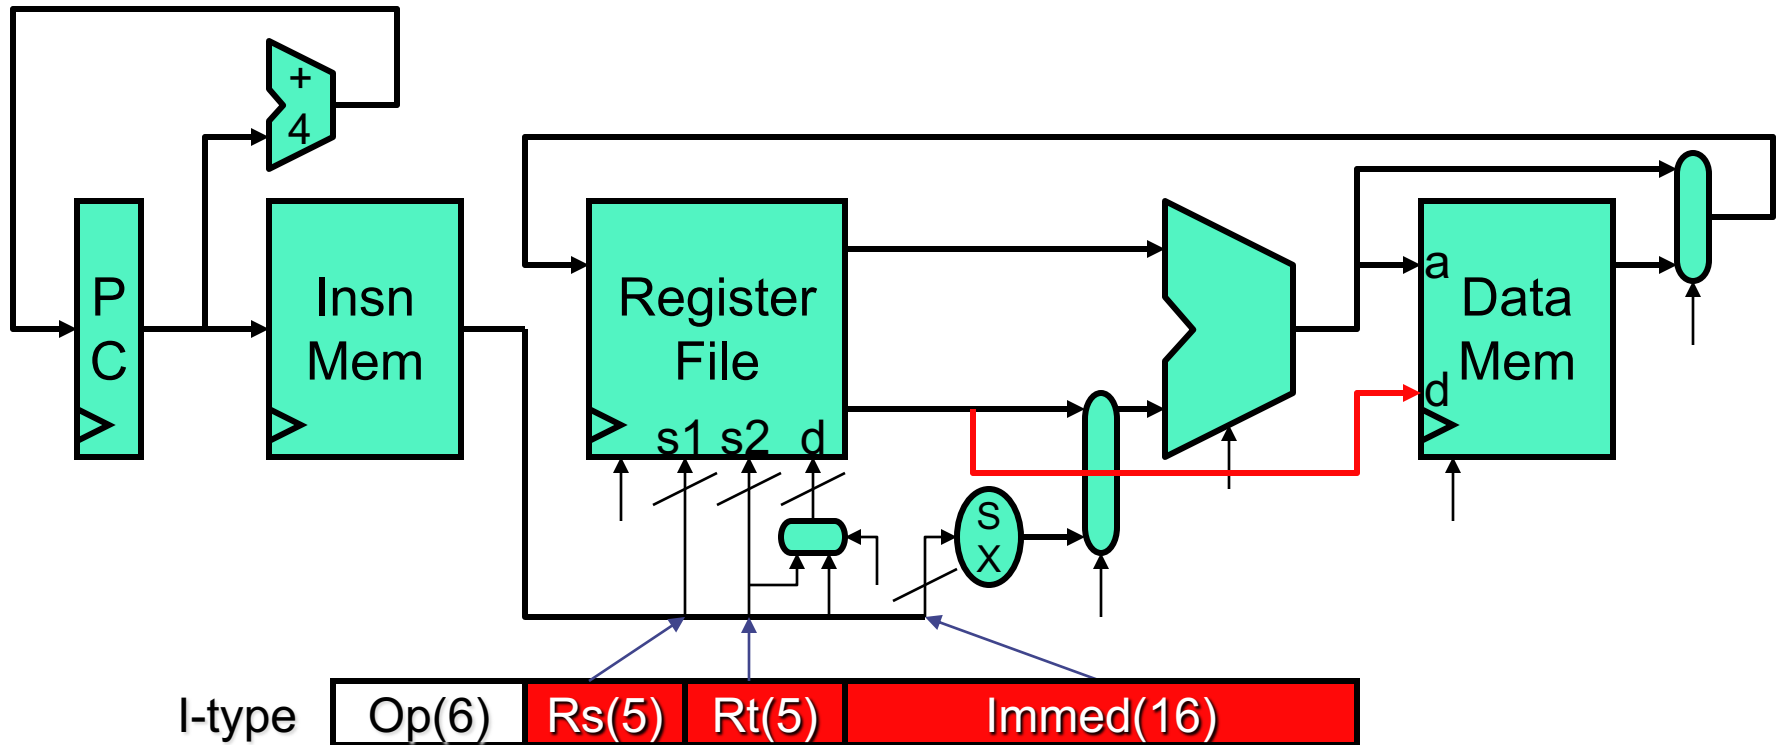
- Add path from second input register to data memory data input

# Fifth Instruction: **beq**



I-type

| Op(6) | Rs(5) | Rt(5) | Immed(16) |
|-------|-------|-------|-----------|
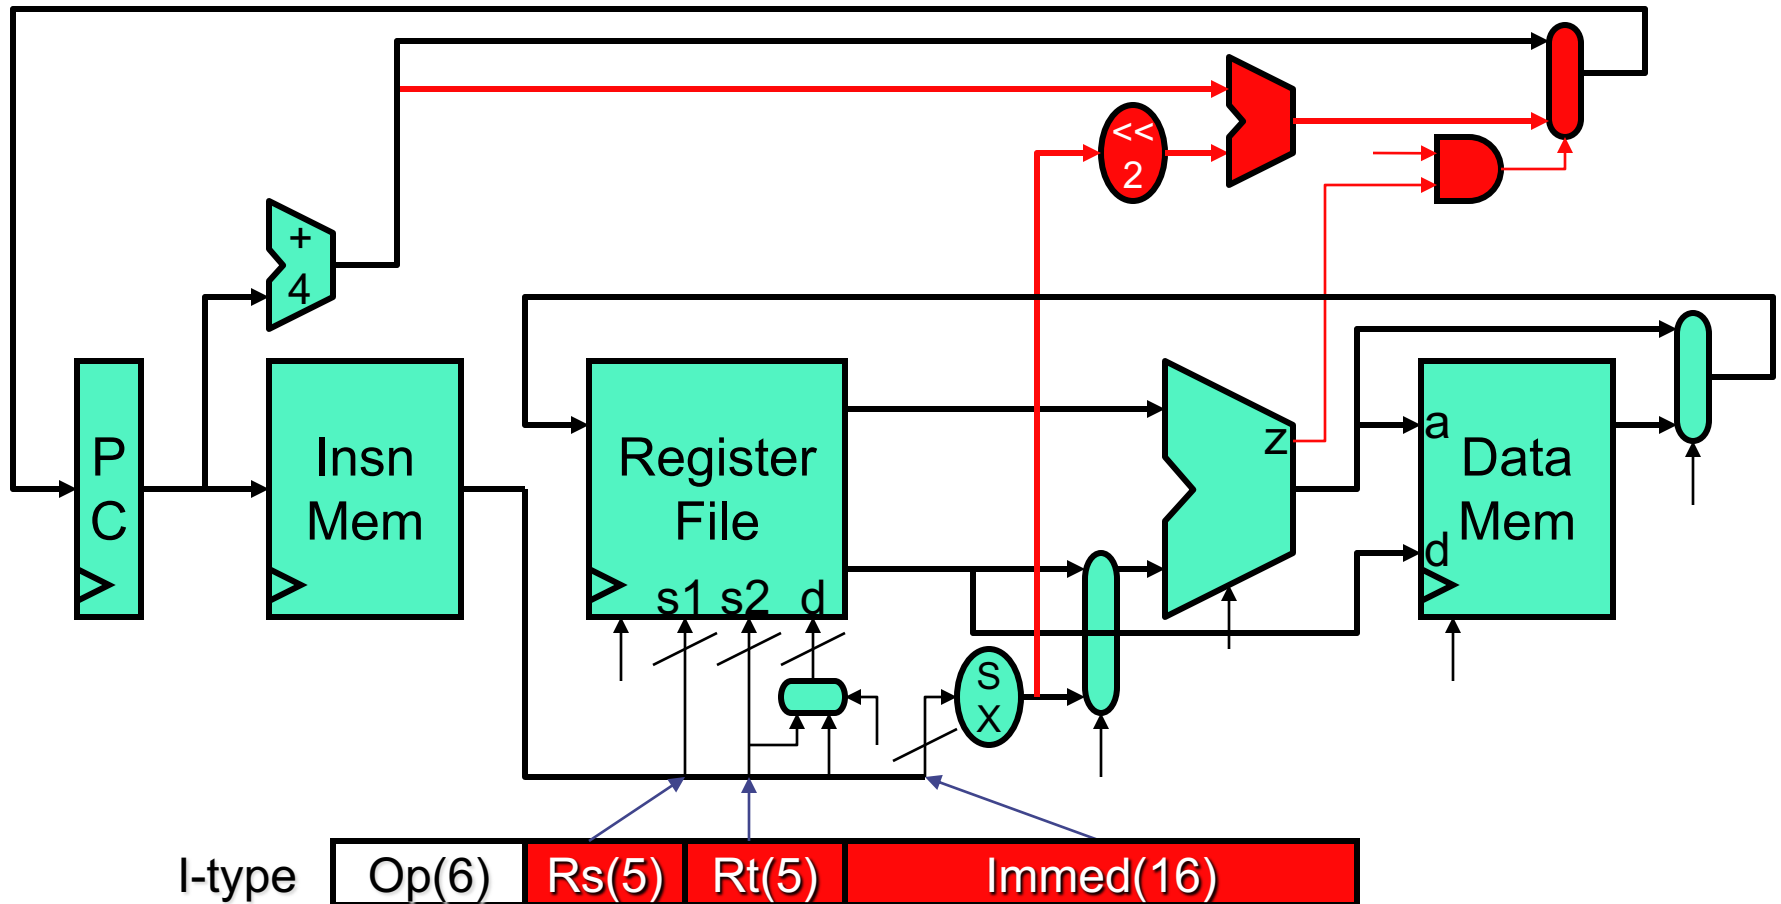
- Add left shift unit and adder to compute PC-relative branch target
- Add PC input mux to select PC+4 or branch target

# Another Use of Wire Concatenation

- How do you do <<2? **Wire concatenation**

```
wire [31:0] insn;
wire [25:0] imm26 = insn[25:0]
wire [31:0] imm26_shifted_by_2 = {4'b0000, imm26, 2'b00};
```

| Op(6) | Immed(26) |
|-------|-----------|

# Sixth Instruction: **j**



J-type | Op(6) | Immed(26)
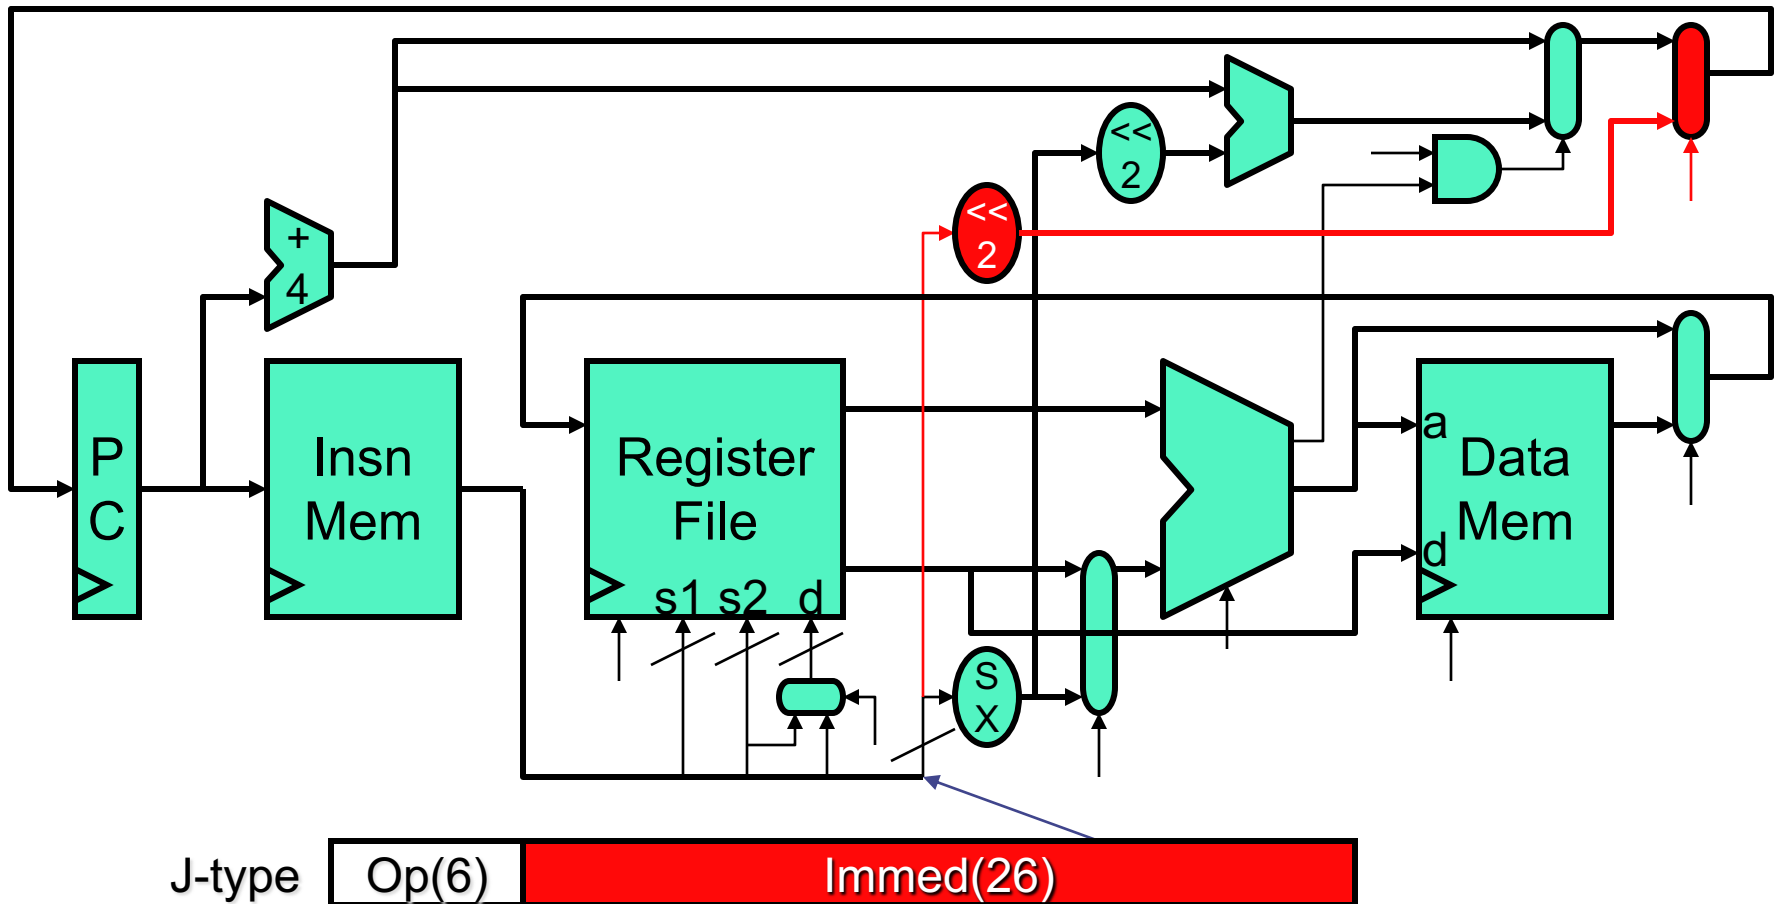
- Add shifter to compute left shift of 26-bit immediate
- Add additional PC input mux for jump target

# MIPS Control

# What Is Control?
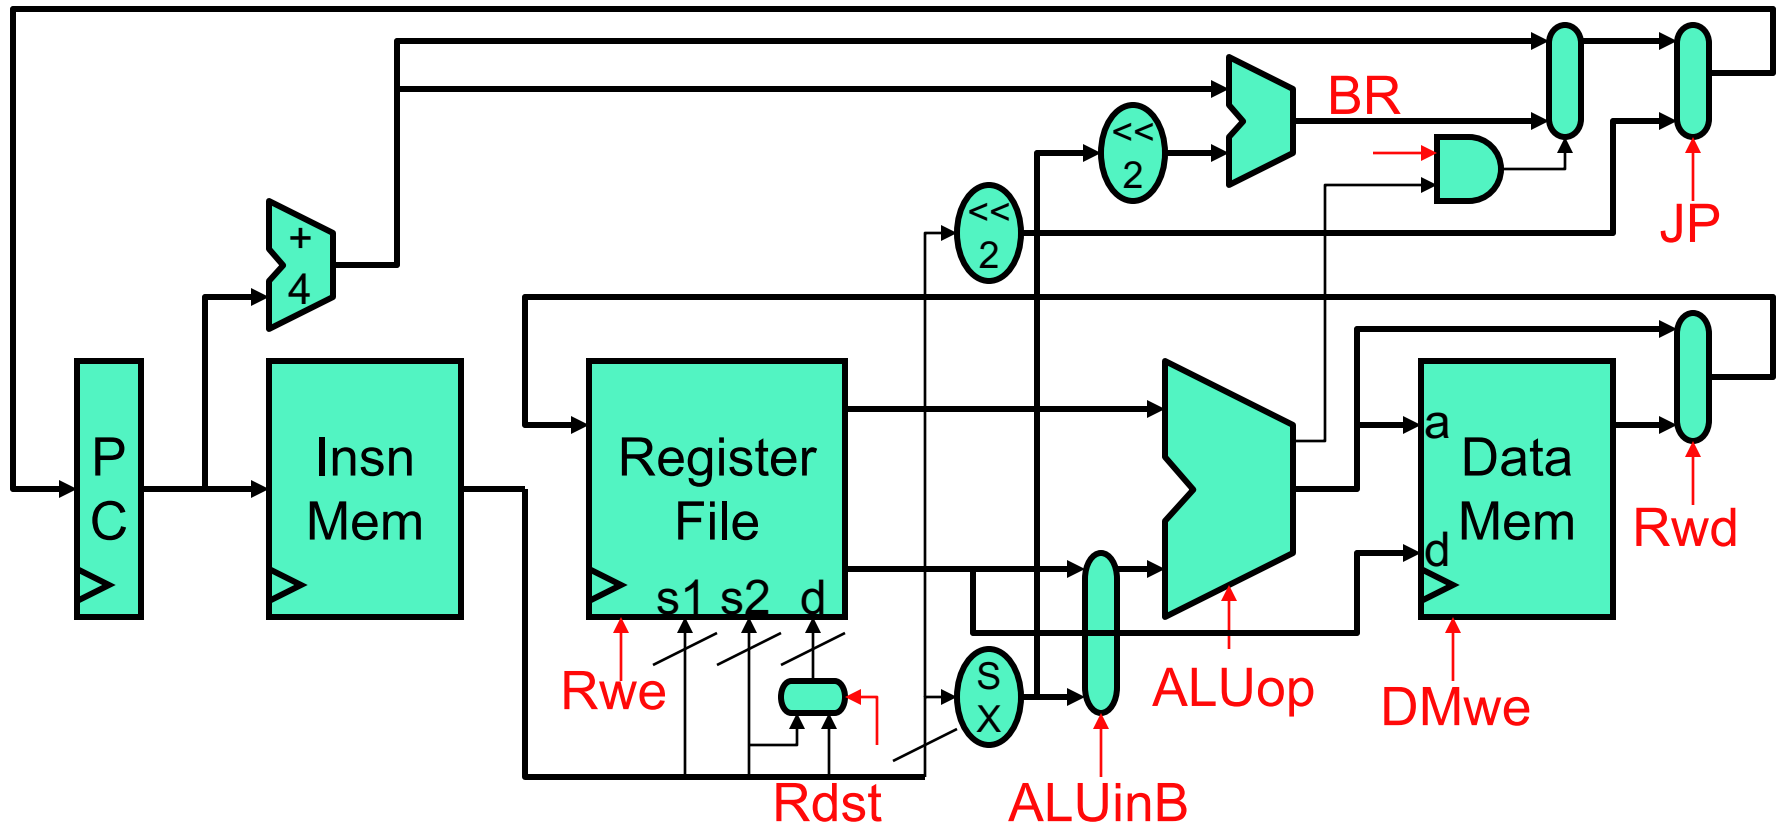


- 8 signals control flow of data through this datapath
  - MUX selectors, or register/memory write enable signals
  - A real datapath has 300-500 control signals

# Example: Control for **add**
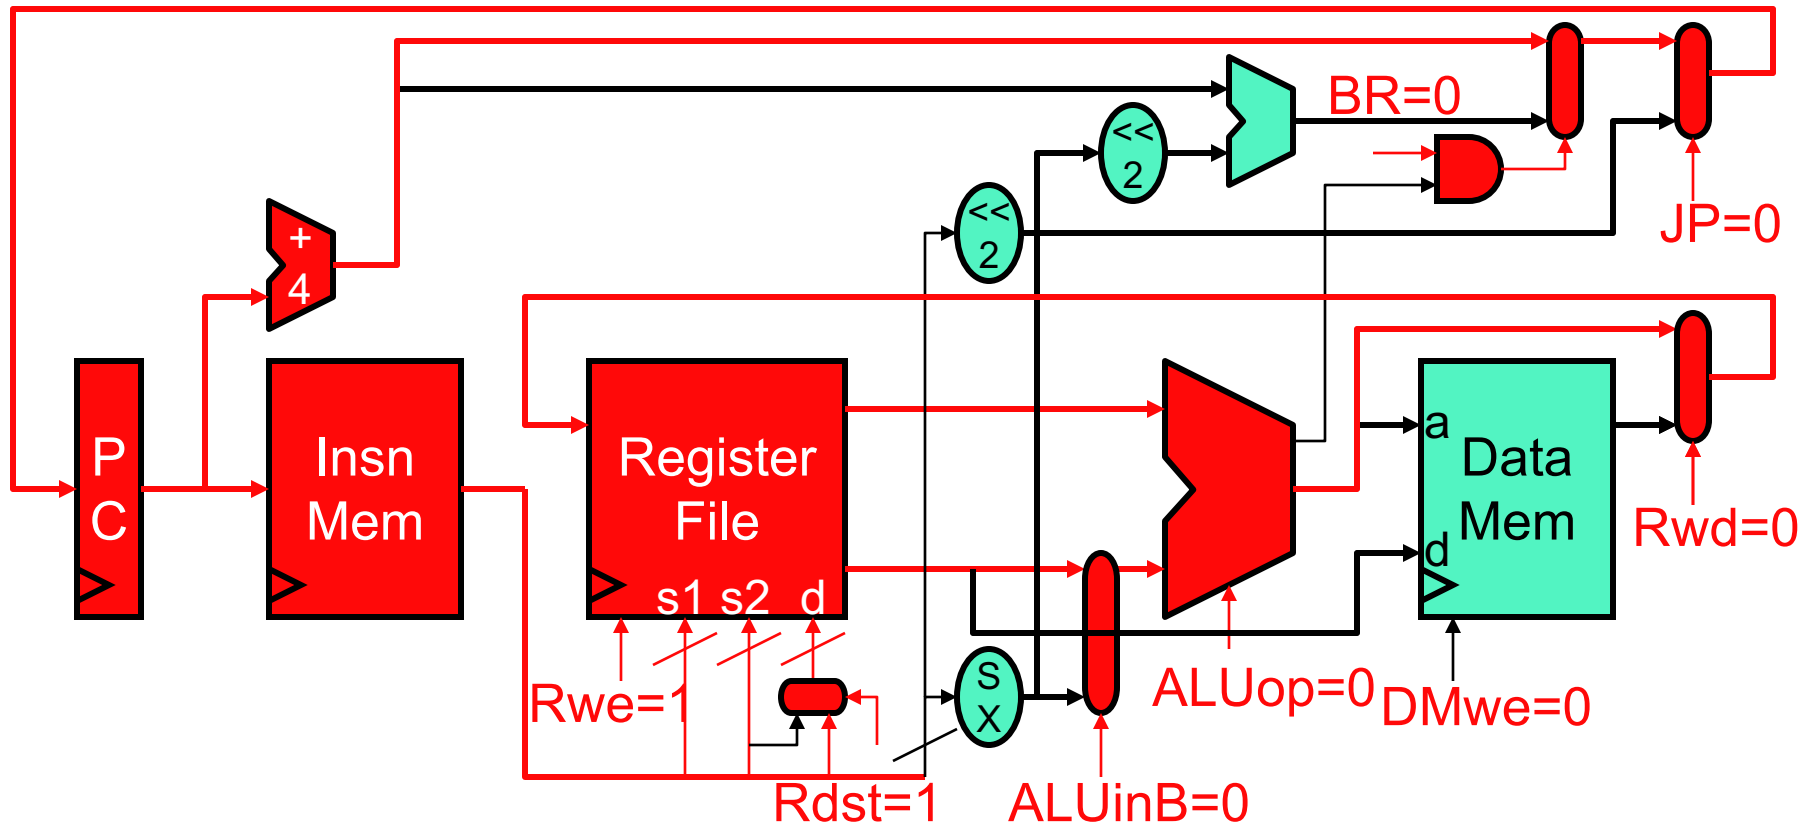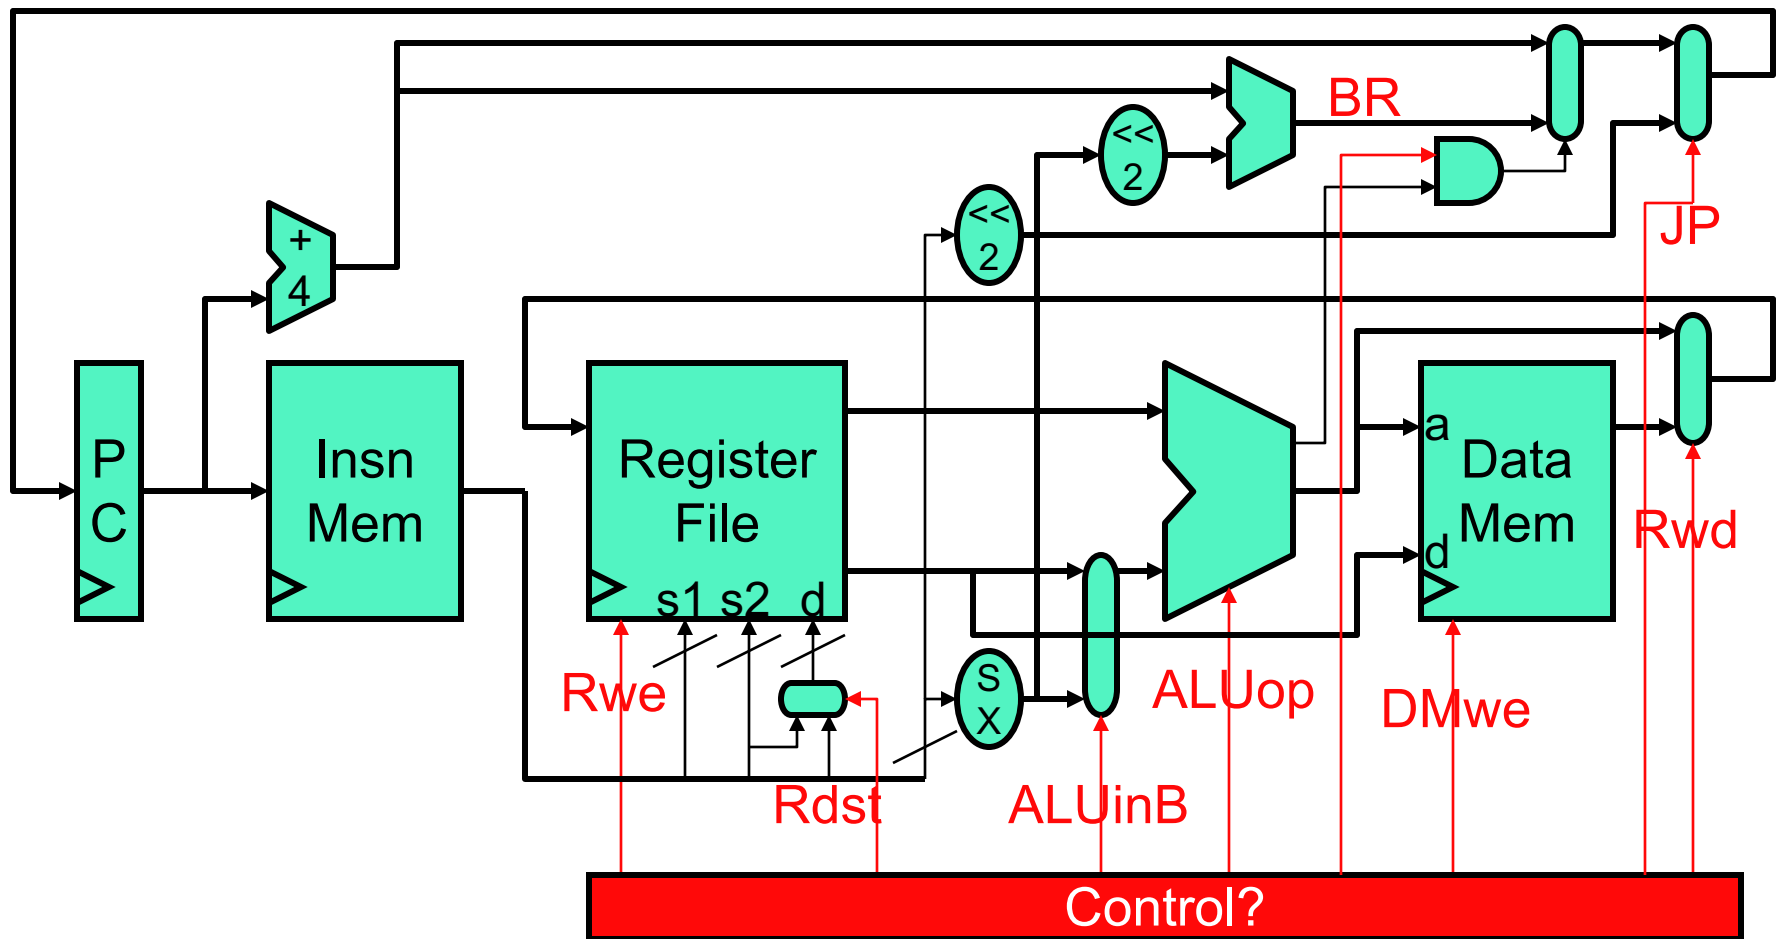
# Example: Control for **sw**
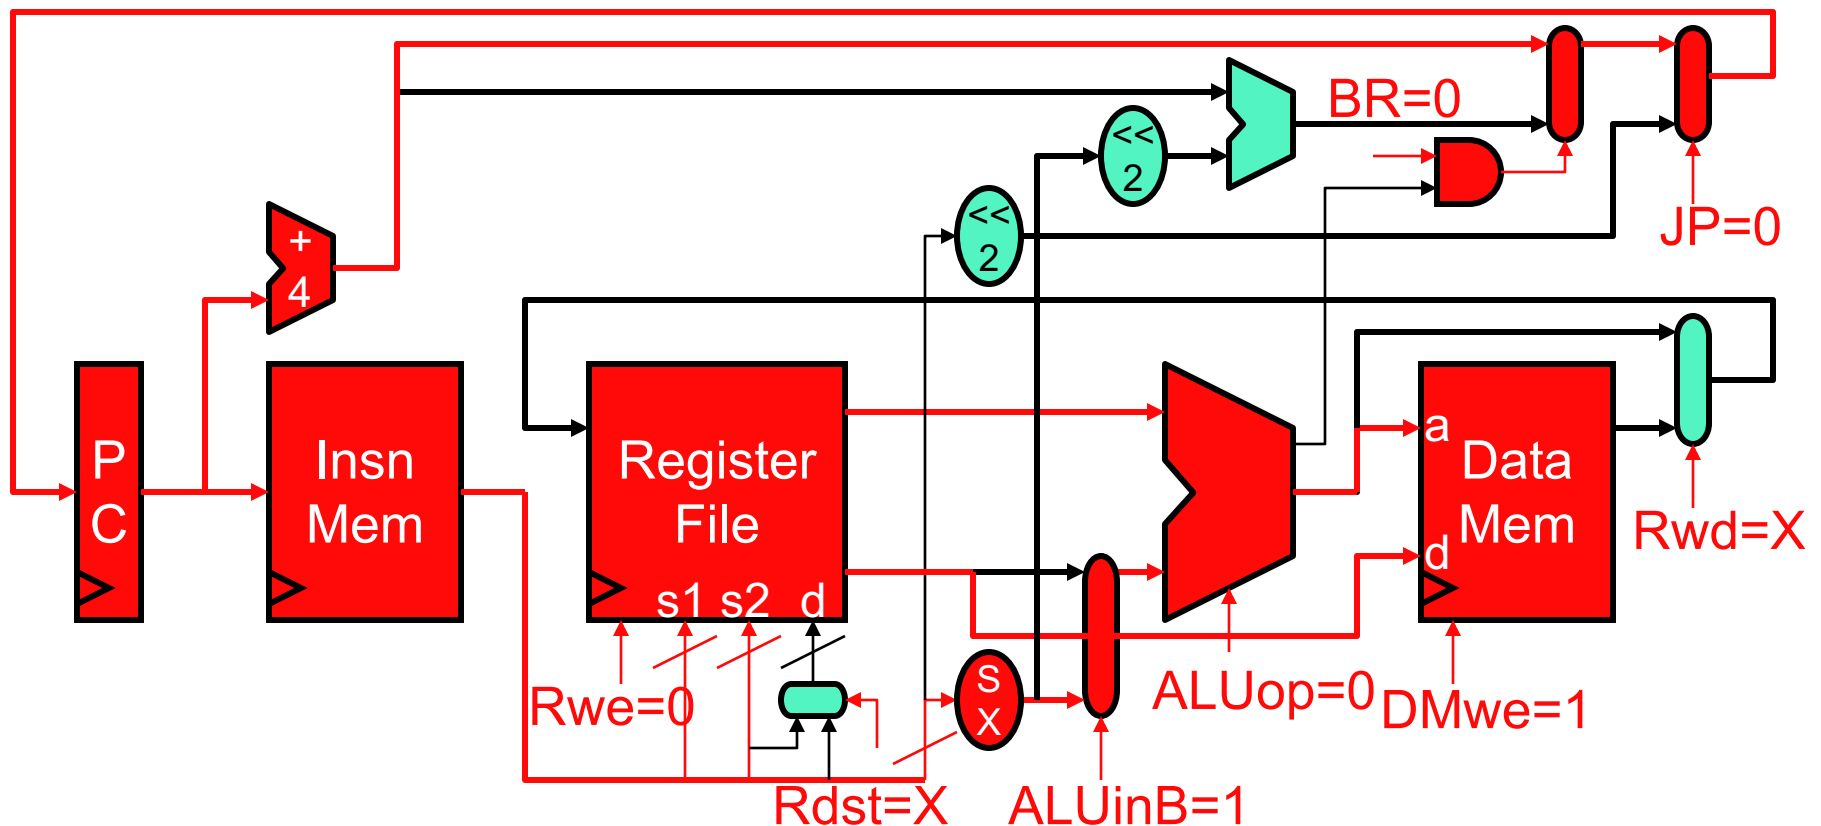


- Difference between **sw** and **add** is 5 signals
  - 3 if you don't count the X (don't care) signals

# Example: Control for **beq**
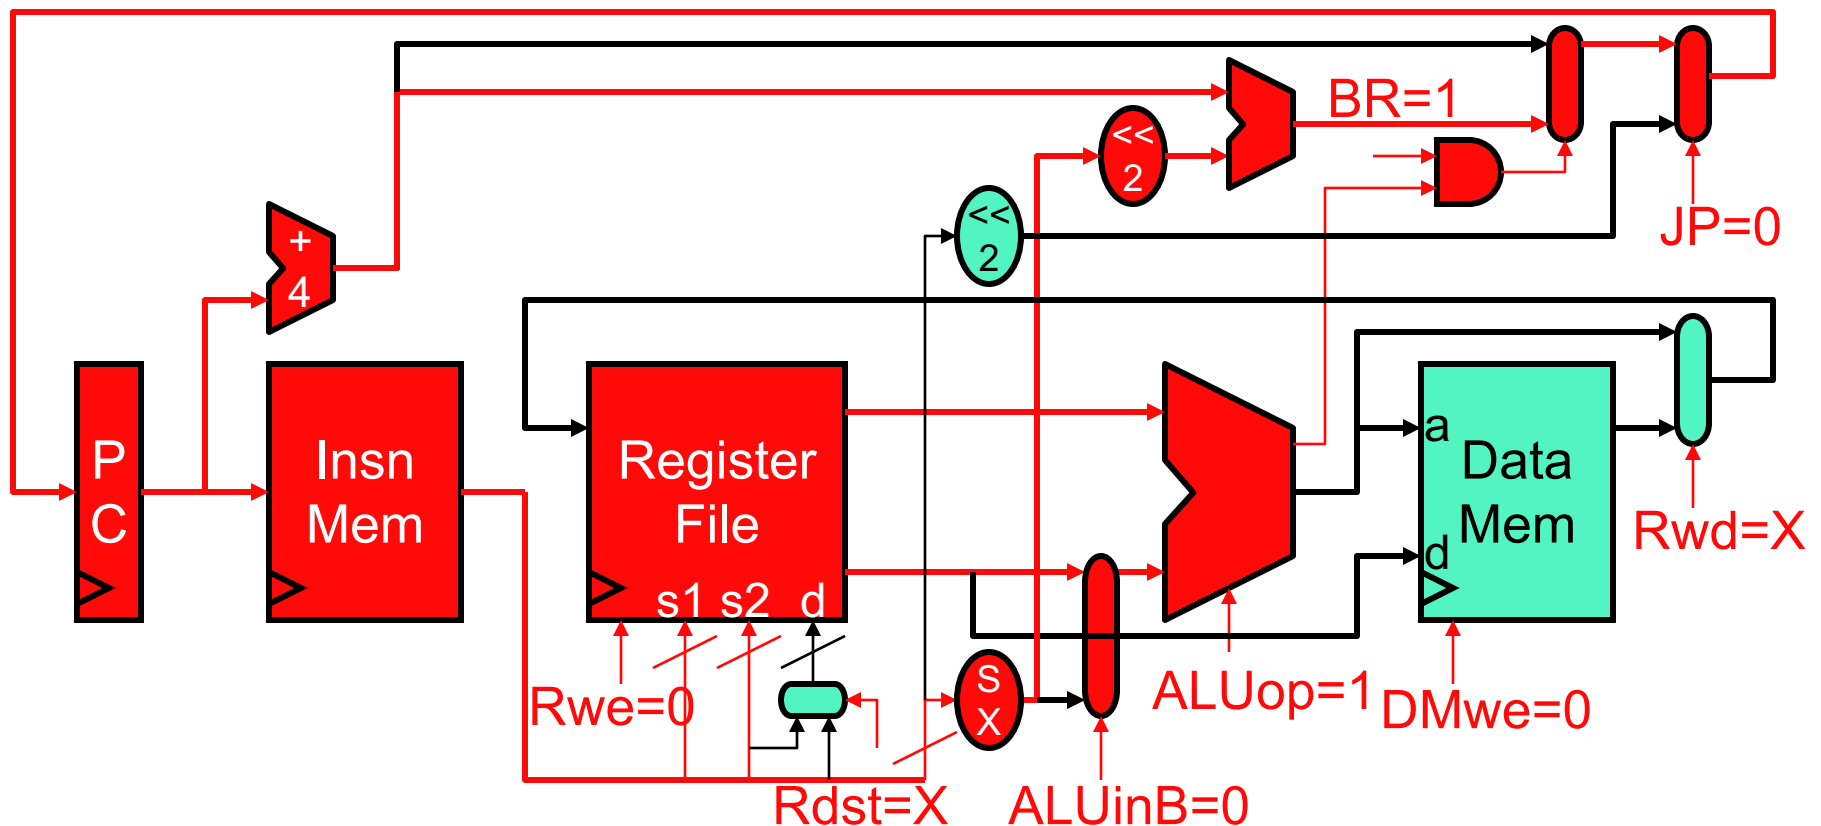


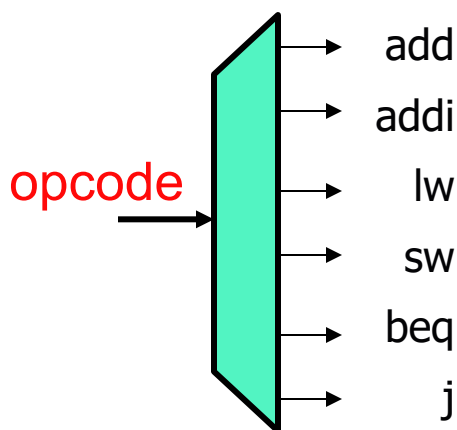- Difference between **sw** and **beq** is only 4 signals

# How Is Control Implemented?

# Implementing Control

- Each instruction has a unique set of control signals
  - Most are function of opcode
  - Some may be encoded in the instruction itself
    - E.g., the ALUop signal is some portion of the MIPS Func field
    + Simplifies controller implementation
    - Requires careful ISA design

# Control Implementation: ROM

- **ROM (read only memory)**: like a RAM but unwritable
  - Bits in data words are control signals
  - Lines indexed by opcode
  - Example: ROM control for 6-insn MIPS datapath
  - X is "don't care"

opcode

| | BR | JP | ALUinB | ALUop | DMwe | Rwe | Rdst | Rwd |
|------|----|----|--------|-------|------|-----|------|-----|
| add | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 |
| addi | 0 | 0 | 1 | 0 | 0 | 1 | 1 | 0 |
| lw | 0 | 0 | 1 | 0 | 0 | 1 | 1 | 1 |
| sw | 0 | 0 | 1 | 0 | 1 | 0 | X | X |
| beq | 1 | 0 | 0 | 1 | 0 | 0 | X | X |
| j | 0 | 1 | 0 | 0 | 0 | 0 | X | X |

# Control Implementation: Logic

- Real machines have 100+ insns 300+ control signals
  - 30,000+ control bits (~4KB)
  - Not huge, but hard to make faster than datapath (important!)
- Alternative: **logic gates** or "random logic" (unstructured)
  - Exploits the observation: many signals have few 1s or few 0s
  - Example: random logic control for 6-insn MIPS datapath
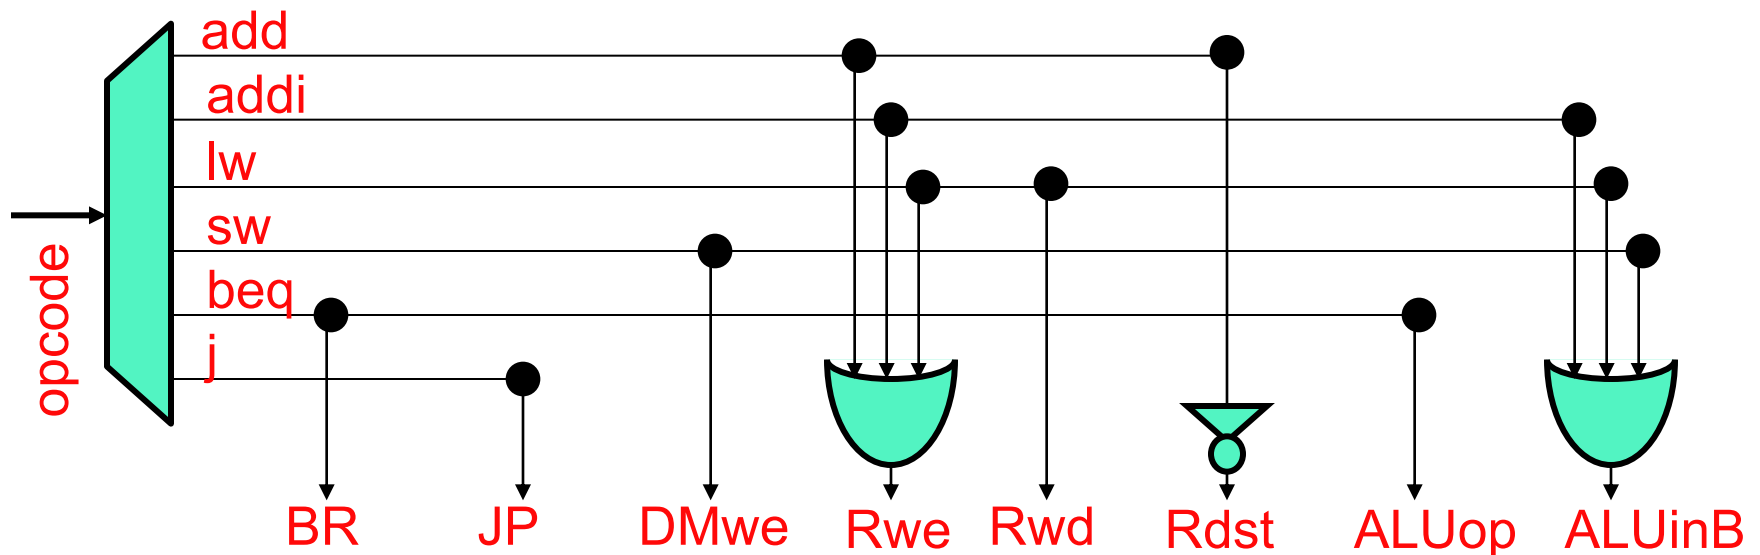
# Control Logic in Verilog

```verilog
wire [31:0] insn;
wire [5:0] func = insn[5:0]
wire [5:0] opcode = insn[31:26];
wire is_add = ((opcode == 6'h00) & (func == 6'h20));
wire is_addi = (opcode == 6'h0F);
wire is_lw = (opcode == 6'h23);
wire is_sw = (opcode == 6'h2A);
wire ALUinB = is_addi | is_lw | is_sw;
wire Rwe = is_add | is_addi | is_lw;
wire Rwd = is_lw;
wire Rdst = ~is_add;
wire DMwe = is_sw;
```

add
addi
lw
sw

opcode

DMwe   Rwe   Rwd   Rdst   ALUinB
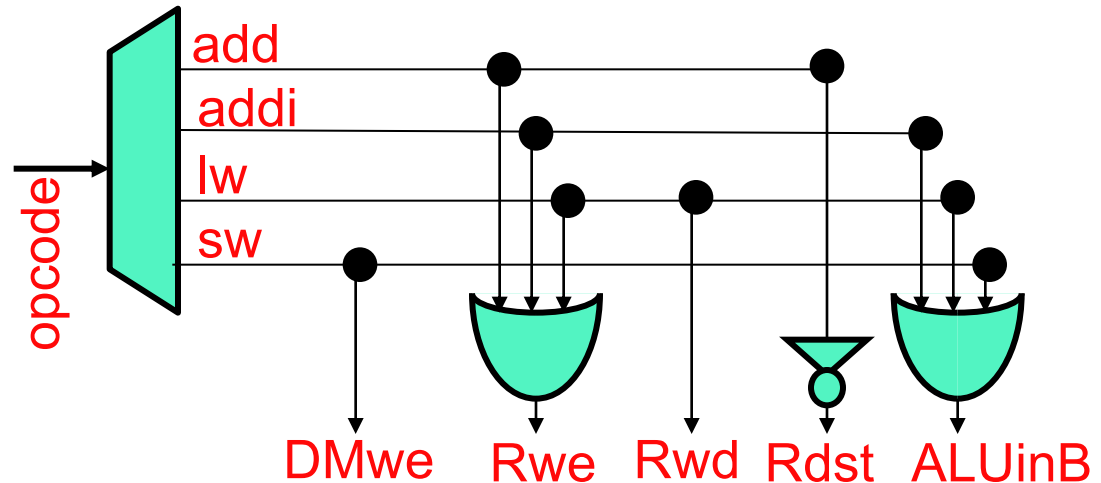
# Datapath Storage Elements

# Register File
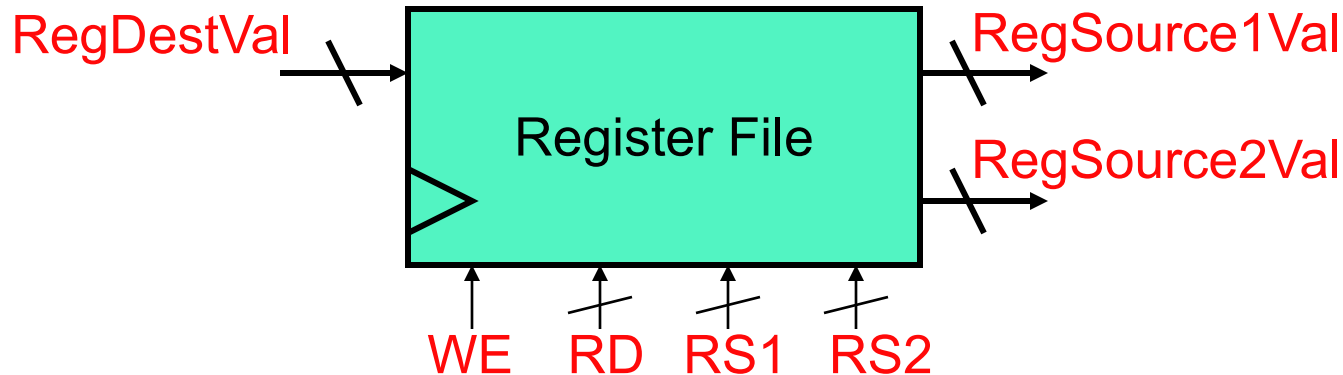


- **Register file**: M N-bit storage words
  - Multiplexed input/output: data buses write/read "random" word
- **"Port"**: set of buses for accessing a random word in array
  - Data bus (N-bits) + address bus ($\log_2$M-bits) + optional WE bit
  - P ports = P parallel and independent accesses
- MIPS integer register file
  - 32 32-bit words, two read ports + one write port (why?)

# Decoder

- **Decoder**: converts binary integer to "1-hot" representation
  - Binary representation of $0 \ldots 2^N - 1$: N bits
  - 1 hot representation of $0 \ldots 2^N - 1$: $2^N$ bits
    - J represented as $J^{th}$ bit 1, all other bits zero
  - Example below: 2-to-4 decoder
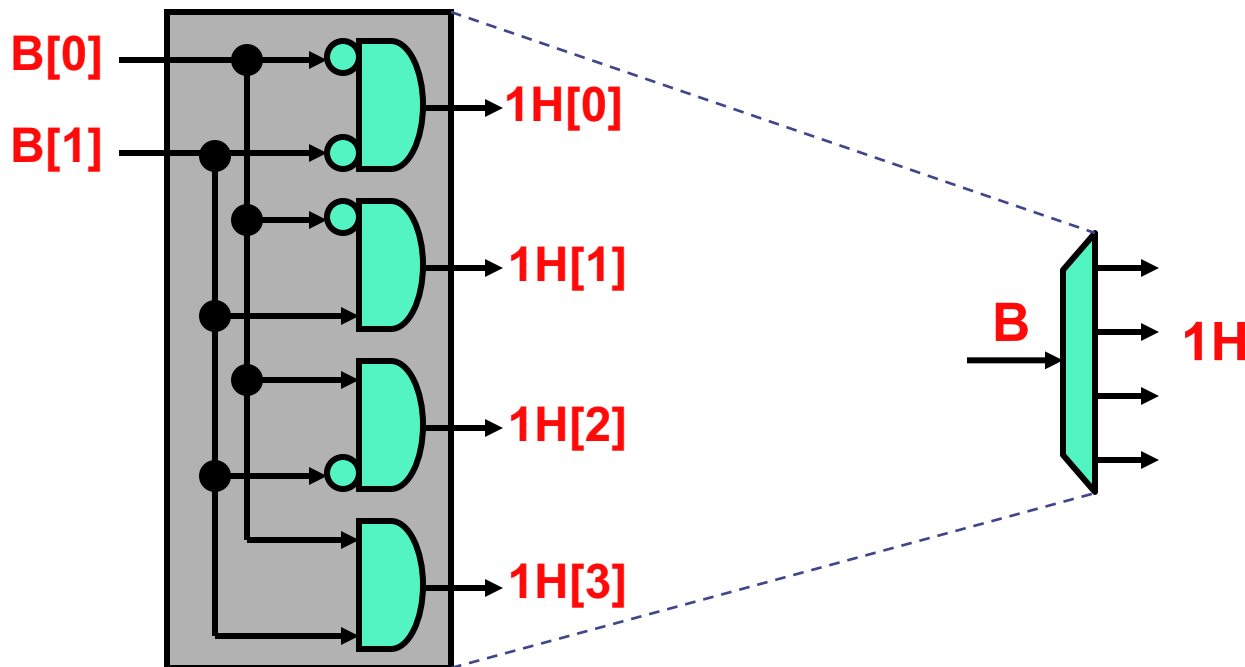
# Decoder in Verilog (1 of 2)

```verilog
module decoder_2_to_4 (binary_in, onehot_out);
    input [1:0] binary_in;
    output [3:0] onehot_out;
    assign onehot_out[0] = (~binary_in[0] & ~binary_in[1]);
    assign onehot_out[1] = (~binary_in[0] & binary_in[1]);
    assign onehot_out[2] = (binary_in[0] & ~binary_in[1]);
    assign onehot_out[3] = (binary_in[0] & binary_in[1]);
endmodule
```

- Is there a simpler way?

```
module decoder_2_to_4 (binary_in, onehot_out);
    input [1:0] binary_in;
    output [3:0] onehot_out;
    assign onehot_out[0] = (binary_in == 2'd0);
    assign onehot_out[1] = (binary_in == 2'd1);
    assign onehot_out[2] = (binary_in == 2'd2);
    assign onehot_out[3] = (binary_in == 2'd3);
endmodule
```

- How is "a == b" implemented for vectors?
    - ~|(a ^ b)     (this is a "nor" reduction of bitwise "a xor b")
    - When one of the inputs to "==" is a constant
        - Simplifies to simpler inverter on bits with "one" in constant
        - Exactly what was on previous slide! (apply DeMorgan's law)

# Register File Interface

RDestVal
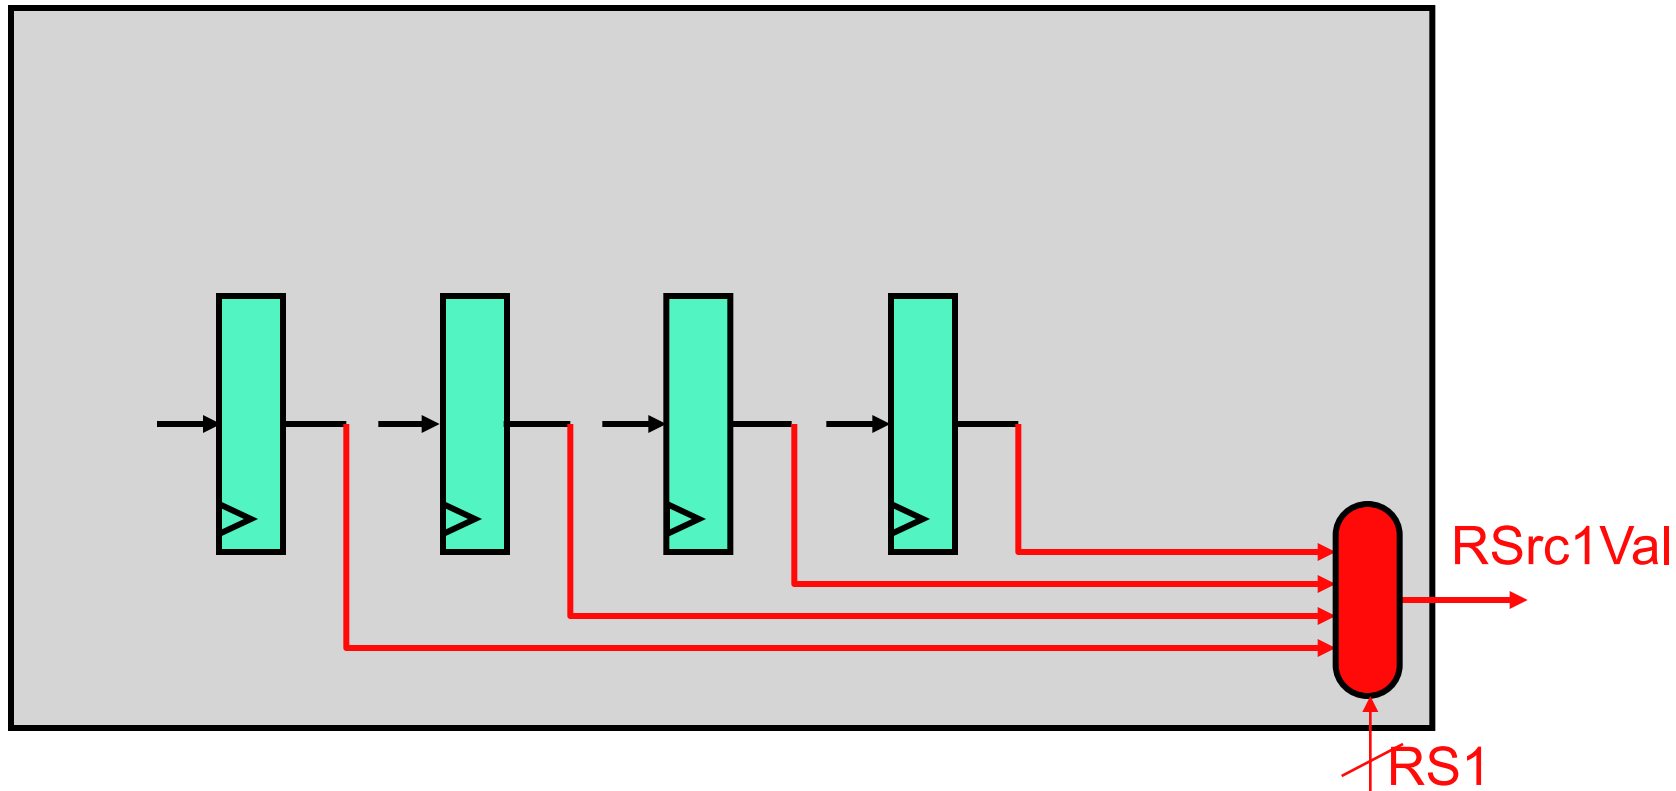
RSrc2Val

RSrc1Val

WE　RD

RS2　RS1

- Inputs:
  - RS1, RS2 (reg. sources to read), RD (reg. destination to write)
  - WE (write enable), RDestVal (value to write)
- Outputs: RSrc1Val, RSrc2Val (value of RS1 & RS2 registers)
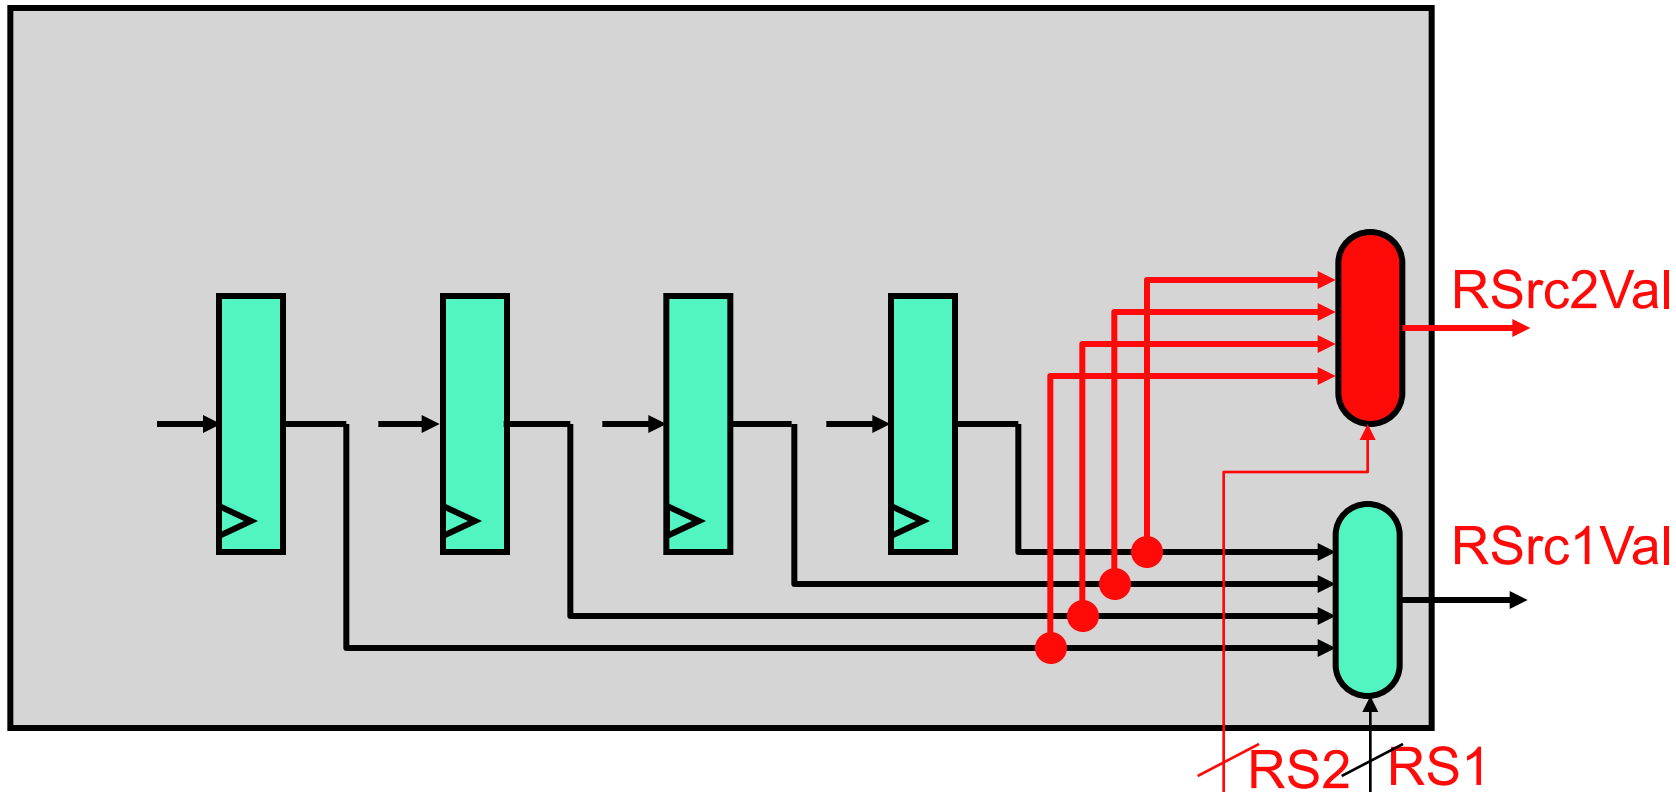
# Register File: Four Registers
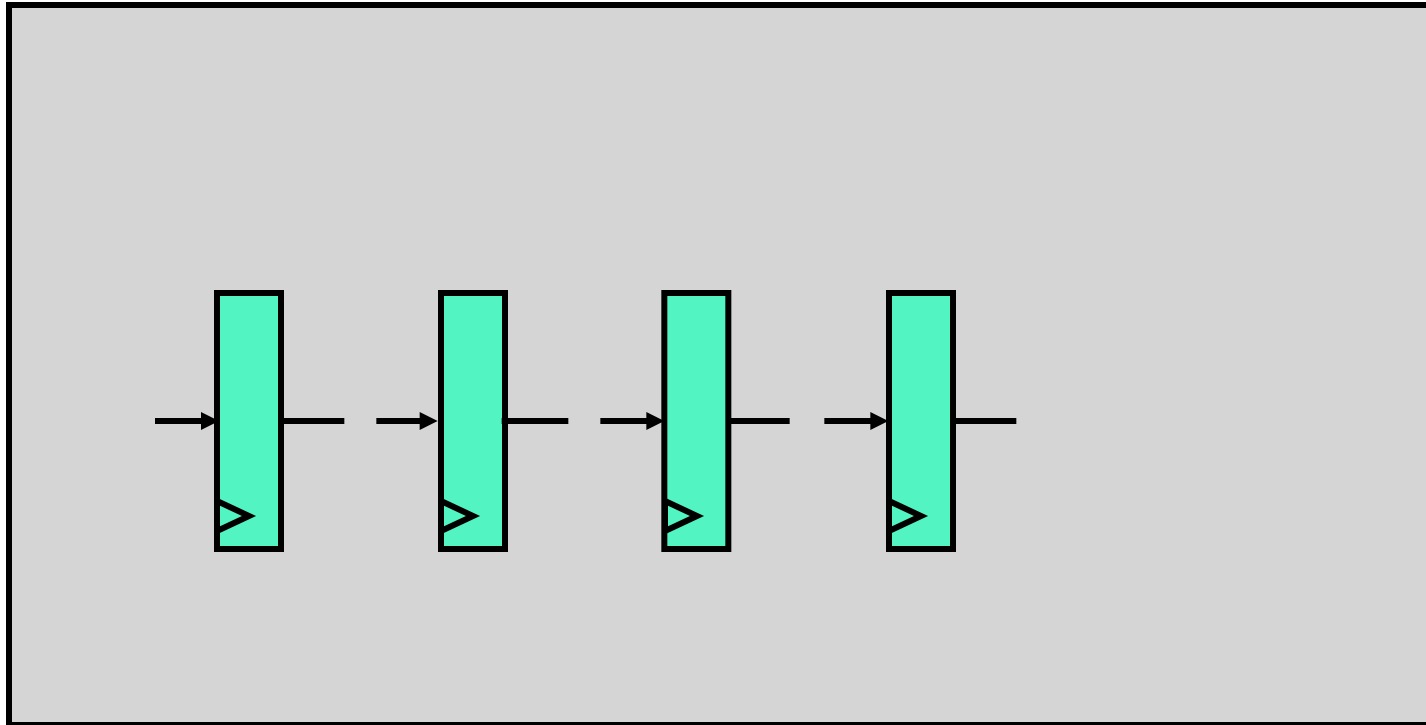


- Register file with four registers
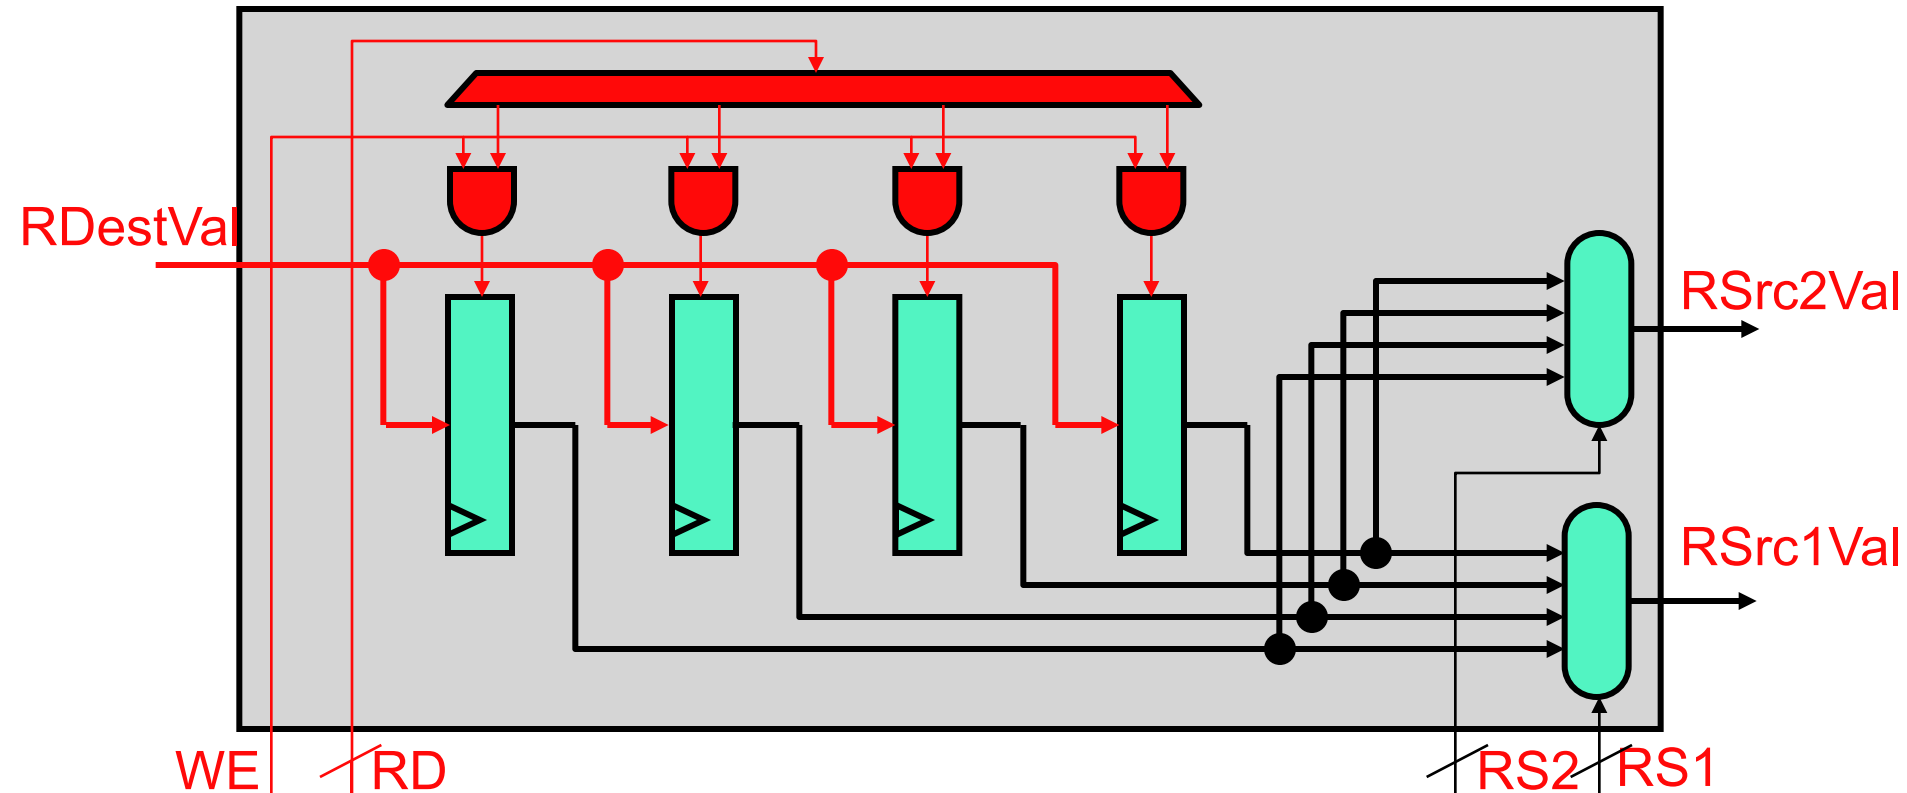
# Add a Read Port



- Output of each register into 4to1 mux (RSrc1Val)
  - RS1 is select input of RSrc1Val mux

# Add Another Read Port



- Output of each register into another 4to1 mux (RSrc2Val)
  - RS2 is select input of RSrc2Val mux

# Add a Write Port



- Input RegDestVal into each register
  - Enable only one register's WE: (Decoded RD) & (WE)
- What if we needed two write ports?

# Register File Interface (Verilog)

```verilog
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
   parameter n = 1;
   input [1:0] rs1, rs2, rd;
   input we, rst, clk;
   input [n-1:0] rdval;
   output [n-1:0] rs1val, rs2val;

   …

endmodule
```

- Building block modules:
  - **module register (out, in, wen, rst, clk);**
  - **module decoder_2_to_4 (binary_in, onehot_out)**
  - **module Nbit_mux4to1 (sel, a, b, c, d, out);**

# Register File Interface (Verilog)

```
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
    input [1:0] rs1, rs2, rd;
    input we, rst, clk;
    input [15:0] rdval;
    output [15:0] rs1val, rs2val;




endmodule
```

- Warning: this code not tested, may contain typos, do not blindly trust!

# Register File Interface (Verilog)

```verilog
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
   parameter n = 1;
   input [1:0] rs1, rs2, rd;
   input we, rst, clk;
   input [n-1:0] rdval;
   output [n-1:0] rs1val, rs2val;




   endmodule
```

- Warning: this code not tested, may contain typos, do not blindly trust!

# Register File: Four Registers (Verilog)

```verilog
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
   parameter n = 1;
   input [1:0] rs1, rs2, rd;
   input we, rst, clk;
   input [n-1:0] rdval;
   output [n-1:0] rs1val, rs2val;
   wire [n-1:0] r0v, r1v, r2v, r3v;


   Nbit_reg #(n) r0 (r0v,      ,              , rst, clk);
   Nbit_reg #(n) r1 (r1v,      ,              , rst, clk);
   Nbit_reg #(n) r2 (r2v,      ,              , rst, clk);
   Nbit_reg #(n) r3 (r3v,      ,              , rst, clk);


endmodule
```
- Warning: this code not tested, may contain typos, do not blindly trust!

# Add a Read Port (Verilog)

```verilog
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
   parameter n = 1;
   input [1:0] rs1, rs2, rd;
   input we, rst, clk;
   input [n-1:0] rdval;
   output [n-1:0] rs1val, rs2val;
   wire [n-1:0] r0v, r1v, r2v, r3v;


   Nbit_reg #(n) r0 (r0v,       ,                  , rst, clk);
   Nbit_reg #(n) r1 (r1v,       ,                  , rst, clk);
   Nbit_reg #(n) r2 (r2v,       ,                  , rst, clk);
   Nbit_reg #(n) r3 (r3v,       ,                  , rst, clk);
   Nbit_mux4to1 #(n) mux1 (rs1, r0v, r1v, r2v, r3v, rs1val);

endmodule
```

- Warning: this code not tested, may contain typos, do not blindly trust!

# Add Another Read Port (Verilog)

```verilog
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
   parameter n = 1;
   input [1:0] rs1, rs2, rd;
   input we, rst, clk;
   input [n-1:0] rdval;
   output [n-1:0] rs1val, rs2val;
   wire [n-1:0] r0v, r1v, r2v, r3v;


   Nbit_reg #(n) r0 (r0v,       ,                    , rst, clk);
   Nbit_reg #(n) r1 (r1v,       ,                    , rst, clk);
   Nbit_reg #(n) r2 (r2v,       ,                    , rst, clk);
   Nbit_reg #(n) r3 (r3v,       ,                    , rst, clk);
   Nbit_mux4to1 #(n) mux1 (rs1, r0v, r1v, r2v, r3v, rs1val);
   Nbit_mux4to1 #(n) mux2 (rs2, r0v, r1v, r2v, r3v, rs2val);
endmodule
```

- Warning: this code not tested, may contain typos, do not blindly trust!

# Add a Write Port (Verilog)

```verilog
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
  parameter n = 1;
  input [1:0] rs1, rs2, rd;
  input we, rst, clk;
  input [n-1:0] rdval;
  output [n-1:0] rs1val, rs2val;
  wire [n-1:0] r0v, r1v, r2v, r3v;
  wire [3:0] rd_select;
  decoder_2_to_4 dec (rd, rd_select);
  Nbit_reg #(n) r0 (r0v, rdval, rd_select[0] & we, rst, clk);
  Nbit_reg #(n) r1 (r1v, rdval, rd_select[1] & we, rst, clk);
  Nbit_reg #(n) r2 (r2v, rdval, rd_select[2] & we, rst, clk);
  Nbit_reg #(n) r3 (r3v, rdval, rd_select[3] & we, rst, clk);
  Nbit_mux4to1 #(n) mux1 (rs1, r0v, r1v, r2v, r3v, rs1val);
  Nbit_mux4to1 #(n) mux2 (rs2, r0v, r1v, r2v, r3v, rs2val);
endmodule
```

- Warning: this code not tested, may contain typos, do not blindly trust!

# Final Register File (Verilog)

```verilog
module regfile4(rs1, rs1val, rs2, rs2val, rd, rdval, we, rst, clk);
  parameter n = 1;
  input [1:0] rs1, rs2, rd;
  input we, rst, clk;
  input [n-1:0] rdval;
  output [n-1:0] rs1val, rs2val;
  wire [n-1:0] r0v, r1v, r2v, r3v;


  Nbit_reg #(n) r0 (r0v, rdval, (rd == 2'd0) & we, rst, clk);
  Nbit_reg #(n) r1 (r1v, rdval, (rd == 2'd1) & we, rst, clk);
  Nbit_reg #(n) r2 (r2v, rdval, (rd == 2'd2) & we, rst, clk);
  Nbit_reg #(n) r3 (r3v, rdval, (rd == 2'd3) & we, rst, clk);
  Nbit_mux4to1 #(n) mux1 (rs1, r0v, r1v, r2v, r3v, rs1val);
  Nbit_mux4to1 #(n) mux2 (rs2, r0v, r1v, r2v, r3v, rs2val);
endmodule
```

- **Warning: this code not tested, may contain typos, do not blindly trust!**

# Another Useful Component: Memory
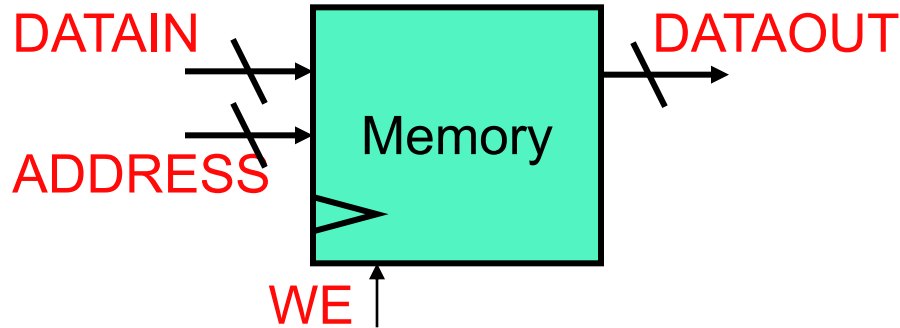


- Register file: M N-bit storage words
  - Few words (< 256), many ports, dedicated read and write ports
- **Memory**: M N-bit storage words, yet not a register file
  - Many words (> 1024), few ports (1, 2), shared read/write ports
- Leads to different implementation choices
  - Lots of circuit tricks and such
  - Larger memories typically only 6 transistors per bit
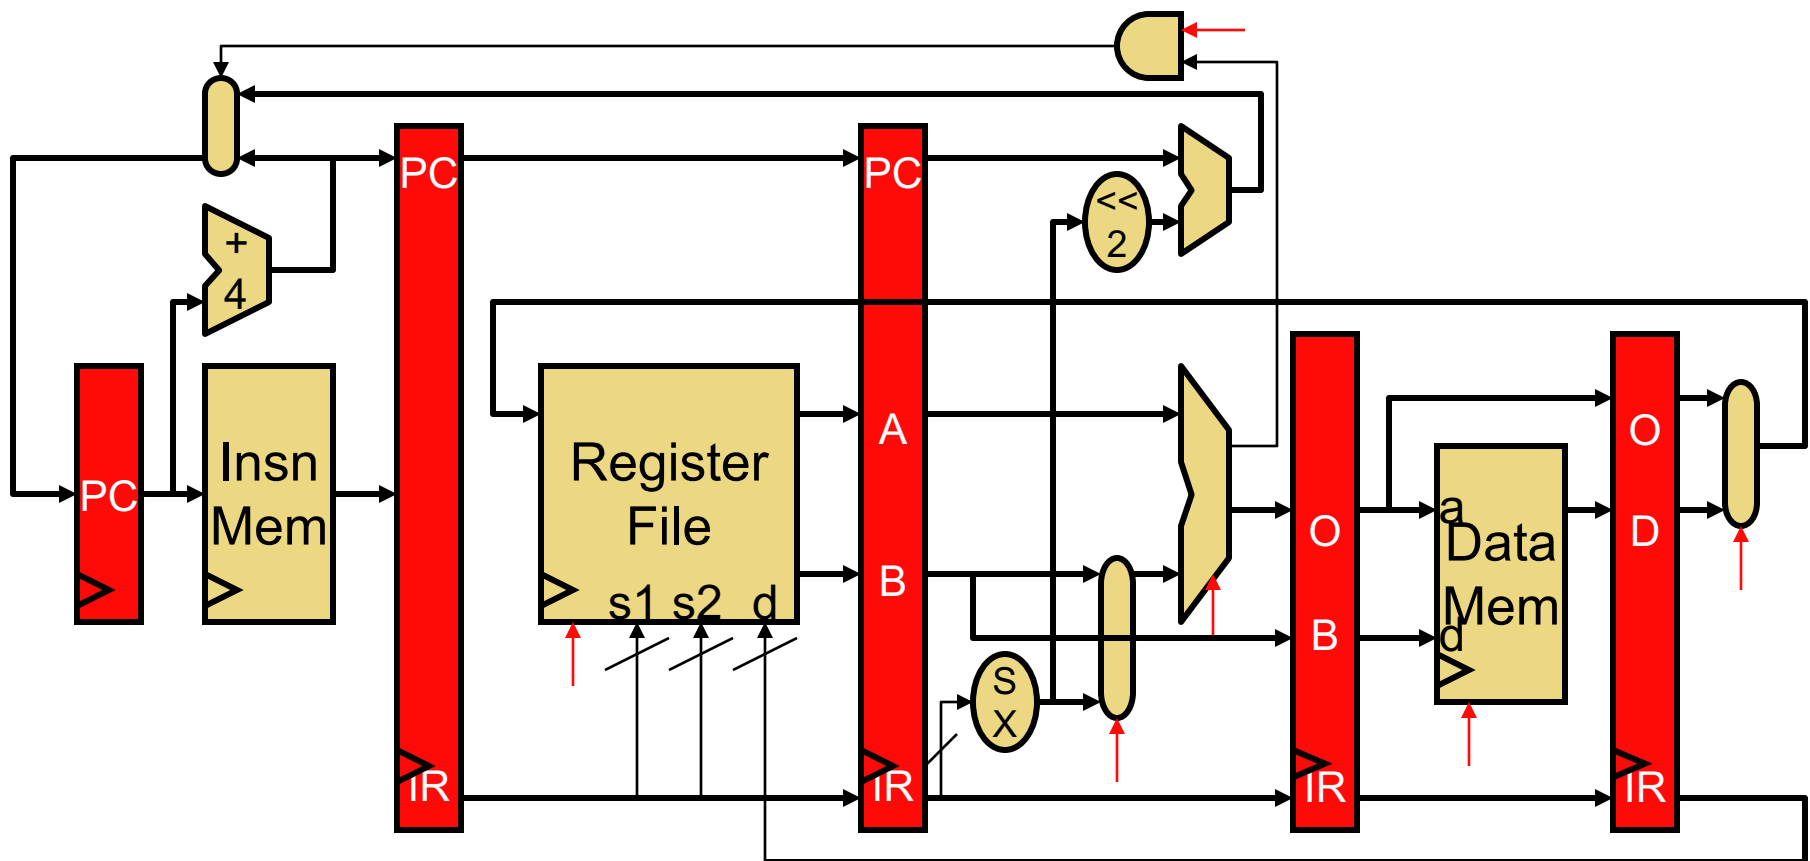- In Verilog?  We'll give you the code for memories

# Single-Cycle Performance
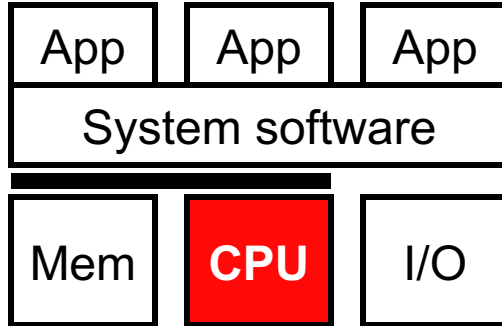
# Single-Cycle Datapath Performance
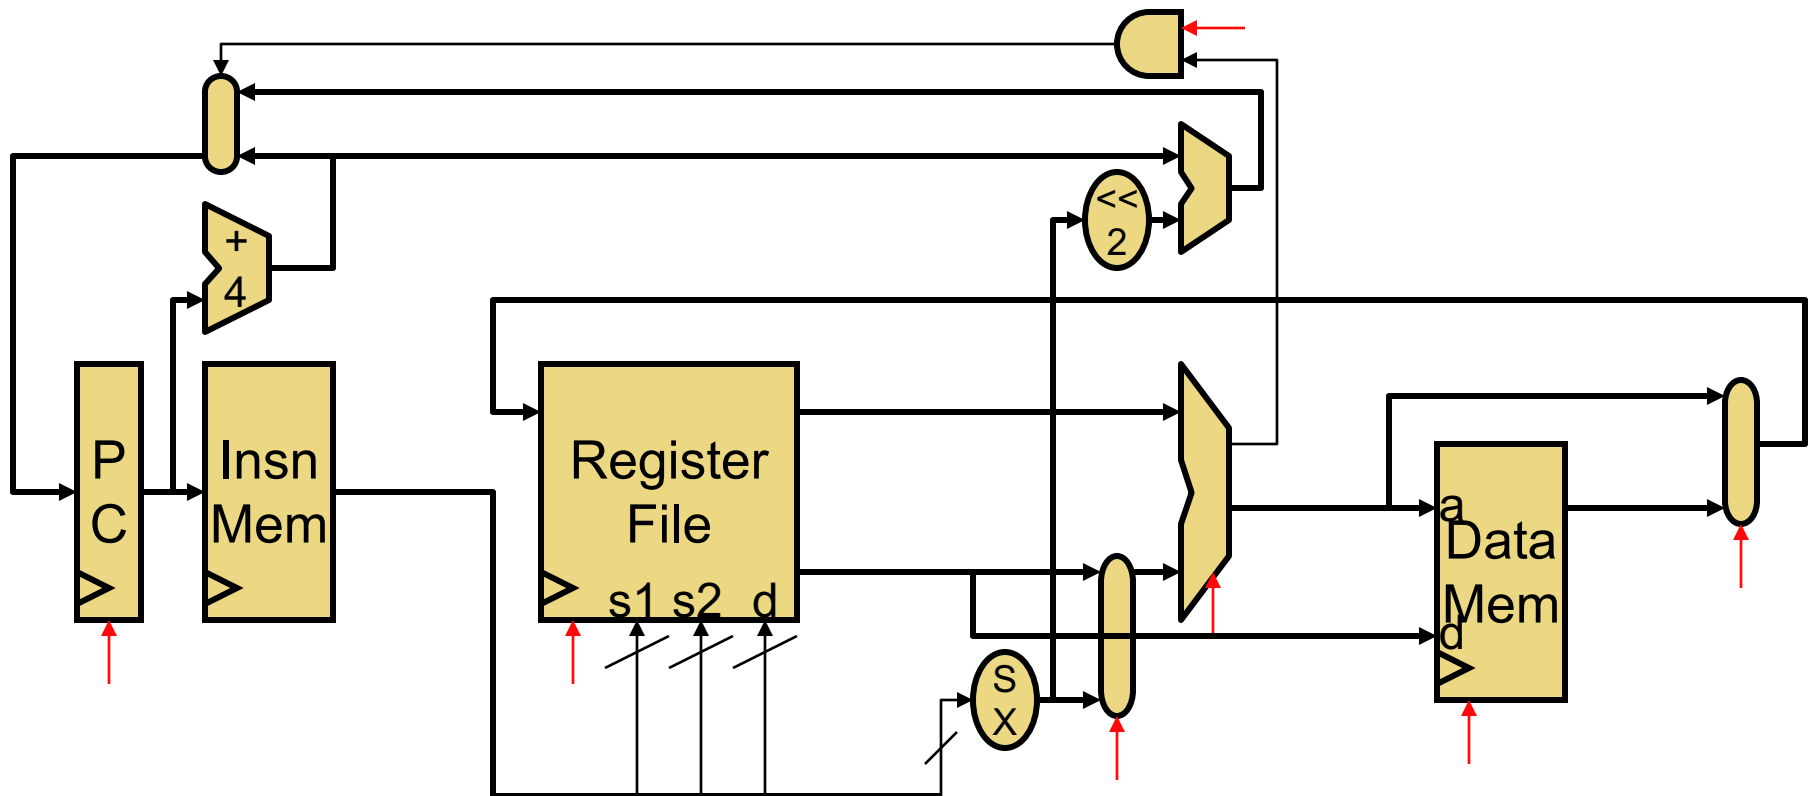


- One cycle per instruction (CPI)
- **Clock cycle time proportional to worst-case logic delay**
  - In this datapath: insn fetch, decode, register read, ALU, data memory access, write register
  - Can we do better?

# Foreshadowing: Pipelined Datapath



- Split datapath into multiple stages
  - Assembly line analogy
  - 5 stages results in up to 5x clock & performance improvement

# Summary

| App | App | App |
|-----|-----|-----|
| System software | | |

| Mem | **CPU** | I/O |
|-----|---------|-----|

- Overview of ISAs
- Datapath storage elements
- MIPS Datapath
- MIPS Control