


CIS 5710

Computer Organization and Design

Unit 6: Pipelining

Based on slides by Profs. Amir Roth, Milo Martin & C.J. Taylor

This Unit: Pipelining




- Processor performance
 - Latency vs throughput
- Single-cycle & multi-cycle datapaths
- Basic pipelining
- Data hazards
 - Software interlocks and scheduling
 - Hardware interlocks and stalling
 - Bypassing
 - Load-use stalling
- Pipelined multi-cycle operations
- Control hazards
 - Branch prediction



Readings

- P&H
 - Chapter 4

Welcome to the Laundromat



Henry Ford's Big Idea:



In-Class Exercise

- You have a washer, dryer, and “folding robot”
 - Each takes 1 unit of time per load
 - How long for one load in total?
 - How long for two loads of laundry?
 - How long for 100 loads of laundry?
- Now assume:
 - Washing takes 30 minutes, drying 60 minutes, and folding 15 min
 - How long for one load in total?
 - How long for two loads of laundry?
 - How long for 100 loads of laundry?

In-Class Exercise Answers

- You have a washer, dryer, and “folding robot”
 - Each takes 1 unit of time per load
 - How long for one load in total?
 - How long for two loads of laundry?
 - How long for 100 loads of laundry?
- Now assume:
 - Washing takes 30 minutes, drying 60 minutes, and folding 15 min
 - How long for one load in total? $30+60+15=$ **1h45m**
 - How long for two loads of laundry? $30+(60*2)+15 =$ **2h45m**
 - How long for 100 loads of laundry? $30+(60*100)+15=$ **100h45m**

Performance

Processor Performance Equation

$$\text{Execution time} = \text{"seconds per program"} = (\text{instructions/program}) * (\text{seconds/cycle}) * (\text{cycles/instruction})$$

$$(1 \text{ billion instructions}) * (1\text{ns per cycle}) * (1 \text{ cycle per insn}) = 1 \text{ second}$$

- Instructions per program: "dynamic instruction count"
 - Runtime count of instructions executed by the program
 - Determined by program, compiler, instruction set architecture (ISA)
- **Cycles per instruction: "CPI"** (typical range: 2 to 0.5)
 - On average, how many *cycles* does an instruction take to execute?
 - Determined by program, compiler, ISA, micro-architecture
- **Sec. per cycle: "clock period"** (typical range: 2ns to 0.25ns)
 - Reciprocal is **frequency**: 0.5 Ghz to 4 Ghz (1 Hertz = 1 cycle per sec)
 - Determined by micro-architecture, technology parameters
- For minimum execution time, minimize each term
 - Difficult: *often pull against one another*

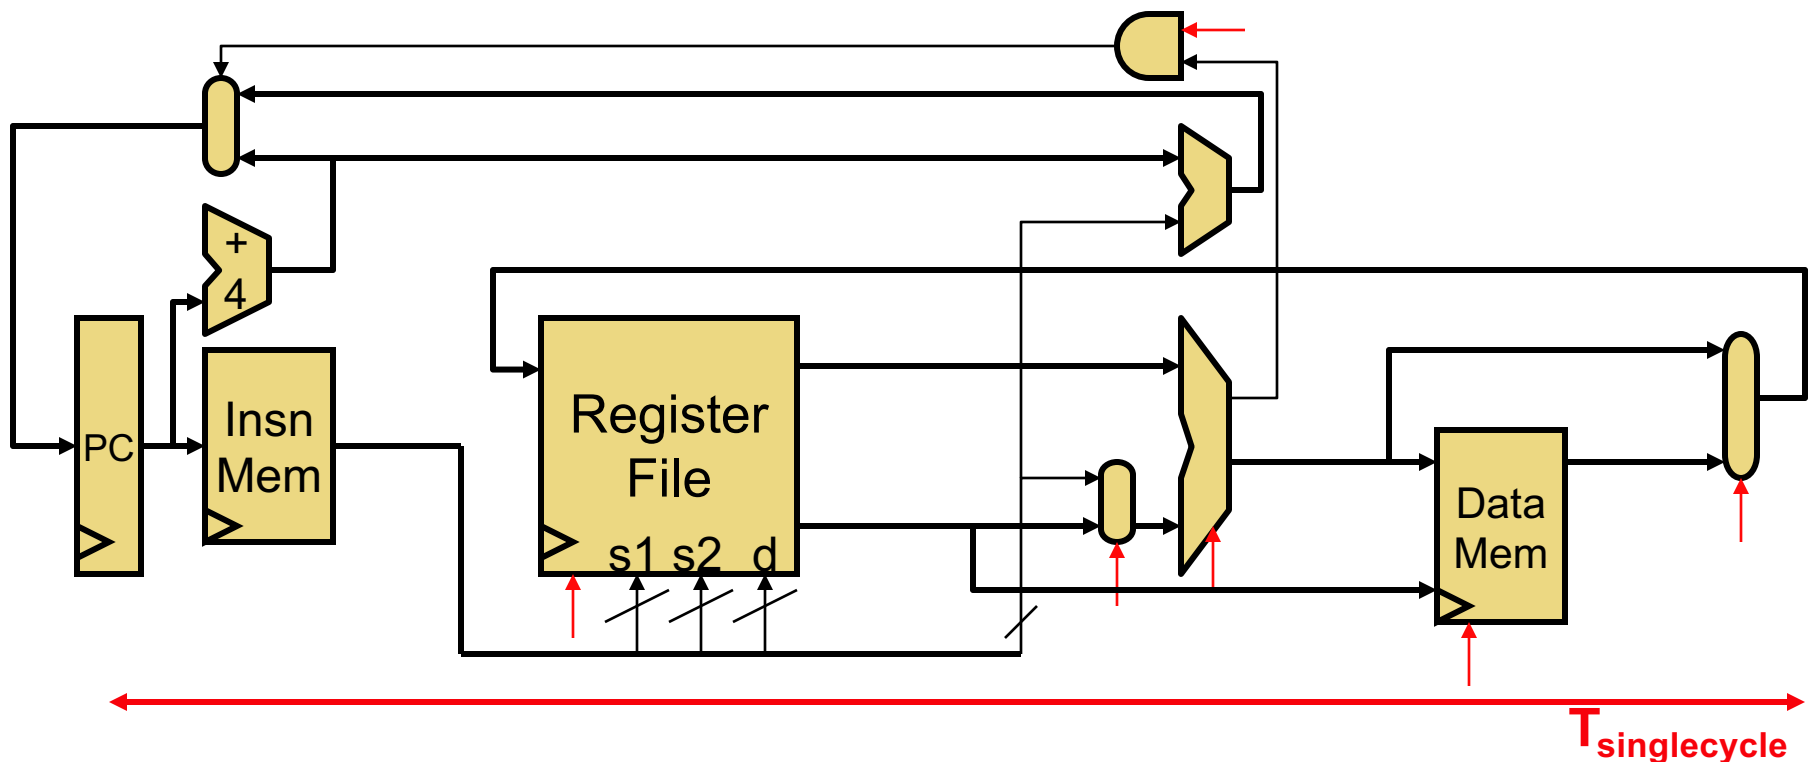
Cycles per Instruction (CPI)

- **CPI**: Cycle/instruction **on average**
 - **IPC** = $1/\text{CPI}$
 - Used more frequently than CPI
 - Favored because “bigger is better”, but harder to compute with
 - Different instructions have different cycle costs
 - E.g., “add” typically takes 1 cycle, “divide” takes >10 cycles
 - Depends on relative instruction frequencies
- CPI example
 - A program executes equal: integer, floating point (FP), memory ops
 - Cycles per instruction type: integer = 1, memory = 2, FP = 3
 - What is the CPI? $(33\% * 1) + (33\% * 2) + (33\% * 3) = 2$
 - **Caveat**: this sort of calculation ignores many effects
 - Back-of-the-envelope arguments only

Improving Clock Frequency

- **Faster transistors**
- Micro-architectural techniques
 - **Multi-cycle processors**
 - Break each instruction into small bits
 - Less logic delay -> improved clock frequency
 - Different instructions take different number of cycles
 - $CPI > 1$
 - **Pipelined processors**
 - As above, but overlap parts of instruction (parallelism!)
 - Faster clock, but CPI can still be around 1

Single-Cycle Datapath



- **Single-cycle datapath:** true “atomic” fetch/execute loop
 - Fetch, decode, execute one complete instruction every cycle
 - + Takes 1 cycle to execution any instruction by definition (“CPI” is 1)
 - Long clock period: to accommodate slowest instruction (worst-case delay through circuit, must wait this long *every* time)

Latency versus Throughput

- **Latency (execution time)**: time to finish a fixed task
- **Throughput (bandwidth)**: number of tasks in fixed time
 - Different: exploit parallelism for throughput, not latency (e.g., bread)
 - Often contradictory (latency **vs.** throughput)
 - Will see many examples of this
 - Choose definition of performance that matches your goals
 - Scientific program? Latency, web server: throughput?
- Example: move people 10 miles
 - Car: capacity = 5, speed = 60 miles/hour
 - Bus: capacity = 60, speed = 20 miles/hour
 - Latency: **car = 10 min**, bus = 30 min
 - Throughput: car = 15 PPH (count return trip), **bus = 60 PPH**
- Fastest way to send 10TB of data? (at 1+ gbits/second)

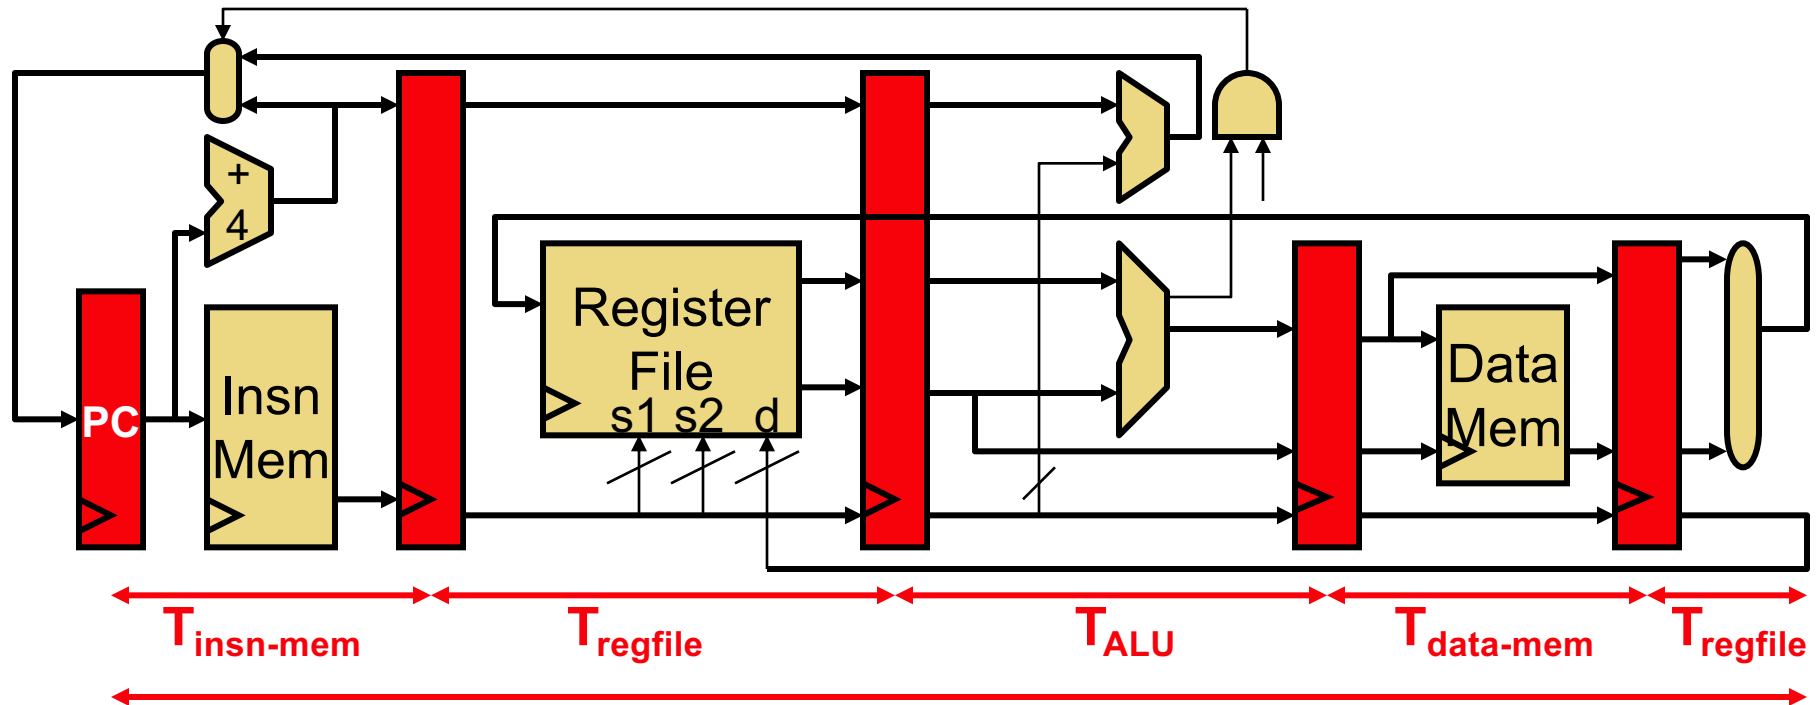
Pipelined Datapath

Pipelining



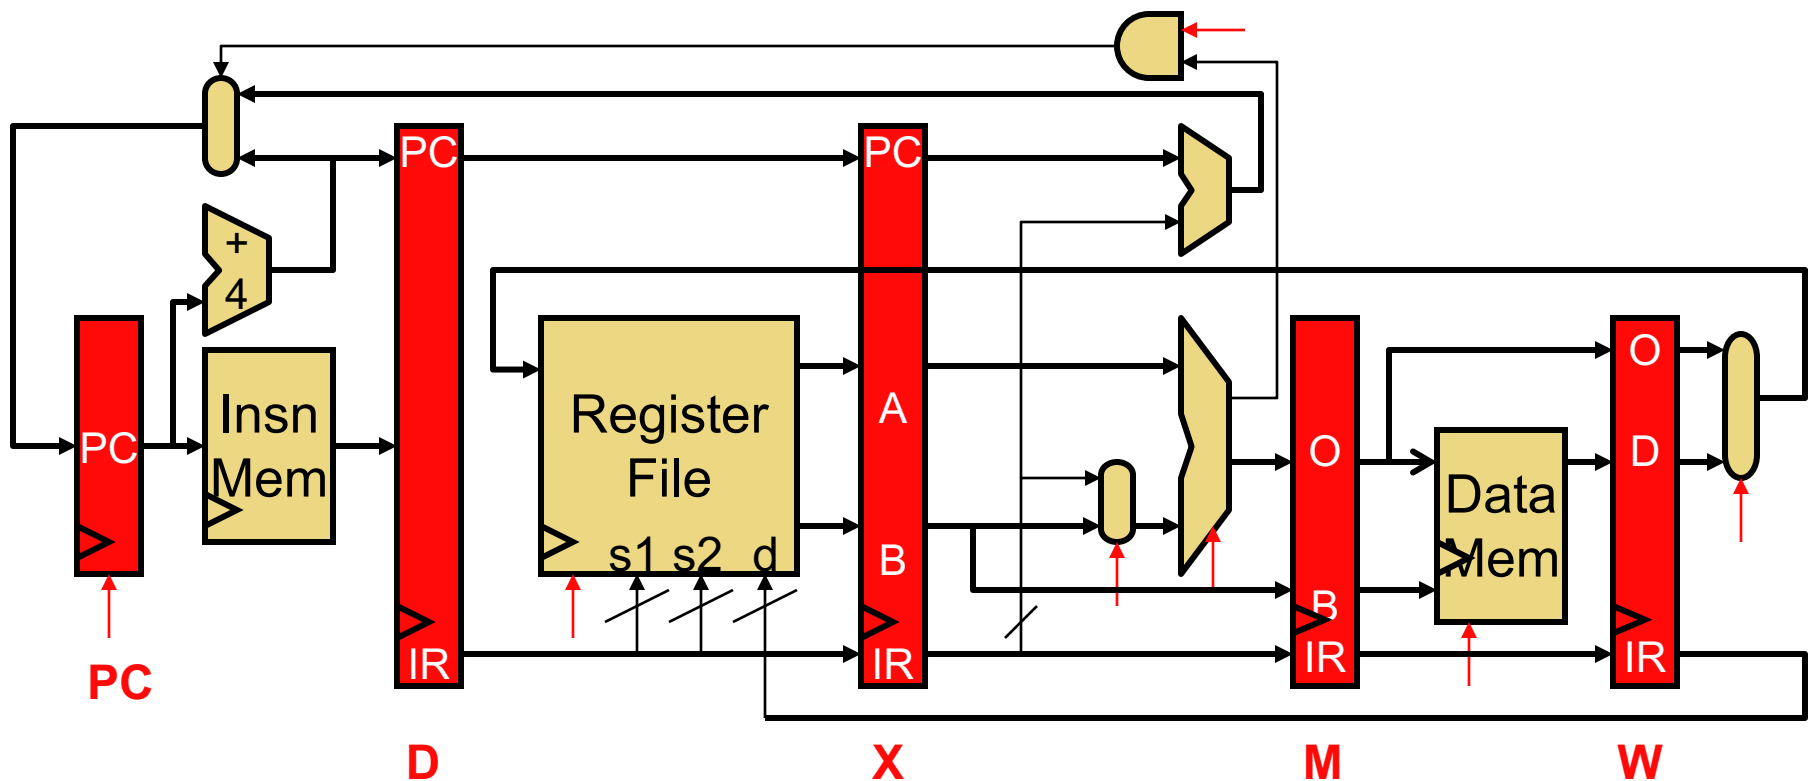
- Important performance technique
 - **Improves instruction throughput, not instruction latency**
- Break instruction execution into stages
 - When insn advances from stage 1 to 2, next insn enters at stage 1
 - Form of parallelism: “insn-stage parallelism”
 - Maintains illusion of sequential fetch/execute loop
 - Individual instruction takes the same number of stages
 - + **But instructions enter and leave at a much faster rate**
- Just like our laundromat

5 Stage Pipeline: Inter-Insn Parallelism



- **Pipelining:** cut datapath into N stages (here 5) $T_{\text{singlecycle}}$
 - One insn in each stage in each cycle
 - + Clock period = $\text{MAX}(T_{\text{insn-mem}}, T_{\text{regfile}}, T_{\text{ALU}}, T_{\text{data-mem}}, T_{\text{writeback}})$
 - + Base CPI = 1: insn enters and leaves every cycle
 - Actual CPI > 1: pipeline must often “stall”
 - Individual insn latency **increases** (pipeline overhead)

5 Stage Pipelined Datapath



- Five stage: **F**etch, **D**ecode, e**X**ecute, **M**emory, **W**riteback
 - Nothing magical about 5 stages (Pentium 4 had 22 stages!)
- Pipeline registers named by stages they begin
 - **PC, D, X, M, W**

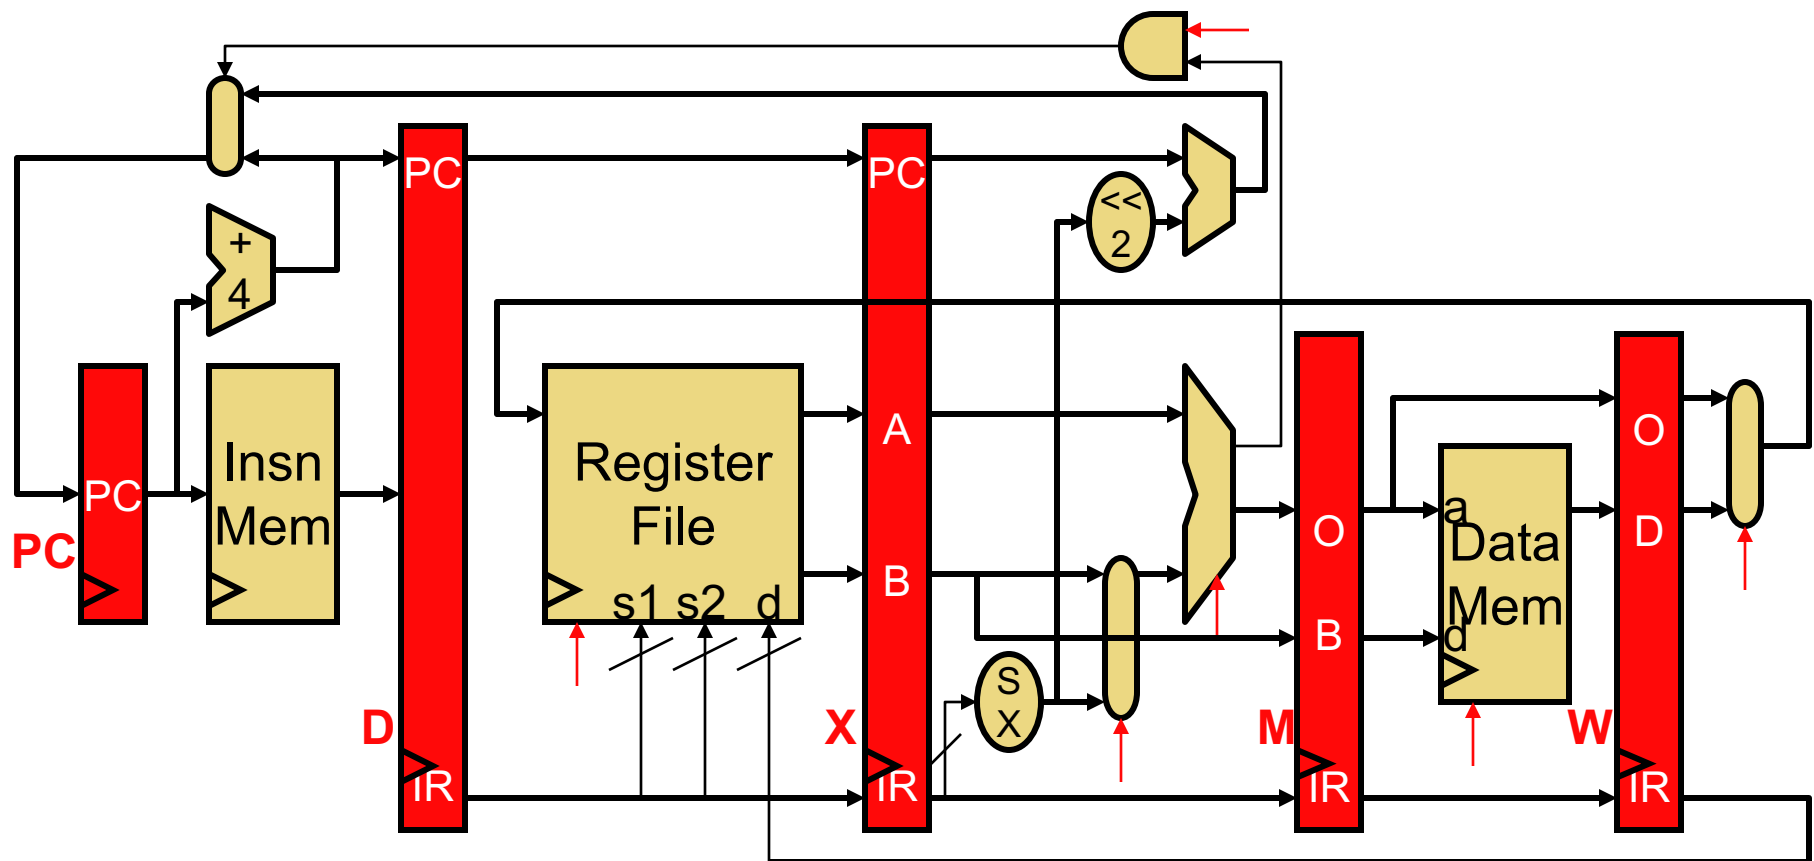
Pipeline Terminology

- **Scalar pipeline**: one insn per stage per cycle
 - Alternative: “superscalar” (later)
- **In-order pipeline**: insns enter execute stage in order
 - Alternative: “out-of-order” (later)
- **Pipeline depth**: number of pipeline stages
 - Nothing magical about five
 - Contemporary high-performance cores have ~15 stage pipelines

Instruction Convention

- Different ISAs use inconsistent register orders
- Some ISAs (for example MIPS)
 - Instruction destination (i.e., output) **on the left**
 - `add $1,$2,$3` means $\$1 \leftarrow \$2 + \$3$
- Other ISAs
 - Instruction destination (i.e., output) **on the right**
 - `add r1,r2,r3` means $r1 + r2 \rightarrow r3$
 - `ld 8(r5),r4` means $\text{mem}[r5+8] \rightarrow r4$
 - `st r4,8(r5)` means $r4 \rightarrow \text{mem}[r5+8]$
- Will try to specify to avoid confusion, next slides MIPS style

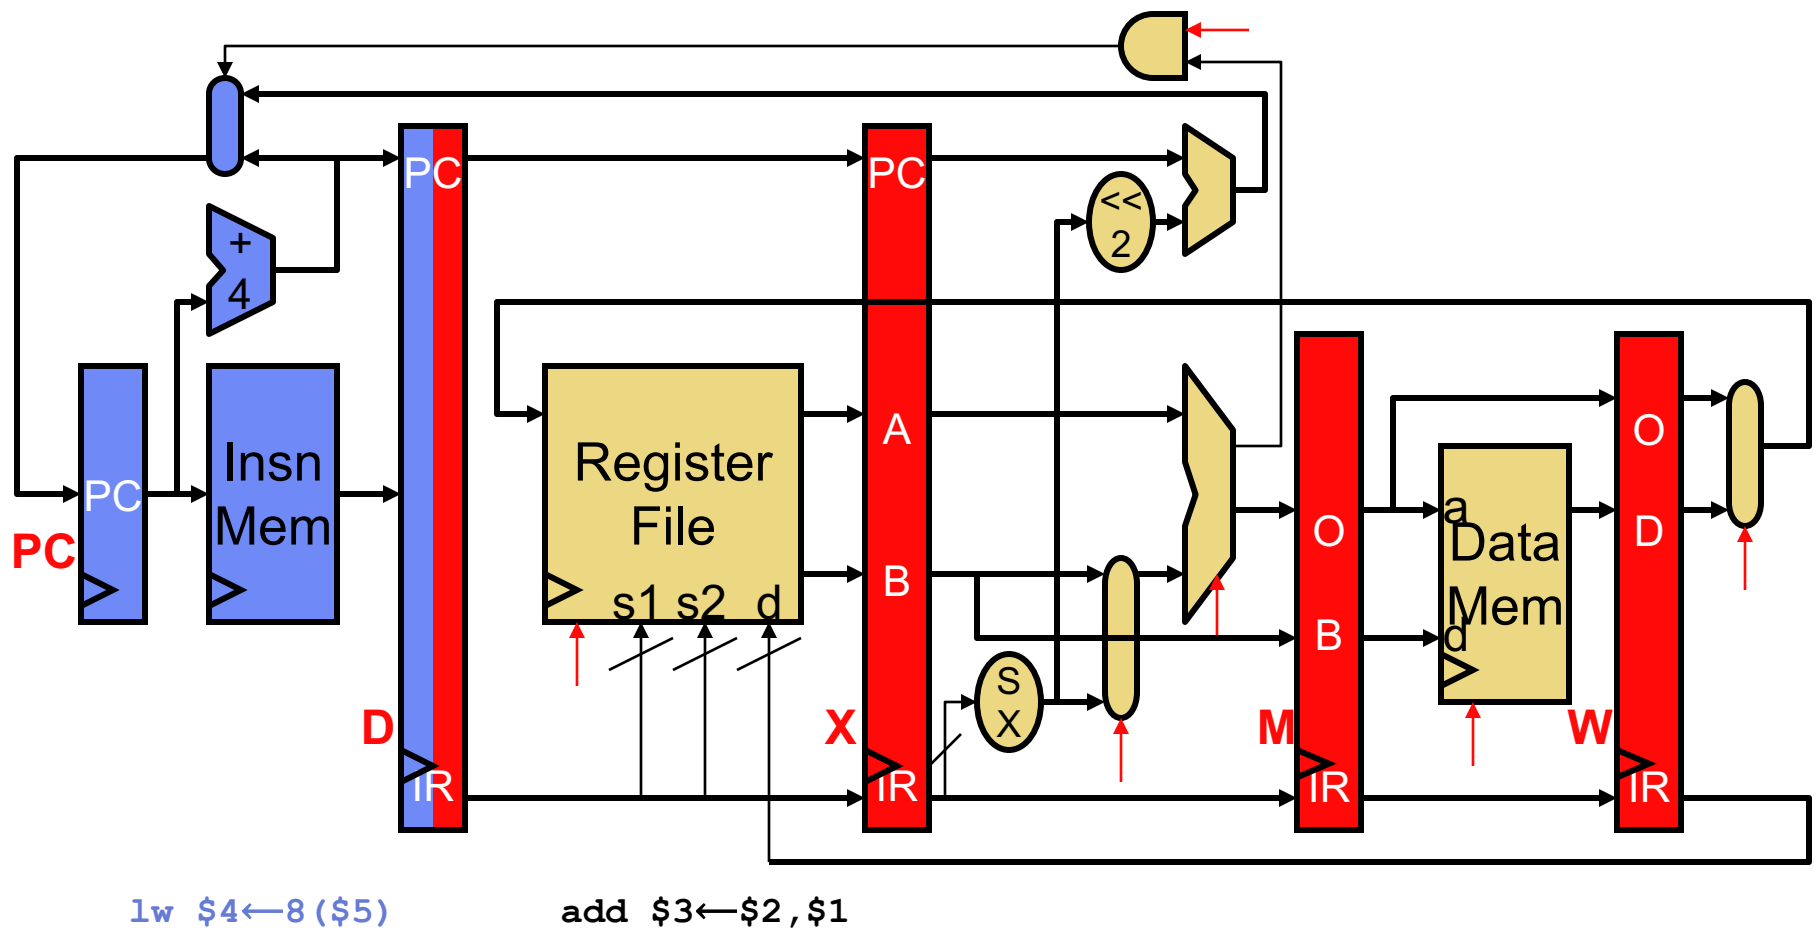
Pipeline Example: Cycle 1



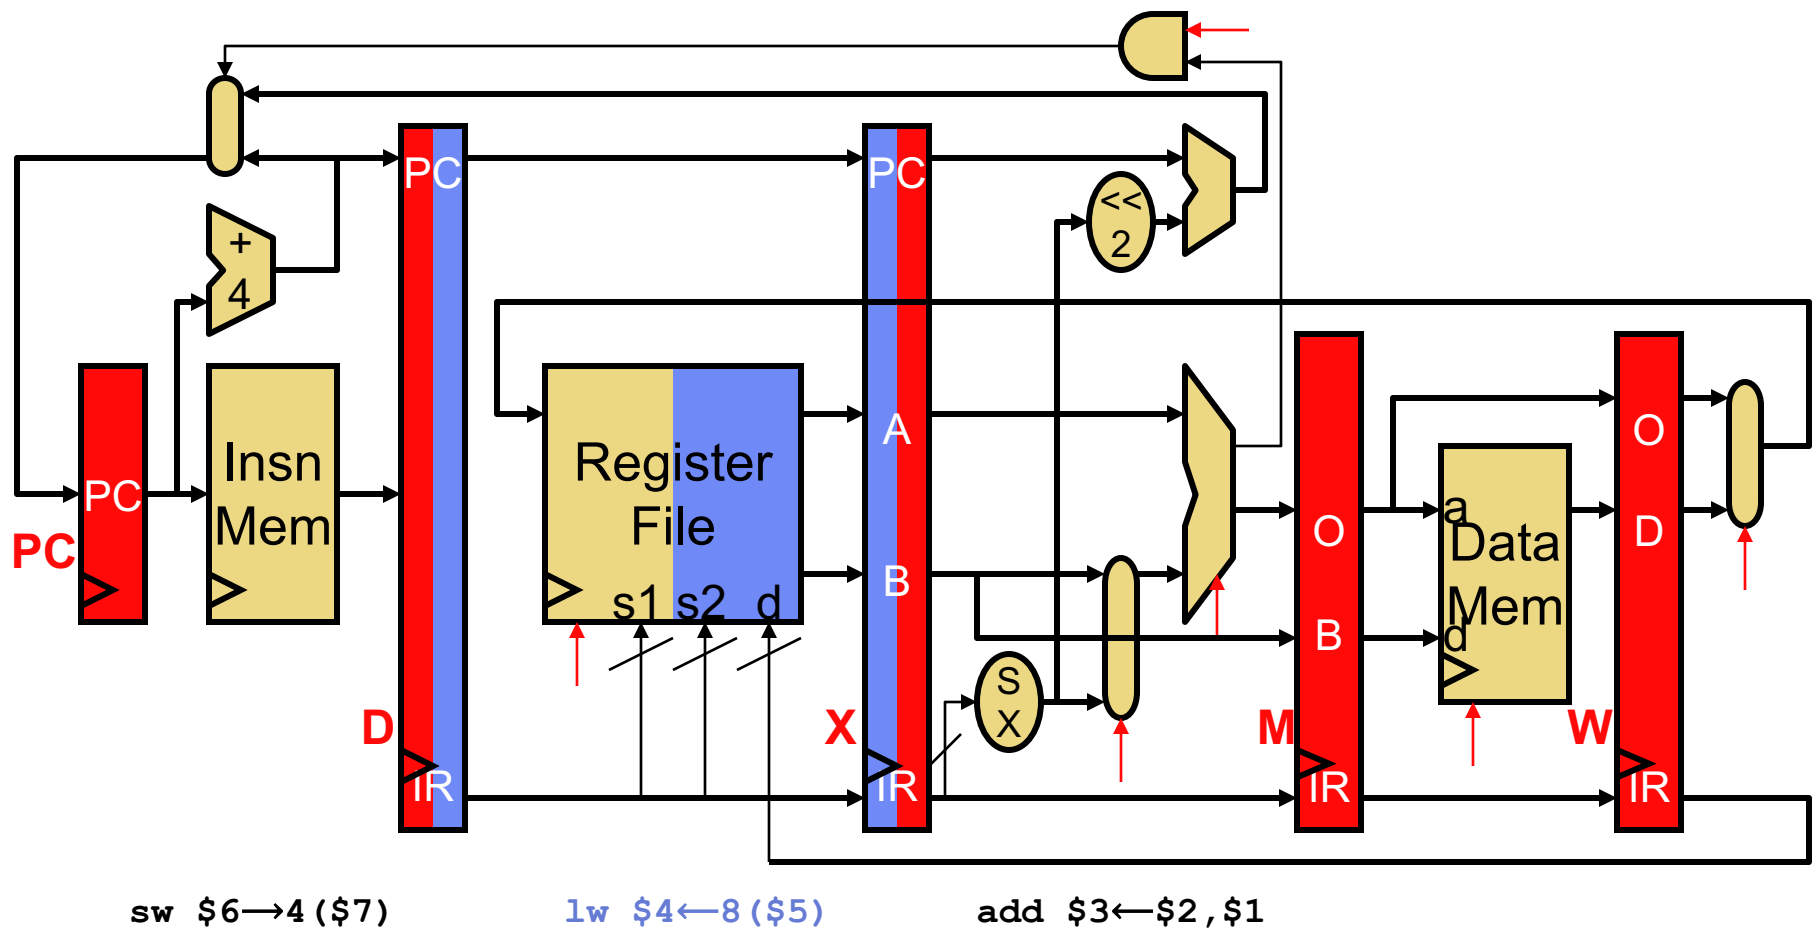
add \$3←\$2,\$1

- 3 instructions

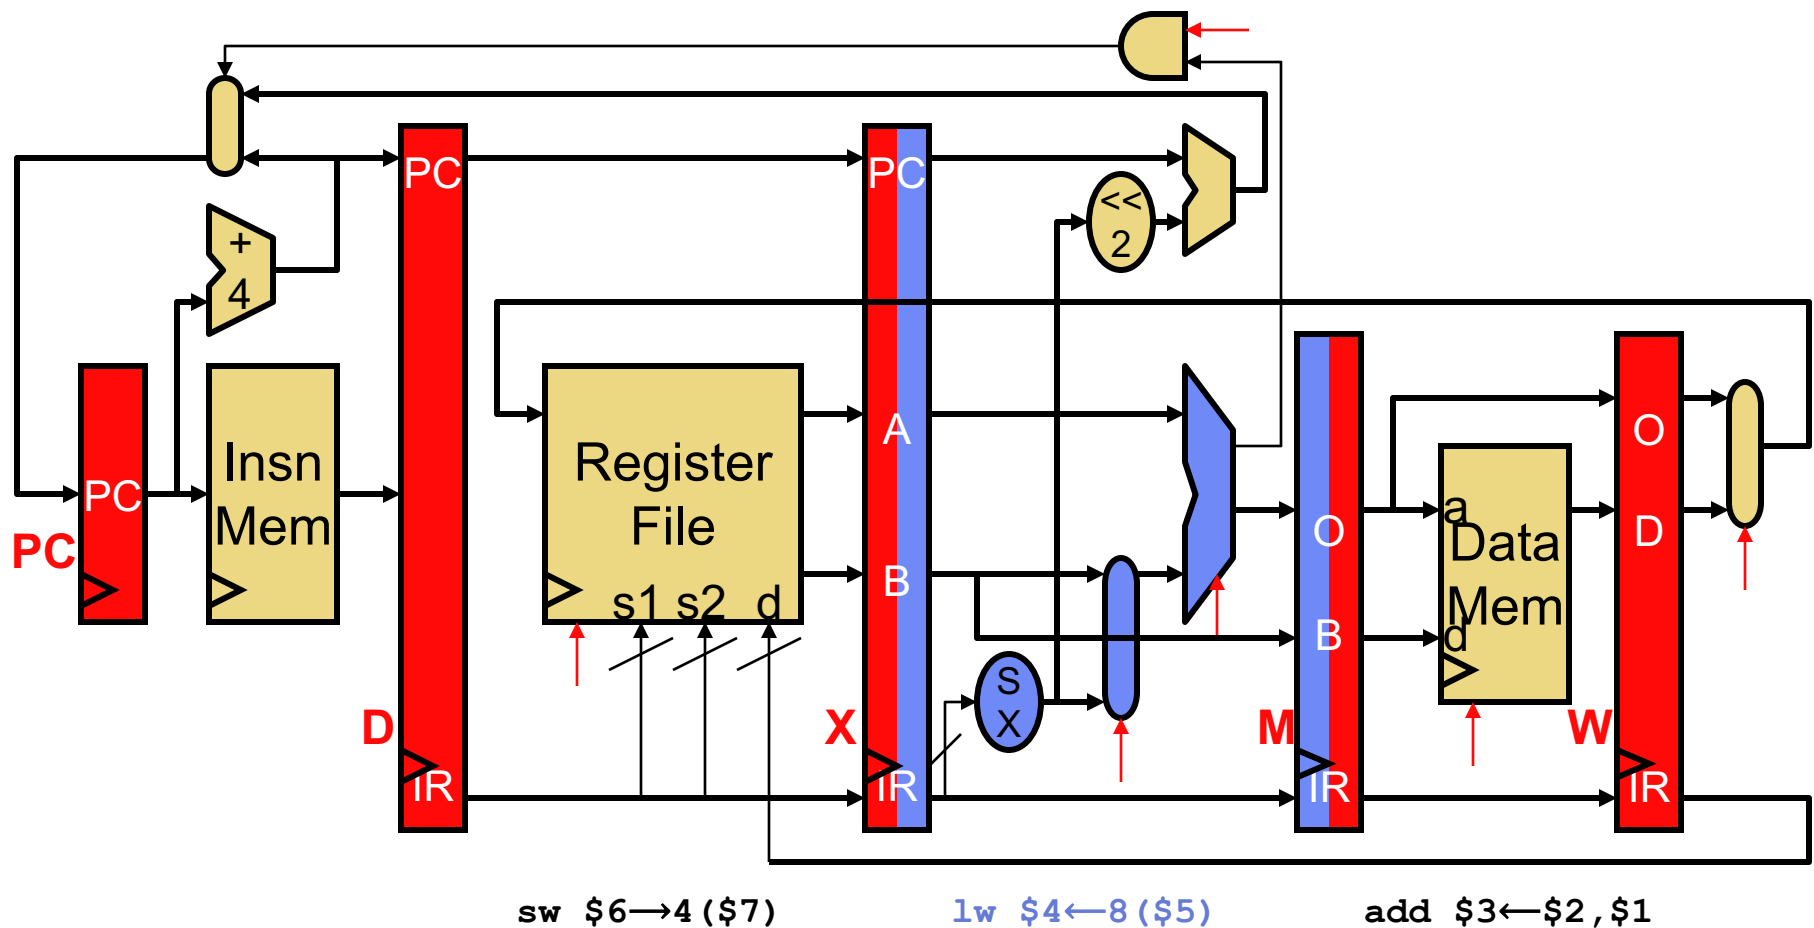
Pipeline Example: Cycle 2



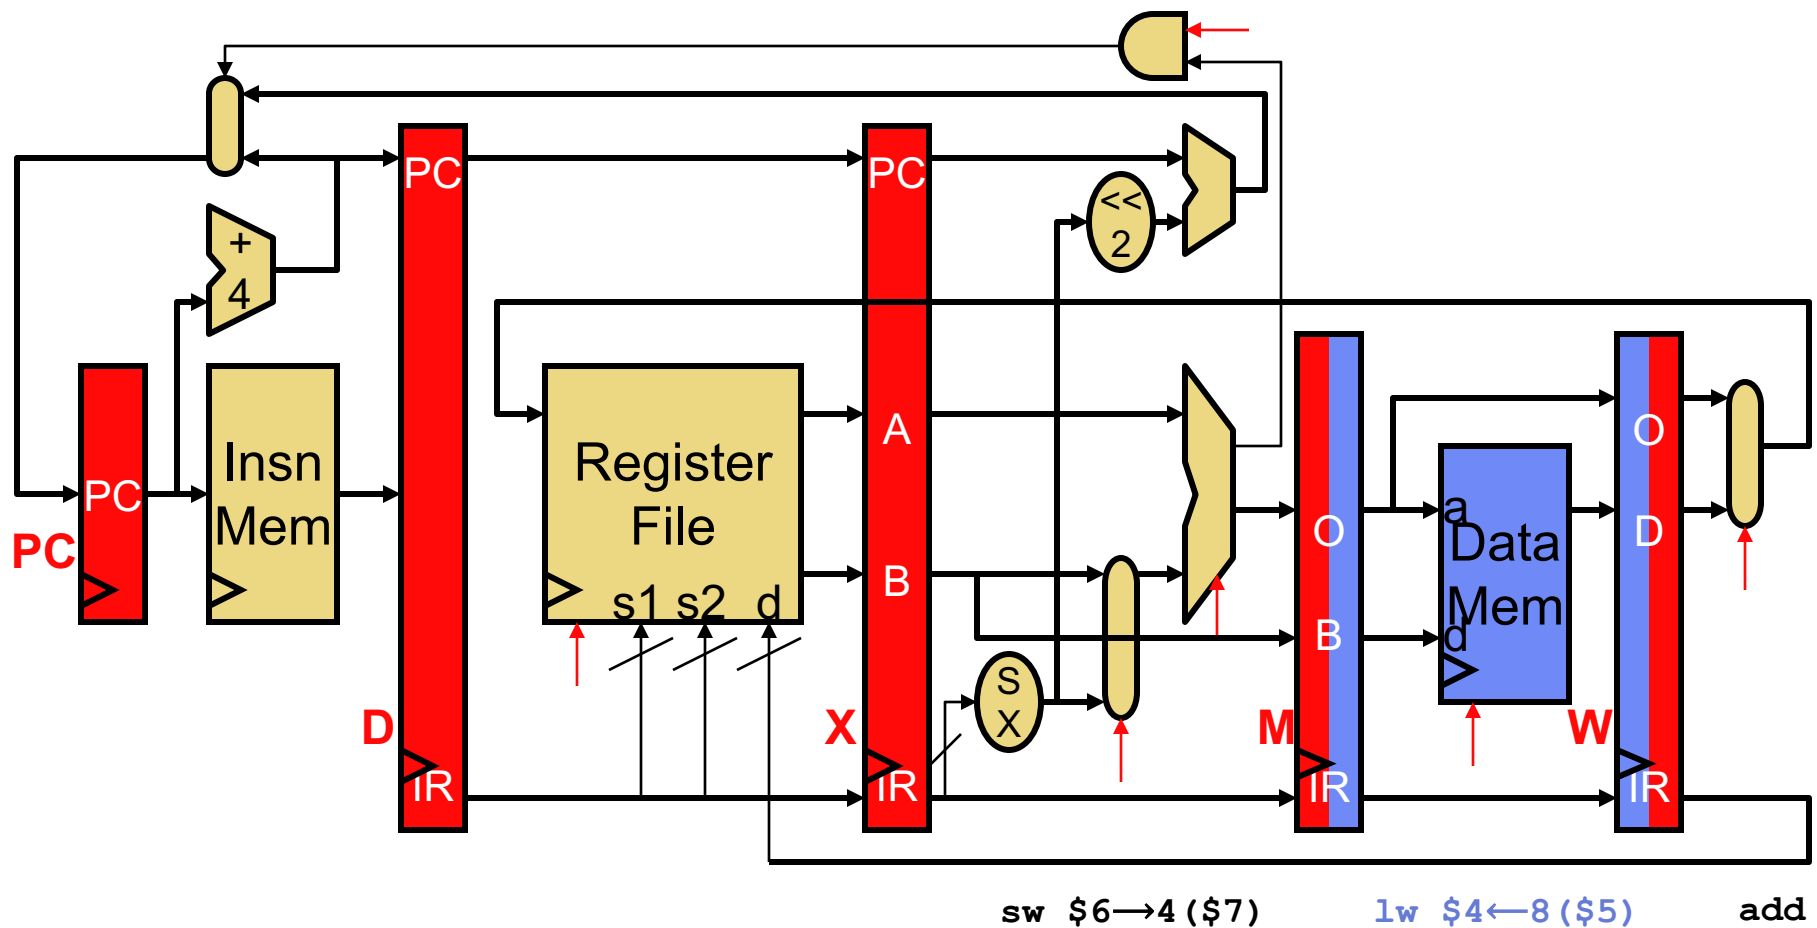
Pipeline Example: Cycle 3



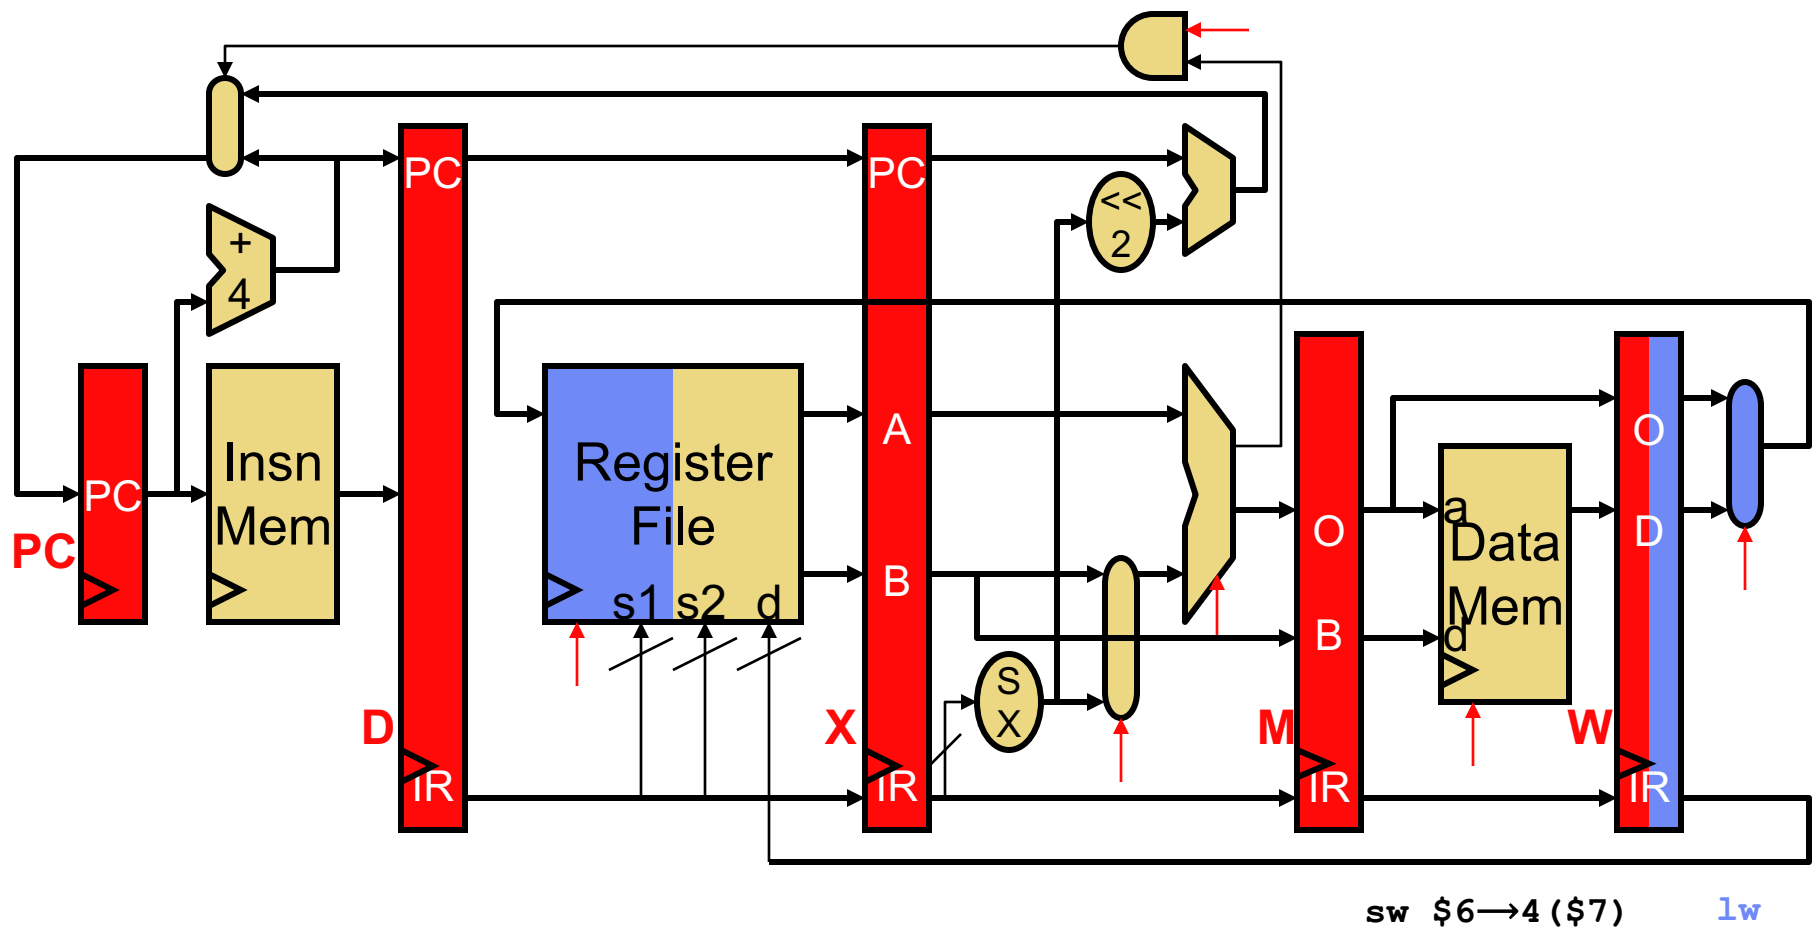
Pipeline Example: Cycle 4



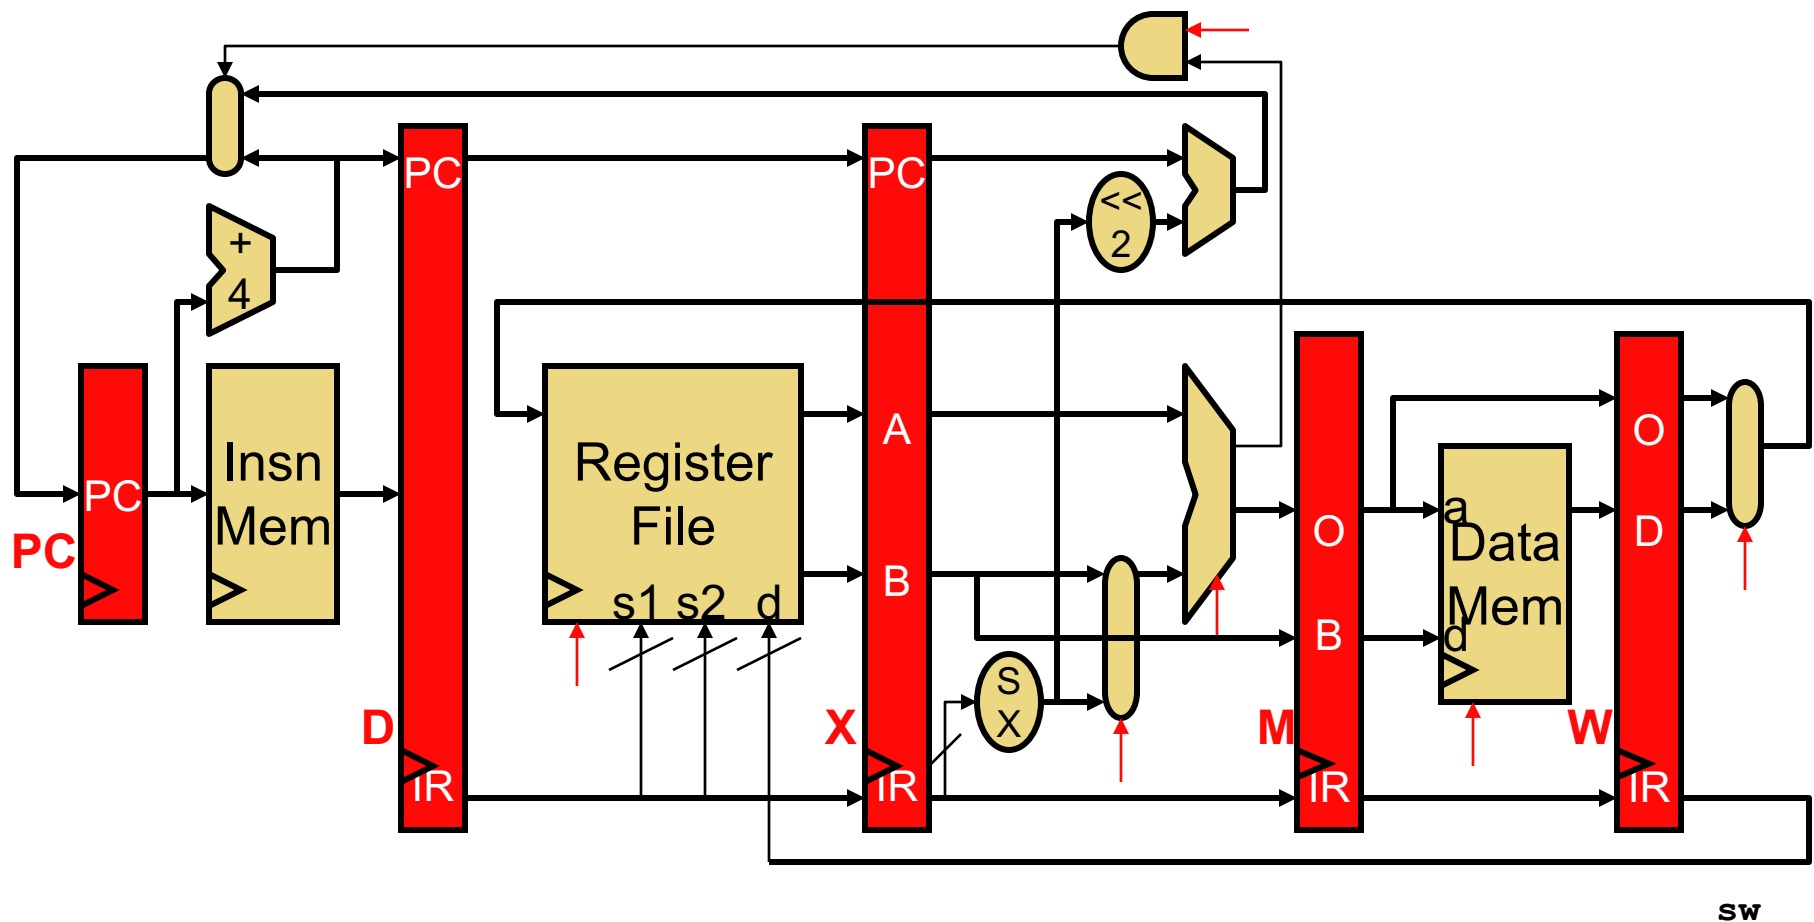
Pipeline Example: Cycle 5



Pipeline Example: Cycle 6



Pipeline Example: Cycle 7



Pipeline Diagram

- **Pipeline diagram**: shorthand for what we just saw
 - Across: cycles
 - Down: insns
 - Convention: **X** means `lw $4 ← 8 ($5)` finishes execute stage and writes into M register at **end** of cycle 4
 - assuming no *stalls* (discussed later)

	1	2	3	4	5	6	7	8	9
<code>add \$3, \$2, \$1</code>	F	D	X	M	W				
<code>lw \$4, 8 (\$5)</code>		F	D	X	M	W			
<code>sw \$6, 4 (\$7)</code>			F	D	X	M	W		

Example Pipeline Perf. Calculation

- Single-cycle
 - Clock period = 50ns, CPI = 1
 - Performance = 50ns/insn
- 5-stage pipelined
 - Clock period = **12ns** approx. (50ns / 5 stages) + overheads
 - + CPI = **1** (each insn takes 5 cycles, but 1 completes each cycle)
 - + Performance = **12ns/insn**
 - Well actually ... CPI = 1 + some penalty for pipelining (next)
 - CPI = **1.5** (on average insn completes every 1.5 cycles)
 - Performance = **18ns/insn**
 - Much higher performance than single-cycle

Q1: Why Is Pipeline Clock Period ...

- ... > (delay thru datapath) / (number of pipeline stages)?
 - Three reasons:
 - Registers add delay
 - Pipeline stages have different delays, clock period is **max** delay
 - Extra datapaths for pipelining (bypassing paths)
 - These factors have implications for ideal number pipeline stages
 - Diminishing clock frequency gains for longer (deeper) pipelines

Q2: Why Is Pipeline CPI...


- ... > 1?
 - CPI for scalar in-order pipeline is 1 + **stall penalties**
 - Stalls used to resolve hazards
 - **Hazard**: condition that jeopardizes sequential illusion
 - **Stall**: pipeline delay introduced to restore sequential illusion
- Calculating pipeline CPI
 - **Frequency of stall * stall cycles**
 - Penalties add (stalls typically don't overlap in in-order pipelines)
 - $1 + (\text{stall-freq}_1 * \text{stall-cyc}_1) + (\text{stall-freq}_2 * \text{stall-cyc}_2) + \dots$
- Correctness/performance/make common case fast
 - Long penalties OK if they are rare, e.g., $1 + (0.01 * 10) = 1.1$
 - Stalls also have implications for ideal number of pipeline stages

Data Dependences, Pipeline Hazards, and Bypassing

Dependences and Hazards

- **Dependence**: relationship between two insns
 - **Data**: two insns use same storage location
 - **Control**: one insn affects whether another executes at all
 - Not a bad thing, programs would be boring without them
 - Enforced by making older insn go before younger one
 - Happens naturally in single-/multi-cycle designs
 - But not in a pipeline
- **Hazard**: dependence & possibility of wrong insn order
 - Effects of wrong insn order must not be externally visible
 - **Stall**: for order by keeping younger insn in same stage
 - Hazards are a bad thing: stalls reduce performance

Data Hazards



- Let's forget about branches for now
- The three insn sequence we saw earlier executed fine...
 - But it wasn't a real program
 - Real programs have **data dependences**
 - They pass values via registers and memory

Dependent Operations

- Independent operations

```
add $3←$2,$1  
add $6←$5,$4
```


- Would this program execute correctly on a pipeline?

```
add $3←$2,$1  
add $6←$5,$3
```

- What about this program?

```
add $3←$2,$1  
lw $4←8($3)  
addi $6←1,$3  
sw $3→8($7)
```

Data Hazards



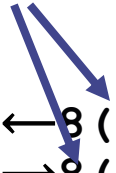
- Would this “program” execute correctly on this pipeline?
 - Which insns would execute with correct inputs?
 - `add` is writing its result into `$3` in current cycle
 - `lw` read `$3` two cycles ago → got wrong value
 - `addi` read `$3` one cycle ago → got wrong value
 - `sw` is reading `$3` this cycle → maybe (depending on regfile design)

Fixing Register Data Hazards

- Can only read register value three cycles after writing it
- **Option #1: make sure programs don't do it**
 - Compiler puts two independent insns between write/read insn pair
 - If they aren't there already
 - Independent means: "do not interfere with register in question"
 - Do not write it: otherwise meaning of program changes
 - Do not read it: otherwise create new data hazard
 - **Code scheduling**: compiler moves existing insns to do this
 - If none can be found, must use **nops** (no-operation)
- This is called **software interlocks**
 - **MIPS**: **M**icroprocessor w/out **I**nterlocking **P**ipeline **S**tages

Software Interlock Example

```
add  $3←$2, $1
nop
nop
lw   $4←8 ($3)
sw   $7→8 ($3)
sub  $6←$2, $8
addi $3←$5, 4
```



- Can any of last three insns be scheduled between first two
 - `sw $7→8 ($3)`? No, creates hazard with `add $3←$2, $1`
 - `sub $6←$2, $8`? Okay
 - `addi $3←$5, 4`? No, `lw` would read \$3 from it
 - Still need one more insn, use `nop`

```
add  $3←$2, $1
sub  $6←$2, $8
nop
lw   $4←8 ($3)
sw   $7→8 ($3)
addi $3←$5, 4
```

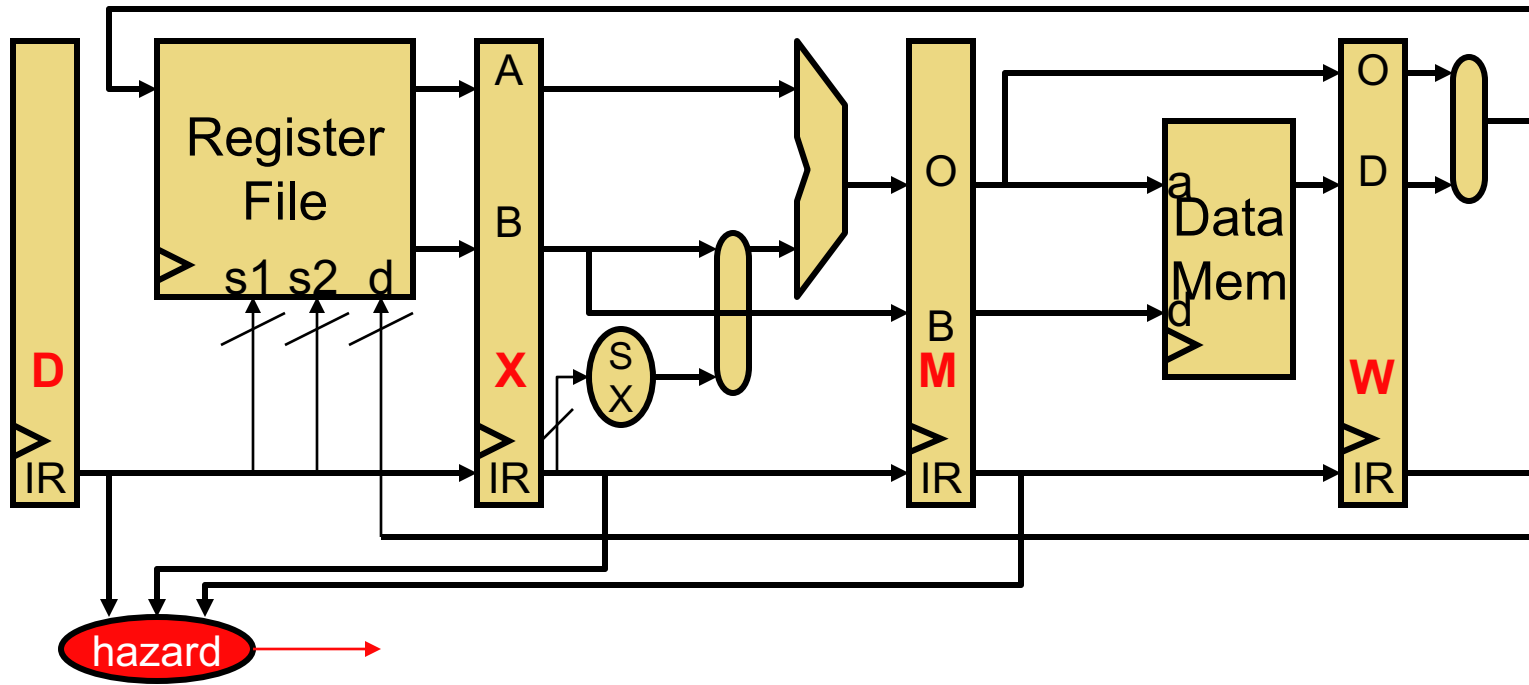
Software Interlock Performance

- Assume
 - Branch: 20%, load: 20%, store: 10%, other: 50%
- For software interlocks, let's assume:
 - 20% of insns require insertion of 1 `nop`
 - 5% of insns require insertion of 2 `nops`
- Result:
 - CPI is still 1 technically
 - But now there are more insns
 - $\#insns = 1 + 0.20*1 + 0.05*2 = \mathbf{1.3}$
 - **30% more insns (30% slowdown) due to data hazards**

Hardware Interlocks

- Problem with software interlocks? Not compatible
 - Where does **3** in “read register 3 cycles after writing” come from?
 - From structure (depth) of pipeline
 - What if next MIPS version uses a 7-stage pipeline?
 - Programs compiled assuming 5-stage pipeline won't work!
- **Option #2: hardware interlocks**
 - Processor detects data hazards and fixes them
 - Resolves the above compatibility concern
 - Two aspects to this
 - Detecting hazards
 - Fixing hazards

Detecting Data Hazards

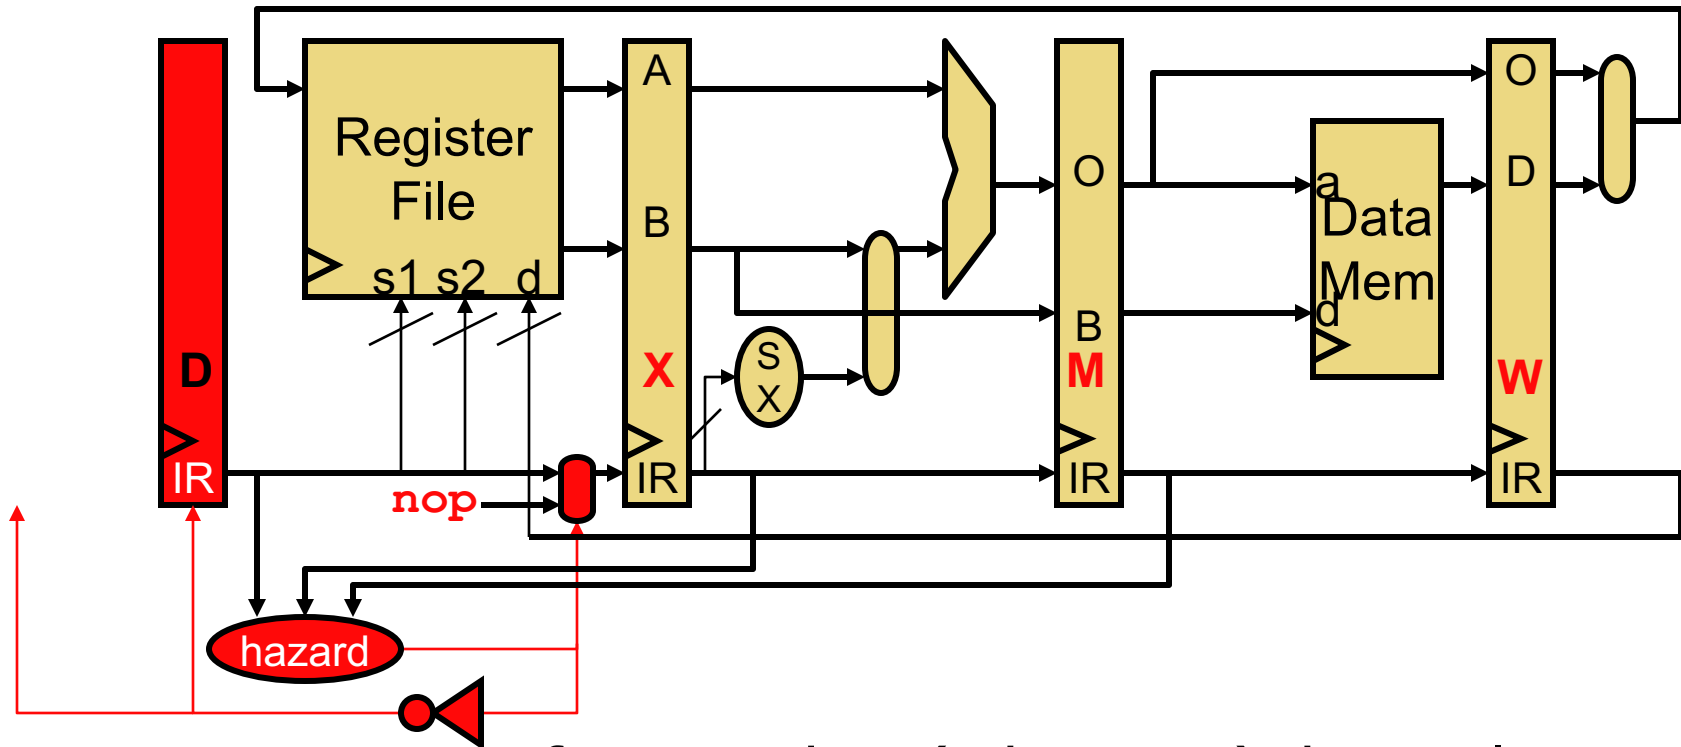


- Compare input register names of insn in D stage with output register names of older insns in pipeline

Stall =

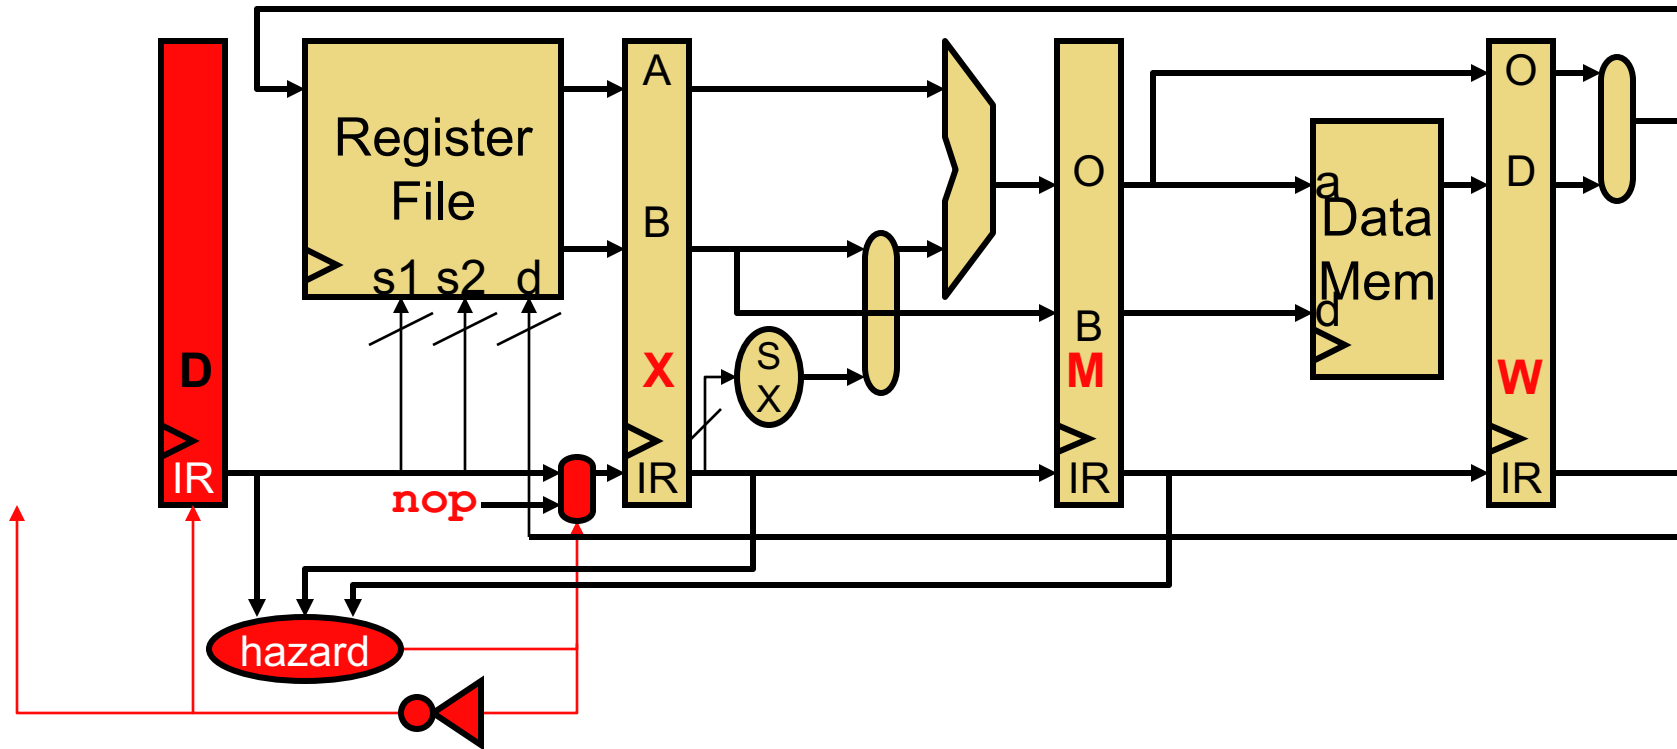
```
(D.IR.RegSrc1 == X.IR.RegDest) ||
(D.IR.RegSrc2 == X.IR.RegDest) ||
(D.IR.RegSrc1 == M.IR.RegDest) ||
(D.IR.RegSrc2 == M.IR.RegDest)
```

Fixing Data Hazards



- Prevent D insn from reading (advancing) this cycle
 - Write **nop** into X.IR (effectively, insert **nop** in hardware)
 - Also reset (clear) the datapath control signals
 - Disable D register and PC write enables (why?)
- Re-evaluate situation next cycle

Hardware Interlock Example: cycle 1



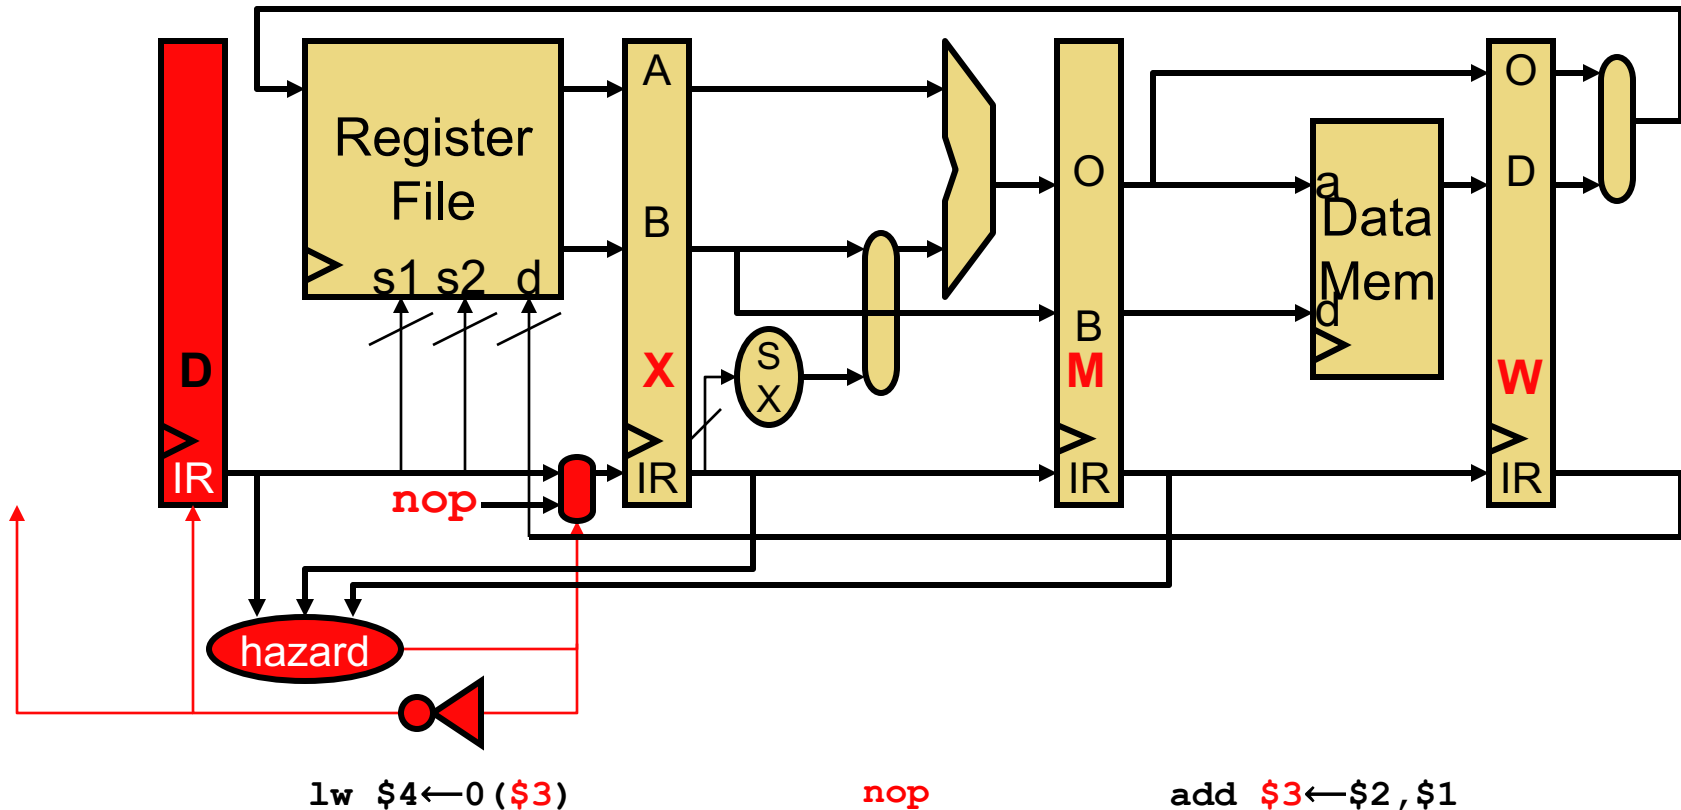
lw \$4 ← 0 (\$3)

add \$3 ← \$2, \$1

Stall =

$(D.IR.RegSrc1 == X.IR.RegDest) ||$
 $(D.IR.RegSrc2 == X.IR.RegDest) ||$
 $(D.IR.RegSrc1 == M.IR.RegDest) ||$
 $(D.IR.RegSrc2 == M.IR.RegDest) = 1$

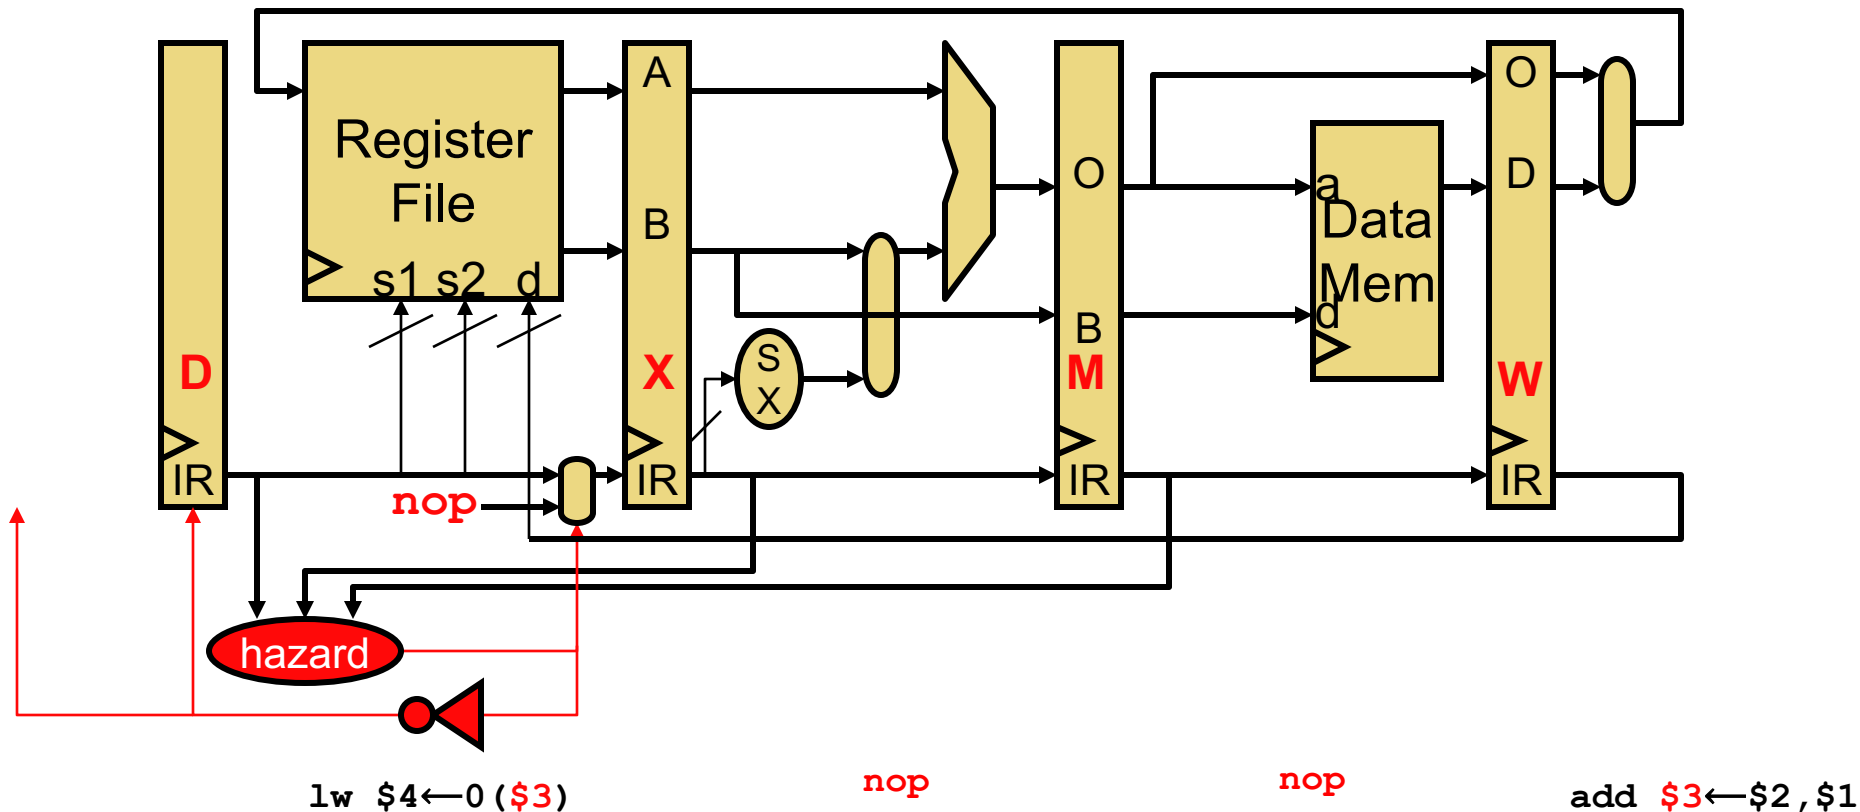
Hardware Interlock Example: cycle 2



Stall =

```
(D.IR.RegSrc1 == X.IR.RegDest) ||
(D.IR.RegSrc2 == X.IR.RegDest) ||
(D.IR.RegSrc1 == M.IR.RegDest) ||
(D.IR.RegSrc2 == M.IR.RegDest) = 1
```


Hardware Interlock Example: cycle 3



Stall =

```
(D.IR.RegSrc1 == X.IR.RegDest) ||
(D.IR.RegSrc2 == X.IR.RegDest) ||
(D.IR.RegSrc1 == M.IR.RegDest) ||
(D.IR.RegSrc2 == M.IR.RegDest) = 0
```


Pipeline Control Terminology

- Hardware interlock maneuver is called **stall** or **bubble**
- Mechanism is called **stall logic**
- Part of more general **pipeline control** mechanism
 - Controls the advance of insns through pipeline
- Distinguish from **pipelined datapath control**
 - Controls datapath **within** each stage
 - **Pipeline control** controls advancement of **datapath control**

Hardware Interlock Performance

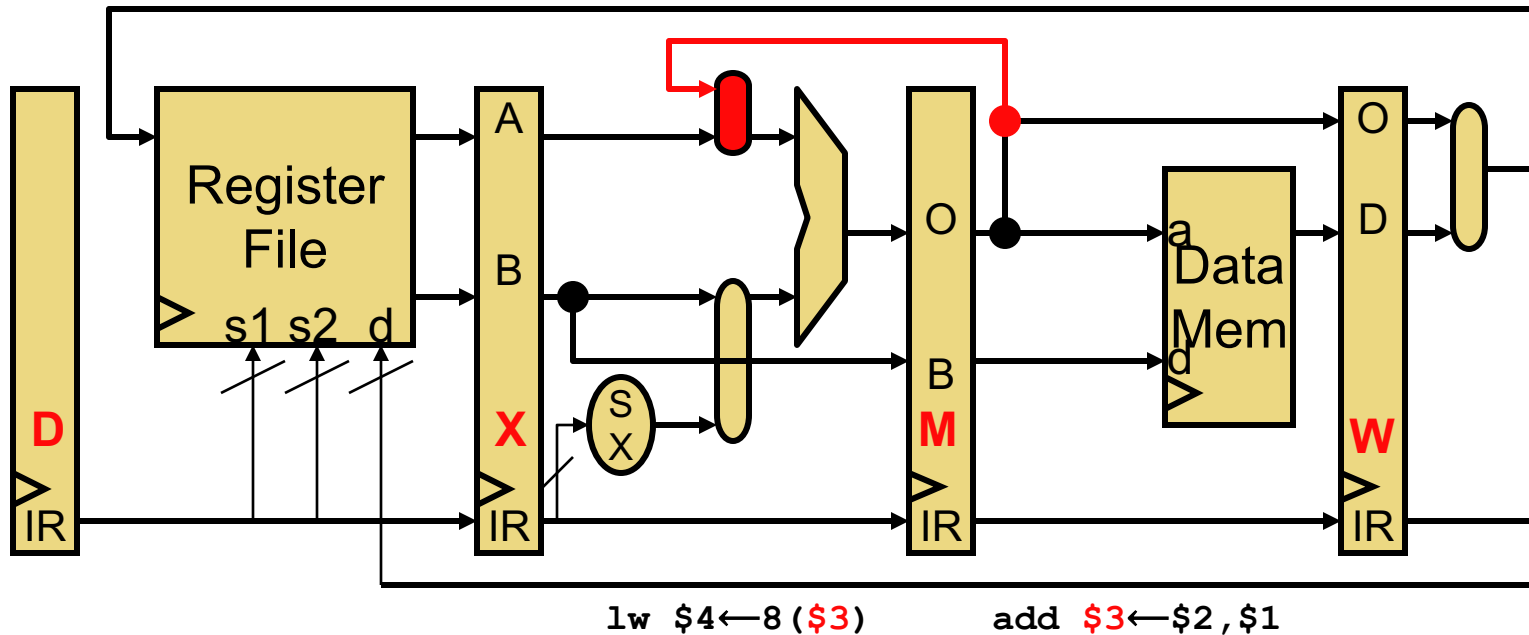
- As before:
 - Branch: 20%, load: 20%, store: 10%, other: 50%
- Hardware interlocks: same as software interlocks
 - 20% of insns require 1 cycle stall (I.e., insertion of 1 `nop`)
 - 5% of insns require 2 cycle stall (I.e., insertion of 2 `nops`)
 - $\text{CPI} = 1 + 0.20 \cdot 1 + 0.05 \cdot 2 = \mathbf{1.3}$
 - So, either CPI stays at 1 and #insns increases 30% (software)
 - Or, #insns stays at 1 (relative) and CPI increases 30% (hardware)
 - Same performance ☹ but better software compatibility ☺
- Anyway, we can do better

Observation!



- Technically, this situation is broken
 - `lw $4 ← 8($3)` has already read **\$3** from regfile
 - `add $3 ← $2, $1` hasn't yet written **\$3** to regfile
- But fundamentally, everything is OK
 - `lw $4 ← 8($3)` hasn't actually used **\$3** yet
 - `add $3 ← $2, $1` has already computed **\$3**

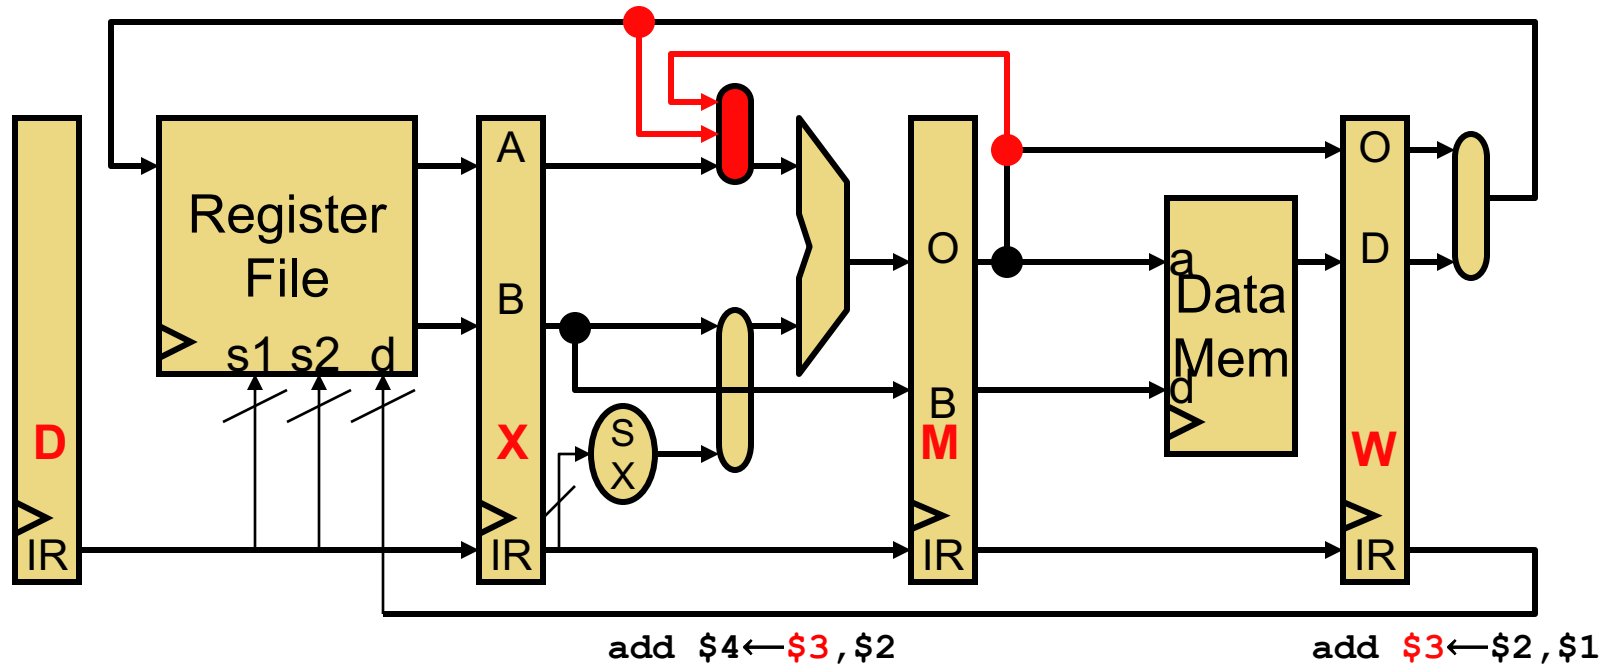
Bypassing



- **Bypassing**

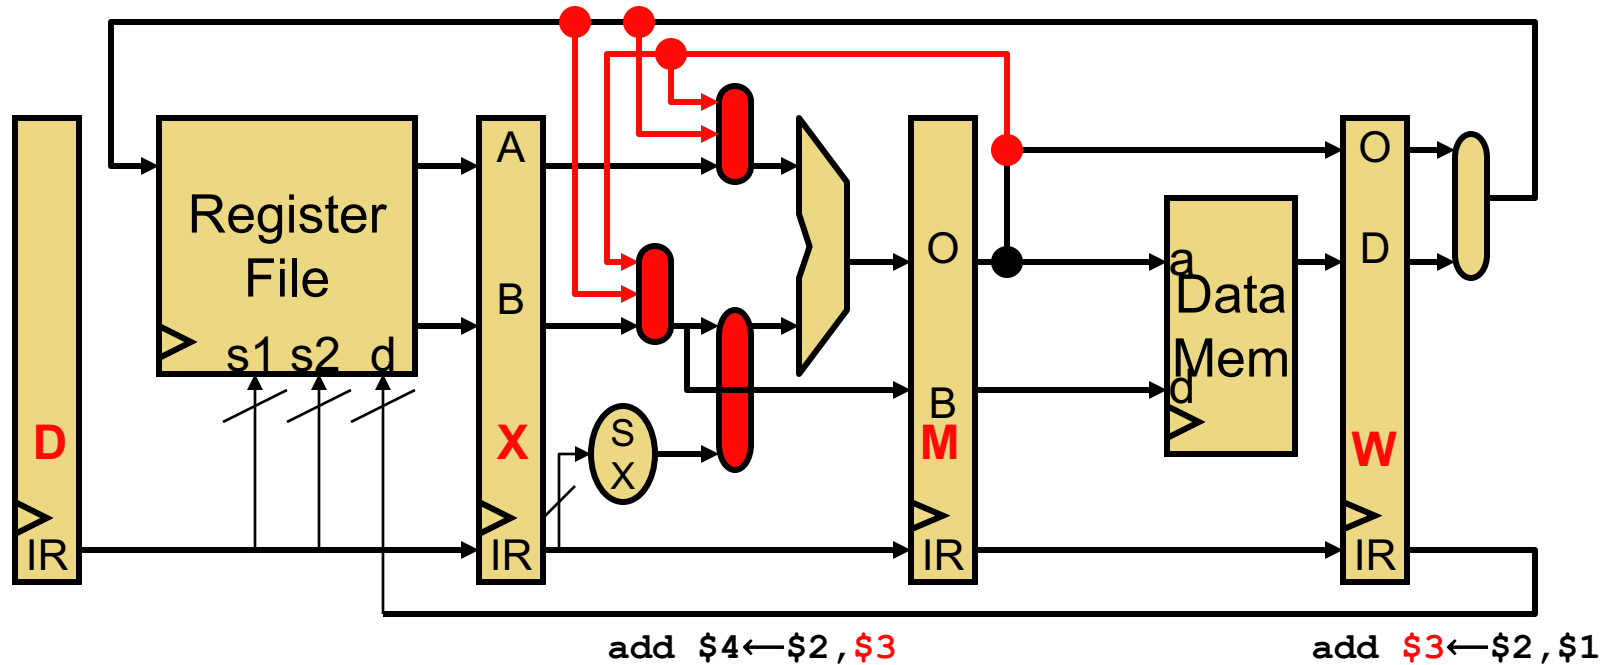
- Reading a value from an intermediate (μ architectural) source
- Not waiting until it is available from primary source
- Here, we are bypassing the register file
- Also called **forwarding**

WX Bypassing




- What about this combination?
 - Add another bypass path and MUX (multiplexor) input
 - First one was an **MX** bypass
 - This one is a **WX** bypass

ALUinB Bypassing



- Can also bypass to ALU input B

WM Bypassing?



- Does WM bypassing make sense?
 - Not to the address input (why not?)

`sw $4 -> 4 ($3)` `lw $3 <- 8 ($2)`

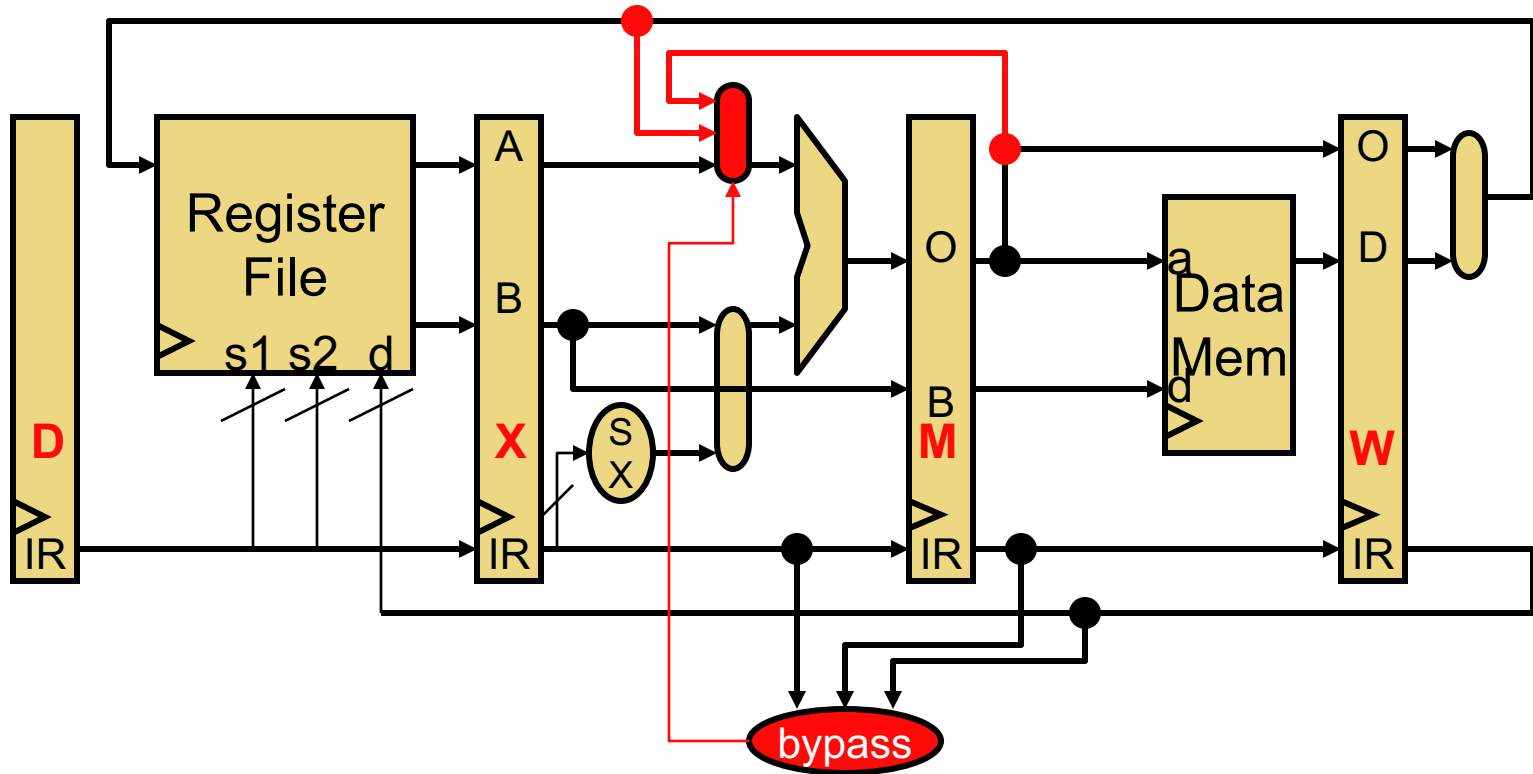


- But to the store data input, yes

`sw $3 -> 4 ($4)` `lw $3 <- 8 ($2)`



Bypass Logic



- Each multiplexor has its own, here it is for "ALUinA"
 $(X.IR.RegSrc1 == M.IR.RegDest) \Rightarrow 0$
 $(X.IR.RegSrc1 == W.IR.RegDest) \Rightarrow 1$
Else $\Rightarrow 2$

Pipeline Diagrams with Bypassing

- If bypass exists, "from"/"to" stages execute in same cycle

- Example: MX bypass

	1	2	3	4	5	6	7	8	9	10
add r1 ← r2, r3	F	D	X	M	W					
sub r2 ← r1 , r4		F	D	X	M	W				

- Example: WX bypass

	1	2	3	4	5	6	7	8	9	10
add r1 ← r2, r3	F	D	X	M	W					
ld r5 ← 4(r7)		F	D	X	M	W				
sub r2 ← r1 , r4			F	D	X	M	W			

- Example: WM bypass

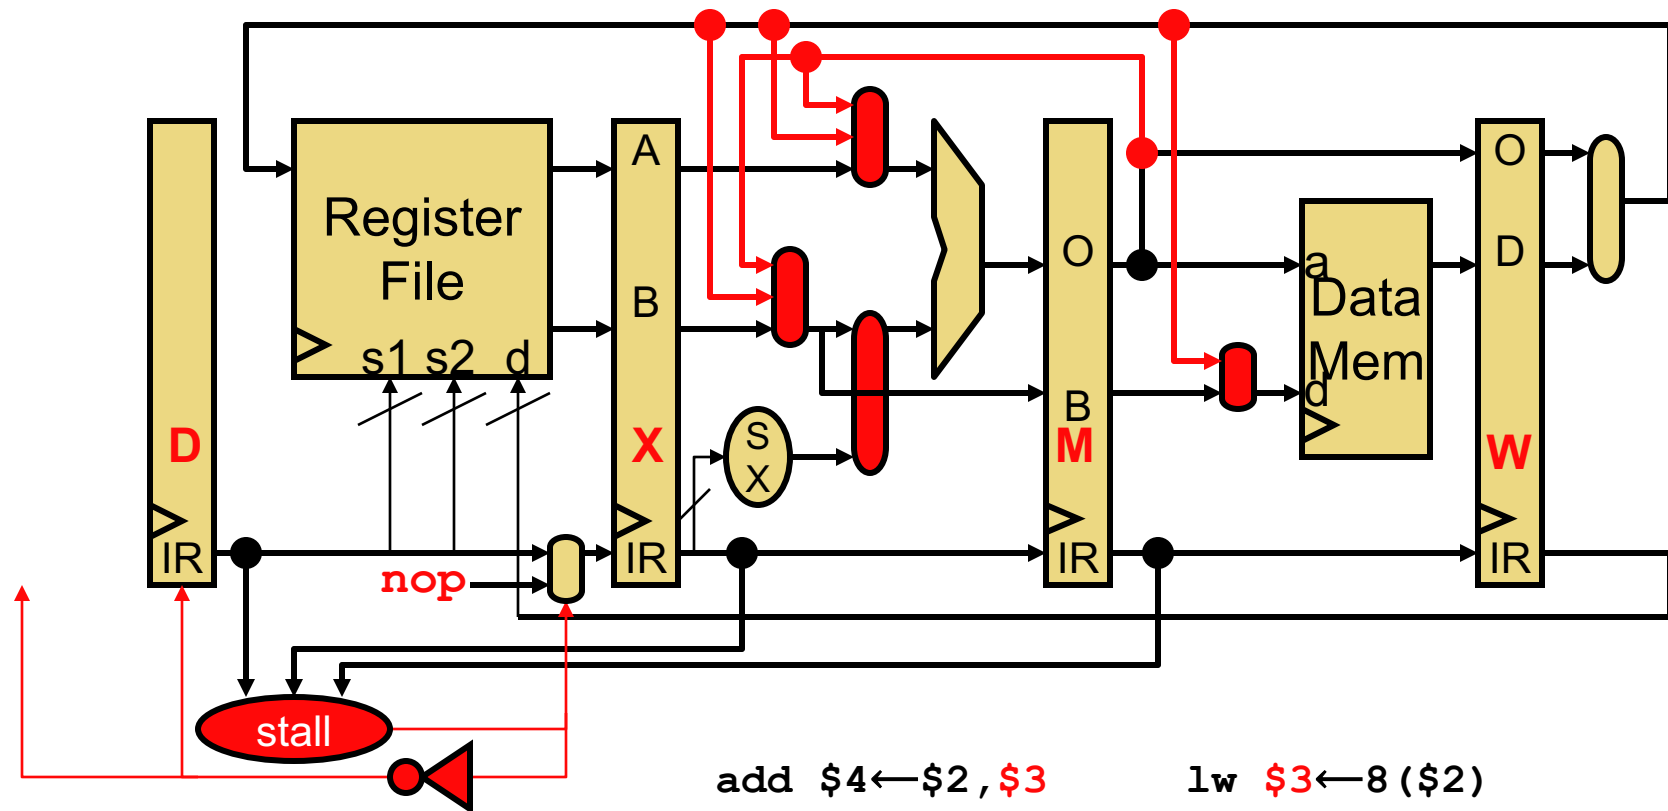
	1	2	3	4	5	6	7	8	9	10
add r1 ← r2, r3	F	D	X	M	W					
?		F	D	X	M	W				

- Can you think of a code example that uses the WM bypass?

Bypass and Stall Logic

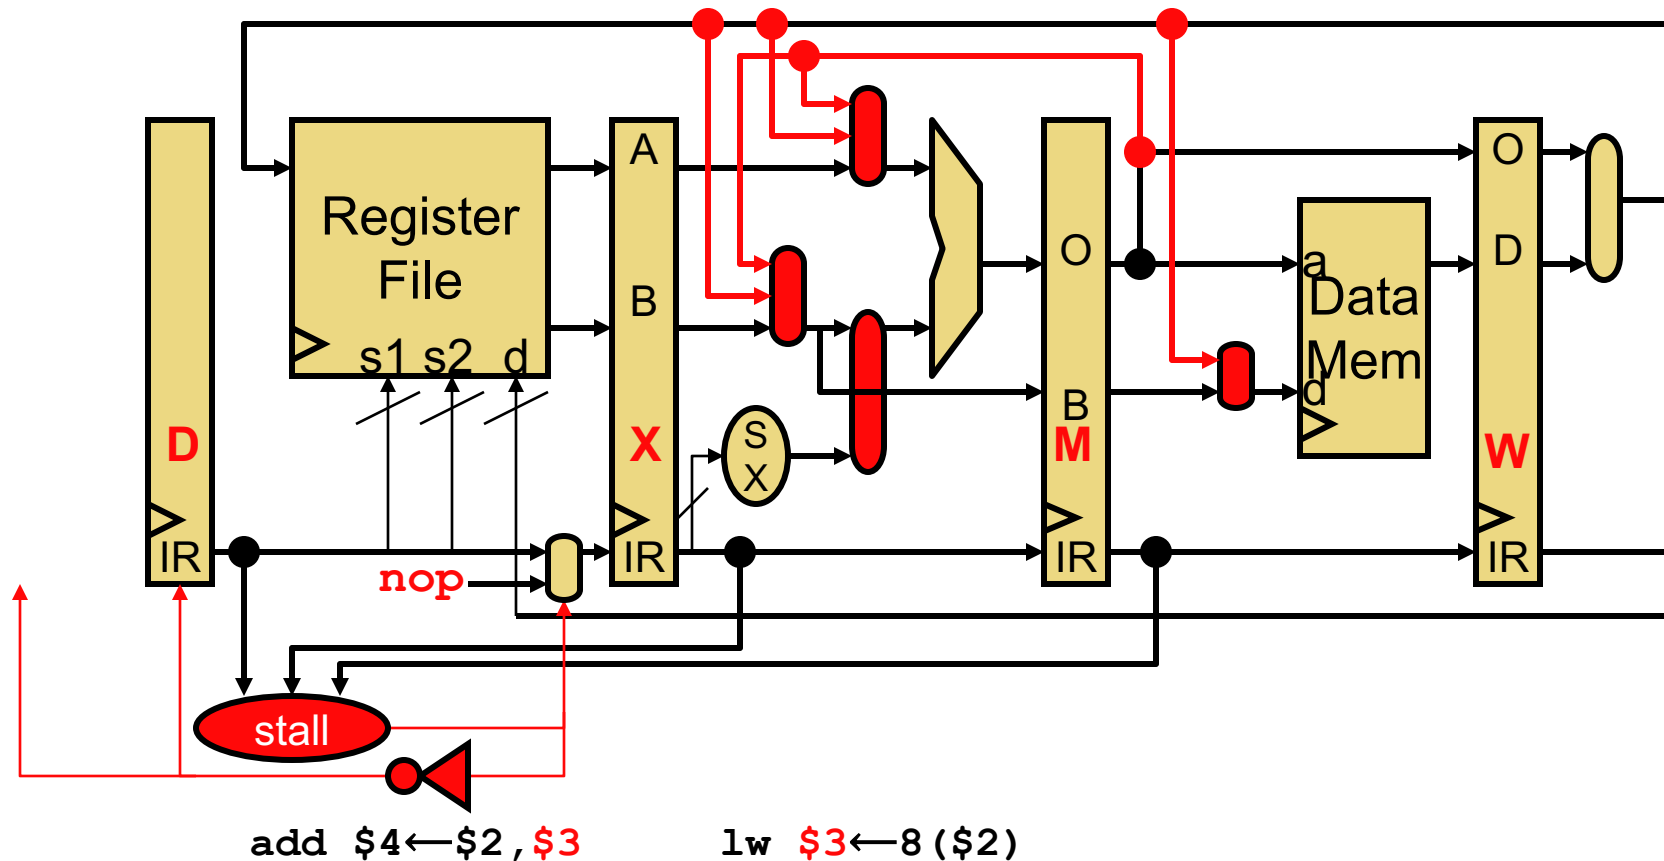
- Two separate things
 - Stall logic controls pipeline registers
 - Bypass logic controls multiplexors
- But complementary
 - For a given data hazard: if you can't bypass, must stall
- Previous slide shows **full bypassing**: all bypasses possible
 - Have we prevented all data hazards? (Thus obviating stall logic)

Have We Prevented All Data Hazards?



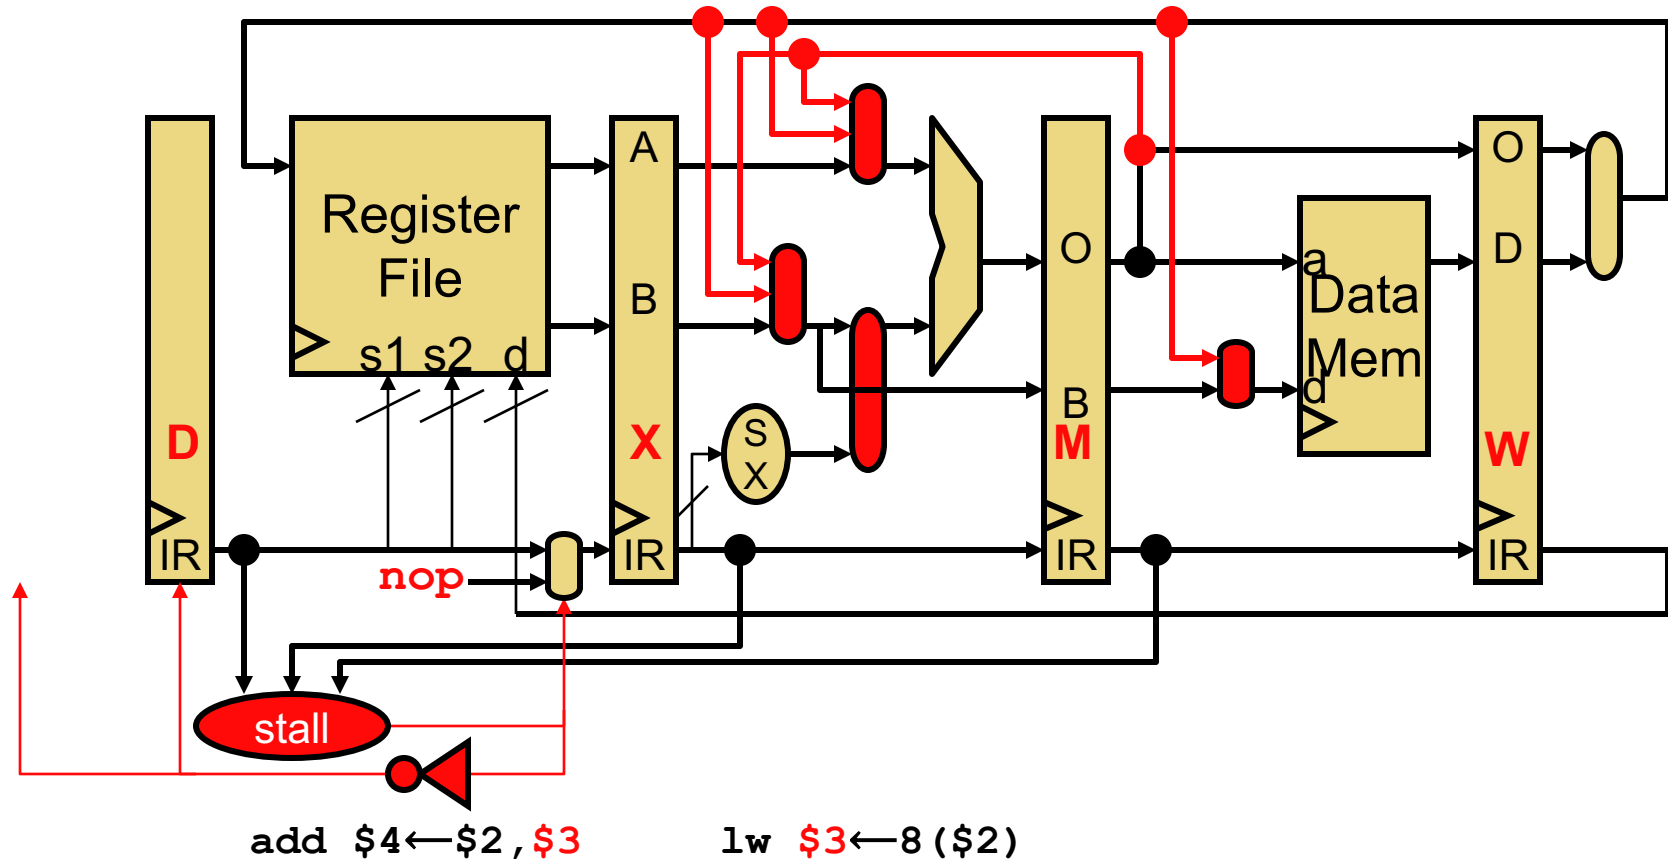
- No. Consider a "load" followed by a dependent "add" insn
- Bypassing alone isn't sufficient!
- **Hardware solution:** detect this situation and inject a stall cycle
- **Software solution:** ensure compiler doesn't generate such code

Stalling on Load-To-Use Dependences



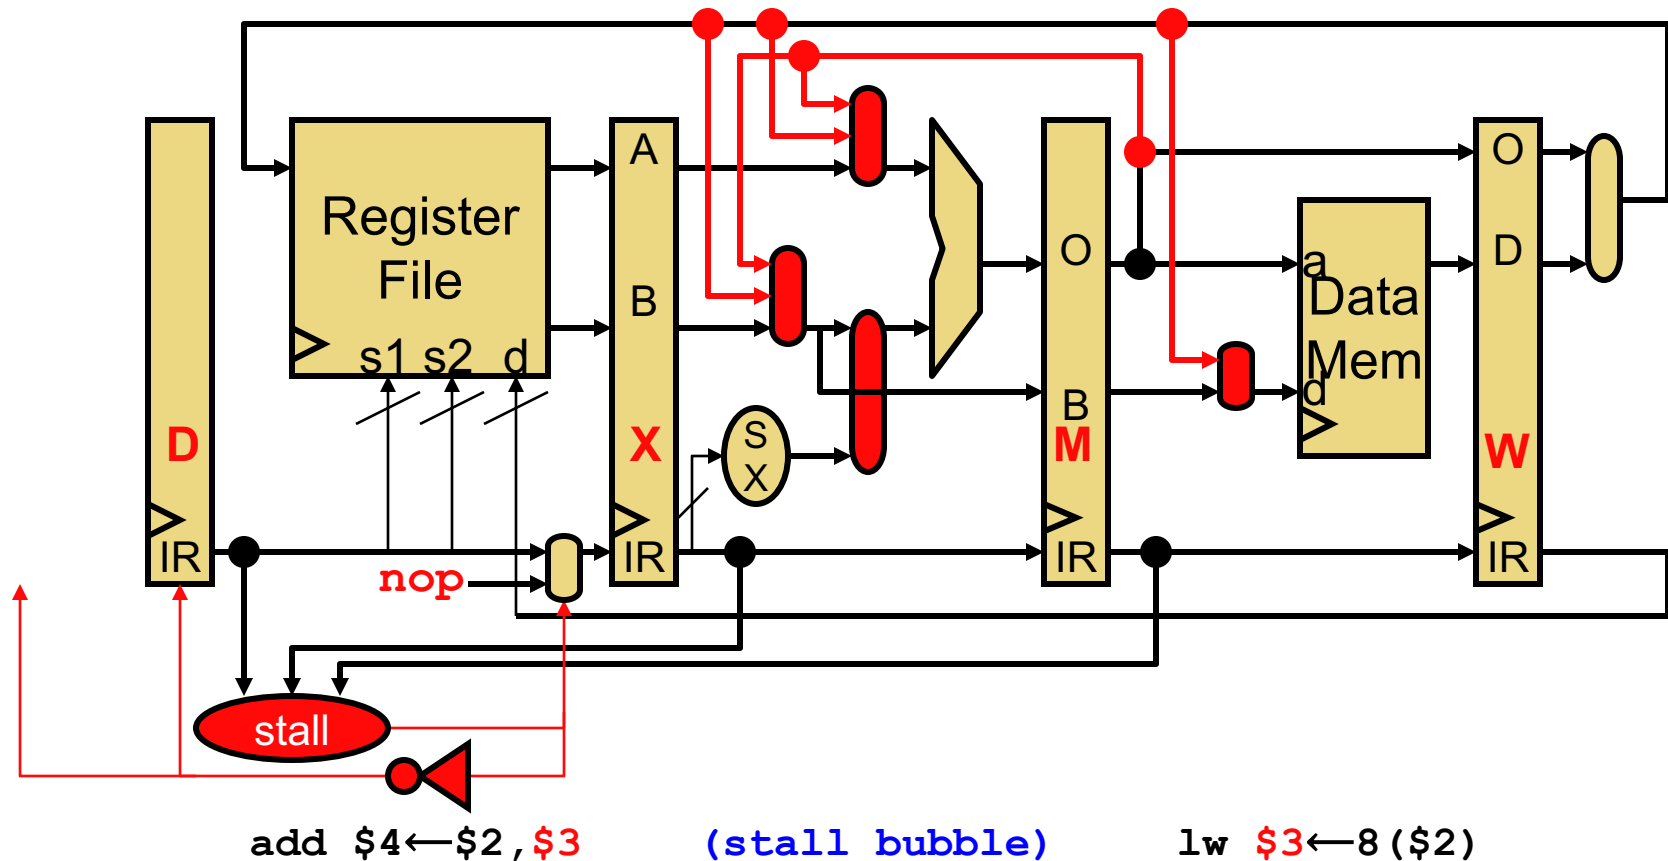
- Prevent "D insn" from advancing this cycle
 - Write **nop** into X.IR (effectively, insert **nop** in hardware)
 - Keep same "D insn", same PC next cycle
- Re-evaluate situation next cycle

Stalling on Load-To-Use Dependences



Stall = (X.IR.Operation == LOAD) &&
 ((D.IR.RegSrc1 == X.IR.RegDest) ||
 ((D.IR.RegSrc2 == X.IR.RegDest) && (D.IR.Op != STORE))
)

Stalling on Load-To-Use Dependencies



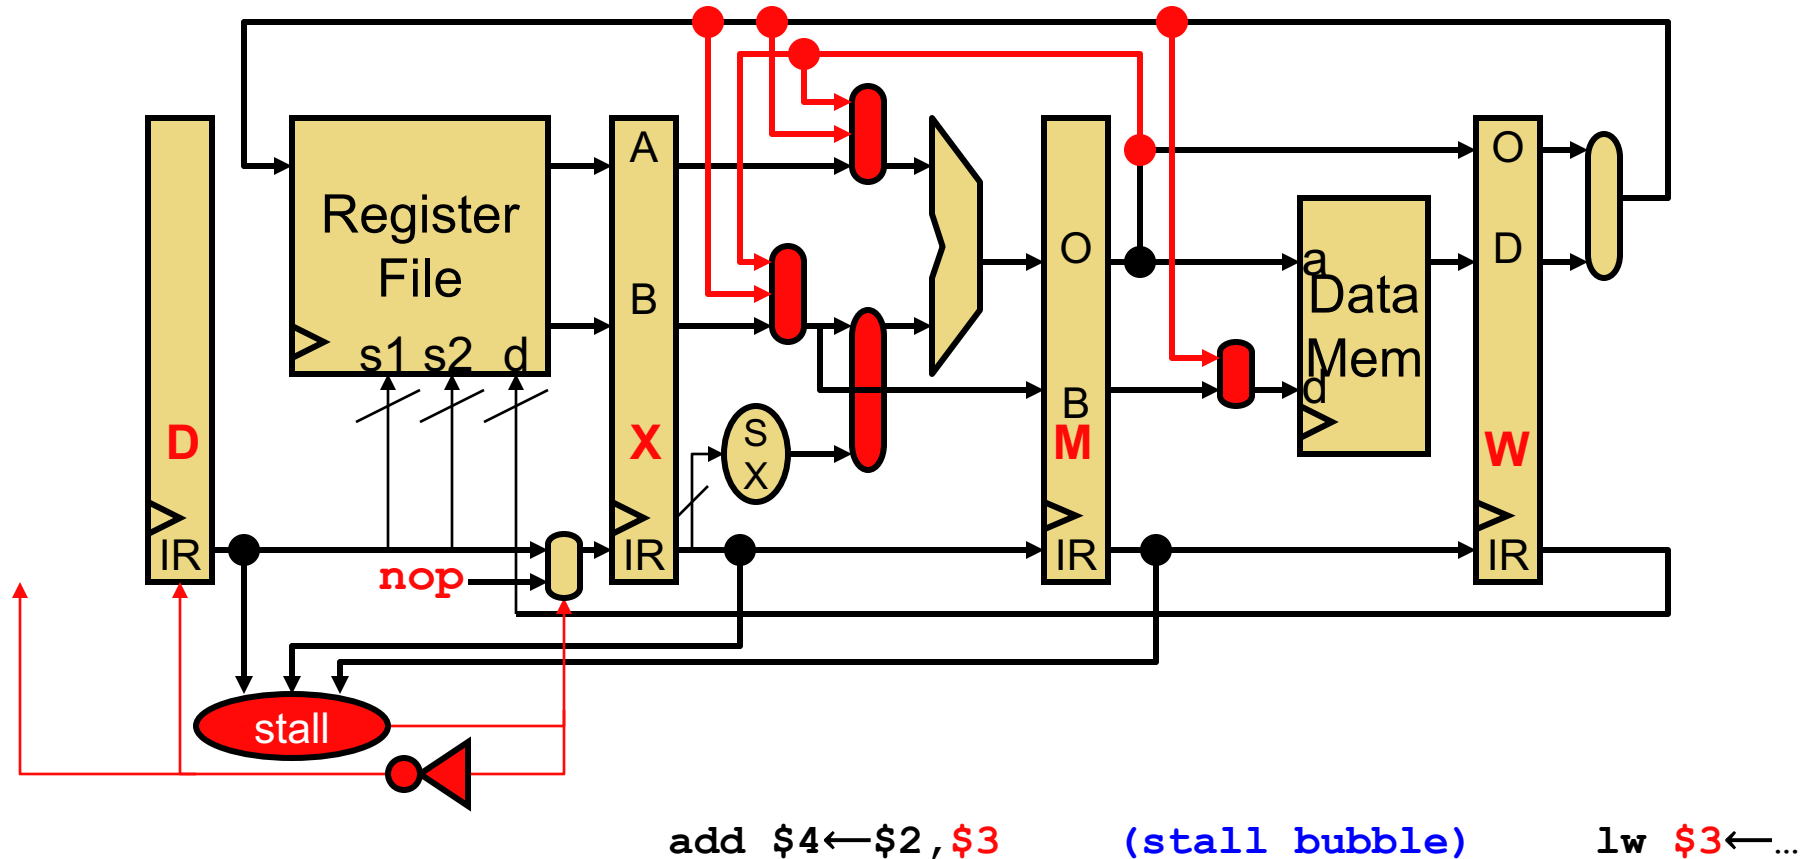
$$\text{Stall} = (\text{X.IR.Operation} == \text{LOAD}) \ \&\&$$

$$(\ (\text{D.IR.RegSrc1} == \text{X.IR.RegDest}) \ ||$$

$$((\text{D.IR.RegSrc2} == \text{X.IR.RegDest}) \ \&\& (\text{D.IR.Op} != \text{STORE}))$$

$$)$$

Stalling on Load-To-Use Dependences



Stall = (X.IR.Operation == LOAD) &&
 ((D.IR.RegSrc1 == X.IR.RegDest) ||
 ((D.IR.RegSrc2 == X.IR.RegDest) && (D.IR.Op != STORE))
)

Performance Impact of Load/Use Penalty

- Assume
 - Branch: 20%, load: 20%, store: 10%, other: 50%
 - 50% of loads are followed by dependent instruction
 - require 1 cycle stall (I.e., insertion of 1 `nop`)
- Calculate CPI
 - $\text{CPI} = 1 + (1 * 20\% * 50\%) = \mathbf{1.1}$


Reducing Load-Use Stall Frequency

	1	2	3	4	5	6	7	8	9
add \$3 , \$2, \$1	F	D	X	M	W				
lw \$4 , 4(\$3)		F	D	X	M	W			
addi \$6, \$4 , 1			F	D	d*	X	M	W	
sub \$8, \$3 , \$1				F	d*	D	X	M	W

- **d*** shows stall, addi writes into X register only in cycle 5
- Use compiler scheduling to reduce load-use stall frequency
 - As done for software interlocks, but for performance not correctness

	1	2	3	4	5	6	7	8	9
add \$3 , \$2, \$1	F	D	X	M	W				
lw \$4 , 4(\$3)		F	D	X	M	W			
sub \$8, \$3 , \$1			F	D	X	M	W		
addi \$6, \$4 , 1				F	D	X	M	W	

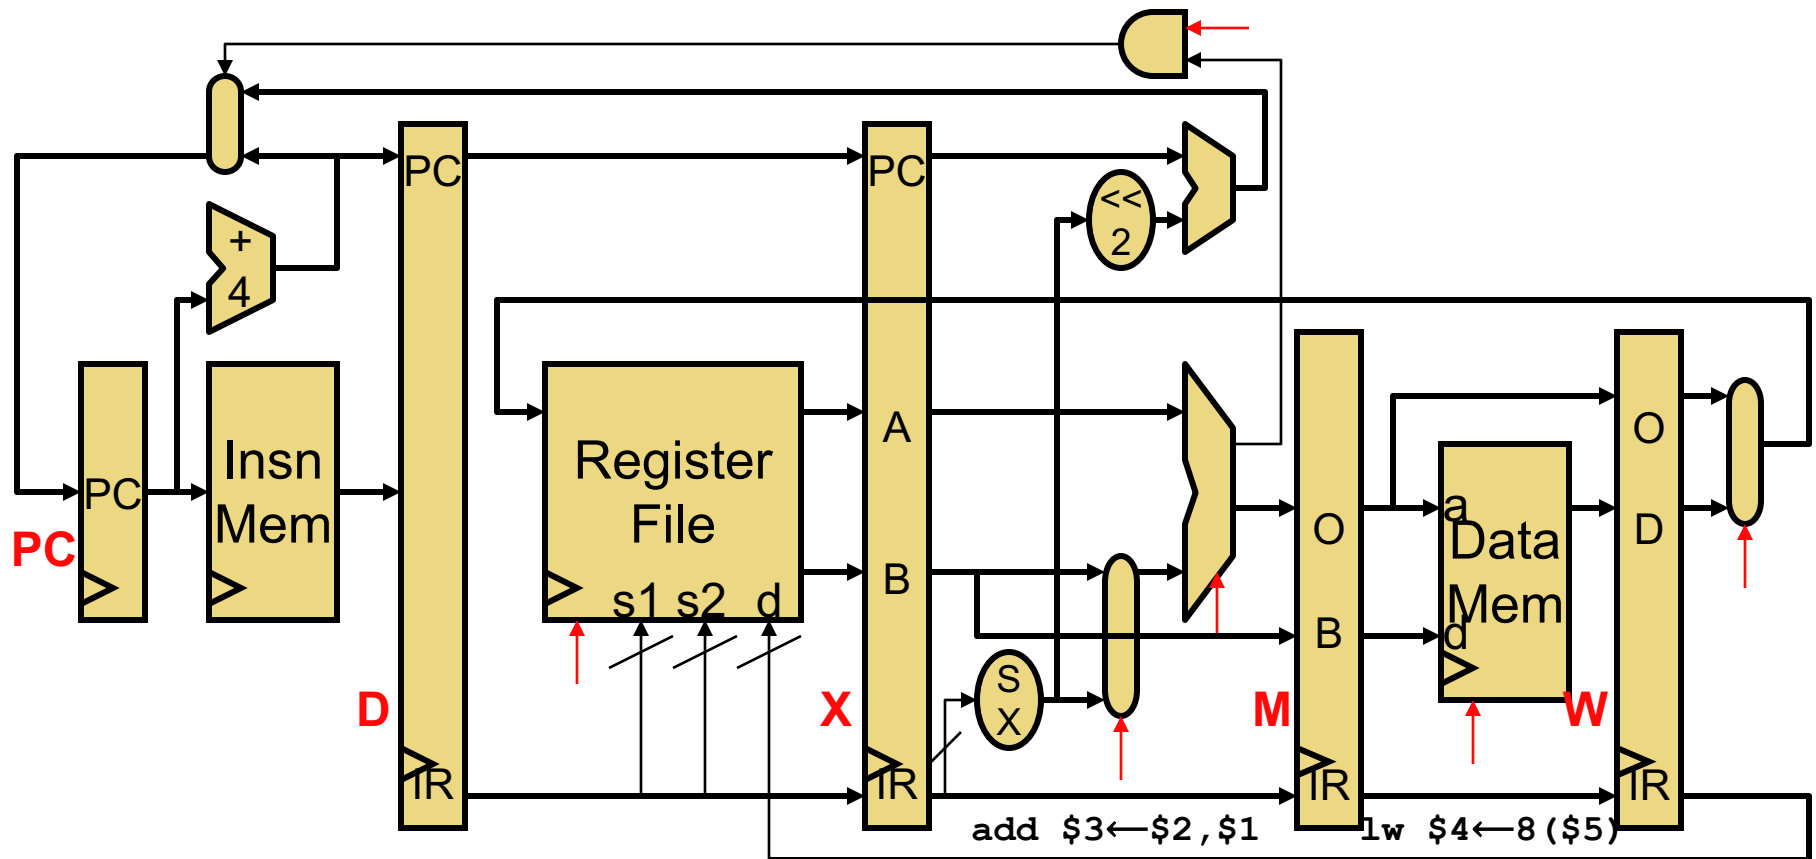
Dependencies Through Memory



- Are “load to store” memory dependencies a problem?
 - Nope
 - **lw** following **sw** to same address in next cycle, gets right value
 - Why? Data mem read/write always take place in same stage

Multi-Cycle Operations

Why Does Every Insn Take 5 Cycles?

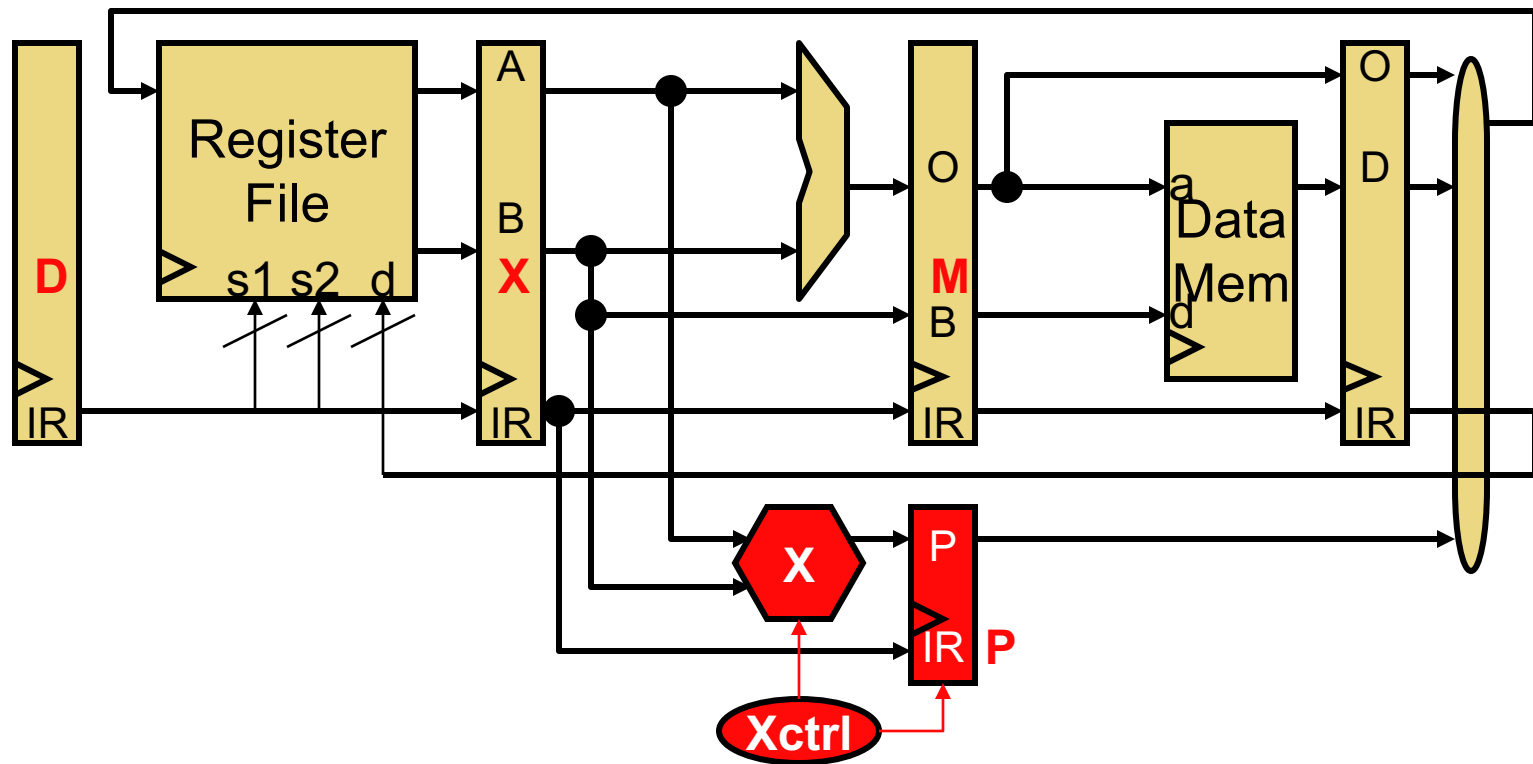


- Could/should we allow **add** to skip M and go to W? No
 - It wouldn't help: peak fetch still only 1 insn per cycle
 - **Structural hazards**: imagine **add** after **lw** (only 1 reg. write port)

Structural Hazards

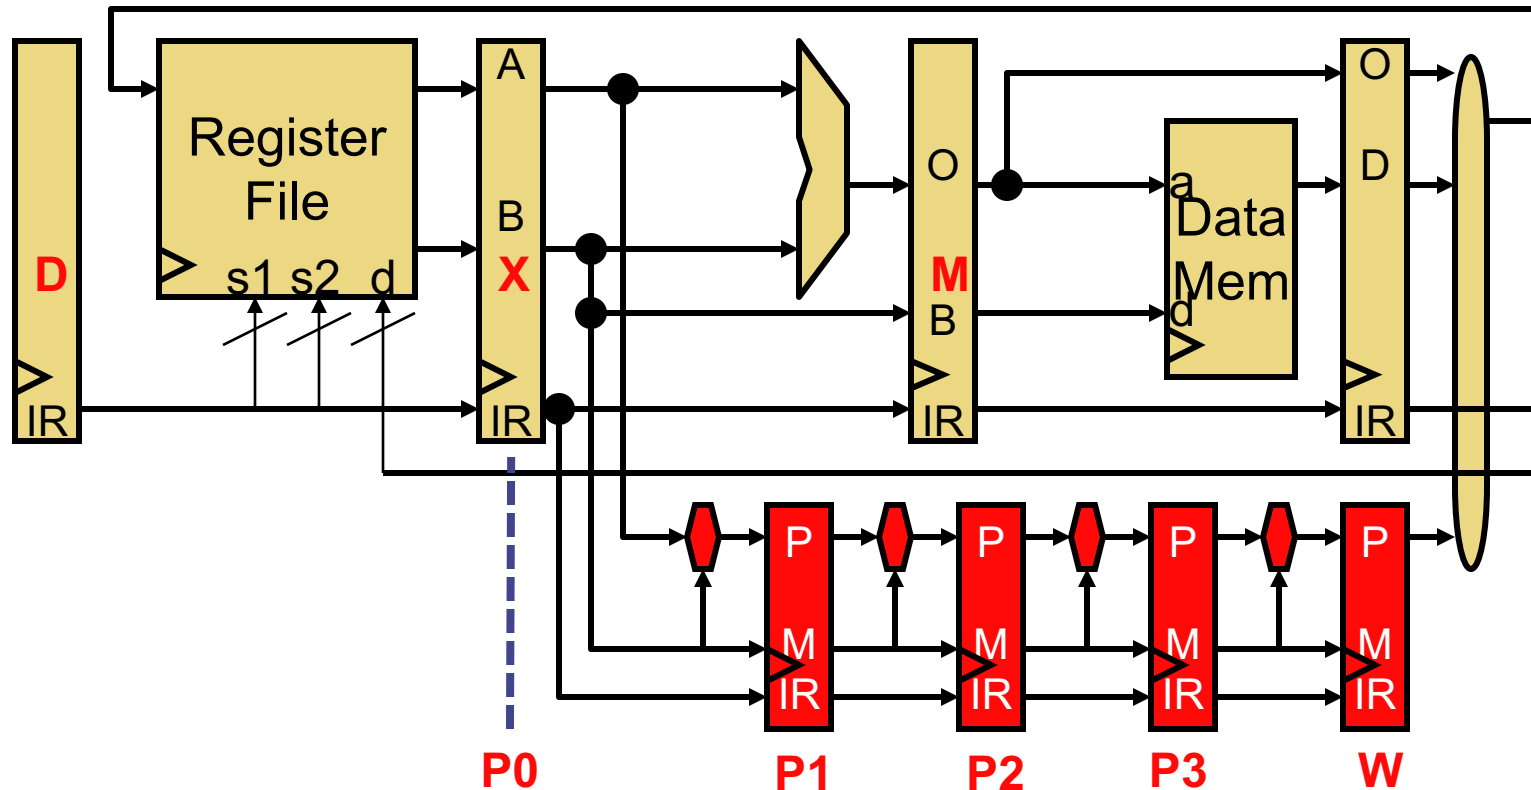
- **Structural hazards**
 - Two insns trying to use same circuit at same time
 - E.g., structural hazard on register file write port
- **To avoid structural hazards**
 - Avoided if:
 - Each insn uses every structure exactly once
 - For at most one cycle
 - All instructions travel through all stages
 - Add more resources:
 - Example: two memory accesses per cycle (Fetch & Memory)
 - Split instruction & data memories allows simultaneous access
- **Tolerate structure hazards**
 - Add stall logic to stall pipeline when hazards occur

Pipelining and Multi-Cycle Operations



- What if you wanted to add a multi-cycle operation?
 - E.g., 4-cycle multiply
 - **P**: separate output register connects to W stage
 - Controlled by pipeline control finite state machine (FSM)

A Pipelined Multiplier



- Multiplier itself is often pipelined, what does this mean?
 - Product/multiplicand register/ALUs replicated
 - Can start different multiply operations in consecutive cycles
 - **But still takes 4 cycles to generate output value**

Pipeline Diagram with Multiplier

- Allow **independent** instructions

	1	2	3	4	5	6	7	8	9
mul \$4 ← \$3, \$5	F	D	P0	P1	P2	P3	W		
addi \$6 ← \$7, 1		F	D	X	M	W			

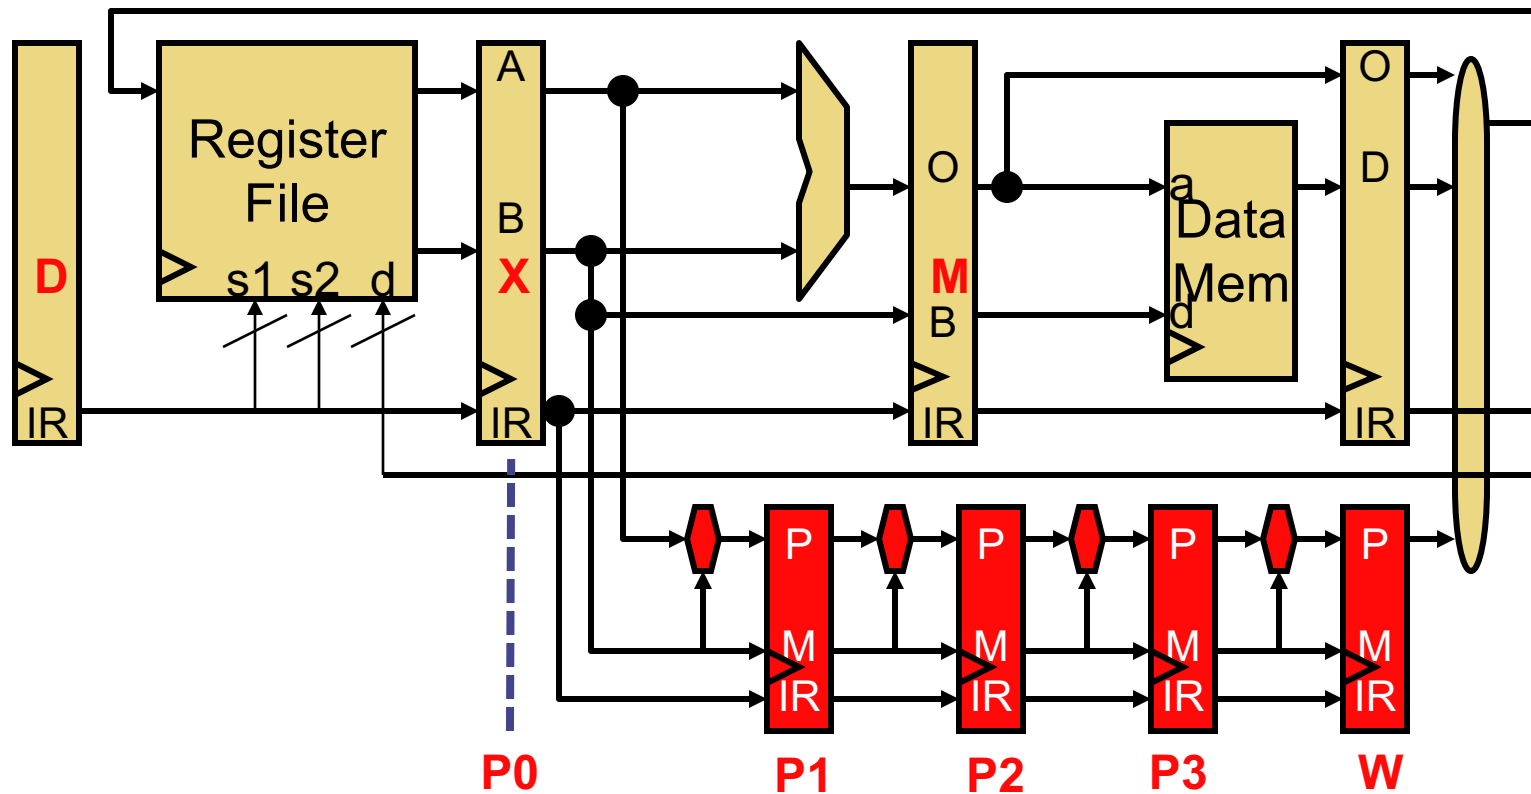
- Even allow **independent multiply** instructions

	1	2	3	4	5	6	7	8	9
mul \$4 ← \$3, \$5	F	D	P0	P1	P2	P3	W		
mul \$6 ← \$7, \$8		F	D	P0	P1	P2	P3	W	

- But must stall subsequent **dependent** instructions:

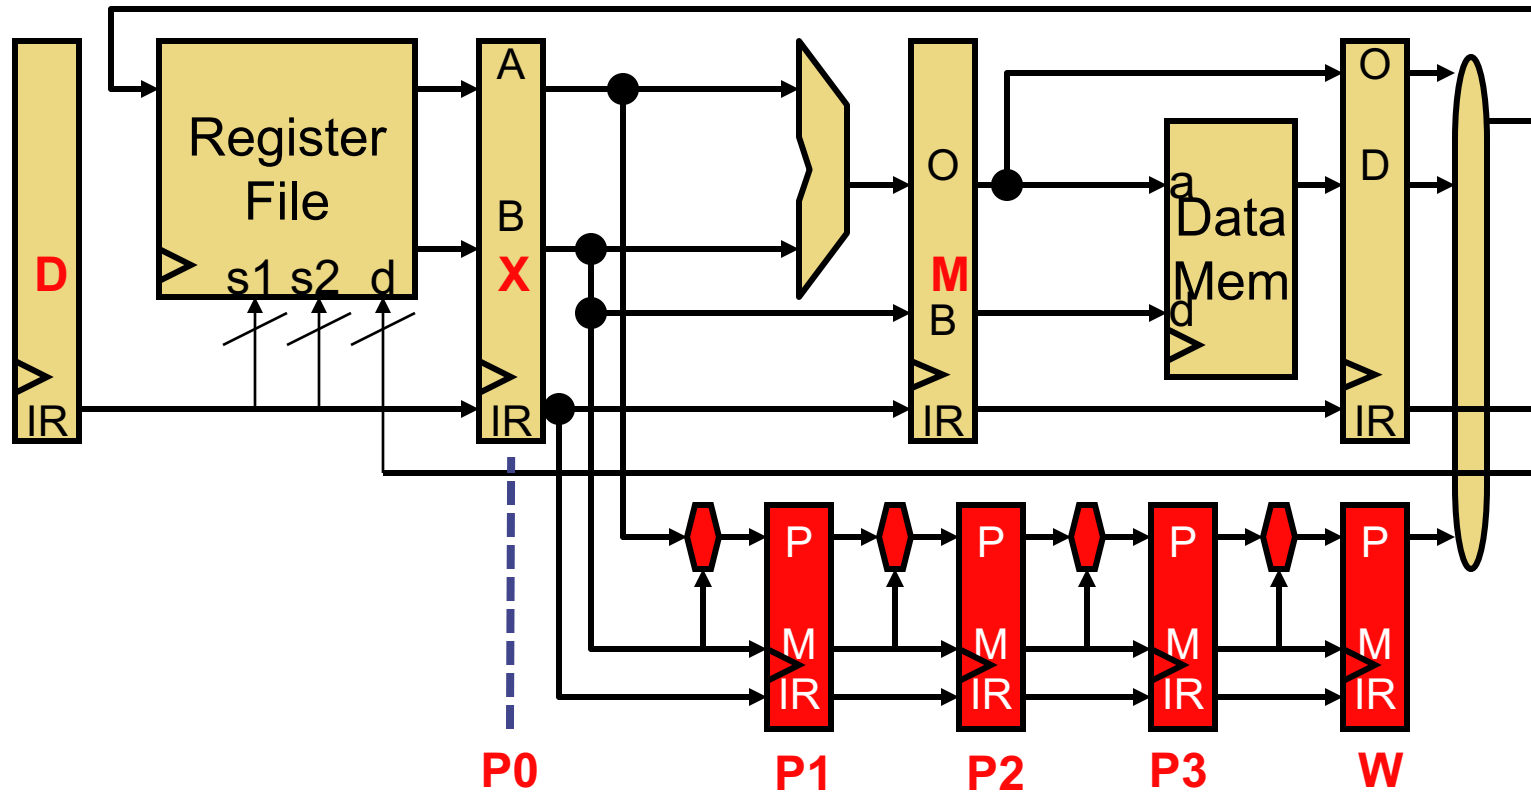
	1	2	3	4	5	6	7	8	9
mul \$4 ← \$3, \$5	F	D	P0	P1	P2	P3	W		
addi \$6 ← \$4 , 1		F	D	d*	d*	d*	↓ X	M	W

What about Stall Logic?



	1	2	3	4	5	6	7	8	9
mul $\$4 \leftarrow \$3, \$5$	F	D	P0	P1	P2	P3	W		
addi $\$6 \leftarrow \$4, 1$		F	D	d*	d*	d*	X	M	W

What about Stall Logic?



```
Stall = (OldStallLogic) ||
```

```
(D.IR.RegSrc1 == P0.IR.RegDest) || (D.IR.RegSrc2 == P0.IR.RegDest) ||
(D.IR.RegSrc1 == P1.IR.RegDest) || (D.IR.RegSrc2 == P1.IR.RegDest) ||
(D.IR.RegSrc1 == P2.IR.RegDest) || (D.IR.RegSrc2 == P2.IR.RegDest)
```

Multiplier Write Port Structural Hazard

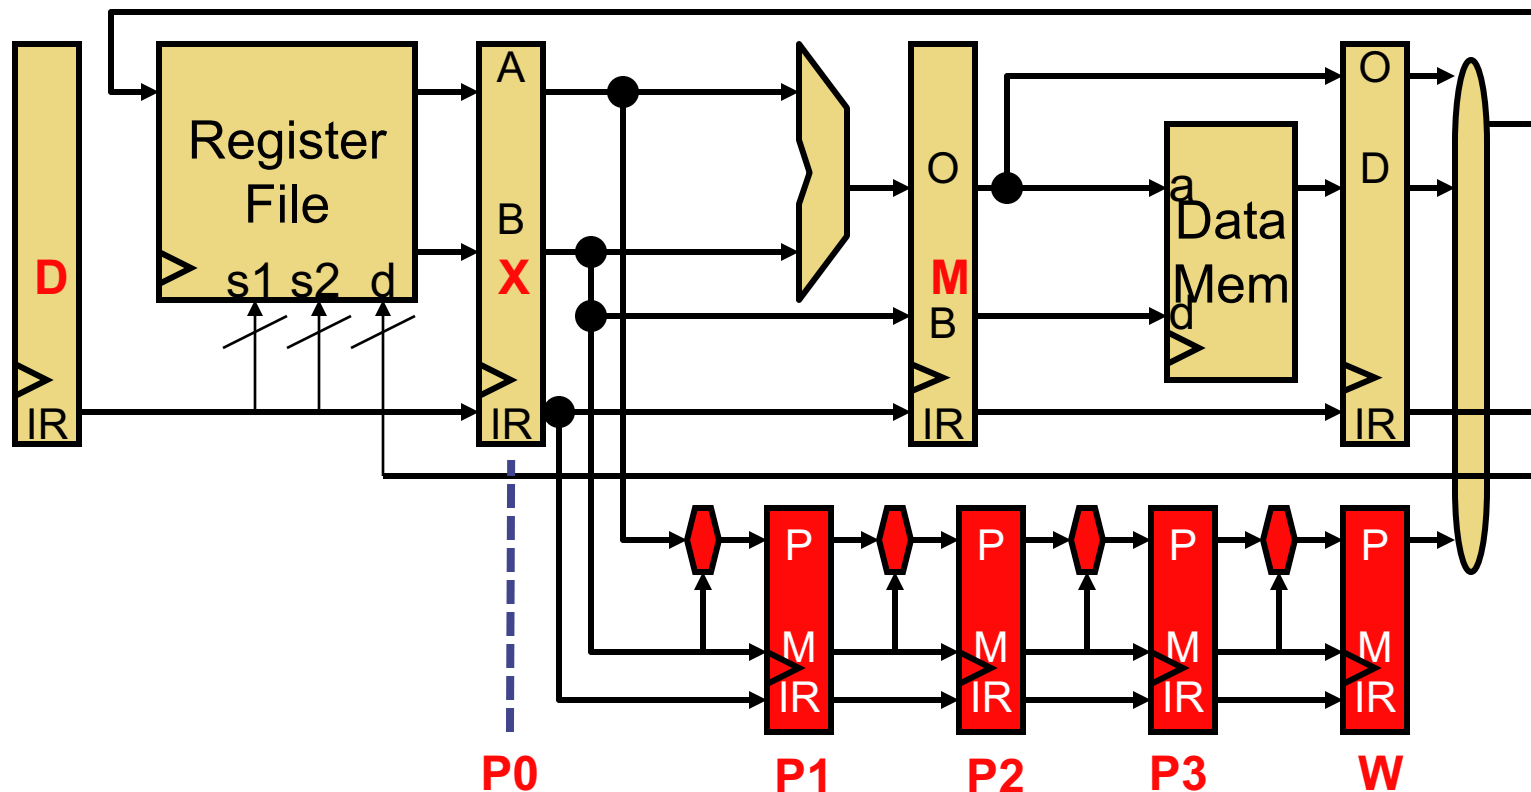
- What about...
 - Two instructions trying to write register file in same cycle?
 - Structural hazard!
- Must prevent:

	1	2	3	4	5	6	7	8	9
<code>mul \$4←\$3,\$5</code>	F	D	P0	P1	P2	P3	W		
<code>addi \$6←\$1,1</code>		F	D	X	M	W			
<code>add \$5←\$6,\$10</code>			F	D	X	M	W		

- Solution? stall the subsequent instruction

	1	2	3	4	5	6	7	8	9
<code>mul \$4←\$3,\$5</code>	F	D	P0	P1	P2	P3	W		
<code>addi \$6←\$1,1</code>		F	D	X	M	W			
<code>add \$5←\$6,\$10</code>			F	D	d*	X	M	W	

Preventing Structural Hazard



- Fix to problem on previous slide:

$$\text{Stall} = (\text{OldStallLogic}) \vee$$

$$(\text{D.IR.RegDest "is valid"} \wedge \wedge$$

$$\text{D.IR.Operation} \neq \text{MULT} \wedge \wedge \text{P1.IR.RegDest "is valid"})$$


More Multiplier Nasties

- What about...
 - Mis-ordered writes to the same register
 - Software thinks `add` gets `$4` from `addi`, actually gets it from `mul`

	1	2	3	4	5	6	7	8	9
<code>mul \$4 ← \$3, \$5</code>	F	D	P0	P1	P2	P3	W		
<code>addi \$4 ← \$1, 1</code>		F	D	X	M	W			
...									
...									
<code>add \$10 ← \$4, \$6</code>					F	D	X	M	W

- Common? Not for a 4-cycle multiply with 5-stage pipeline
 - More common with deeper pipelines
 - In any case, must be corrected

Preventing Mis-Ordered Reg. Write



- Fix to problem on previous slide:

Stall = (OldStallLogic) ||

**((D.IR.RegDest == P0.IR.RegDest) ||
(D.IR.RegDest == P1.IR.RegDest))**

Corrected Pipeline Diagram

- With the correct stall logic
 - Prevent mis-ordered writes to the same register
 - Why two cycles of delay?

	1	2	3	4	5	6	7	8	9
mul $\$4 \leftarrow \$3, \$5$	F	D	P0	P1	P2	P3	W		
addi $\$4 \leftarrow \$1, 1$		F	D	d*	d*	X	M	W	
...									
...									
add $\$10 \leftarrow \$4, \$6$					F	D	X	M	W

- **Multi-cycle operations complicate pipeline logic**

Pipelined Functional Units

- Almost all multi-cycle functional units are pipelined
 - Each operation takes N cycles
 - But can start initiate a new (independent) operation every cycle
 - Requires internal registers and some hardware replication
- + A cheaper way to add bandwidth than multiple non-pipelined units

	1	2	3	4	5	6	7	8	9	10	11
<code>mul f0 ← f1, f2</code>	F	D	E*	E*	E*	E*	W				
<code>mul f3 ← f4, f5</code>		F	D	E*	E*	E*	E*	W			

- One exception: int/FP divide: difficult to pipeline and not worth it

	1	2	3	4	5	6	7	8	9	10	11
<code>div f0 ← f1, f2</code>	F	D	E/	E/	E/	E/	W				
<code>div f3 ← f4, f5</code>		F	D	s*	s*	s*	E/	E/	E/	E/	W

- **s*** = structural hazard, two insns need same structure
 - ISAs and pipelines designed to have few of these
 - Canonical example: all insns forced to go through M stage


Pipeline Depth

Pipelining: Clock Frequency vs. IPC


- Increase number of pipeline stages (“pipeline depth”)
 - Keep cutting datapath into finer pieces
 - + Increases clock frequency (decreases clock period)
 - **Register overhead & unbalanced stages** cause sub-linear scaling
 - Double the number of stages won’t quite double the frequency
 - Increases CPI (decreases IPC)
 - More pipeline “hazards”, higher branch penalty
 - Memory latency relatively higher (same absolute lat., more cycles)
 - Result: after some point, deeper pipelining can decrease performance
 - “Optimal” pipeline depth is program and technology specific

Pipeline Depth

data from <http://cpudb.stanford.edu/>



Summary



- Processor performance
 - Latency vs throughput
- Single-cycle & multi-cycle datapaths
- Basic pipelining
- Data hazards
 - Software interlocks and scheduling
 - Hardware interlocks and stalling
 - Bypassing
 - Load-use stalling
- Pipelined multi-cycle operations