

# CIS 5710

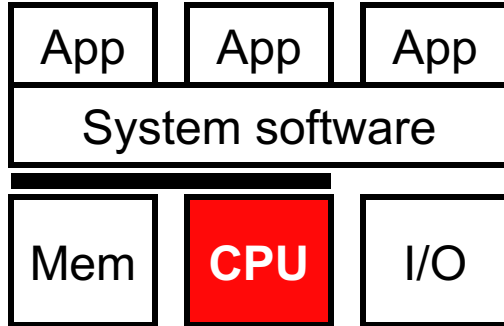
## Computer Organization and Design

### Unit 7: Branch Prediction

Based on slides by Profs. Amir Roth, Milo Martin & C.J. Taylor

# This Unit: Branch Prediction

---



- Control hazards
  - Branch prediction



# Readings


---

- P&H
  - Chapter 4

---

# Control Dependences and Branch Prediction

# What About Branches?



- Wait for branch outcome (two-cycle penalty)
- **Fetch past BR insn before BR outcome is known**
  - Default: assume “**not-taken**” (at fetch, can’t tell it’s a branch)

# Big Idea: Speculative Execution

---


- Speculation: “risky transactions on chance of profit”
- **Speculative execution**
  - Execute before all parameters known with certainty
  - **Correct speculation**
    - + Avoid stall, improve performance
  - **Incorrect speculation (mis-speculation)**
    - Must abort/flush/squash incorrect insns
    - Must undo incorrect changes (recover pre-speculation state)
- **Control speculation**: speculation aimed at control hazards
  - Unknown parameter: are these the correct insns to execute next?

# Control Speculation Mechanics

---

- Guess branch target, start fetching at guessed position
    - Doing nothing is implicitly guessing target is the next sequential PC
      - We were **already** speculating before!
    - Can actively guess other targets: **dynamic branch prediction**
  - Execute branch to verify (check) guess
    - Correct speculation? keep going
    - Mis-speculation? Flush mis-speculated insns
      - Hopefully haven't modified permanent state (Regfile, DMem)
- + Happens naturally in in-order 5-stage pipeline

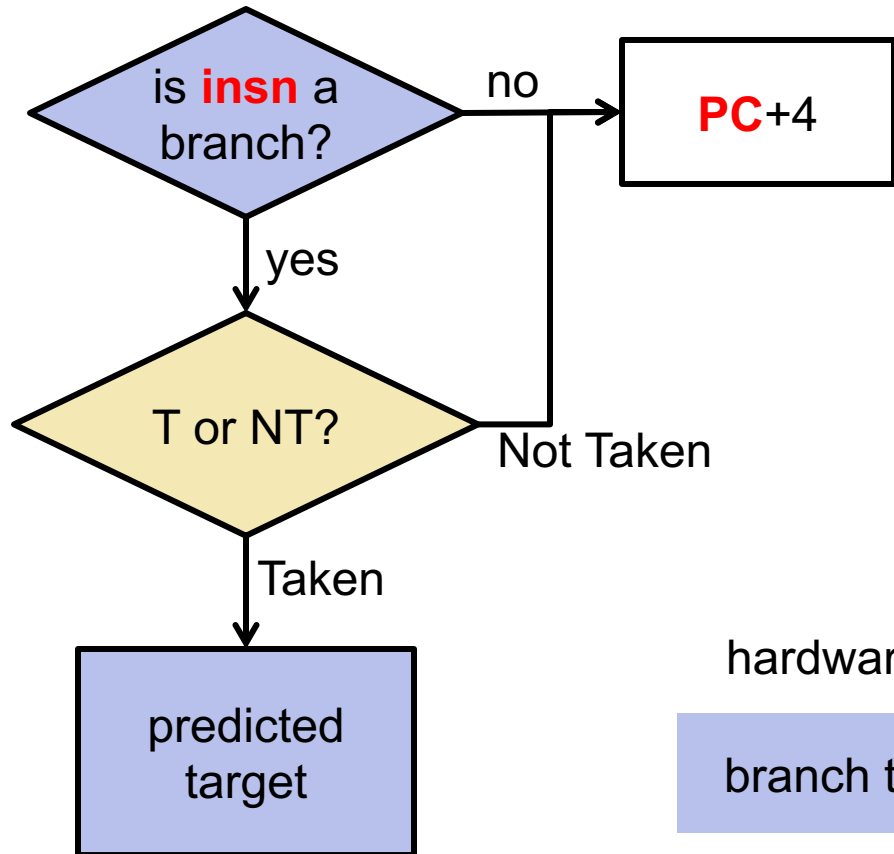
# Dynamic Branch Prediction Components



- Step #1: is it a branch?
  - Easy after decode...
- Step #2: is the branch taken or not taken?
  - Direction predictor (applies to conditional branches only)
  - Predicts taken/not-taken
- Step #3: if the branch is taken, where does it go?
  - Easy after decode...



# Branch Prediction Steps



- Which **insn**'s behavior are we trying to predict?
- Where does **PC** come from?


hardware structure:

branch target buffer

direction predictor

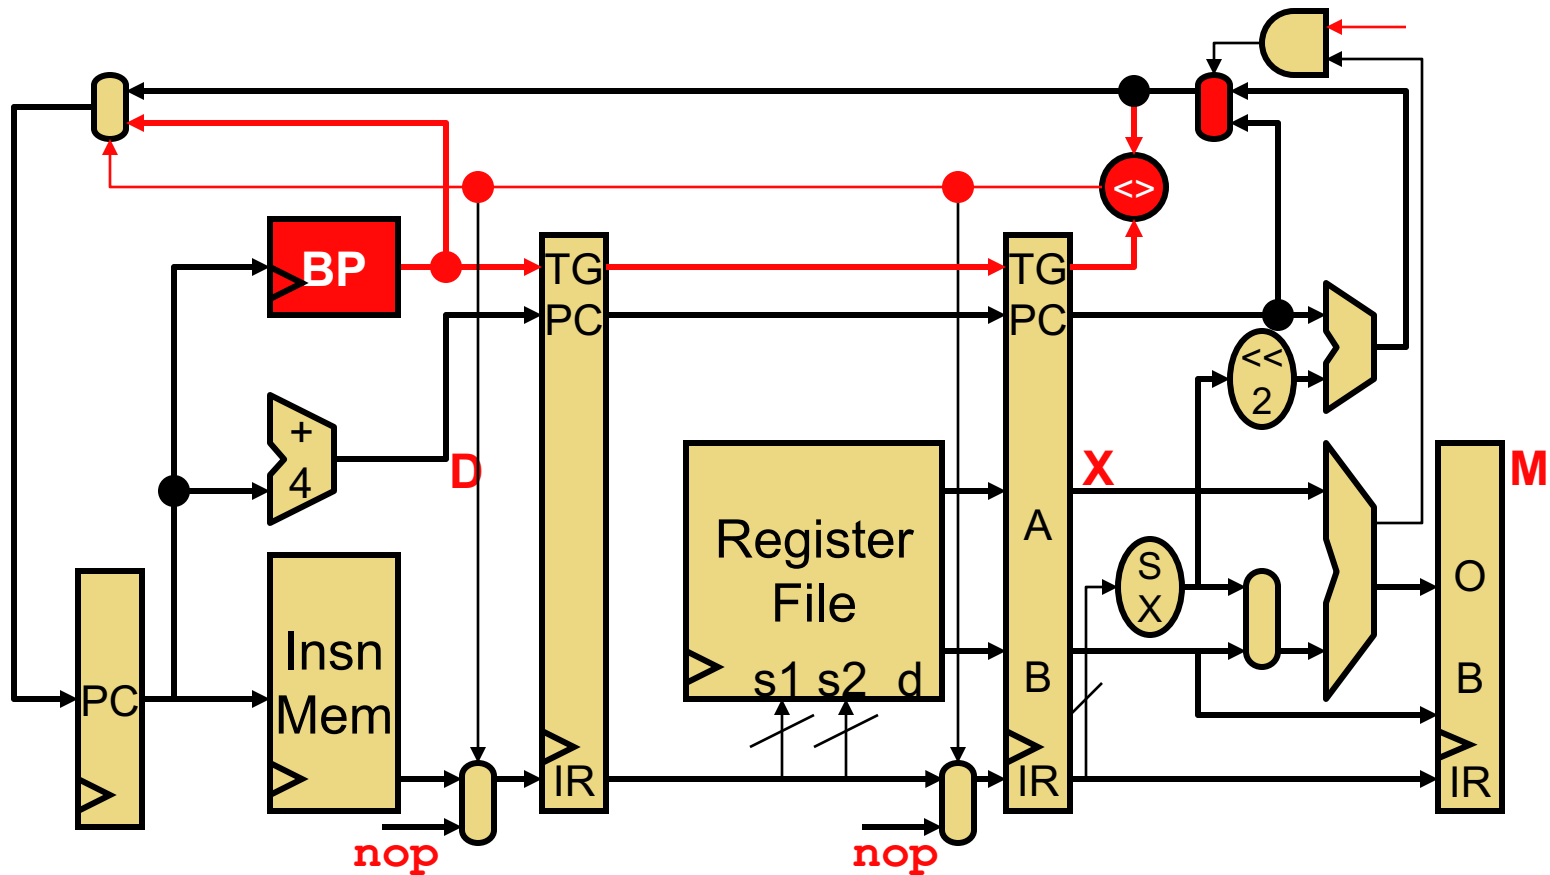
# When to Perform Branch Prediction?

- Option #1: During Decode
  - Look at instruction opcode to determine branch instructions
  - Can calculate next PC from instruction (for PC-relative branches)
    - One cycle “mis-fetch” penalty **even if branch predictor is correct**



- Option #2: During Fetch?
  - How do we do that?

# Dynamic Branch Prediction

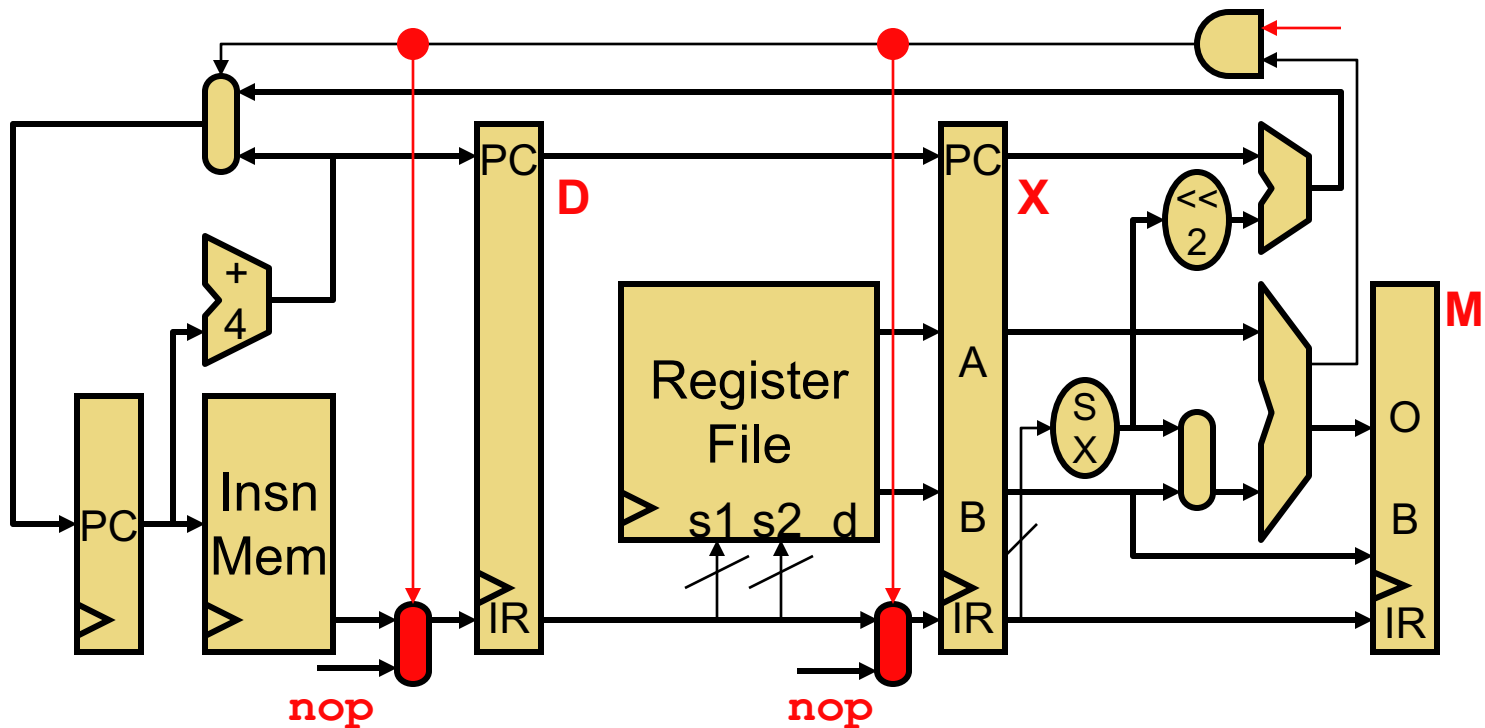


- **Dynamic branch prediction:** hardware guesses outcome
  - Start fetching from guessed address
  - Flush on **mis-prediction**

---

# MISPREDICTION RECOVERY

# Branch Recovery



- **Branch recovery**: what to do when branch is actually taken
  - Insns that are in F and D are wrong
  - **Flush them**, i.e., replace them with **nops**
  - + They haven't written permanent state yet (regfile, DMem)
  - Two cycle penalty for taken branches

# Branch Speculation and Recovery

**Correct:**

	1	2	3	4	5	6	7	8	9
<code>addi r3←r1,1</code>	F	D	X	M	W				
<code>bnez r3,targ</code>		F	D	X	M	W			
<code>st r6→[r7+4]</code>			<b>F</b>	<b>D</b>	X	M	W		
<code>mul r10←r8,r9</code>				<b>F</b>	D	X	M	W	

speculative

- **Mis-speculation recovery:** what to do on wrong guess
  - Not too painful in a short, in-order pipeline
  - Branch resolves in X
  - + Younger insns (in F, D) haven't changed permanent state
  - **Flush** insns currently in D and X (i.e., replace with **nops**)


**Recovery:**

	1	2	3	4	5	6	7	8	9
<code>addi r3←r1,1</code>	F	D	X	M	W				
<code>bnez r3,targ</code>		F	D	<b>X</b>	M	W			
<del><code>st r6→[r7+4]</code></del>			<b>F</b>	<b>D</b>	--	--	--		
<del><code>mul r10←r8,r9</code></del>				<b>F</b>	--	--	--	--	
<code>targ:add r4←r4,r5</code>					<b>F</b>	D	X	M	W

---

# IS IT A BRANCH?

# Revisiting Branch Prediction Components

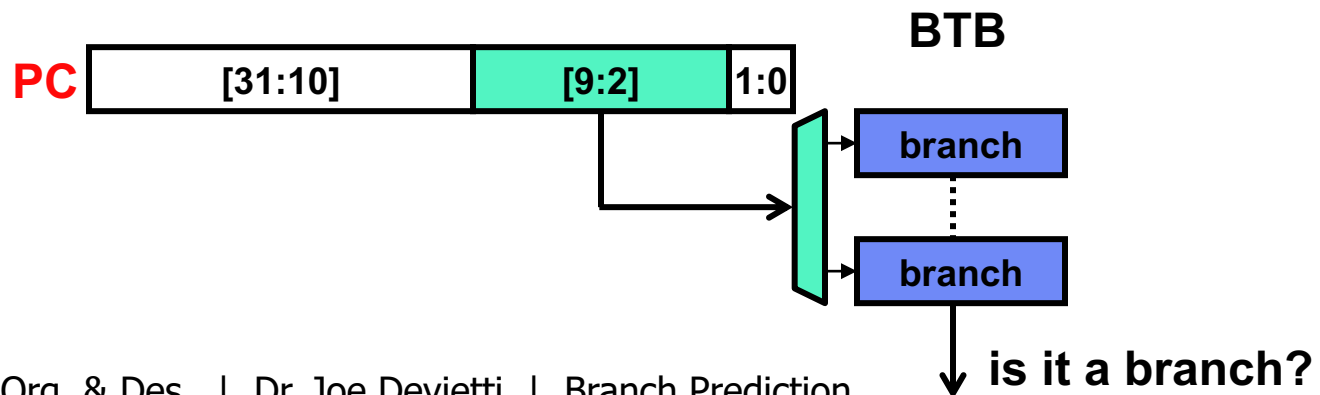


- Step #1: is it a branch?
  - Easy after decode... during fetch: **predictor**
- Step #2: is the branch taken or not taken?
  - Direction predictor (later)
- Step #3: if the branch is taken, where does it go?
  - Branch target predictor (BTB)
  - Supplies target PC if branch is taken




# Branch Target Buffer

- **Learn from the past to predict the future**
  - Record the past in a hardware structure
- **Branch target buffer (BTB):**
  - Record a list of branches we have seen
    - + code doesn't change
  - PC indexes table of bits
    - each entry is 1 bit: is there a branch here?
    - set the bit if we see a branch at that index
    - What about two PCs with the same lower bits...?



# Branch Target Buffer


- BTB entries are too coarse-grained
- + Record only branches that were taken at least once
  - a never-taken branch might as well be a NOP
    - doesn't help enough
- better idea: **Tag each BTB entry**
  - remember **some** things precisely, rather than **everything** imprecisely
  - record a subset of actual taken branches
  - $\text{is\_a\_branch} = (\text{BTB}[\text{PC}].\text{branch} \ \&\& \ \text{BTB}[\text{PC}].\text{tag} == \text{PC})$
  - How big is each tag?



---

# **BRANCH TARGET PREDICTION**


# Revisiting Branch Prediction Components



- Step #1: is it a branch?
  - Easy after decode... during fetch: predictor
- Step #2: is the branch taken or not taken?
  - Direction predictor
- Step #3: if the branch is taken, where does it go?
  - **Branch target predictor (BTB)**
  - Supplies target PC if branch is taken


# Branch Target Buffer, Again

- **Branch target buffer (BTB):**
  - “guess” the future PC based on past behavior
  - “Last time the branch X was taken, it went to address Y”
    - “So, in the future, if address X is fetched, fetch address Y next”
  - Essentially: branch will go to same place it went last time
  - PC indexes table of **target addresses**
    - use tags to precisely remember a subset of branch targets
  - What about aliasing?
    - Two PCs with the same lower bits?
    - No problem, just a prediction!



# Branch Target Buffer (continued)

- At Fetch, how do we know we have a branch? We don't...
  - ...**all insns access BTB in parallel with Imem Fetch**
- **BTB predicts which insn are branches, and targets**
  - tag each entry with its corresponding PC
  - Update BTB on every taken branch insn, record target PC:
    - $\text{BTB}[\text{PC}].\text{tag} = \text{PC}$ ,  $\text{BTB}[\text{PC}].\text{target} = \text{target of branch}$
  - All insns access at Fetch *in parallel* with Imem
    - Check for tag match, indicates insn at that PC is a branch
      - otherwise, assume insn is **not** a branch
    - $\text{Predicted PC} = (\text{BTB}[\text{PC}].\text{tag} == \text{PC}) ? \text{BTB}[\text{PC}].\text{target} : \text{PC}+4$




# Why Does a BTB Work?

---

- Because most control insns use **direct targets**
  - Target encoded in insn itself → same “taken” target every time
- What about **indirect targets**?
  - Target held in a register → can be different each time
  - Two indirect call idioms
    - + Dynamically linked functions (DLLs): target always the same
      - Dynamically dispatched (virtual) functions: hard but uncommon
  - Also two indirect unconditional jump idioms
    - Switches: hard but uncommon
    - Function returns: **hard and common**

# Return Address Stack (RAS)




- **Return address stack (RAS)**
  - Call instruction?  $RAS[TopOfStack++] = PC+4$
  - Return instruction?  $Predicted\text{-}target = RAS[--TopOfStack]$
  - Q: how can you tell if an insn is a call/return before decoding it?
    - mark some BTB entries as “returns”, or use another table



---

# BRANCH DIRECTION PREDICTION

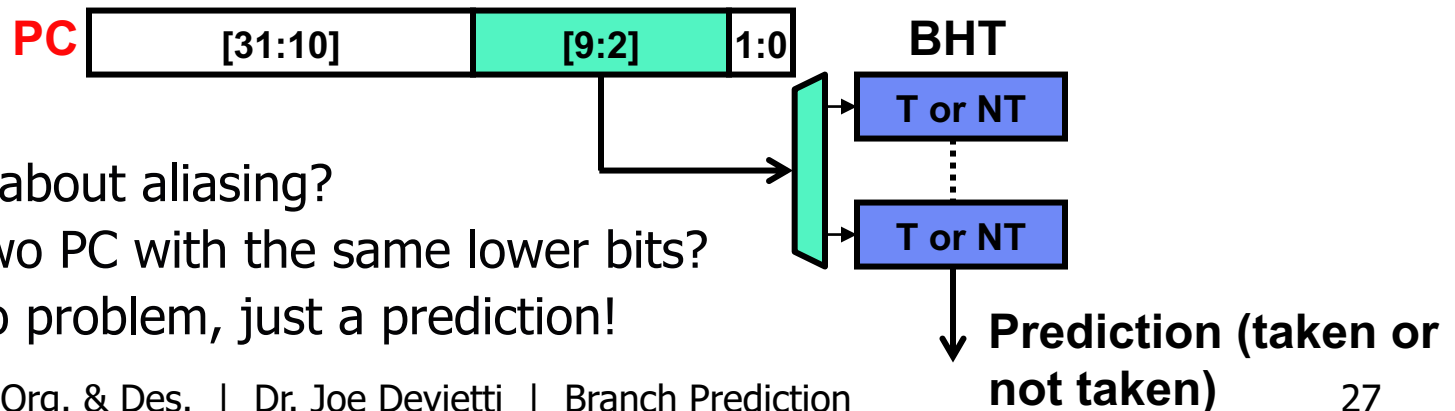
# Revisiting Branch Prediction Components



- Step #1: is it a branch?
  - Easy after decode... during fetch: predictor
- Step #2: is the branch taken or not taken?
  - **Direction predictor**
- Step #3: if the branch is taken, where does it go?
  - Branch target predictor (BTB)
  - Supplies target PC if branch is taken

# Branch Direction Prediction

- **Learn from past, predict the future**
  - Record the past in a hardware structure
- **Direction predictor (DIRP)**
  - Map conditional-branch PC to taken/not-taken (T/N) decision
  - Individual conditional branches often biased or weakly biased
    - 90%+ one way or the other considered **"biased"**
    - Why? Loop back edges, checking for uncommon conditions
- **Bimodal predictor**: simplest predictor
  - PC indexes Branch History Table of bits (0 = N, 1 = T), no tags
  - Essentially: branch will go same way it went last time



# Bimodal Branch Predictor

- simplest direction predictor
  - PC indexes table of bits (0 = N, 1 = T), no tags
  - Essentially: branch will go same way it went last time
  - Problem: **inner loop branch** below

```
for (i=0;i<100;i++)
    for (j=0;j<3;j++)
        // whatever
```

    - Two “built-in” mis-predictions per inner loop iteration
    - Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4	T	T	N	Wrong
5	N	N	T	Wrong
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	N	N	T	Wrong
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong

# Two-Bit Saturating Counters (2bc)

- **Two-bit saturating counters (2bc)**

[Smith 1981]

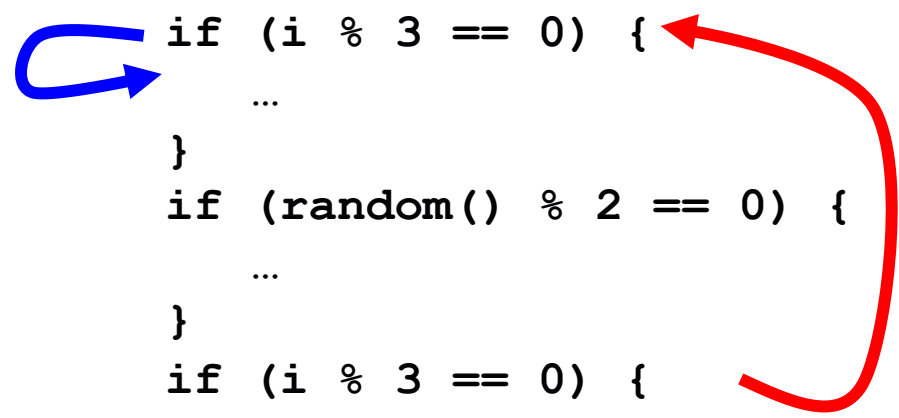
- Replace each single-bit prediction
  - $(0,1,2,3) = (N,n,t,T)$
- Adds "hysteresis"
  - Force predictor to mis-predict twice before "changing its mind"
- One mispredict each loop execution (rather than two)
  - + Fixes this pathology (which is not contrived, by the way)
  - Can we do even better?

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	T	T	Correct
4	T	T	N	Wrong
5	t	T	T	Correct
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	t	T	T	Correct
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong

# Branches may be correlated

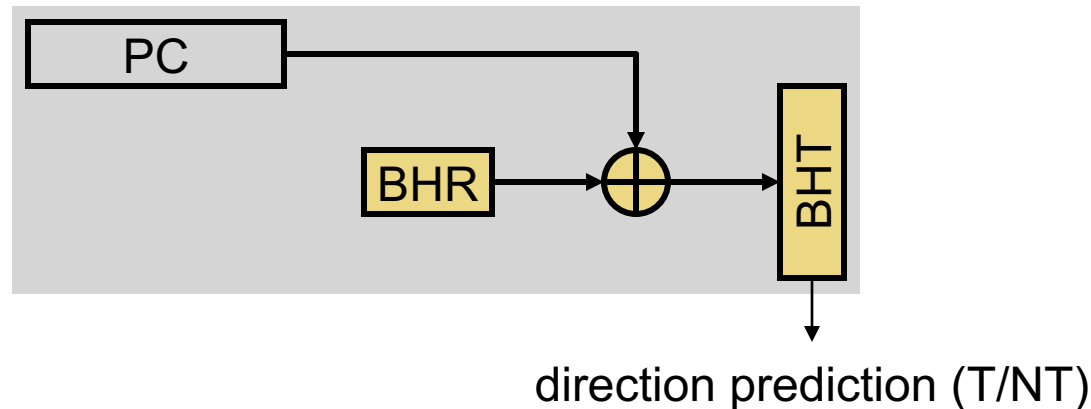
- Consider:

```
for (i=0; i<1000000; i++) {    // Highly biased
    if (i % 3 == 0) {          // Locally correlated
        ...
    }
    if (random() % 2 == 0) {    // Unpredictable
        ...
    }
    if (i % 3 == 0) {          // Globally correlated
        ...
    }
}
```



# Gshare History-Based Predictor

- Exploits observation that branch outcomes are correlated
- Maintains recent branch outcomes in **Branch History Register (BHR)**
  - In addition to BHT of counters (typically 2-bit sat. counters)
- How do we incorporate history into our predictions?
  - Use PC xor BHR to index into BHT. Why?



# Gshare History-based Predictor


- Gshare working example
  - assume program has one branch
  - BHT: one 1-bit DIRP entry
  - **3BHR**: last 3 branch outcomes
  - train counter, **and** update BHR after each branch

Time	State	BHR	Prediction	Outcome	Result?
1	N	NNN	N	T	wrong
2	N	NNT	N	T	wrong
3	N	NTT	N	T	wrong
4	N	TTT	N	N	correct
5	N	TTN	N	T	wrong
6	N	TNT	N	T	wrong
7	T	NTT	T	T	correct
8	N	TTT	N	N	correct
9	T	TTN	T	T	correct
10	T	TNT	T	T	correct
11	T	NTT	T	T	correct
12	N	TTT	N	N	correct



# Hybrid Predictor

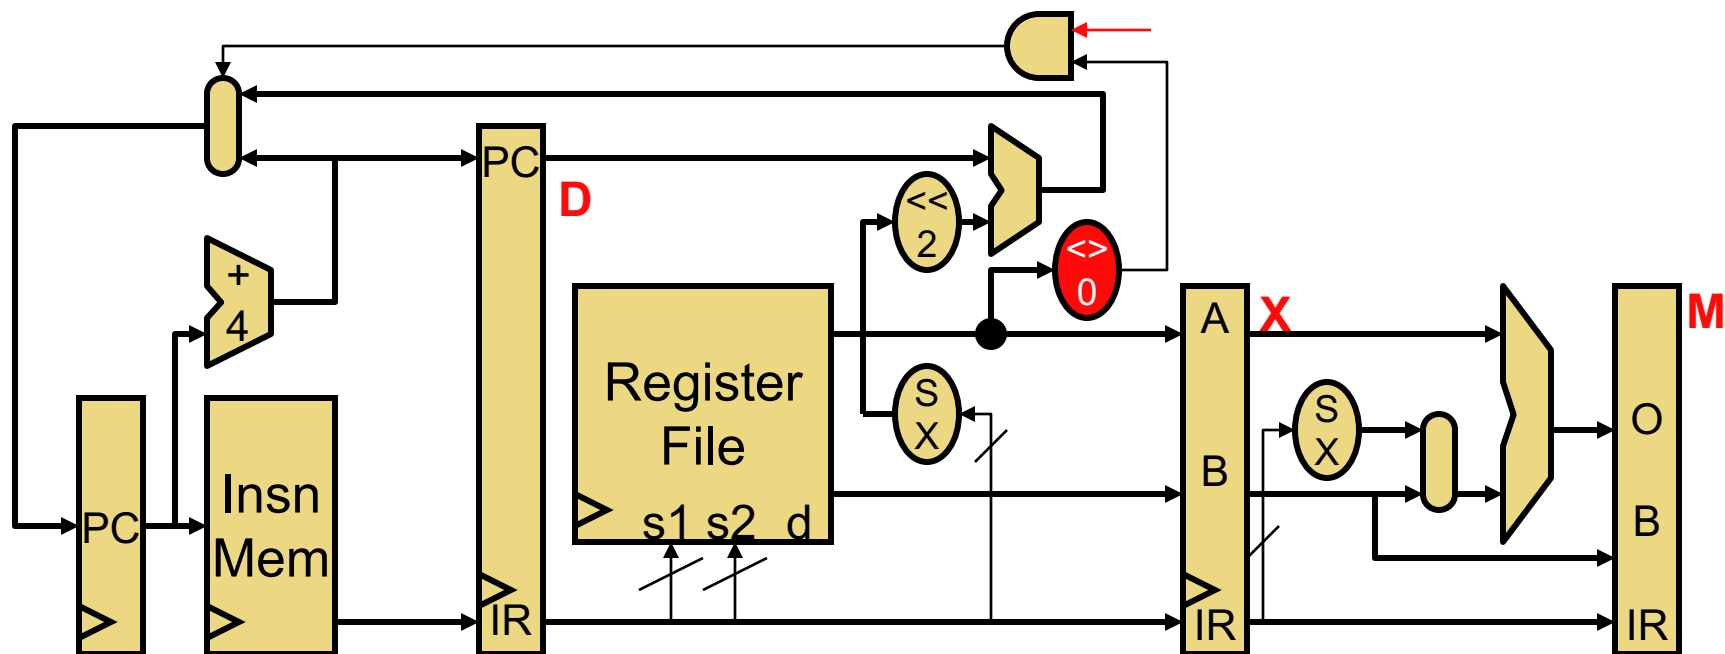
- **Hybrid (tournament) predictor** [McFarling 1993]
  - Attacks correlated predictor BHT capacity problem
  - Idea: combine two predictors
    - **Simple bimodal predictor** for history-independent branches
    - **Correlated predictor** for branches that need history
    - **Chooser** assigns branches to one predictor or the other
    - Branches start in simple BHT, move mis-prediction threshold
- + Correlated predictor can be made **smaller**, handles fewer branches
- + 90–95% accuracy



---

# REDUCING BRANCH PENALTY


# Reducing Penalty: Fast Branches



- **Fast branch**: can decide at D, not X
  - Test must be comparison to zero or equality, no time for ALU
  - + New taken branch penalty is 1
  - Additional insns (**s1t**) for more complex tests, must bypass to D too

# Reducing Penalty: Fast Branches

- **Fast branch**: targets control-hazard penalty
  - Basically, branch insns that can resolve at D, not X
    - Test must be comparison to zero or equality, **no time for ALU**
- + New taken branch penalty is 1
- Additional comparison insns (e.g., `cmplt`, `slt`) for complex tests
- Must bypass into decode stage now, too




# Fast Branch Performance

---

- Assume: Branch: 20%, 75% of branches are taken
  - $\text{CPI} = 1 + 20\% * 75\% * 1 = 1 + 0.20 * 0.75 * 1 = 1.15$
  - **15% slowdown**
- But wait, fast branches assume only simple comparisons
  - Fine for MIPS
  - But not fine for ISAs with “branch if \$1 > \$2” operations
- In such cases, say 25% of branches require an extra insn
  - $\text{CPI} = 1 + (20\% * 75\% * 1) + 20\% * 25\% * 1(\text{extra insn}) = 1.2$
- Example of ISA and micro-architecture interaction
  - Type of branch instructions
  - Another option: “Delayed branch” or “branch delay slot”
  - What about condition codes?

# Putting It All Together

- BTB & branch direction predictor during fetch



- If branch prediction correct, no taken branch penalty

# Branch Prediction Performance

---

- Dynamic branch prediction
  - 20% of instruction branches
  - Simple predictor: branches predicted with 75% accuracy
    - $\text{CPI} = 1 + (20\% * 25\% * 2) = 1.1$
  - More advanced predictor: 95% accuracy
    - $\text{CPI} = 1 + (20\% * 5\% * 2) = 1.02$
- Branch mis-predictions still a big problem though
  - Pipelines are long: typical mis-prediction penalty is 10+ cycles
  - For cores that do more per cycle, predictions more costly (later)

---

# PREDICATION



# Predication

---

- Instead of predicting which way we're going, why not go **both ways**?
  - compute a **predicate bit** indicating a condition
  - ISA includes **predicated instructions**
    - predicated insns either execute as normal or as NOPs, depending on the predicate bit
- Examples
  - x86 cmov performs conditional load/store
  - 32b ARM allows almost all insns to be predicated
    - 64b ARM has predicated reg-reg move, inc, dec, not
  - Nvidia GPU ISA supports predication on most insns
  - predicate bits are like LC4 NZP bits
    - x86 FLAGS, ARM condition codes

# Predication Example

---

- Instead of predicting which way we're going, why not go both ways?
  - compute a predicate bit indicating a condition
  - ISA includes predicated instructions
    - predicated insns either execute as normal or as NOPs, depending on the predicate bits

// C code	; original LC4	; predicated LC4
if (a <= b) {	CMP R1 R2	CMP R1 R2
x += y;	BRn else	ADDzp R3 <- R3 R4
} else {	ADD R3 <- R3 R4	SUBn R3 <- R3 R5
x -= z;	JMP after	
}	else:	
	SUB R3 <- R3 R5	
	after:	

# Predication Performance

---

- Predication overhead is additional insns
  - Sometimes overhead is zero
    - for if-then statement where condition is true
  - Most of the times it isn't
    - if-then-else statement, only one of the paths is useful
- Calculation for a given branch, predicate (vs speculate) if...
  - Average number of additional insns > overall mis-prediction penalty
  - For an individual branch
    - Mis-prediction penalty in a 5-stage scalar pipeline = 2
    - Mis-prediction rate is <50%, and often <20%
    - Overall mis-prediction penalty <1 and often <0.4
  - So when is predication ever worth it?


# Predication Performance

---

- What does predication actually accomplish?
  - In a scalar 5-stage pipeline (penalty = 2): nothing!
  - In a 4-way superscalar 15-stage pipeline (penalty = 60)?
    - Use when mis-predictions >10% and insn overhead <6
  - In a 4-way out-of-order superscalar (penalty  $\sim 150$ )
    - potentially useful in more situations
  - typically only desirable for branches that mis-predict frequently
- Other predication advantages
  - Low-power: eliminates the need for a large branch predictor
  - Real-time: predicated code performs consistently
- Predication disadvantages
  - wasted time/energy compared to correct prediction
  - doesn't nest well

# Summary

---



- Control hazards
  - Branch target prediction
  - Branch direction prediction