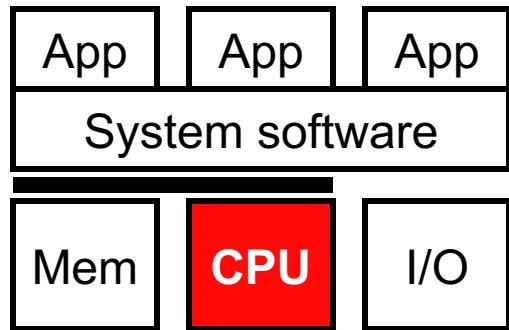


CIS 5710: Computer Organization & Design

Unit 11: Static & Dynamic Scheduling

Slides originally developed by
Drew Hilton, Amir Roth, Milo Martin and Joe Devietti
at University of Pennsylvania

This Unit: Static & Dynamic Scheduling



- Code scheduling
 - To reduce pipeline stalls
 - To increase ILP (insn level parallelism)
- Static scheduling by the compiler
 - Approach & limitations
- Dynamic scheduling in hardware
 - Register renaming
 - Instruction selection
 - Handling memory operations

Code Scheduling & Limitations

Code Scheduling

- Scheduling: act of finding independent instructions
 - “Static” done at compile time by the compiler (software)
 - “Dynamic” done at runtime by the processor (hardware)
- Why schedule code?
 - Scalar pipelines: fill in load-to-use delay slots to improve CPI
 - Superscalar: place independent instructions together
 - As above, load-to-use delay slots
 - Allow multiple-issue decode logic to let them execute at the same time

Compiler Scheduling

- Compiler can schedule (move) instructions to reduce stalls
 - Basic pipeline scheduling**: eliminate back-to-back load-use pairs
 - Example code sequence: $a = b + c; d = f - e;$
 - sp stack pointer, $sp+0$ is “ a ”, $sp+4$ is “ b ”, etc...

Before

```
ld [sp+4]→r2
ld [sp+8]→r3
add r2,r3→r1 //stall
st r1→[sp+0]
ld [sp+16]→r5
ld [sp+20]→r6
sub r6,r5→r4 //stall
st r4→[sp+12]
```

After

```
ld [sp+4]→r2
ld [sp+8]→r3
ld [sp+16]→r5
add r2,r3→r1 //no stall
ld [sp+20]→r6
st r1→[sp+0]
sub r6,r5→r4 //no stall
st r4→[sp+12]
```

Compiler Scheduling Requires

- **Large scheduling scope**

- Independent instruction to put between load-use pairs
 - + Original example: large scope, two independent computations
 - This example: small scope, one computation

Before

```
ld [sp+4]→r2
ld [sp+8]→r3
add r2,r3→r1 //stall
st r1→[sp+0]
```

After (same!)

```
ld [sp+4]→r2
ld [sp+8]→r3
add r2,r3→r1 //stall
st r1→[sp+0]
```

- Compiler can create **larger** scheduling scopes
 - For example: loop unrolling & function inlining

Scheduling Scope Limited by Branches

r1 and r2 are inputs

loop:

 jz r1, not_found

 ld [r1+0] → r3

 sub r2, r3 → r4

 jz r4, found

 ld [r1+4] → r1

 jmp loop

Aside: what does this code do?



Legal to move load up past branch?

Compiler Scheduling Requires

- **Enough registers**

- To hold additional “live” values
- Example code contains 7 different values (including `sp`)
- Before: max 3 values live at any time → 3 registers enough
- After: max 4 values live → 3 registers not enough

Original

```
ld [sp+4]→r2
ld [sp+8]→r1
add r1,r2→r1 //stall
st r1→[sp+0]
ld [sp+16]→r2
ld [sp+20]→r1
sub r2,r1→r1 //stall
st r1→[sp+12]
```

Wrong!

```
ld [sp+4]→r2
ld [sp+8]→r1
ld [sp+16]→r2
add r1,r2→r1 // wrong r2
ld [sp+20]→r1
st r1→[sp+0] // wrong r1
sub r2,r1→r1
st r1→[sp+12]
```

Compiler Scheduling Requires

- **Alias analysis**

- Ability to tell whether load/store reference same memory locations
 - Effectively, whether load/store can be rearranged
- Previous example: easy, loads/stores use same base register (**sp**)
- New example: can compiler tell that **r8 != r9**?
- Must be **conservative**

Before

```
ld [r9+4]→r2
ld [r9+8]→r3
add r3,r2→r1 //stall
st r1→[r9+0]
ld [r8+0]→r5
ld [r8+4]→r6
sub r5,r6→r4 //stall
st r4→[r8+8]
```

Wrong(?)

```
ld [r9+4]→r2
ld [r9+8]→r3
ld [r8+0]→r5 //does r8==r9?
add r3,r2→r1
ld [r8+4]→r6 //does r8+4==r9?
st r1→[r9+0]
sub r5,r6→r4
st r4→[r8+8]
```

A Good Case: Static Scheduling of SAXPY

- **SAXPY** (Single-precision A X Plus Y)
 - Linear algebra routine (used in solving systems of equations)

```
for (i=0;i<N;i++)  
    Z[i]=(A*X[i])+Y[i];
```

```
0: ldf [X+r1]→f1      // loop  
1: mulf f0,f1→f2      // A in f0  
2: ldf [Y+r1]→f3      // X,Y,Z are constant addresses  
3: addf f2,f3→f4  
4: stf f4→[Z+r1]  
5: addi r1,4→r1        // i in r1  
6: blt r1,r2,0          // N*4 in r2
```

- Static scheduling works great for SAXPY
 - All loop iterations independent
 - Use loop unrolling to increase scheduling scope
 - Aliasing analysis is tractable (just ensure X, Y, Z are independent)
 - Still limited by number of registers

Unrolling & Scheduling SAXPY

- Fuse two (in general, K) iterations of loop
 - Fuse loop control: induction variable (*i*) increment + branch
 - Adjust register names & induction uses (constants → constants+4)
 - Reorder operations to reduce stalls

```
ldf [X+r1]→f1  
mulf f0,f1→f2  
ldf [Y+r1]→f3  
addf f2,f3→f4  
stf f4→[Z+r1]  
addi r1,4→r1  
blt r1,r2,0
```



```
ldf [X+r1]→f1  
mulf f0,f1→f2  
ldf [Y+r1]→f3  
addf f2,f3→f4  
stf f4→[Z+r1]
```



```
ldf [X+r1]→f1  
mulf f0,f1→f2  
ldf [Y+r1]→f3  
addf f2,f3→f4  
stf f4→[Z+r1]  
addi r1,4→r1  
blt r1,r2,0
```

```
ldf [X+r1]→f1  
ldf [X+r2+4]→f5  
mulf f0,f1→f2  
mulf f0,f5→f6  
ldf [Y+r1]→f3  
ldf [Y+r1+4]→f7  
addf f2,f3→f4  
addf f6,f7→f8  
stf f4→[Z+r1]  
stf f8→[Z+r1+4]  
addi r1,8→r1  
blt r1,r2,0
```

Compiler Scheduling Limitations

- Scheduling scope
 - Example: can't generally move memory operations past branches
- Limited number of registers (set by ISA)
- Inexact “memory aliasing” information
 - Often prevents reordering of loads above stores by compiler
- Caches misses (or any runtime event) confound scheduling
 - How can the compiler know which loads will miss vs hit?
 - Can impact the compiler’s scheduling decisions

Dynamic (Hardware) Scheduling

Can Hardware Overcome These Limits?

- **Dynamically-scheduled processors**
 - Also called “out-of-order” processors
 - Hardware re-schedules insns...
 - ...within a sliding window of VonNeumann insns
 - As with pipelining and superscalar, ISA unchanged
 - Same hardware/software interface, appearance of in-order
- Increases scheduling scope
 - Does loop unrolling transparently!
 - Uses branch prediction to “unroll” branches
- Examples:
 - Pentium Pro/II/III (3-wide), Core 2 (4-wide), Alpha 21264 (4-wide), MIPS R10000 (4-wide), Power5 (5-wide)

Example: In-Order Limitations #1

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [r1] → r2	F	D	X	M ₁	M ₂	W							
add r2 + r3 → r4	F	D	d*	d*	d*	X	M ₁	M ₂	W				
xor r4 ^ r5 → r6		F	D	d*	d*	d*	X	M ₁	M ₂	W			
ld [r7] → r4		F	D	p*	p*	p*	X	M ₁	M ₂	W			

- In-order pipeline, three-cycle load-use penalty
 - 2-wide
- Why not the following?

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [r1] → r2	F	D	X	M ₁	M ₂	W							
add r2 + r3 → r4	F	D	d*	d*	d*	X	M ₁	M ₂	W				
xor r4 ^ r5 → r6		F	D	d*	d*	d*	X	M ₁	M ₂	W			
ld [r7] → r4		F	D	X	M₁	M₂	W						

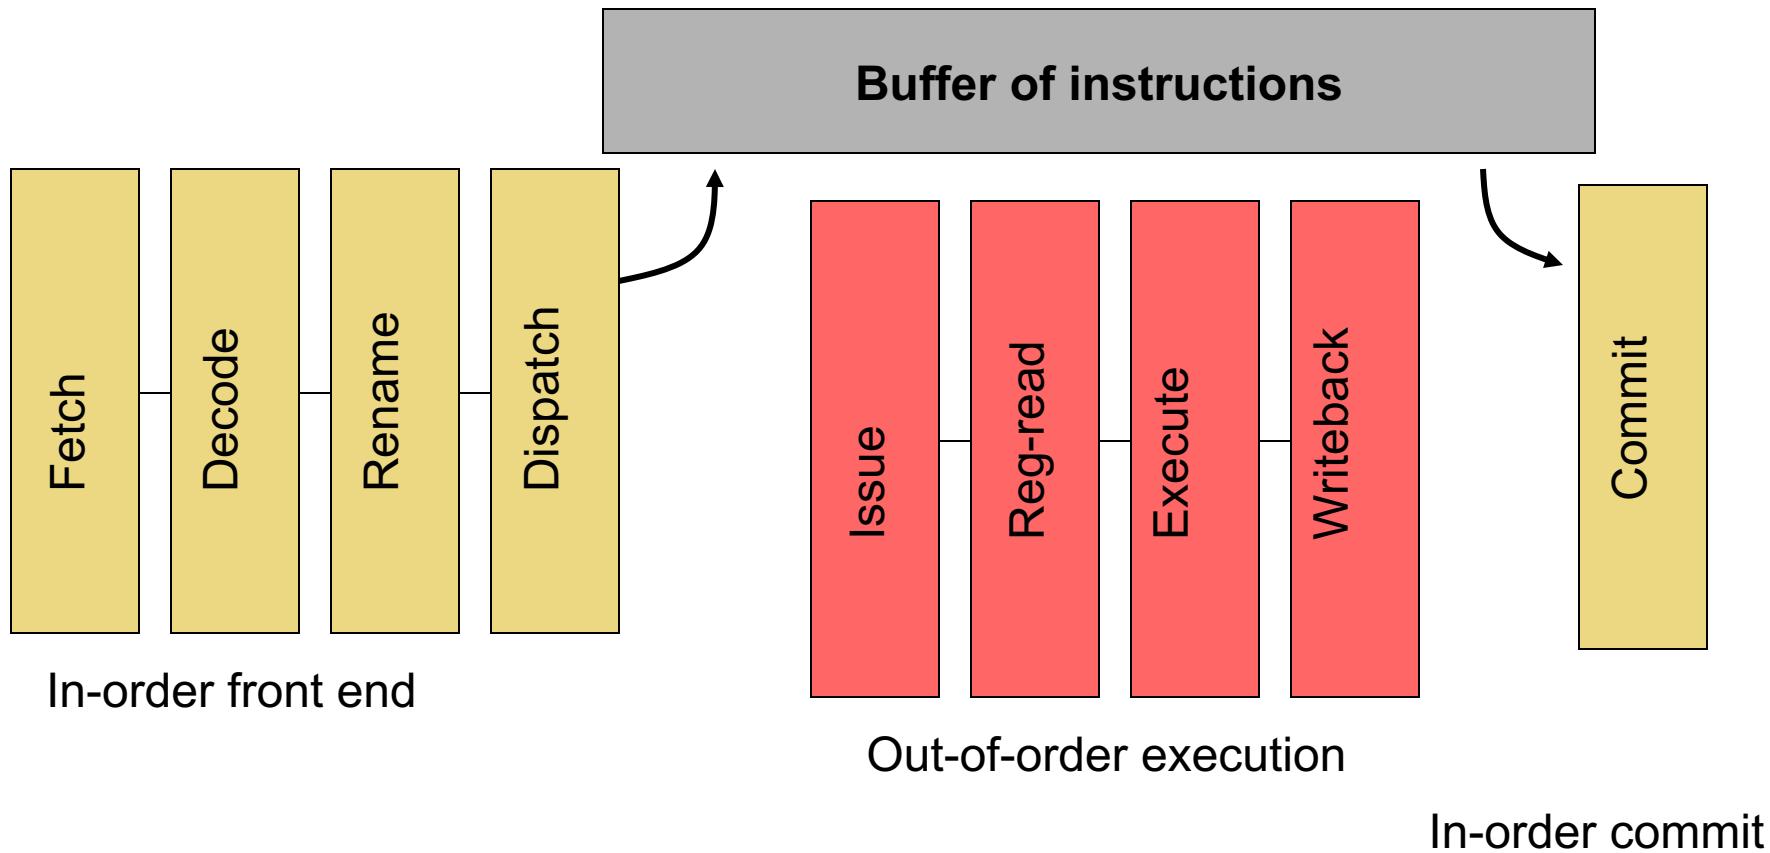
Example: In-Order Limitations #2

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [p1] → p2	F	D	X	M ₁	M ₂	W							
add p2 + p3 → p4	F	D	d*	d*	d*	X	M ₁	M ₂	W				
xor p4 ^ p5 → p6		F	D	d*	d*	d*	X	M ₁	M ₂	W			
ld [p7] → p8		F	D	p*	p*	p*	X	M ₁	M ₂	W			

- In-order pipeline, three-cycle load-use penalty
 - 2-wide
- Why not the following:

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [p1] → p2	F	D	X	M ₁	M ₂	W							
add p2 + p3 → p4	F	D	d*	d*	d*	X	M ₁	M ₂	W				
xor p4 ^ p5 → p6		F	D	d*	d*	d*	X	M ₁	M ₂	W			
ld [p7] → p8		F	D	X	M₁	M₂	W						

Out-of-Order Pipeline

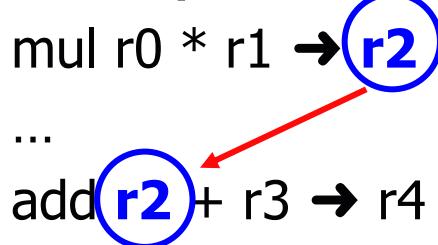


Out-of-Order Execution

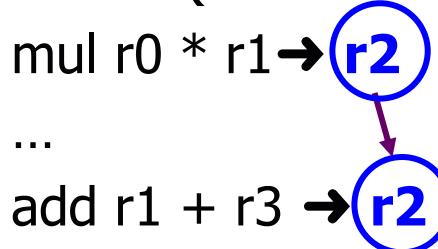
- Also called “Dynamic scheduling”
 - Done by the hardware on-the-fly during execution
- Looks at a “window” of instructions waiting to execute
 - Each cycle, picks the next ready instruction(s)
- Two steps to enable out-of-order execution:
 - Step #1: Register renaming – to avoid “false” dependencies
 - Step #2: Dynamically schedule – to enforce “true” dependencies
- Key to understanding out-of-order execution:
 - **Data dependencies**

Dependence types

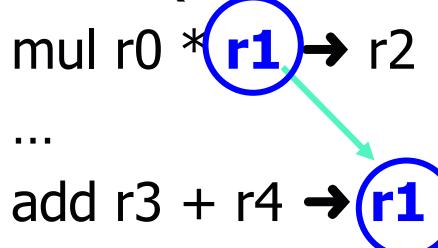
- **RAW** (Read After Write) = “true dependence” (true)



- **WAW** (Write After Write) = “output dependence” (false)



- **WAR** (Write After Read) = “anti-dependence” (false)



- WAW & WAR are “false”, Can be **totally eliminated** by “renaming”

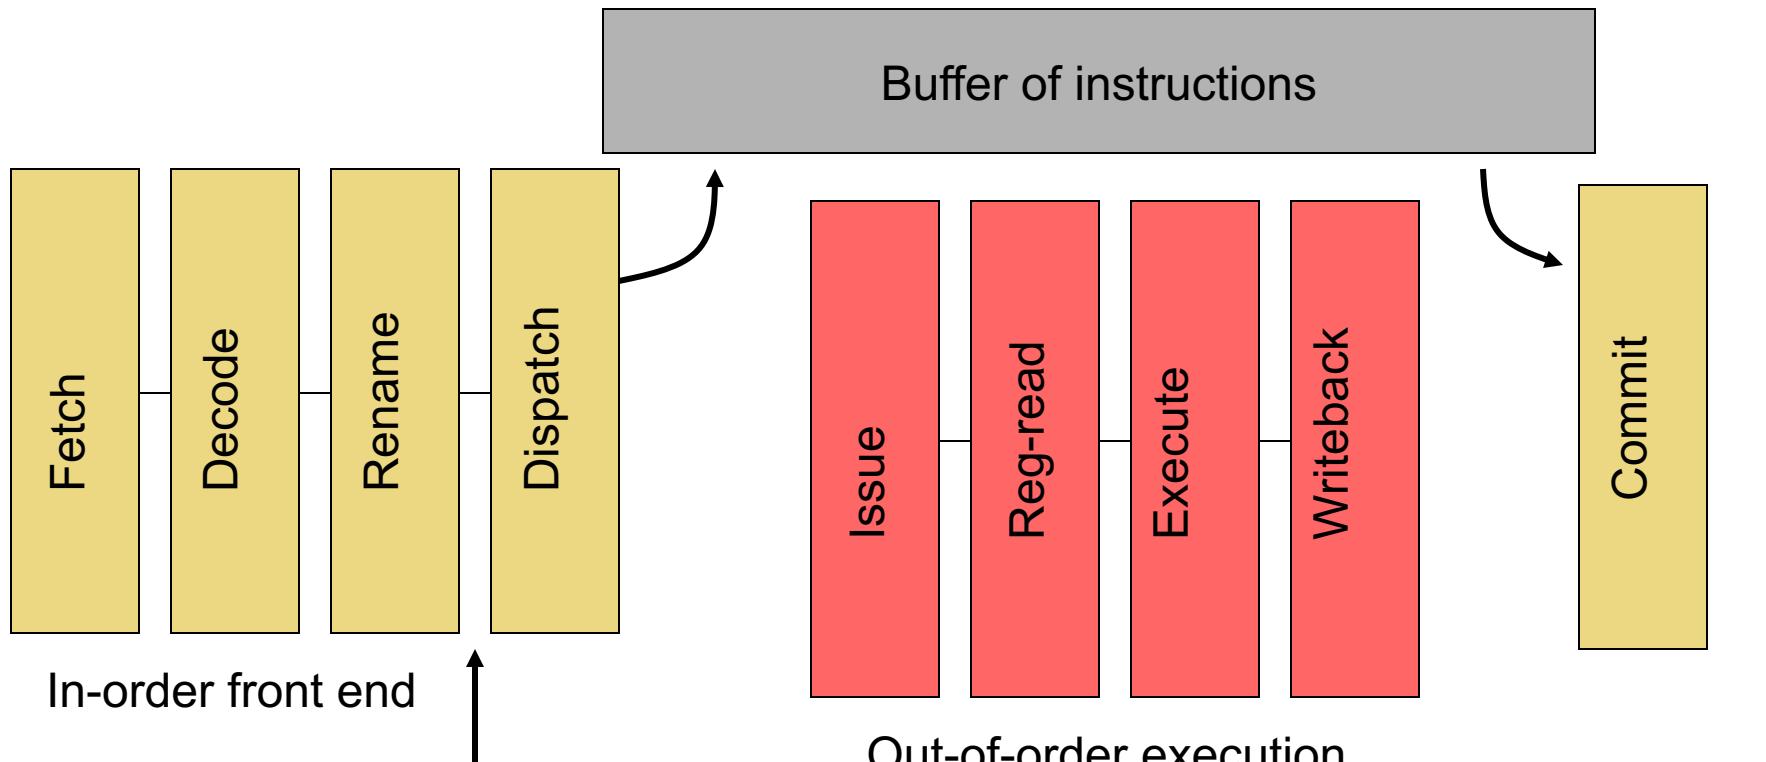
Register Renaming Algorithm

- Two key data structures:
 - `maptable[architectural_reg] → physical_reg`
 - Free list: allocate (new) & free registers (implemented as a queue)
- Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]
insn.phys_input2 = maptable[insn.arch_input2]
insn.old_phys_output = maptable[insn.arch_output]
new_reg = new_phys_reg()
maptable[insn.arch_output] = new_reg
insn.phys_output = new_reg
```

- At “commit”
 - Once all older instructions have committed, free register
`free_phys_reg(insn.old_phys_output)`

Out-of-order Pipeline



Have unique register names

Now put into out-of-order execution structures

In-order commit

Register Renaming

Register Renaming Algorithm (Simplified)

- Two key data structures:
 - `maptable[architectural_reg] → physical_reg`
 - Free list: allocate (new) & free registers (implemented as a queue)
 - ignore freeing of registers for now
- Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]  
insn.phys_input2 = maptable[insn.arch_input2]
```

```
new_reg = new_phys_reg()  
maptable[insn.arch_output] = new_reg  
insn.phys_output = new_reg
```

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 →

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

Renaming example

$\text{xor } r1 \wedge r2 \rightarrow r3$ \longrightarrow $\text{xor } p1 \wedge p2 \rightarrow p6$
 $\text{add } r3 + r4 \rightarrow r4$
 $\text{sub } r5 - r2 \rightarrow r3$
 $\text{addi } r3 + 1 \rightarrow r1$

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

Renaming example

$\text{xor } r1 \wedge r2 \rightarrow r3$
 $\text{add } r3 + r4 \rightarrow r4$
 $\text{sub } r5 - r2 \rightarrow r3$
 $\text{addi } r3 + 1 \rightarrow r1$

\longrightarrow $\text{xor } p1 \wedge p2 \rightarrow p6$

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

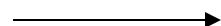
Map table

p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add **r3** + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add **p6** + **p4** →

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → **p7**

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

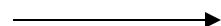
Map table

p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → p7

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 →

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → **p8**

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

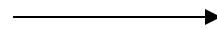
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi **r3** + 1 → r1

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi **p8** + 1 →



r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

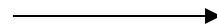
p9
p10

Free-list

Renaming example

$\text{xor } r1 \wedge r2 \rightarrow r3$
 $\text{add } r3 + r4 \rightarrow r4$
 $\text{sub } r5 - r2 \rightarrow r3$
 $\text{addi } r3 + 1 \rightarrow r1$

$\text{xor } p1 \wedge p2 \rightarrow p6$
 $\text{add } p6 + p4 \rightarrow p7$
 $\text{sub } p5 - p2 \rightarrow p8$
 $\text{addi } p8 + 1 \rightarrow \mathbf{p9}$



r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

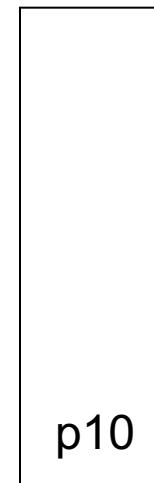
Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

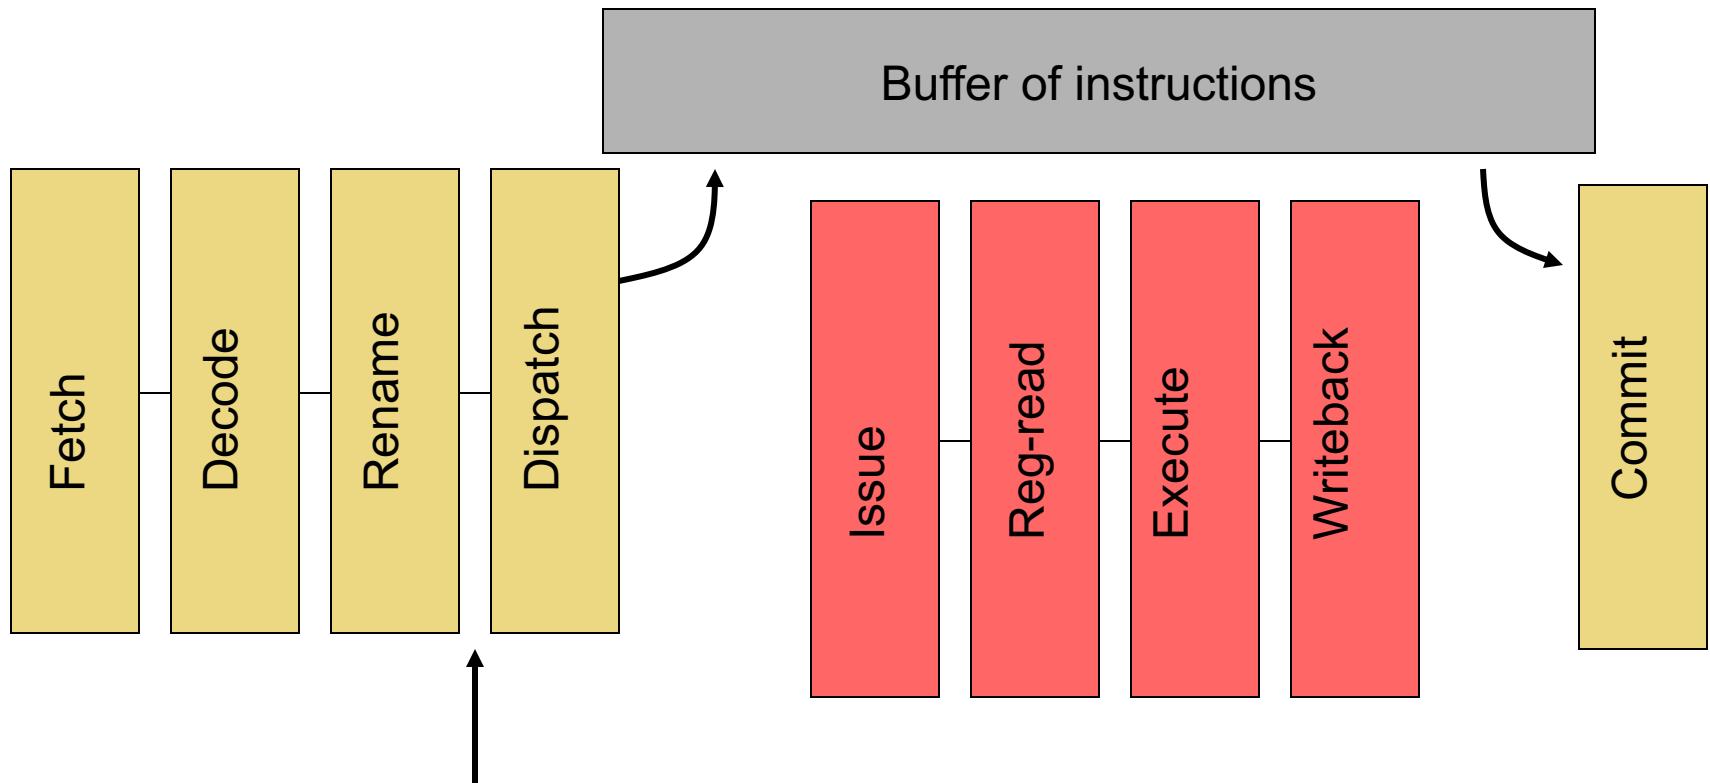
Register Renaming Overview

- To eliminate register conflicts/hazards
- “Architected” vs “Physical” registers – level of indirection
 - Names: r_1, r_2, r_3
 - Locations: $p_1, p_2, p_3, p_4, p_5, p_6, p_7$
 - Original mapping: $r_1 \rightarrow p_1, r_2 \rightarrow p_2, r_3 \rightarrow p_3, p_4-p_7$ are “available”

	MapTable	FreeList	Original insns	Renamed insns																											
Time	<table border="1"><tr><td>r_1</td><td>r_2</td><td>r_3</td></tr><tr><td>p_1</td><td>p_2</td><td>p_3</td></tr><tr><td>p_4</td><td>p_2</td><td>p_3</td></tr><tr><td>p_4</td><td>p_2</td><td>p_5</td></tr><tr><td>p_4</td><td>p_2</td><td>p_6</td></tr></table>	r_1	r_2	r_3	p_1	p_2	p_3	p_4	p_2	p_3	p_4	p_2	p_5	p_4	p_2	p_6	<table border="1"><tr><td>p_4, p_5, p_6, p_7</td></tr><tr><td>p_5, p_6, p_7</td></tr><tr><td>p_6, p_7</td></tr><tr><td>p_7</td></tr></table>	p_4, p_5, p_6, p_7	p_5, p_6, p_7	p_6, p_7	p_7	<table><tr><td>add $r_2, r_3 \rightarrow r_1$</td></tr><tr><td>sub $r_2, r_1 \rightarrow r_3$</td></tr><tr><td>mul $r_2, r_3 \rightarrow r_3$</td></tr><tr><td>div $r_1, 4 \rightarrow r_1$</td></tr></table>	add $r_2, r_3 \rightarrow r_1$	sub $r_2, r_1 \rightarrow r_3$	mul $r_2, r_3 \rightarrow r_3$	div $r_1, 4 \rightarrow r_1$	<table><tr><td>add $p_2, p_3 \rightarrow p_4$</td></tr><tr><td>sub $p_2, p_4 \rightarrow p_5$</td></tr><tr><td>mul $p_2, p_5 \rightarrow p_6$</td></tr><tr><td>div $p_4, 4 \rightarrow p_7$</td></tr></table>	add $p_2, p_3 \rightarrow p_4$	sub $p_2, p_4 \rightarrow p_5$	mul $p_2, p_5 \rightarrow p_6$	div $p_4, 4 \rightarrow p_7$
r_1	r_2	r_3																													
p_1	p_2	p_3																													
p_4	p_2	p_3																													
p_4	p_2	p_5																													
p_4	p_2	p_6																													
p_4, p_5, p_6, p_7																															
p_5, p_6, p_7																															
p_6, p_7																															
p_7																															
add $r_2, r_3 \rightarrow r_1$																															
sub $r_2, r_1 \rightarrow r_3$																															
mul $r_2, r_3 \rightarrow r_3$																															
div $r_1, 4 \rightarrow r_1$																															
add $p_2, p_3 \rightarrow p_4$																															
sub $p_2, p_4 \rightarrow p_5$																															
mul $p_2, p_5 \rightarrow p_6$																															
div $p_4, 4 \rightarrow p_7$																															

- Renaming – conceptually write each register once
 - + Removes **false** dependences
 - + Leaves **true** dependences intact!
- When to reuse a physical register? After overwriting insn done

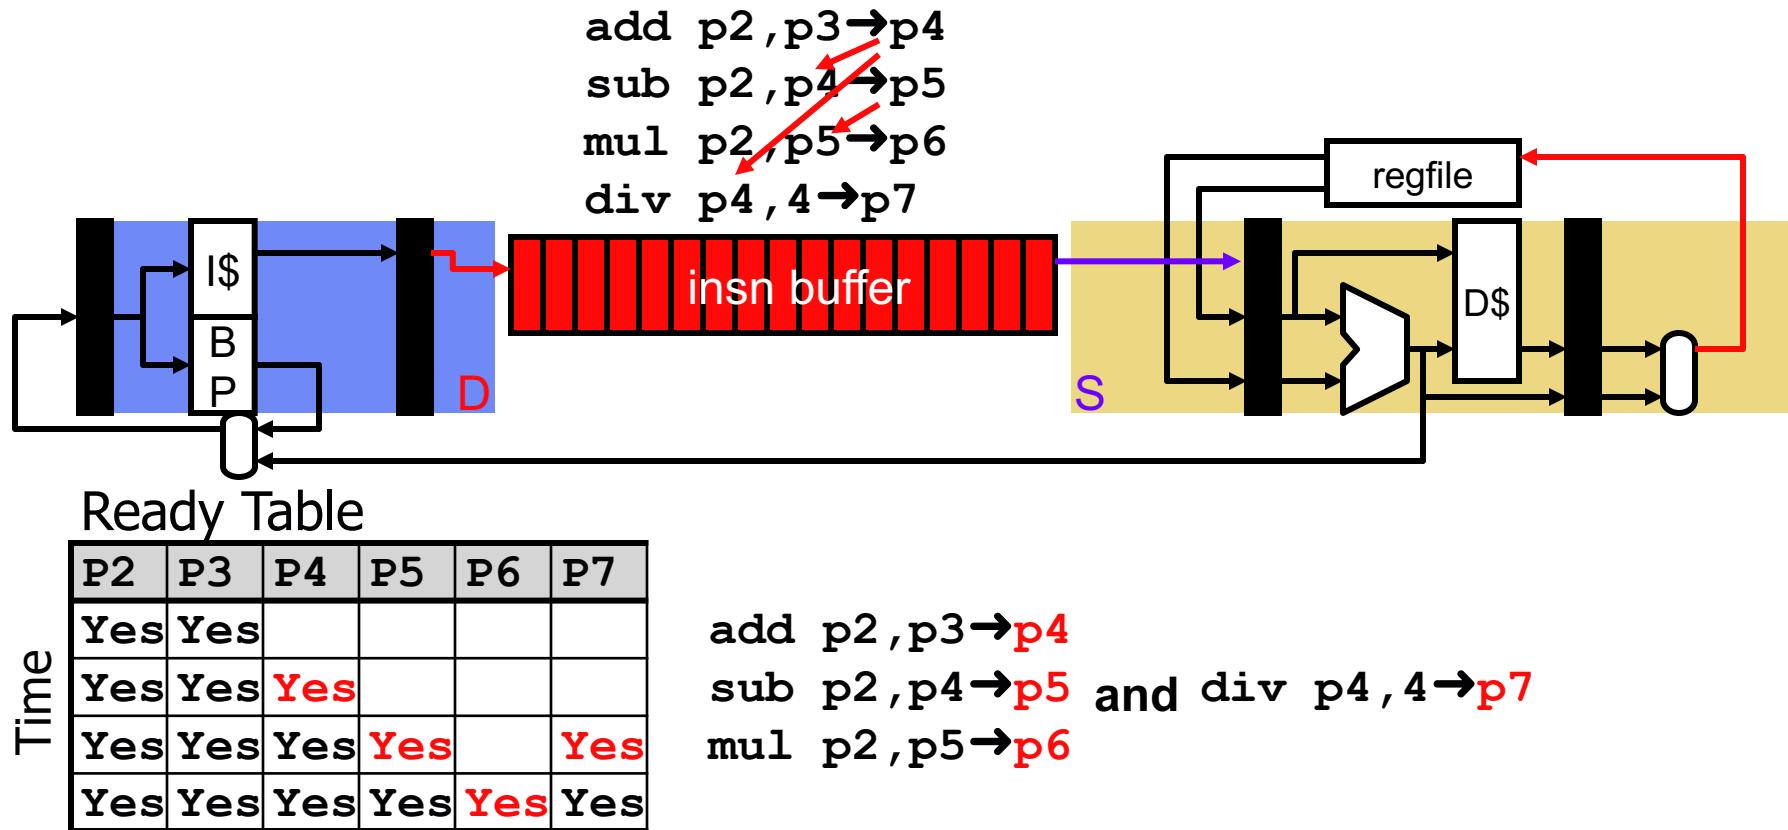
Out-of-order Pipeline



Have unique register names
Now put into out-of-order execution structures

Dynamic Scheduling Mechanisms

Dynamic Scheduling Overview



- Instructions fetch/decoded/renamed into *Instruction Buffer*
 - Also called “instruction window” or “instruction scheduler”
- Instructions (conceptually) check ready bits every cycle
 - Execute oldest “ready” instruction, set output as “ready”

Dynamic Scheduling/Issue Algorithm

- Data structures:
 - Ready table[phys_reg] → yes/no (part of “issue queue”)
- Algorithm at “issue” stage (prior to read registers):

foreach instruction:

```
if table[insn.phys_input1] == ready &&
    table[insn.phys_input2] == ready then
        insn is “ready”
```

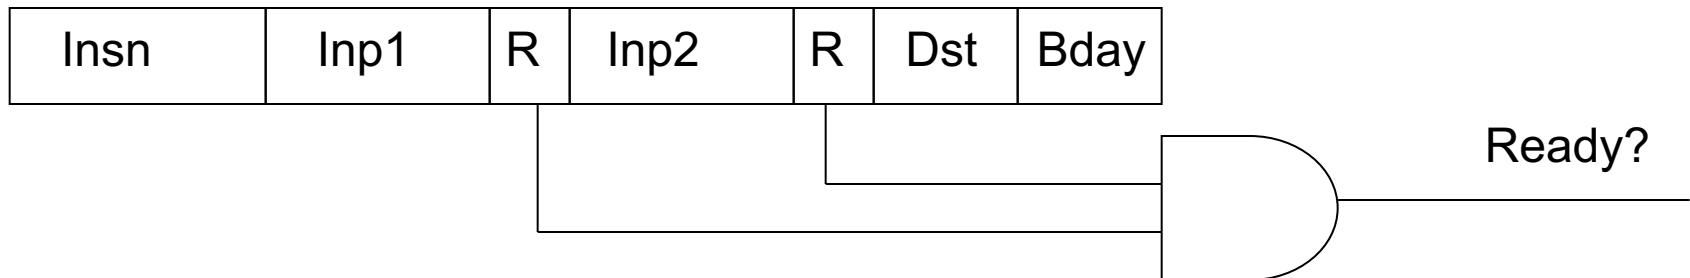
select the oldest “ready” instruction

```
table[insn.phys_output] = ready
```

- Multiple-cycle instructions? (such as loads)
 - For an insn with latency of N, set “ready” bit N-1 cycles in future

Dispatch

- Put renamed instructions into out-of-order structures
- Re-order buffer (ROB)
 - Holds instructions from Fetch through Commit
- Issue Queue
 - Central piece of scheduling logic
 - Holds instructions from Dispatch through Issue
 - Tracks ready inputs
 - Physical register names + ready bit
 - “AND” the bits to tell if ready



Dispatch Steps

- Allocate Issue Queue (IQ) slot
 - Full? Stall
- Read **ready bits** of inputs
 - 1-bit per physical reg
- Clear **ready bit** of output in table
 - Instruction has not produced value yet
- Write data into Issue Queue (IQ) slot

Dispatch Example

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	y
p8	y
p9	y

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday

Dispatch Example

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	y
p8	y
p9	y

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0

Dispatch Example

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	y
p9	y

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1

Dispatch Example

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	y

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2

Dispatch Example

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

Ready bits

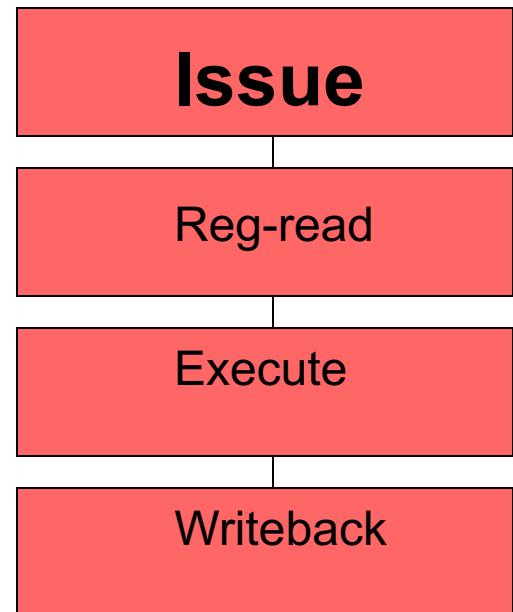
p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	n

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	n	---	y	p9	3

Out-of-order pipeline

- Execution (out-of-order) stages
- **Select** ready instructions
 - Send for execution
- **Wakeup** dependents



Dynamic Scheduling/Issue Algorithm

- Data structures:
 - Ready table[phys_reg] → yes/no (part of issue queue)
- Algorithm at “schedule” stage (prior to read registers):

foreach instruction:

```
if table[insn.phys_input1] == ready &&
    table[insn.phys_input2] == ready then
        insn is "ready"
```

select the oldest “ready” instruction

```
table[insn.phys_output] = ready
```

Issue = Select + Wakeup

- **Select** oldest of “ready” instructions
 - “xor” is the oldest ready instruction below
 - “xor” and “sub” are the two oldest ready instructions below
 - Note: may have resource constraints: i.e. load/store/floating point

Insn	Inp1	R	Inp2	R	Dst	Bday	
xor	p1	y	p2	y	p6	0	Ready!
add	p6	n	p4	y	p7	1	
sub	p5	y	p2	y	p8	2	Ready!
addi	p8	n	---	y	p9	3	

Issue = Select + Wakeup

- Wakeup dependent instructions
 - Search for destination (Dst) in inputs & set “ready” bit
 - Implemented with a special memory array circuit called a Content Addressable Memory (CAM)
 - Also update ready-bit table for future instructions

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	y	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	y	---	y	p9	3

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	n
p8	y
p9	n

- For multi-cycle operations (loads, floating point)
 - Wakeup deferred a few cycles
 - Include checks to avoid structural hazards

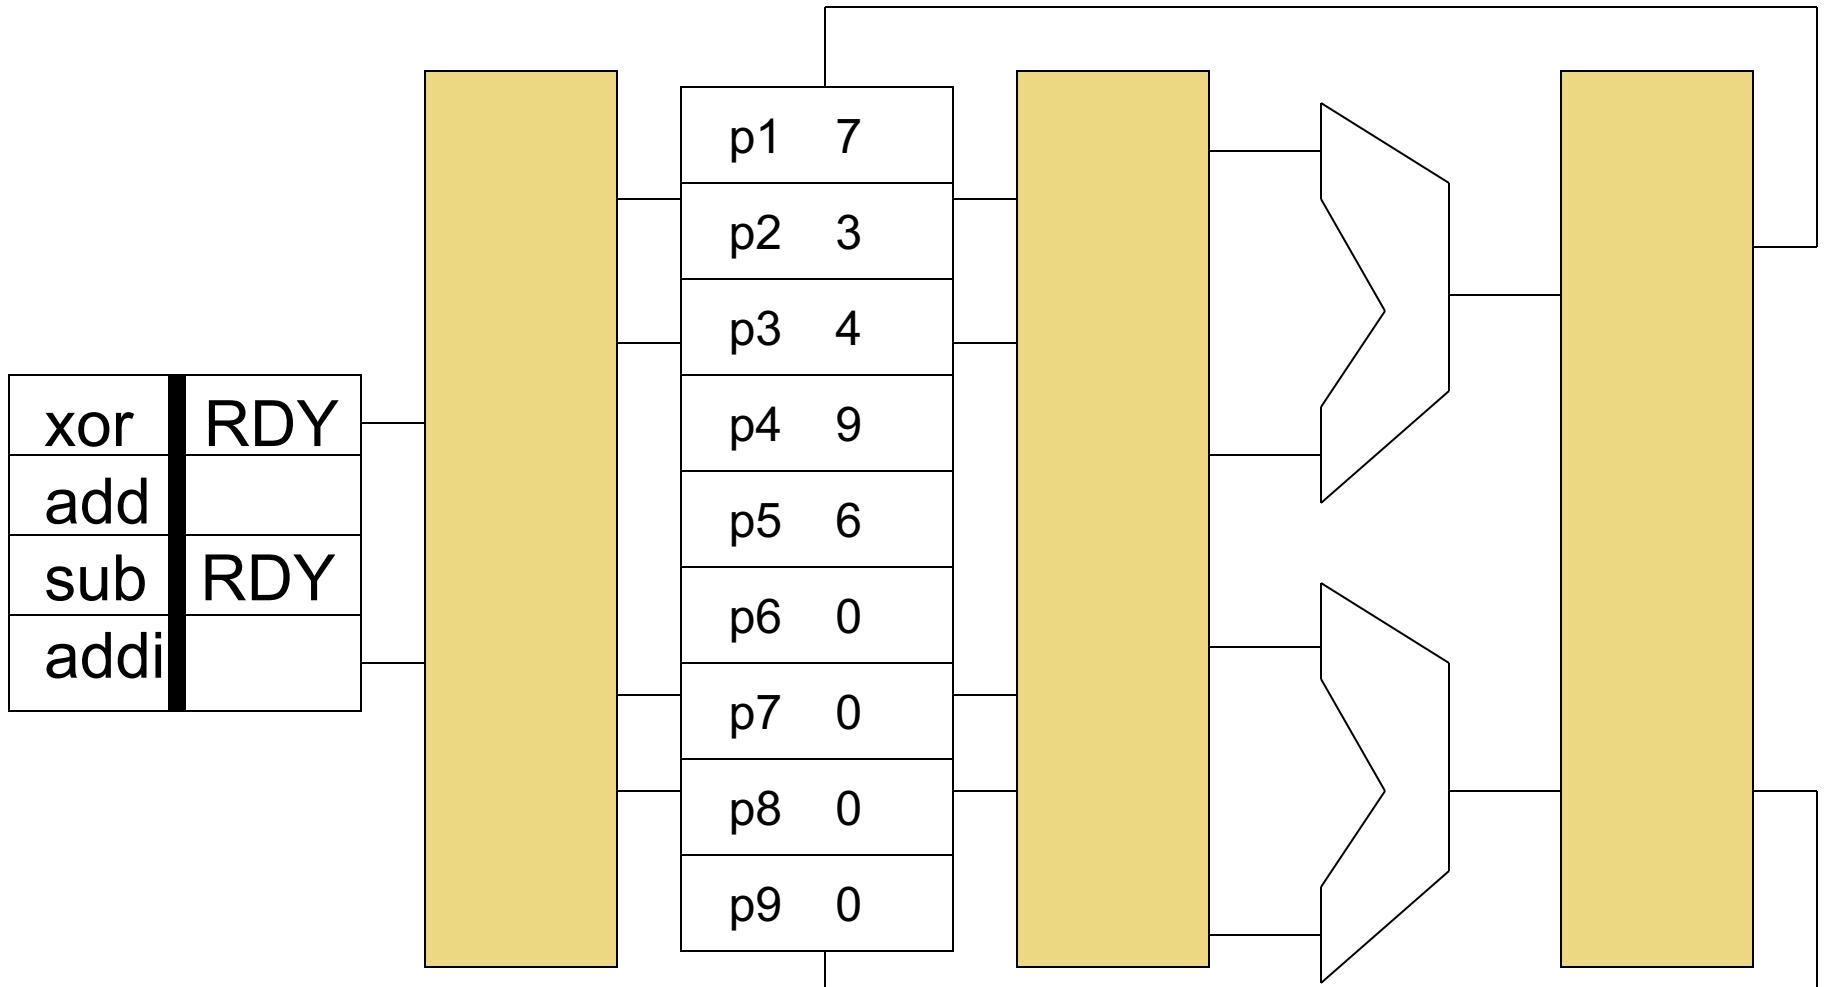
Issue

- **Select/Wakeup** one cycle
- Dependent instructions execute on back-to-back cycles
 - Next cycle: add/addi are ready:

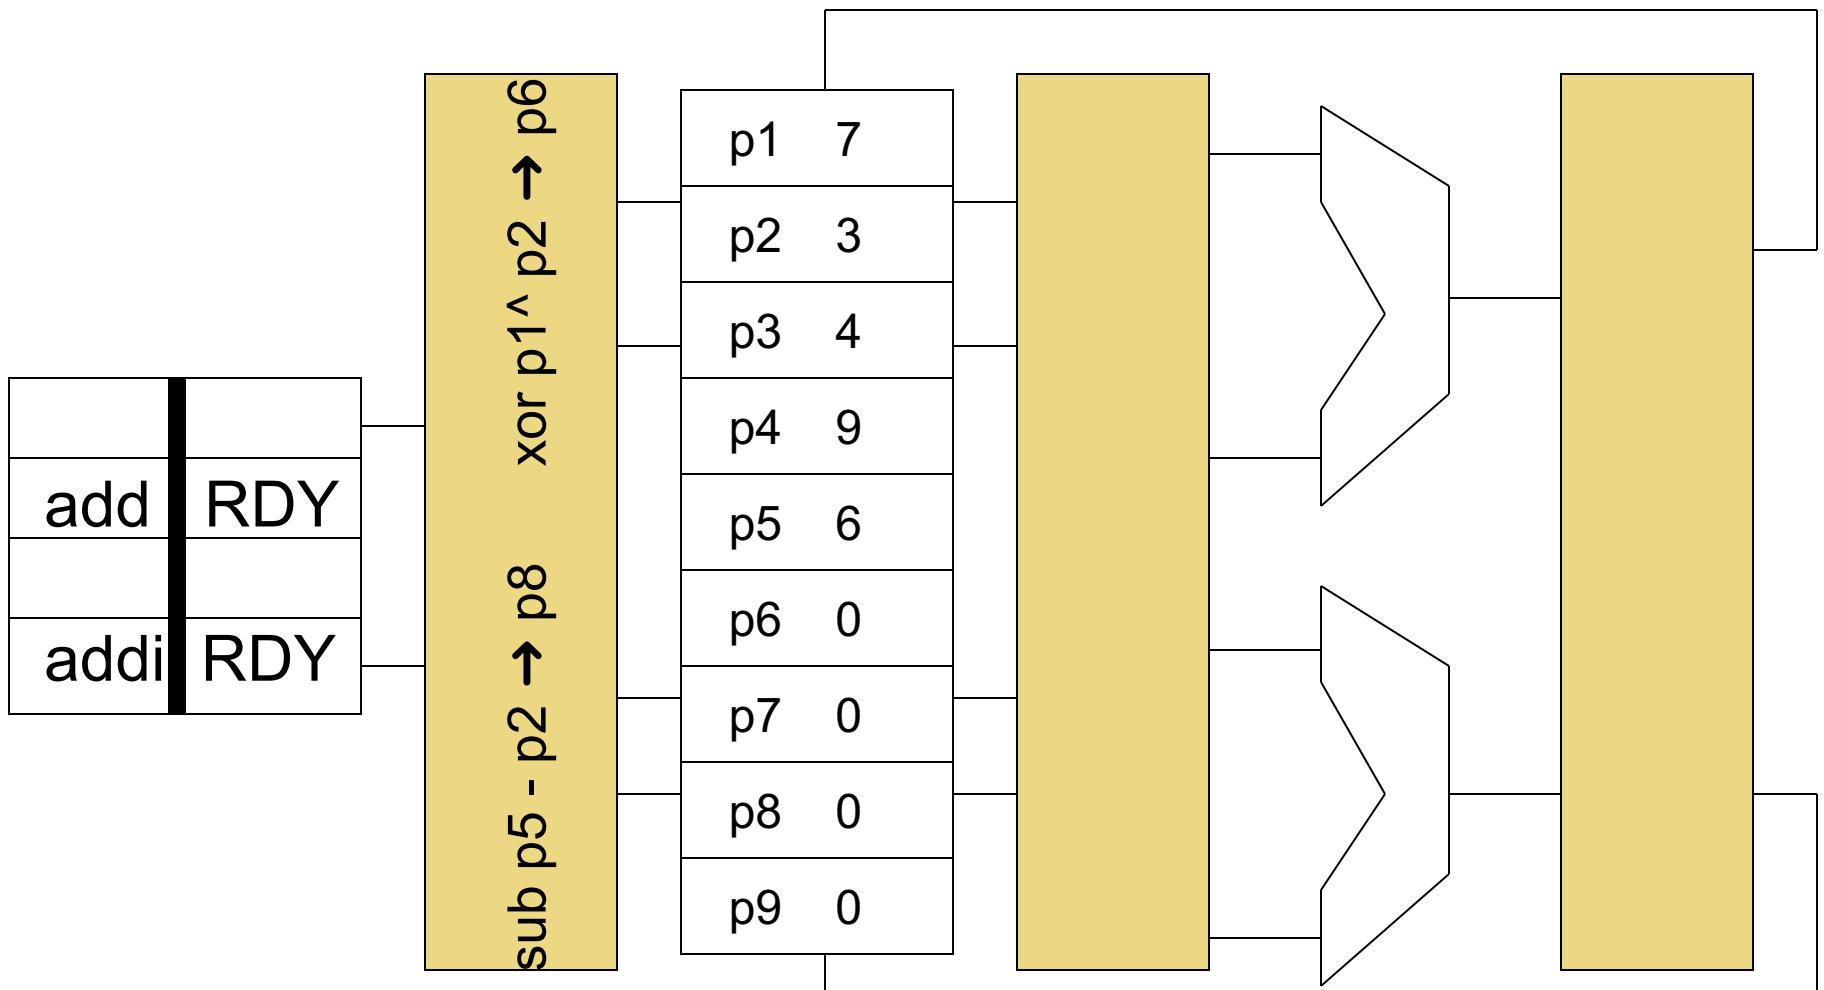
Insn	Inp1	R	Inp2	R	Dst	Bday
add	p6	y	p4	y	p7	1
addi	p8	y	---	y	p9	3

- Issued instructions are removed from issue queue
 - Free up space for subsequent instructions

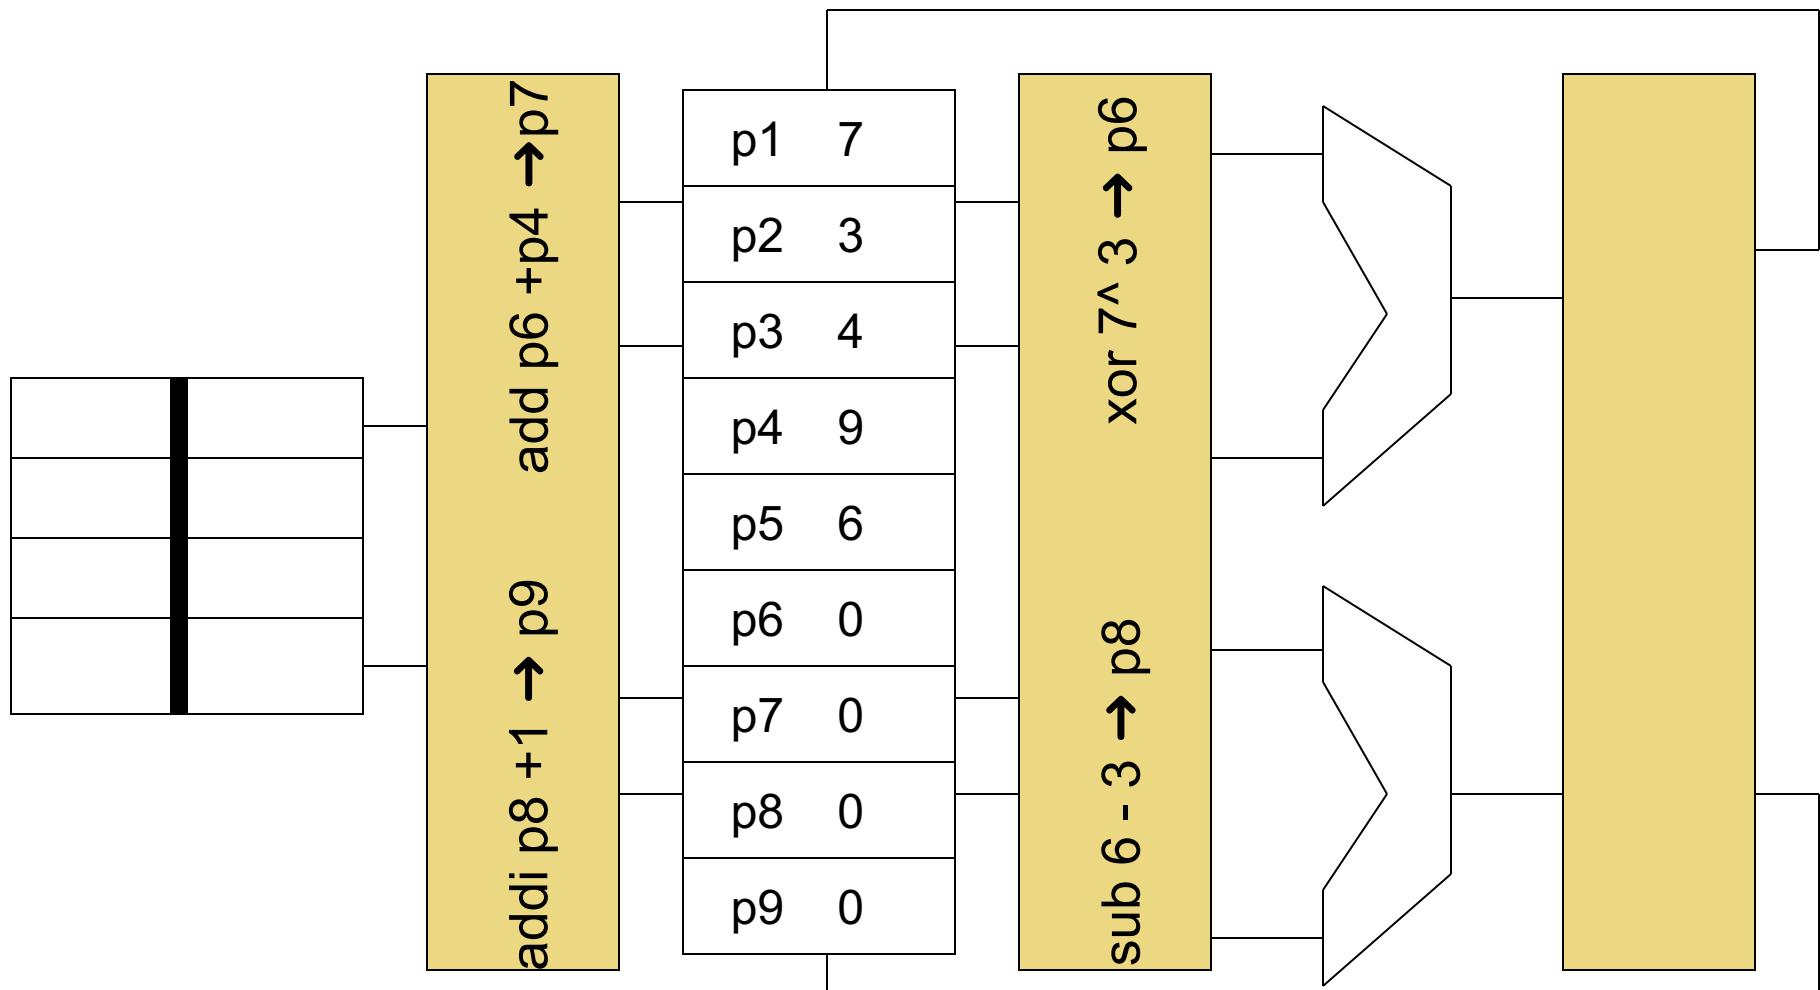
000 execution (2-wide)



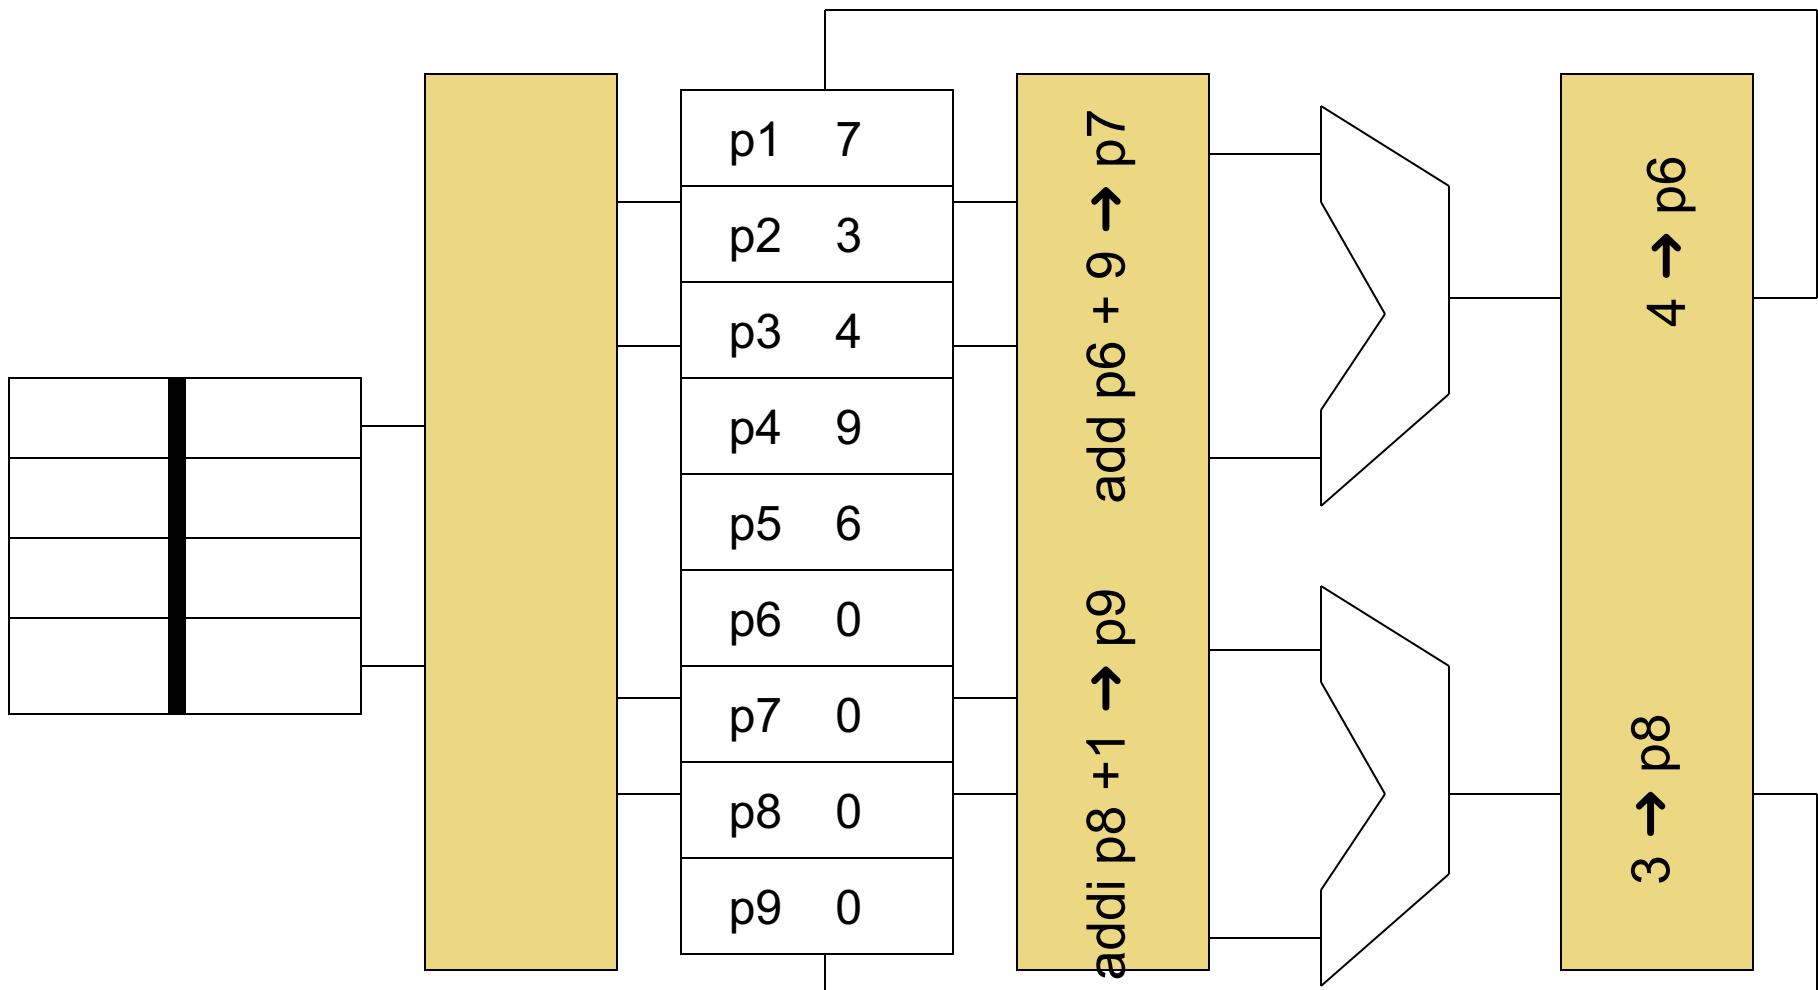
000 execution (2-wide)



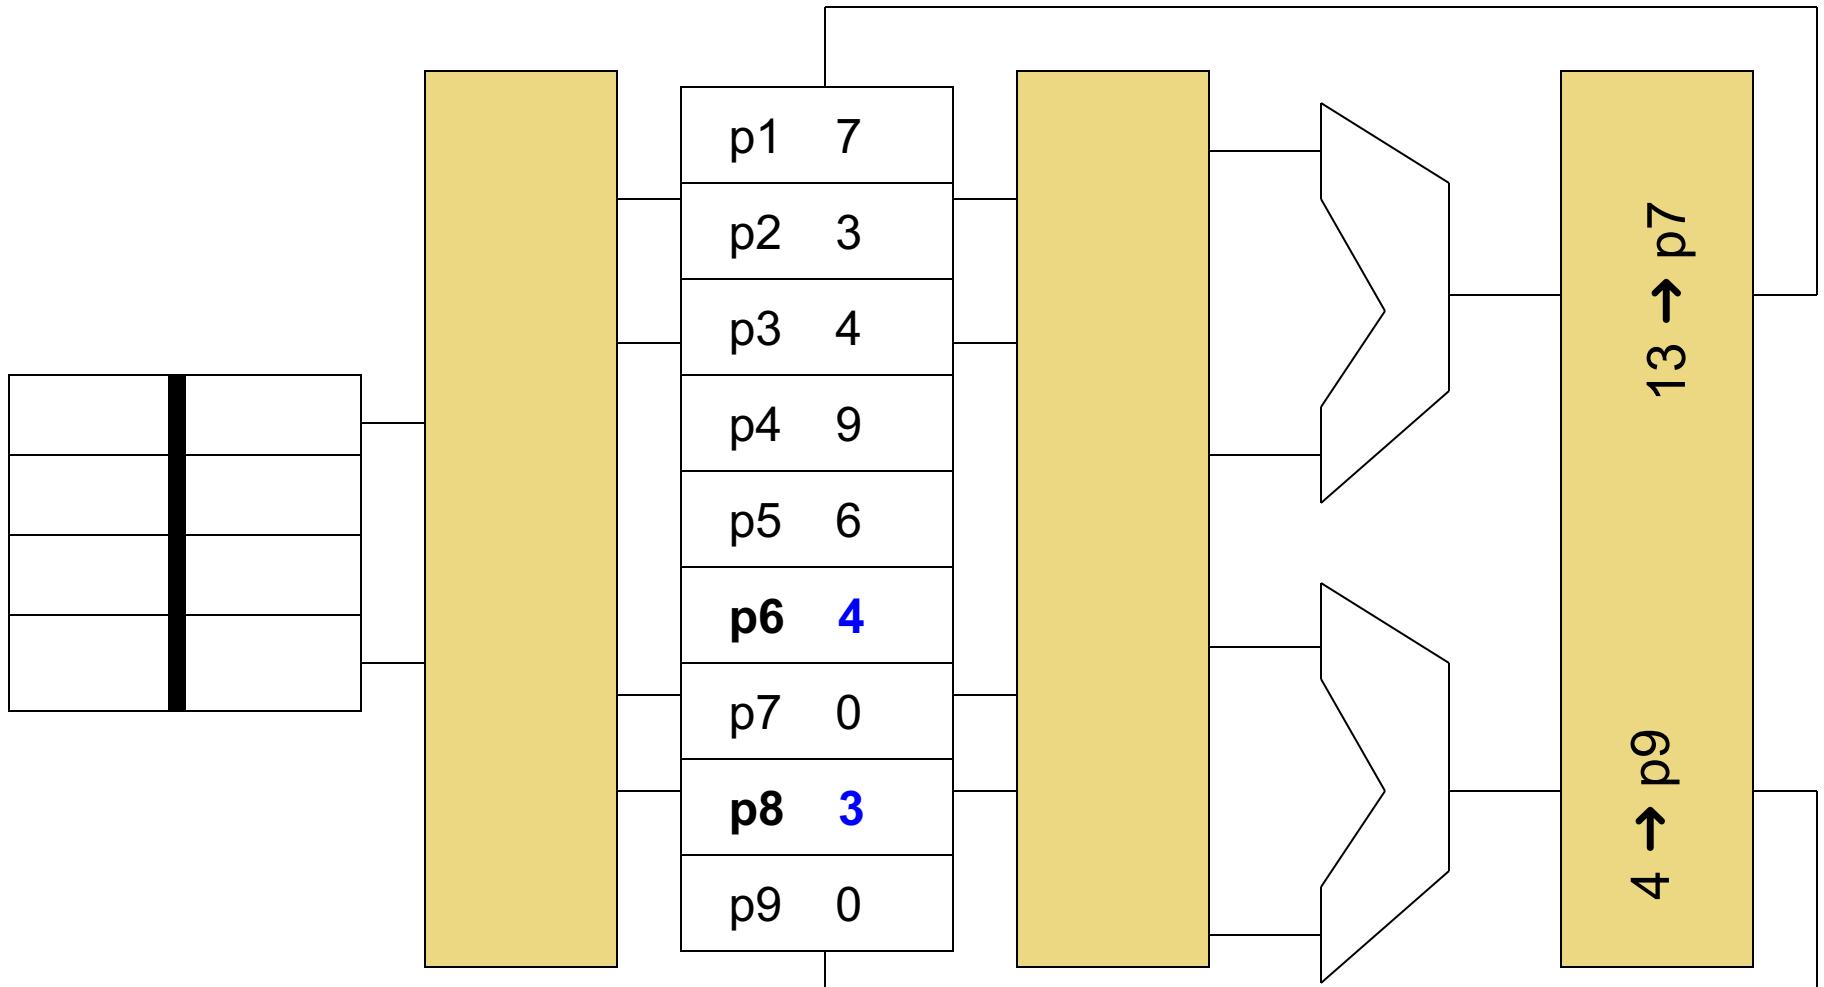
000 execution (2-wide)



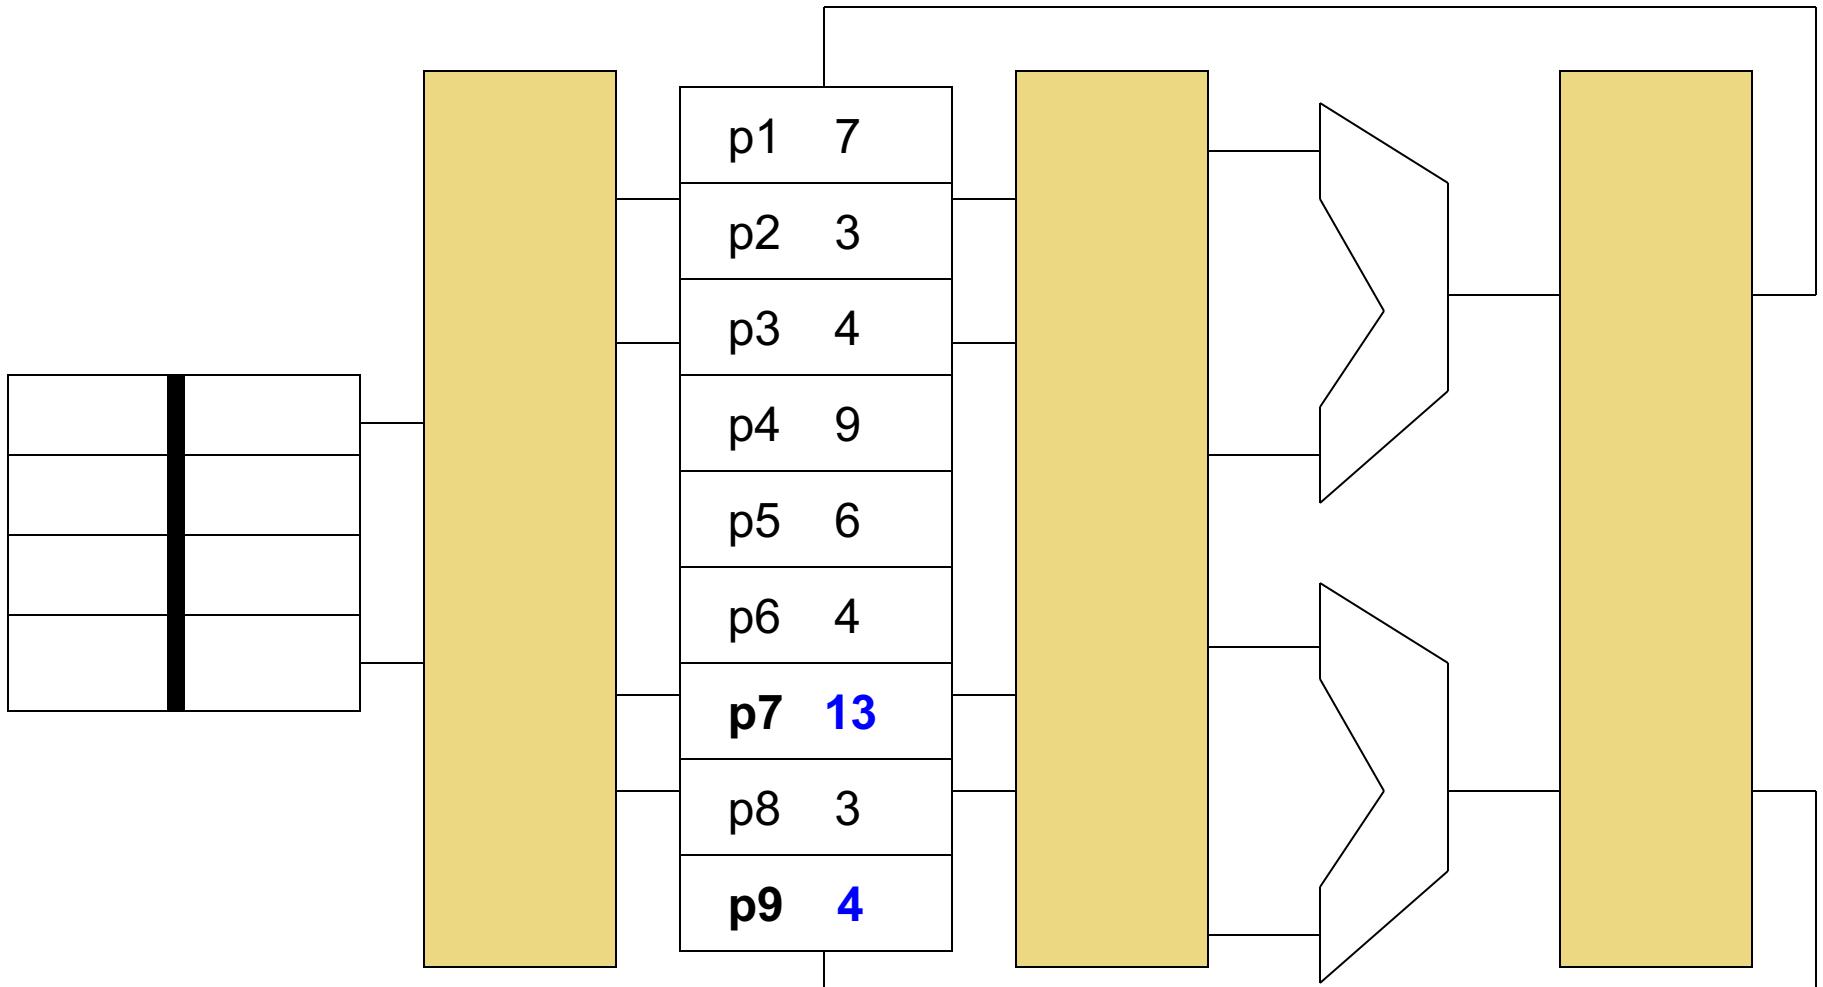
000 execution (2-wide)



000 execution (2-wide)

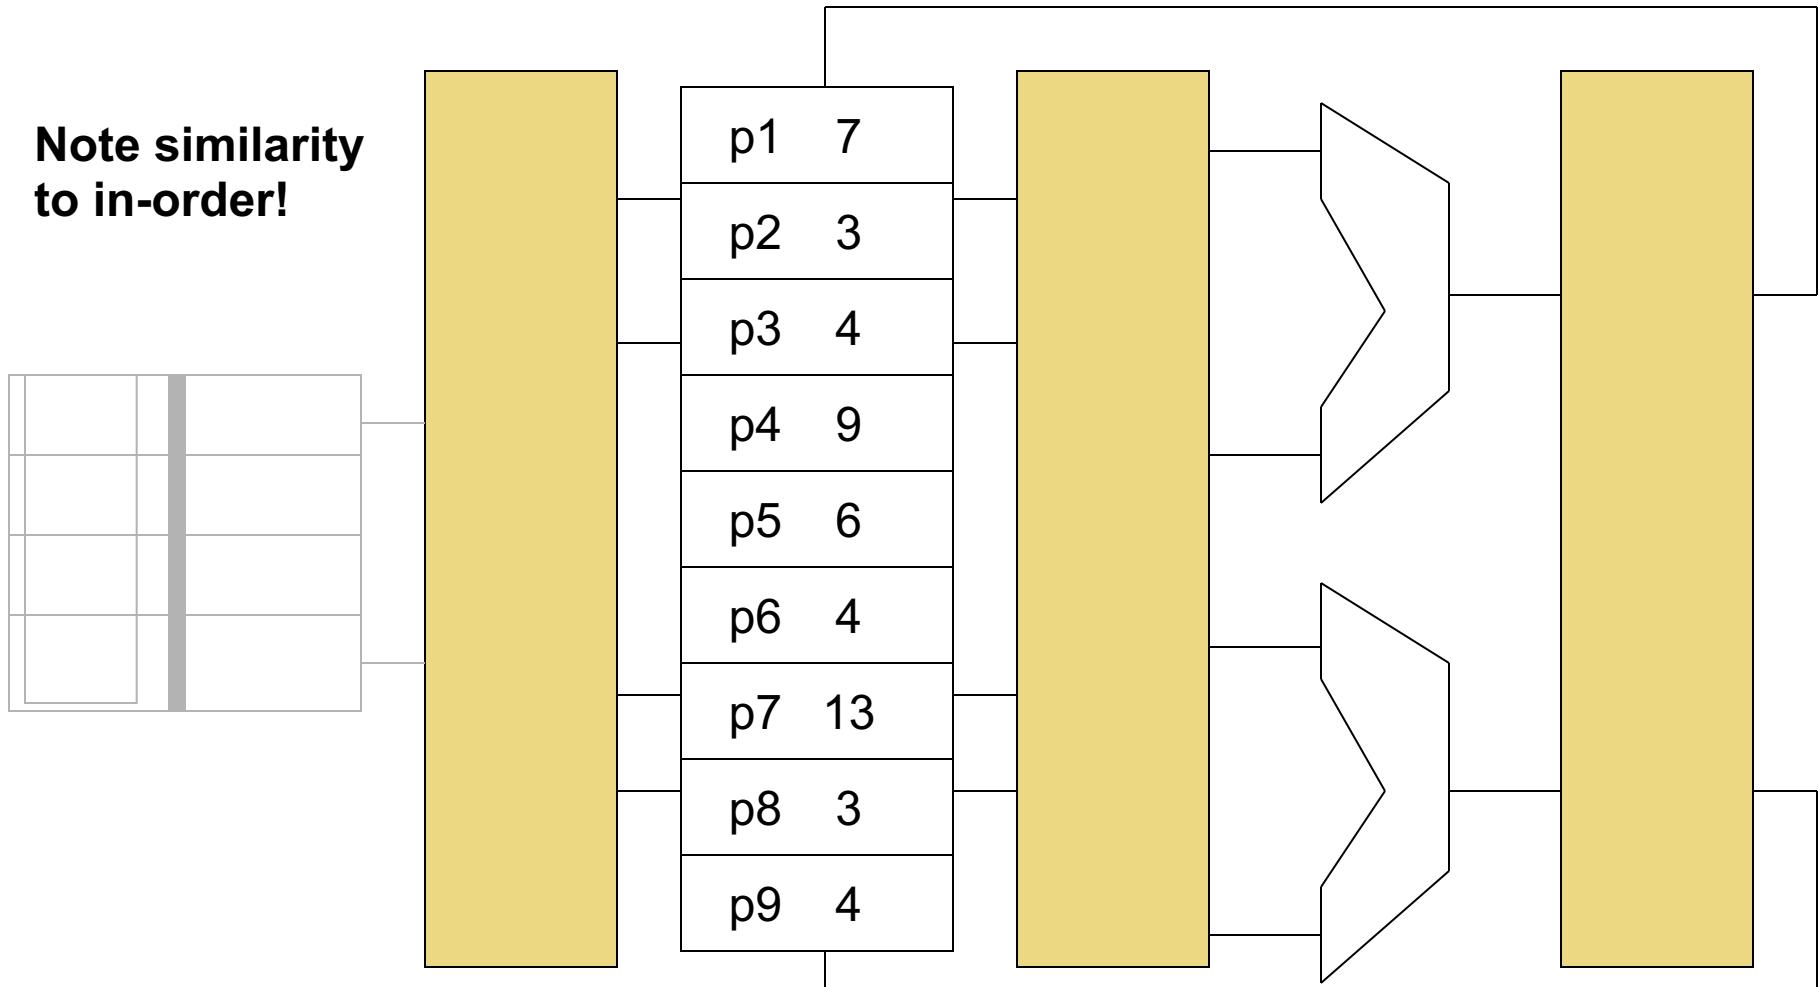


000 execution (2-wide)



OOO execution (2-wide)

Note similarity
to in-order!



Renaming Revisited

Re-order Buffer (ROB)

- ROB entry holds all info for recovery/commit
 - **All instructions** & in order
 - Architectural register names, physical register names, insn type
 - Not removed until very last thing ("commit")
- Operation
 - Fetch: insert at tail (if full, stall)
 - Commit: remove from head (if not yet done, stall)
- Purpose: tracking for in-order commit
 - Maintain appearance of in-order execution
 - Needed to support:
 - **Misprediction recovery**
 - **Freeing of physical registers**

Renaming revisited

- Track (or “log”) the “overwritten register” in ROB
 - Free this register at commit
 - Also used to restore the map table on “recovery”
 - Used for branch misprediction recovery

Register Renaming Algorithm (Full)

- Two key data structures:
 - `maptable[architectural_reg] → physical_reg`
 - Free list: allocate (new) & free registers (implemented as a queue)
- Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]
insn.phys_input2 = maptable[insn.arch_input2]
insn.old_phys_output = maptable[insn.arch_output]
new_reg = new_phys_reg()
maptable[insn.arch_output] = new_reg
insn.phys_output = new_reg
```

- At “commit”
 - **Once all older instructions have committed, free register `free_phys_reg(insn.old_phys_output)`**

Recovery

- Completely remove wrong path instructions
 - Flush from IQ
 - Remove from ROB
 - Restore map table to before misprediction
 - Free destination registers
- How to restore map table?
 - Option #1: log-based reverse renaming to recover each instruction
 - Tracks the old mapping to allow it to be reversed
 - Done sequentially for each instruction (slow)
 - See next slides
 - Option #2: checkpoint-based recovery
 - Checkpoint state of maptable and free list each cycle
 - Faster recovery, but requires more state
 - Option #3: hybrid (checkpoint for branches, unwind for others)

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

→ xor p1 ^ p2 → [p3]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

→ xor p1 ^ p2 → p6 [p3]

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

→ xor p1 ^ p2 → p6
add p6 + p4 → [p3]
[p4]

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

→ xor p1 ^ p2 → p6
add p6 + p4 → p7
[p3]
[p4]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

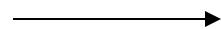
Map table

p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 →
[p3]
[p4]
[p6]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

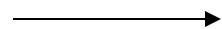
Map table

p8
p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p3 + 1 → p1

[p3]
[p4]
[p6]

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

Renaming example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1



xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 →

[p3]
[p4]
[p6]
[p1]

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

Renaming example

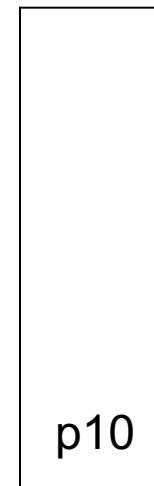
xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

[p3]
[p4]
[p6]
[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

Recovery Example

Now, let's use this info. to recover from a branch misprediction

bnz r1 loop

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

bnz p1, loop

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

[]

[p3]

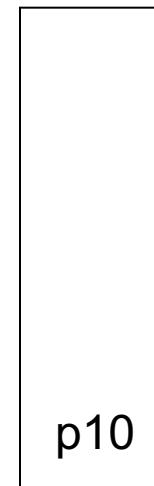
[p4]

[p6]

[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

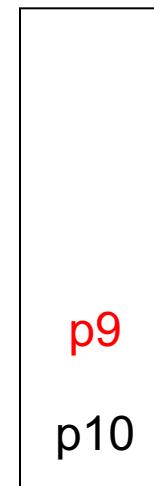
Recovery Example

bnz r1 loop
xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

bnz p1, loop []
xor p1 ^ p2 → p6 [p3]
add p6 + p4 → p7 [p4]
sub p5 - p2 → p8 [p6]
addi p8 + 1 → p9 [p1]

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

Recovery Example

bnz r1 loop
xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3

bnz p1, loop
xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8

[]
[p3]
[p4]
[p6]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

Recovery Example

bnz r1 loop
xor r1 ^ r2 → r3
add r3 + r4 → r4

bnz p1, loop
xor p1 ^ p2 → p6
add p6 + p4 → p7

[]
[p3]
[p4]

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

Recovery Example

bnz r1 loop

xor r1 ^ r2 → r3

bnz p1, loop

xor p1 ^ p2 → p6

[]
[p3]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

Recovery Example

bnz r1, loop

bnz p1, loop

[]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

Commit

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

xor p1 ^ p2 → p6 [p3]
add p6 + p4 → p7 [p4]
sub p5 - p2 → p8 [p6]
addi p8 + 1 → p9 [p1]

- Commit: instruction becomes **architected state**
 - In-order, only when instructions are finished
 - Free overwritten register (why?)

Freeing over-written register

```
xor r1 ^ r2 -> r3  
add r3 + r4 -> r4  
sub r5 - r2 -> r3  
addi r3 + 1 -> r1
```

```
xor p1 ^ p2 -> p6  
add p6 + p4 -> p7  
sub p5 - p2 -> p8  
addi p8 + 1 -> p9
```

[p3]
[p4]
[p6]
[p1]

- P3 was r3 **before** xor
- P6 is r3 **after** xor
 - Anything older than xor should read p3
 - Anything younger than xor should read p6 (until another insn writes r3)
- At commit of xor, **no older instructions exist**

Commit Example

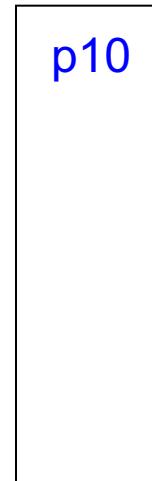
xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

xor p1 ^ p2 → p6
add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

[p3]
[p4]
[p6]
[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

Commit Example

xor r1 ^ r2 → r3
add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

xor p1 ^ p2 → p6 [p3]
add p6 + p4 → p7 [p4]
sub p5 - p2 → p8 [p6]
addi p8 + 1 → p9 [p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3

Free-list

Commit Example

add r3 + r4 → r4
sub r5 - r2 → r3
addi r3 + 1 → r1

add p6 + p4 → p7
sub p5 - p2 → p8
addi p8 + 1 → p9

[p4]
[p6]
[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4

Free-list

Commit Example

sub r5 - r2 → r3
addi r3 + 1 → r1

sub p5 - p2 → p8
addi p8 + 1 → p9

[p6]
[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6

Free-list

Commit Example

addi r3 + 1 → r1

addi p8 + 1 → p9

[p1]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6
p1

Free-list

Commit Example

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6
p1

Free-list

Dynamic Scheduling Example

Dynamic Scheduling Example

- The following slides are a detailed but concrete example
- Yet, it contains enough detail to be overwhelming
 - Try not to worry about the details
- Focus on the big picture:

**Hardware can reorder instructions
to extract instruction-level parallelism**

Recall: Motivating Example

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [p1] → p2	F	Di	I	RR	X	M ₁	M ₂	W	C				
add p2 + p3 → p4	F	Di				I	RR	X	W	C			
xor p4 ^ p5 → p6		F	Di			I	RR	X	W	C			
Id [p7] → p8		F	Di	I	RR	X	M ₁	M ₂	W		C		

- How would this execution occur cycle-by-cycle?
- Execution latencies assumed in this example:
 - Loads have two-cycle load-to-use penalty
 - Three cycle total execution latency
 - All other instructions have single-cycle execution latency
- “Issue queue”: hold all waiting (un-executed) instructions
 - Holds ready/not-ready status
 - Faster than looking up in ready table each cycle

Out-of-Order Pipeline – Cycle 0

Map Table		Ready Table	
r1	p8	p1	yes
r2	p7	p2	yes
r3	p6	p3	yes
r4	p5	p4	yes
r5	p4	p5	yes
r6	p3	p6	yes
r7	p2	p7	yes
r8	p1	p8	yes
		p9	---
		p10	---
		p11	---
		p12	---

Reorder Buffer	Insn	To Free	Done?
	ld		no
	add		no

Insn	Src1	R?	Src2	R?	Dest	Bdy

Out-of-Order Pipeline – Cycle 1a

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di											
add r2 + r3 → r4	F												
xor r4 ^ r5 → r6													
Id [r7] → r4													

Map Table	Ready Table
r1 p8	p1 yes
r2 p9	p2 yes
r3 p6	p3 yes
r4 p5	p4 yes
r5 p4	p5 yes
r6 p3	p6 yes
r7 p2	p7 yes
r8 p1	p8 yes
	p9 no
	p10 ---
	p11 ---
	p12 ---

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes	---	yes	p9	0

Out-of-Order Pipeline – Cycle 1b

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di											
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6													
Id [r7] → r4													

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p10
r5	p4
r6	p3
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Insn	To Free	Done?
	Id	p7	no
	add	p5	no

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes	---	yes	p9	0
add	p9	no	p6	yes	p10	1

Out-of-Order Pipeline – Cycle 1c

Map Table		Ready Table	
r1	p8	p1	yes
r2	p9	p2	yes
r3	p6	p3	yes
r4	p10	p4	yes
r5	p4	p5	yes
r6	p3	p6	yes
r7	p2	p7	yes
r8	p1	p8	yes
		p9	no
		p10	no
		p11	---
		p12	---

Reorder Buffer	Insn	To Free		Done?
	ld	p7		no
	add	p5		no
	xor			no
	ld			no

Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes	---	yes	p9	0
add	p9	no	p6	yes	p10	1

Out-of-Order Pipeline – Cycle 2a

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I										
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6			F										
Id [r7] → r4			F										

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p10
r5	p4
r6	p3
r7	p2
r8	p1

Reorder Buffer		Insn	To Free	Done?
		Id	p7	no
		add	p5	no
		xor		no
		Id		no

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes	---	yes	p9	0
add	p9	no	p6	yes	p10	1

Out-of-Order Pipeline – Cycle 2b

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I										
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F											

Map Table
r1 p8
r2 p9
r3 p6
r4 p10
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 no
p10 no
p11 no

Reorder Buffer						
Insn	To Free	Done?				
Id	p7	no				
add	p5	no				
xor	p3	no				
Id		no				

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	no	p6	yes	p10	1
xor	p10	no	p4	yes	p11	2

Out-of-Order Pipeline – Cycle 2c

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I										
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F	Di										

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?				
	Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p7						
add	p5						
xor	p3						
Id	p10						

Issue Queue

Out-of-Order Pipeline – Cycle 3

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR									
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F	Di	I									

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?				
	Insn	Src1	R?	Src2	R?	Dest	Bdy
	Id	p7	no				
	add	p5	no				
	xor	p3	no				
	Id	p10	no				
Issue Queue							
	Insn	Src1	R?	Src2	R?	Dest	Bdy
	Id	p8	yes		yes	p9	0
	add	p9	no	p6	yes	p10	1
	xor	p10	no	p4	yes	p11	2
	Id	p2	yes	---	yes	p12	3

Out-of-Order Pipeline – Cycle 4

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X								
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F	Di	I	RR								

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?				
	Insn	Src1	R?	Src2	R?	Dest	Bdy
	Id	p8	yes		yes	p9	0
	add	p9	yes	p6	yes	p10	1
	xor	p10	no	p4	yes	p11	2
	Id	p2	yes		yes	p12	3

Issue Queue

Out-of-Order Pipeline – Cycle 5a

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁							
add r2 + r3 → r4	F	Di				I							
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F	Di	I	RR	X							

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer		Insn	To Free	Done?
		Id	p7	no
		add	p5	no
		xor	p3	no
		Id	p10	no

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes		yes	p12	3

Out-of-Order Pipeline – Cycle 5b

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁							
add r2 + r3 → r4	F	Di				I							
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F	Di	I	RR	X							

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Id	p7	no
	add	p5	no
	xor	p3	no
	Id	p10	no

Issue Queue	Insn	Src1	R?	Src2	R?	Dest	Bdy
	Id	p8	yes		yes	p9	0
	add	p9	yes	p6	yes	p10	1
	xor	p10	yes	p4	yes	p11	2
	Id	p2	yes		yes	p12	3

Out-of-Order Pipeline – Cycle 6

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂						
add r2 + r3 → r4	F	Di				I	RR						
xor r4 ^ r5 → r6		F	Di					I					
Id [r7] → r4		F	Di	I	RR	X	M ₁						

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Id	p7	no
	add	p5	no
	xor	p3	no
	Id	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes		yes	p12	3

Out-of-Order Pipeline – Cycle 7

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂	W					
add r2 + r3 → r4	F	Di				I	RR	X					
xor r4 ^ r5 → r6		F	Di				I	RR					
Id [r7] → r4		F	Di	I	RR	X	M ₁	M ₂					

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Id	p7	yes
	add	p5	no
	xor	p3	no
	Id	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes		yes	p12	3

Out-of-Order Pipeline – Cycle 8a

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂	W	C				
add r2 + r3 → r4	F	Di				I	RR	X					
xor r4 ^ r5 → r6		F	Di				I	RR					
Id [r7] → r4		F	Di	I	RR	X	M ₁	M ₂					

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Id	p7	yes
	add	p5	no
	xor	p3	no
	Id	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes		yes	p12	3

Out-of-Order Pipeline – Cycle 8b

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W				
xor r4 ^ r5 → r6		F	Di				I	RR	X				
Id [r7] → r4		F	Di	I	RR	X	M ₁	M ₂	W				

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Id	p7	yes
	add	p5	yes
	xor	p3	no
	Id	p10	yes

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes		yes	p12	3

Out-of-Order Pipeline – Cycle 9a

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 ^ r5 → r6		F	Di				I	RR	X				
Id [r7] → r4		F	Di	I	RR	X	M ₁	M ₂	W				

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 ---
p6 yes
p7 ---
p8 yes
p9 yes
p10 yes
p11 yes
p12 yes

Reorder Buffer		Insn	To Free	Done?
		Id	p7	yes
		add	p5	yes
		xor	p3	no
		Id	p10	yes

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes			p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes			p12	3

Out-of-Order Pipeline – Cycle 9b

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 ^ r5 → r6		F	Di				I	RR	X	W			
Id [r7] → r4		F	Di	I	RR	X	M ₁	M ₂	W				

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 ---
p6 yes
p7 ---
p8 yes
p9 yes
p10 yes
p11 yes
p12 yes

Reorder Buffer						
Insn	To Free	Done?				
Id	p7					yes
add	p5					yes
xor	p3					yes
Id	p10					yes

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes			p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes			p12	3

Out-of-Order Pipeline – Cycle 10

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 ^ r5 → r6		F	Di				I	RR	X	W	C		
Id [r7] → r4		F	Di	I	RR	X	M ₁	M ₂	W		C		

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Id	p7	yes
	add	p5	yes
	xor	p3	yes
	Id	p10	yes

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes		yes	p12	3

Out-of-Order Pipeline – Done!

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR	X	M ₁	M ₂	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 ^ r5 → r6		F	Di				I	RR	X	W	C		
Id [r7] → r4		F	Di	I	RR	X	M ₁	M ₂	W		C		

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 ---
p4 yes
p5 ---
p6 yes
p7 ---
p8 yes
p9 yes
p10 ---
p11 yes
p12 yes

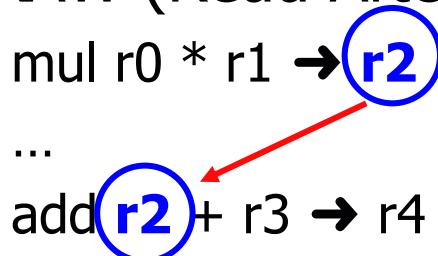
Reorder Buffer		Insn	To Free	Done?
		Id	p7	yes
		add	p5	yes
		xor	p3	yes
		Id	p10	yes

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
Id	p2	yes		yes	p12	3

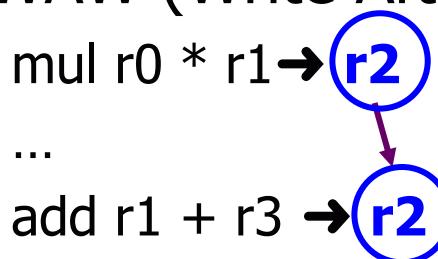
Handling Memory Operations

Recall: Types of Dependencies

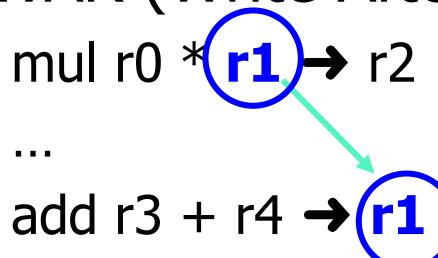
- RAW (Read After Write) = “true dependence”



- WAW (Write After Write) = “output dependence”



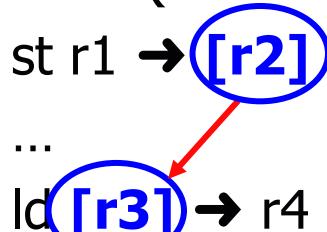
- WAR (Write After Read) = “anti-dependence”



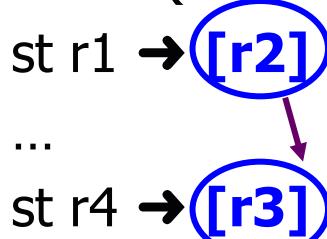
- WAW & WAR are “false”, Can be **totally eliminated** by “renaming”

Also Have Dependencies via Memory

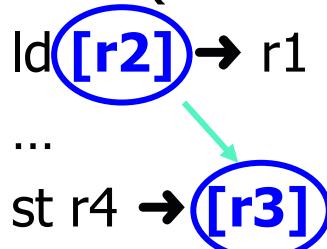
- If value in “r2” and “r3” is the same...
- RAW (Read After Write) – True dependency



- WAW (Write After Write)



- WAR (Write After Read)



WAR/WAW are “false dependencies”

- But can't rename memory in same way as registers
 - Why? Addresses are not known at rename
 - Need to use other tricks

Let's Start with Just Stores

- Stores: Write data cache, not registers
 - Can we rename memory?
- No (at least not easily)
 - Cache writes unrecoverable
- Solution: write stores into cache only when certain
 - When are we certain? At "commit"

Handling Stores

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	M	W	C	
st p4 → [p6+8]		F	Di	I?									

- Can “st p4 → [p6+8]” issue in cycle 3?
 - Its register inputs are ready...

Problem #1: Out-of-Order Stores

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	M	W	C	
st p4 → [p6+8]		F	Di	I?	RR	X	M	W					C

- Can “st p4 → [p6+8]” write the cache in cycle 6?
 - “st p5 → [p3+4]” has not yet executed
- What if p3+4 == p6+8?
 - The two stores write the same address! WAW dependency!
 - Not known until their “X” stages (cycle 5 & 8)
- Unappealing solution: all stores execute in-order
- We can do better...

Problem #2: Speculative Stores

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di			I	RR	X	M	W	C		
st p4 → [p6+8]		F	Di	I?	RR	X	M	W					C

- Can “st p4 → [p6+8]” write the cache in cycle 6?
 - Store is still “speculative” at this point
- What if “jump-not-zero” is mis-predicted?
 - Not known until its “X” stage (cycle 8)
- How does it “undo” the store once it hits the cache?
 - Answer: it can’t; stores write the cache only at **commit**
 - Guaranteed to be non-speculative at that point

Store Queue (SQ)

- Solves two problems
 - Allows for recovery of speculative stores
 - Allows out-of-order stores
- Store Queue (SQ)
 - **At dispatch, each store is given a slot in the Store Queue**
 - First-in-first-out (FIFO) queue
 - Each entry contains: “address”, “value”, and “bday”
- Operation:
 - Dispatch (in-order): allocate entry in SQ (stall if full)
 - Execute (out-of-order): write store value into store queue
 - Commit (in-order): read value from SQ and write into data cache
 - Branch recovery: remove entries from the store queue
- Also solves problems with loads

Store Queue Operation

	0	1	2	3	4	5	6	7	8	9	10	11	12
fdiv p1 / p2 → p9	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	X ₅	X ₆	W	C	
st p4 → [p5+4]	F	Di	I	RR	X	SQ						C	
st p3 → [p6+8]		F	Di	I	RR	X	SQ						C

- Stores write to SQ, not M
 - similar to register renaming, where we allocated a new physical register for each insn
- What if the fdiv triggers a /0 exception?

Memory Forwarding

	0	1	2	3	4	5	6	7	8	9	10	11	12
fdiv p1 / p2 → p9	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	X ₅	X ₆	W	C	
st p4 → [p5+4]	F	Di	I	RR	X	SQ						C	
st p3 → [p6+8]		F	Di	I	RR	X	SQ						C
ld [p7] → p8		F	Di	I?	RR	X	M ₁	M ₂	W				C

- Can “ld [p7] → p8” issue and begin execution?
 - Why or why not?

Memory Forwarding

	0	1	2	3	4	5	6	7	8	9	10	11	12
fdiv p1 / p2 → p9	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	X ₅	X ₆	W	C	
st p4 → [p5+4]	F	Di	I	RR	X	SQ						C	
st p3 → [p6+8]		F	Di	I	RR	X	SQ						C
ld [p7] → p8		F	Di	I?	RR	X	M ₁	M ₂	W				C

- Can “ld [p7] → p8” issue and begin execution?
 - Why or why not?
- If the load reads from either of the stores’ addresses...
 - Load must get correct value, but stores don’t write cache until commit...
- Solution: “memory forwarding”
 - Load also searches the Store Queue (in parallel with cache access)
 - Conceptually like register bypassing, but different implementation
 - Why? Addresses unknown until execute

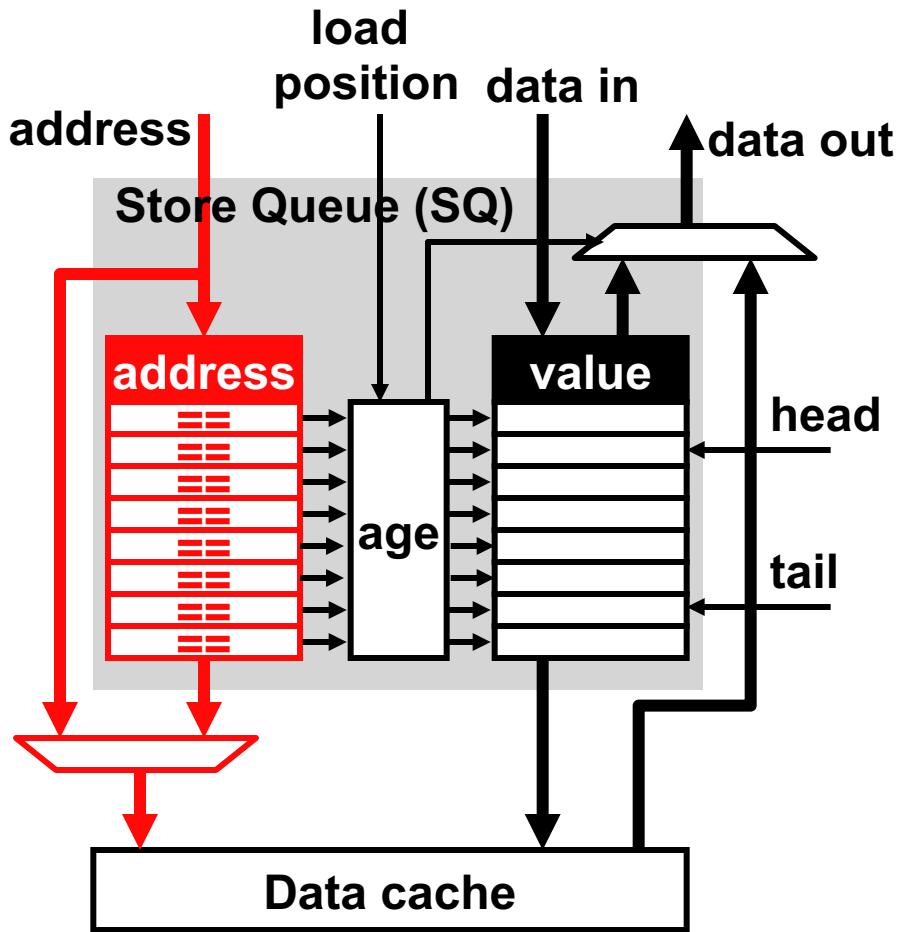
Problem #3: WAR Hazards

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
ld [p3+4] → p5		F	Di				I	RR	X	M ₁	M ₂	W	C
st p4 → [p6+8]		F	Di	I	RR	X	SQ						C

- What if “ $p3+4 == p6 + 8$ ”?
 - WAR: need to make sure that load doesn’t read store’s result
 - Need to get values based on “program order” not “execution order”
- Bad solution: require all stores/loads to execute in-order
- Good solution: add “age” fields to store queue (SQ)
 - Loads read from **youngest older matching** store
 - Another reason the SQ is a FIFO queue

Memory Forwarding via Store Queue

- Store Queue (SQ)
 - Holds all in-flight stores
 - CAM: searchable by address
 - Age logic: determine youngest matching store older than load
- Store rename/dispatch
 - Allocate entry in SQ
- Store execution
 - Update SQ
 - Address + Data
- Load execution
 - Search SQ identify youngest older matching store
 - Match? Read SQ
 - No Match? Read cache



Store Queue (SQ)

- On load execution, select the store that is:
 - To same address as load
 - Older than the load (before the load in program order)
- Of these matching stores, select the youngest
 - The store to the same address that immediately precedes the load

When Can Loads Execute?

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	SQ	C		
ld [p6+8] → p7		F	Di	I?	RR	X	M ₁	M ₂	W				C

- Can “ld [p6+8] → p7” issue in cycle 3
 - Why or why not?

When Can Loads Execute?

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X ₁	X ₂	X ₃	X ₄	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	SQ	C		
ld [p6+8] → p7		F	Di	I?	RR	X	M ₁	M ₂	W				C

- Aliasing! Does $p3+4 == p6+8$?
 - If no, load should get value from memory
 - **Can it start to execute?**
 - If yes, load should get value from store
 - By reading the store queue?
 - **But the value isn't put into the store queue until cycle 9**
- **Key challenge:** don't know addresses until execution!
 - One solution: require all loads to wait for all earlier (prior) stores

Conservative Load Scheduling

- Conservative load scheduling:
 - All older stores have executed
 - Some architectures: split store address / store data
 - Only requires knowing addresses (not the store values)
 - Advantage: always safe
 - Disadvantage: performance (limits ILP)

Conservative Load Scheduling

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
ld [p1] → p4	F	Di	I	Rr	X	M ₁	M ₂	W	C							
ld [p2] → p5	F	Di	I	Rr	X	M ₁	M ₂	W	C							
add p4, p5 → p6		F	Di			I	Rr	X	W	C						
st p6 → [p3]		F	Di			I	Rr	X	SQ	C						
ld [p1+4] → p7			F	Di			I	Rr	X	M ₁	M ₂	W	C			
ld [p2+4] → p8			F	Di			I	Rr	X	M ₁	M ₂	W	C			
add p7, p8 → p9			F	Di						I	Rr	X	W	C		
st p9 → [p3+4]			F	Di						I	Rr	X	SQ	C		

Conservative load scheduling: can't issue ld [p1+4] until cycle 7!

Might as well be an in-order machine on this example

Can we do better? How?

Optimistic Load Scheduling

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
ld [p1] → p4	F	Di	I	Rr	X	M ₁	M ₂	W	C							
ld [p2] → p5	F	Di	I	Rr	X	M ₁	M ₂	W	C							
add p4, p5 → p6		F	Di			I	Rr	X	W	C						
st p6 → [p3]		F	Di			I	Rr	X	SQ	C						
ld [p1+4] → p7			F	Di	I	Rr	X	M ₁	M ₂	W	C					
ld [p2+4] → p8			F	Di	I	Rr	X	M ₁	M ₂	W	C					
add p7, p8 → p9			F	Di			I	Rr	X	W	C					
st p9 → [p3+4]			F	Di			I	Rr	X	SQ	C					

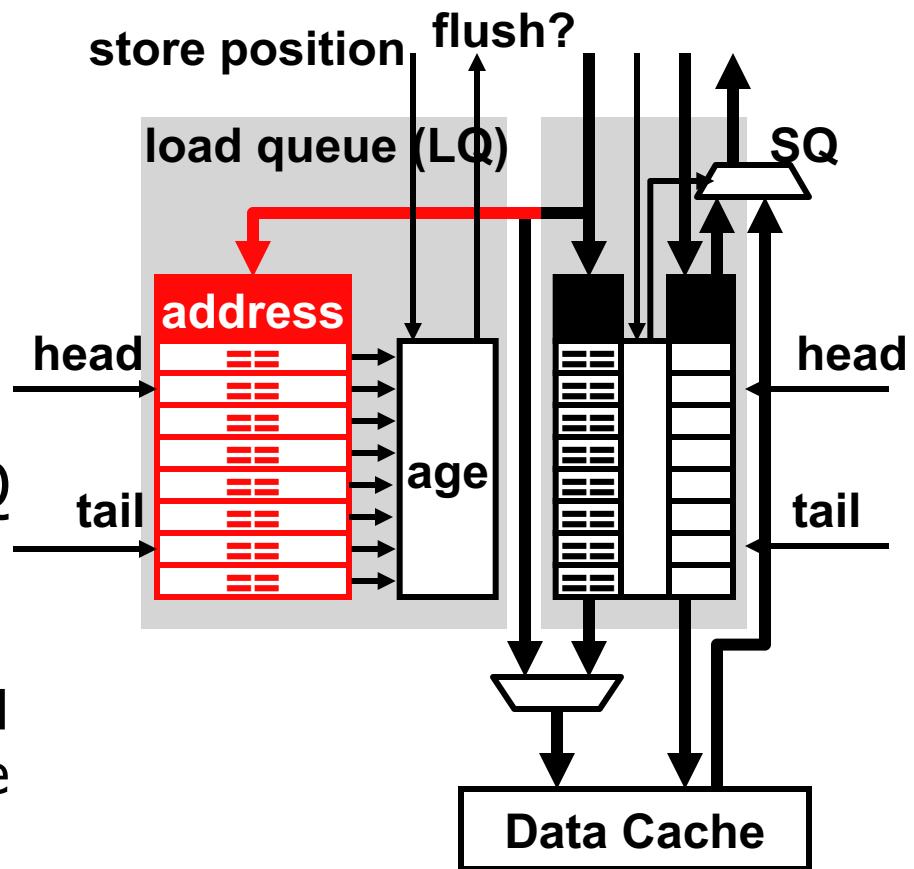
Optimistic load scheduling: can actually benefit from out-of-order!
Let's speculate!

Load Speculation

- Speculation requires three things.....
 - 1. When do we speculate?
 - 2. How do we detect a mis-speculation?
 - 3. How do we recover from mis-speculations?
 - Squash offending load and all newer insns
 - Similar to branch mis-prediction recovery

Load Queue

- Detects load ordering violations
- Load execution: Write address into LQ
 - Also note any store forwarded from
- Store execution: Search LQ
 - Younger load with same addr?
 - Did younger load forward from younger store? [See slide 149]



Store Queue + Load Queue

- Store Queue: handles forwarding, allows OoO stores
 - Entry per store (allocated @ dispatch, deallocated @ commit)
 - Written by stores (@ execute)
 - Searched by loads (@ execute)
 - Read from SQ to write data cache (@ commit)
- Load Queue: detects ordering violations
 - Entry per load (allocated @ dispatch, deallocated @ commit)
 - Written by loads (@ execute)
 - Searched by stores (@ execute)
- Both together
 - Allows aggressive load scheduling
 - Stores don't constrain load execution

Optimistic Load Scheduling Problem

- Allows loads to issue before older stores
 - Increases ILP
 - + Good: When no conflict, increases performance
 - Bad: Conflict => squash => worse performance than waiting
- Can we have our cake AND eat it too?

Predictive Load Scheduling

- Predict which loads must wait for stores
- Fool me once, shame on you-- fool me twice?
 - Loads default to aggressive
 - Keep table of load PCs that have been caused squashes
 - Schedule these conservatively
 - + Simple predictor
 - Makes “bad” loads wait for **all** older stores
- More complex predictors used in practice
 - Predict which stores loads should wait for
 - “Store Sets” paper

Load/Store Queue Examples

Initial State

(Stores to different addresses)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

RegFile		Load Queue	
	Bdy	Addr	
p1	5		
p2	100		
p3	9		
p4	200		
p5	100		
p6	---		
p7	---		
p8	---		

Store Queue

	Bdy	Addr	Val

Cache

Addr	Val
100	13
200	17

RegFile		Load Queue	
	Bdy	Addr	
p1	5		
p2	100		
p3	9		
p4	200		
p5	100		
p6	---		
p7	---		
p8	---		

Store Queue

	Bdy	Addr	Val

Cache

Addr	Val
100	13
200	17

RegFile		Load Queue	
	Bdy	Addr	
p1	5		
p2	100		
p3	9		
p4	200		
p5	100		
p6	---		
p7	---		
p8	---		

Store Queue

	Bdy	Addr	Val

Cache

Addr	Val
100	13
200	17

Good Interleaving

(Shows importance of address check)

- ## 1. St p1 → [p2]

	RegFile	Load Queue
p1	5	
p2	100	
p3	9	
p4	200	
p5	100	
p6	---	
p7	---	
p8	---	

	Store Queue		
	Bdy	Addr	Val
	1	100	5

Addr	Val
100	13
200	17

- ## 2. St p3 → [p4]

RegFile		Load Queue		Store Queue		
	Bdy	Addr		Bdy	Addr	Val
p1	5					
p2	100					
p3	9					
p4	200					
p5	100					
p6	---			1	100	5
p7	---			2	200	9
p8	---					

Addr	Val
100	13
200	17

- 1. St p1 → [p2]
 2. St p3 → [p4]
 3. Ld [p5] → p6

RegFile		Load Queue		Store Queue		
		Bdy	Addr	Bdy	Addr	Val
p1	5					
p2	100		3		100	
p3	9					
p4	200					
p5	100					
p6	5					
p7	---		1	100	5	
p8	---		2	200	9	

Addr	Val
100	13
200	17

Different Initial State

(All to same address)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

RegFile		Load Queue		Store Queue	
	Bdy	Addr		Bdy	Addr
p1	5				
p2	100				
p3	9				
p4	100				
p5	100				
p6	---				
p7	---				
p8	---				

Cache	
Addr	Val
100	13
200	17

RegFile		Load Queue		Store Queue	
	Bdy	Addr		Bdy	Addr
p1	5				
p2	100				
p3	9				
p4	100				
p5	100				
p6	---				
p7	---				
p8	---				

Cache	
Addr	Val
100	13
200	17

RegFile		Load Queue		Store Queue	
	Bdy	Addr		Bdy	Addr
p1	5				
p2	100				
p3	9				
p4	100				
p5	100				
p6	---				
p7	---				
p8	---				

Cache	
Addr	Val
100	13
200	17

Good Interleaving #1

(Program Order)

1. St p1 → [p2]

RegFile	Load Queue	
p1	5	
p2	100	
p3	9	
p4	100	
p5	100	
p6	---	
p7	---	
p8	---	

Store Queue		
Bdy	Addr	Val
1	100	5

2. St p3 → [p4]

RegFile		Load Queue		Store Queue		
p1	5	Bdy	Addr			
p2	100					
p3	9					
p4	100					
p5	100					
p6	---	Bdy	Addr	Val		
p7	---	1	100	5		
p8	---	2	100	9		

3. Ld [p5] → p6

	RegFile	Load Queue		
		Bdy	Addr	
p1	5			
p2	100	3	100	
p3	9			
p4	100			
p5	100			
p6	9			
p7	---	1	100	5
p8	---	2	100	9

1. St p1 → [p2]
 2. St p3 → [p4]
 3. Ld [p5] → p6

Good Interleaving #2

(Stores reordered)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

2. St p3 → [p4]

	RegFile	Load Queue		
	Bdy	Addr		
p1	5			
p2	100			
p3	9			
p4	100			
p5	100			
p6	---			
p7	---			
p8	---			

Store Queue

	Bdy	Addr	Val
	1		
	2	100	9

Cache

	Addr	Val
	100	13
	200	17

1. St p1 → [p2]

	RegFile	Load Queue		
	Bdy	Addr		
p1	5			
p2	100			
p3	9			
p4	100			
p5	100			
p6	---			
p7	---			
p8	---			

Store Queue

	Bdy	Addr	Val
	1		
	2	100	9

Cache

	Addr	Val
	100	13
	200	17

3. Ld [p5] → p6

	RegFile	Load Queue		
	Bdy	Addr		
p1	5			
p2	100			
p3	9			
p4	100			
p5	100			
p6	9			
p7	---			
p8	---			

Store Queue

	Bdy	Addr	Val
	1	100	5
	2	100	9

Cache

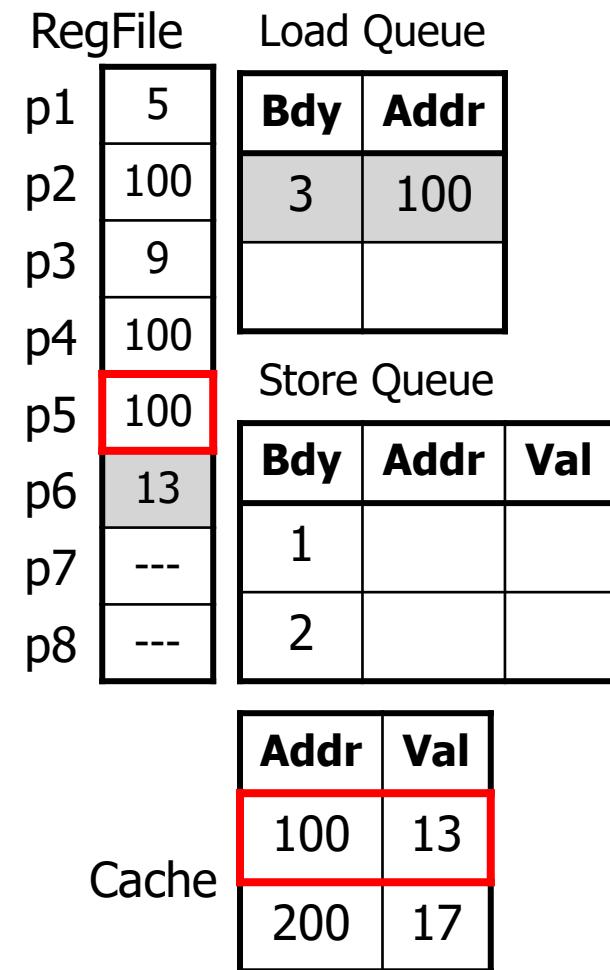
	Addr	Val
	100	13
	200	17

Bad Interleaving #1

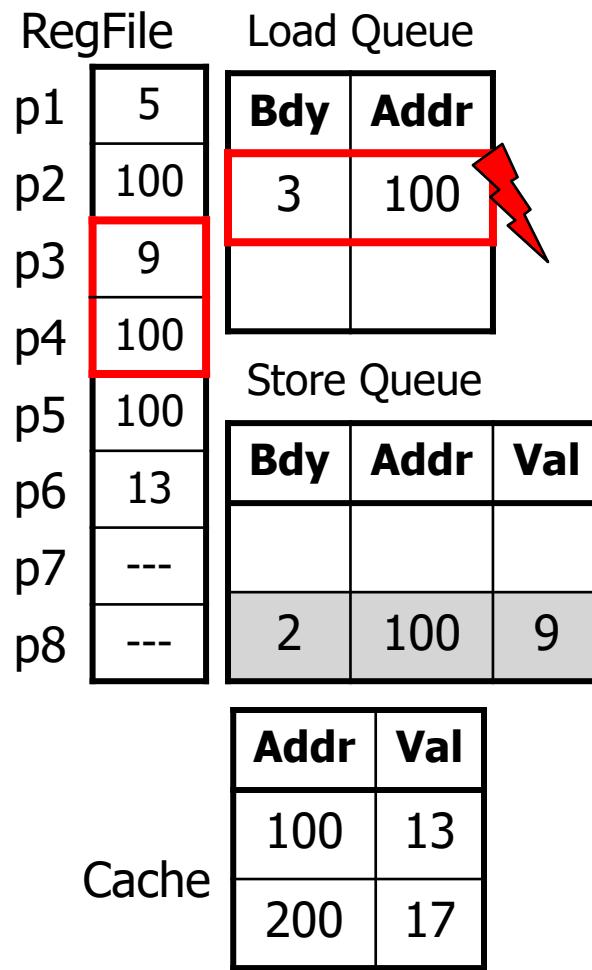
(Load reads the cache)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

3. Ld [p5] → p6



2. St p3 → [p4]



Bad Interleaving #2

(Load gets value from wrong store)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

1. St p1 → [p2]

RegFile		Load Queue			Store Queue		
p1	5	Bdy	Addr	Bdy of Fwd. Store			
p2	100						
p3	9						
p4	100						
p5	100						
p6	---	Bdy	Addr	Val			
p7	---	1	100	5			
p8	---						

Addr	Val
100	13
200	17

3. Ld [p5] → p6

RegFile		Load Queue			Store Queue		
p1	5	Bdy	Addr	Bdy of Fwd. Store			
p2	100	3	100	1			
p3	9						
p4	100						
p5	100						
p6	5	Bdy	Addr	Val			
p7	---	1	100	5			
p8	---	2					

Addr	Val
100	13
200	17

2. St p3 → [p4]

RegFile	Load Queue	Store Queue
p1	5	
p2	100	
p3	9	
p4	100	
p5	100	
p6	5	
p7	---	
p8	---	

Addr	Val
100	13
200	17

Bad/Good Interleaving

(Load gets value from correct store, but does it work?)

2. St p3 → [p4]

	RegFile			Load Queue		
	Bdy	Addr	Bdy of Fwd. Store			
p1	5					
p2	100					
p3	9					
p4	100					
p5	100					
p6	---					
p7	---					
p8	---					

	Store Queue		
	Bdy	Addr	Val
p5	100		
p6	9		
p7	---		
p8	2	100	9

3. Ld [p5] → p6

	RegFile			Load Queue		
	Bdy	Addr	Bdy of Fwd. Store			
p1	5					
p2	100					
p3	9					
p4	100					
p5	100					
p6	9					
p7	---					
p8	---					

	Store Queue		
	Bdy	Addr	Val
p5	100		
p6	9		
p7	---		
p8	2	100	9

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

1. St p1 → [p2]

	RegFile			Load Queue		
	Bdy	Addr	Bdy of Fwd. Store			
p1	5					
p2	100					
p3	9					
p4	100					
p5	100					
p6	9					
p7	---					
p8	---					

	Store Queue		
	Bdy	Addr	Val
p5	100		
p6	9		
p7	---		
p8	2	100	9

Cache

Addr	Val
100	13
200	17

Cache

Addr	Val
100	13
200	17

Cache

Addr	Val
100	13
200	17

Out-of-Order Everywhere

OoO Everywhere

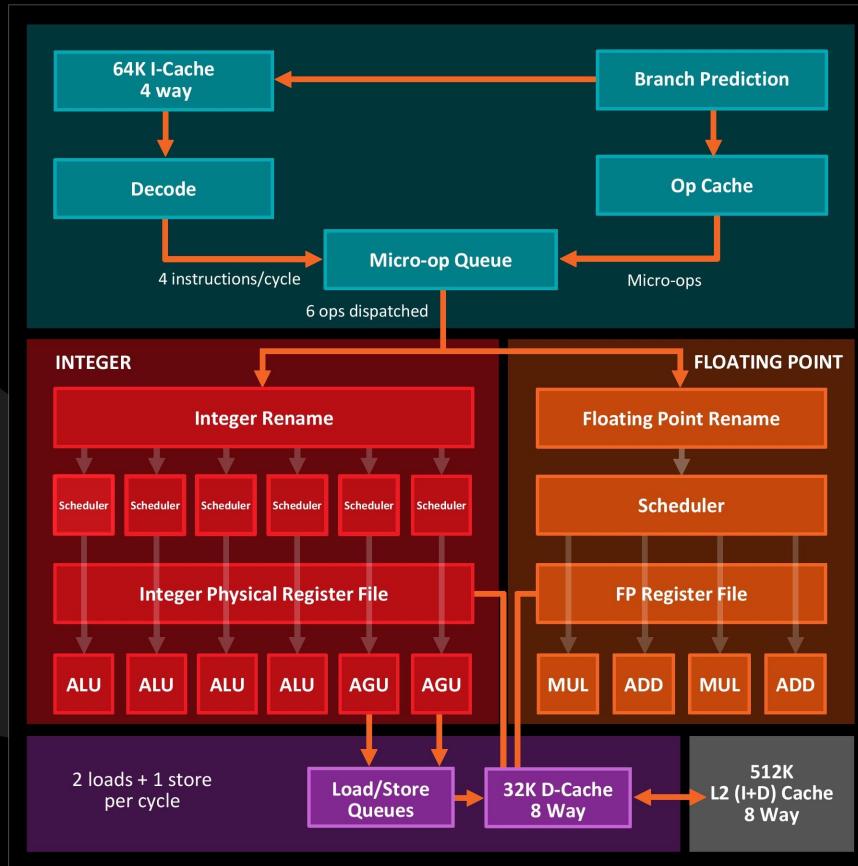
- Joe's phone – iPhone 7
 - Apple A10 processor
 - details very closely guarded
 - definitely OoO
 - some things known about Cyclone (Apple A7 chip)
 - issue 6 insns per cycle
 - 192-entry ROB
 - 4 integer ALUs
 - 2 Load/Store units
 - 14-19 cycle branch misprediction penalty
 - <https://www.anandtech.com/show/7910/apples-cyclone-microarchitecture-detailed>

OoO everywhere for a while now

- Joe's previous phone's CPU: Qualcomm Krait 400 processor
 - based on ARM Cortex A15 processor
 - **out-of-order** 2.5GHz quad-core
 - 3-wide fetch/decode
 - 4-wide issue
 - 11-stage integer pipeline
 - 28nm process technology
 - 4/4KB DM L1\$, 16/16KB 4-way SA L2\$, 2MB 8-way SA L3\$

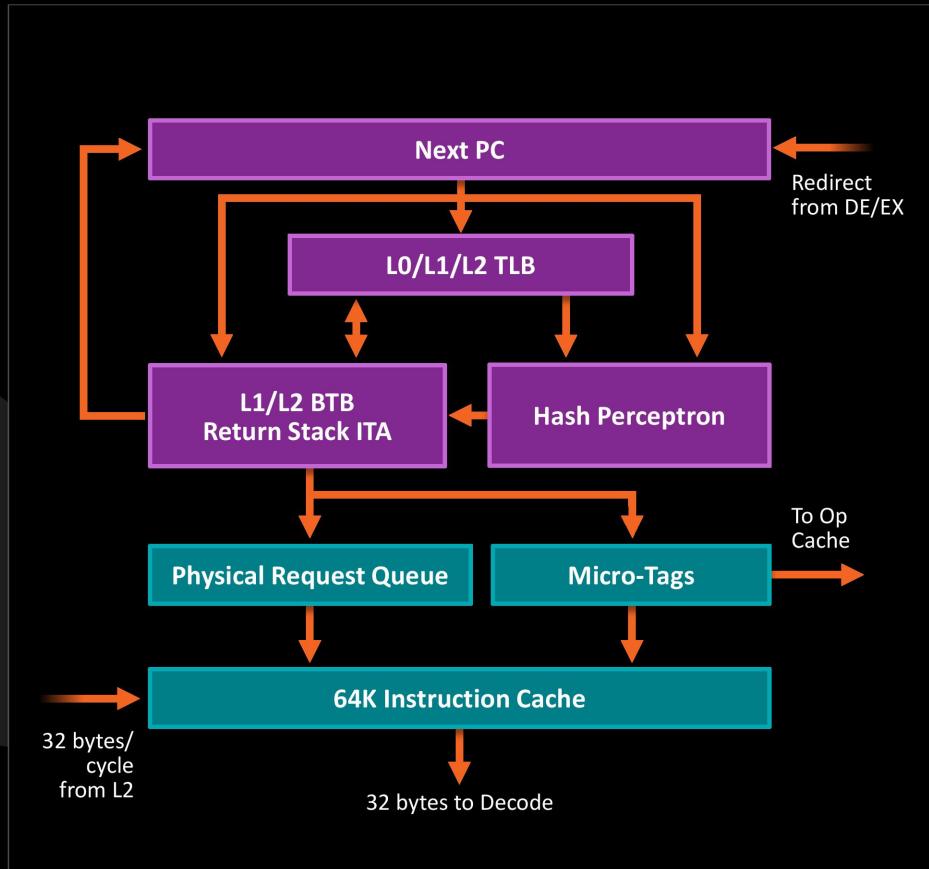
Modern OoO desktop/server CPU

- AMD Zen microarchitecture (2016)
 - AMD slides from HotChips 2016 conference follow



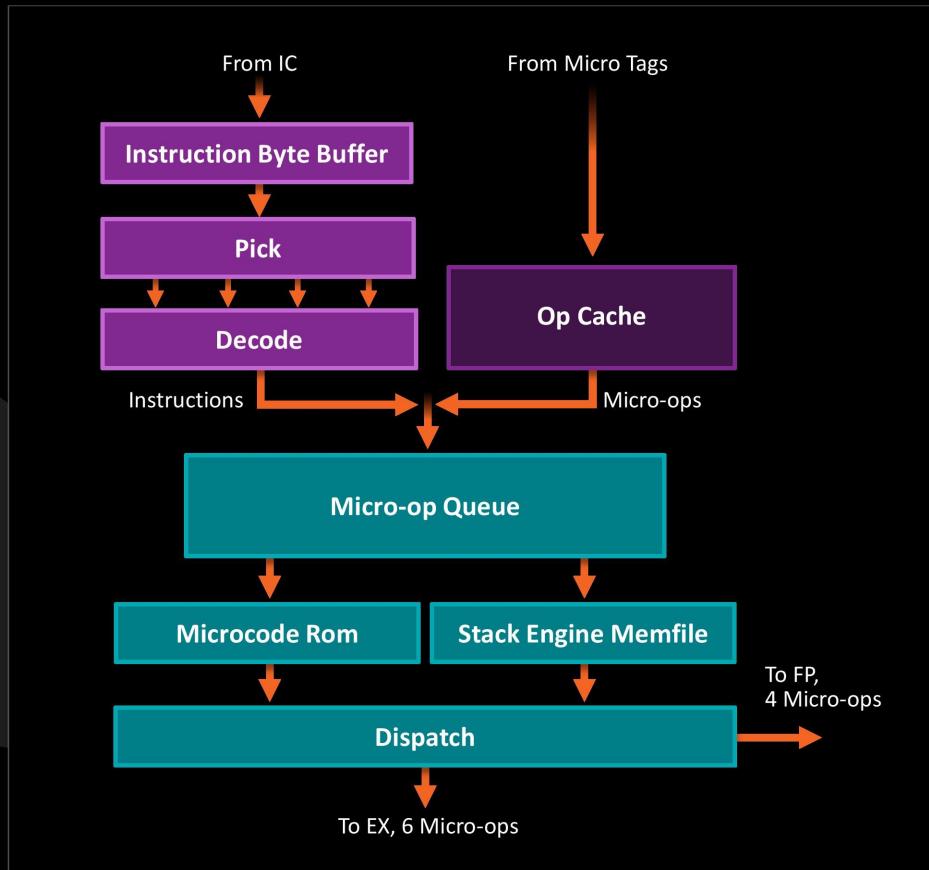
ZEN MICROARCHITECTURE

- ▲ Fetch Four x86 instructions
- ▲ Op Cache instructions
- ▲ 4 Integer units
 - Large rename space – 168 Registers
 - 192 instructions in flight/8 wide retire
- ▲ 2 Load/Store units
 - 72 Out-of-Order Loads supported
- ▲ 2 Floating Point units x 128 FMACs
 - built as 4 pipes, 2 Fadd, 2 Fmul
- ▲ I-Cache 64K, 4-way
- ▲ D-Cache 32K, 8-way
- ▲ L2 Cache 512K, 8-way
- ▲ Large shared L3 cache
- ▲ 2 threads per core



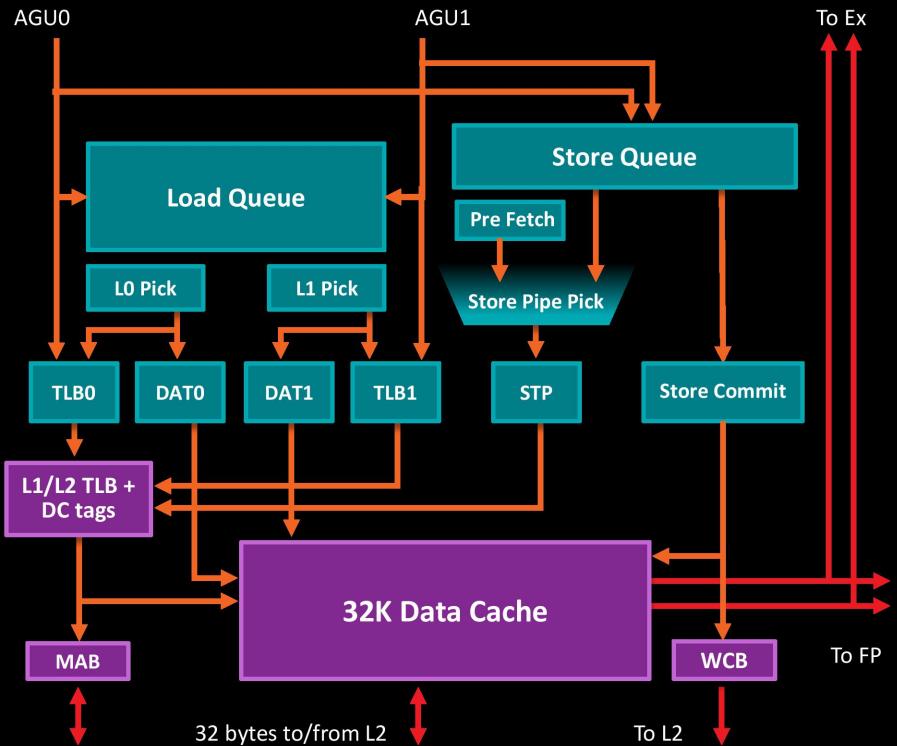
FETCH

- ▲ Decoupled Branch Prediction
- ▲ TLB in the BP pipe
 - 8 entry L0 TLB, all page sizes
 - 64 entry L1 TLB, all page sizes
 - 512 entry L2 TLB, no 1G pages
- ▲ 2 branches per BTB entry
- ▲ Large L1 / L2 BTB
- ▲ 32 entry return stack
- ▲ Indirect Target Array (ITA)
- ▲ 64K, 4-way Instruction cache
- ▲ Micro-tags for IC & Op cache
- ▲ 32 byte fetch



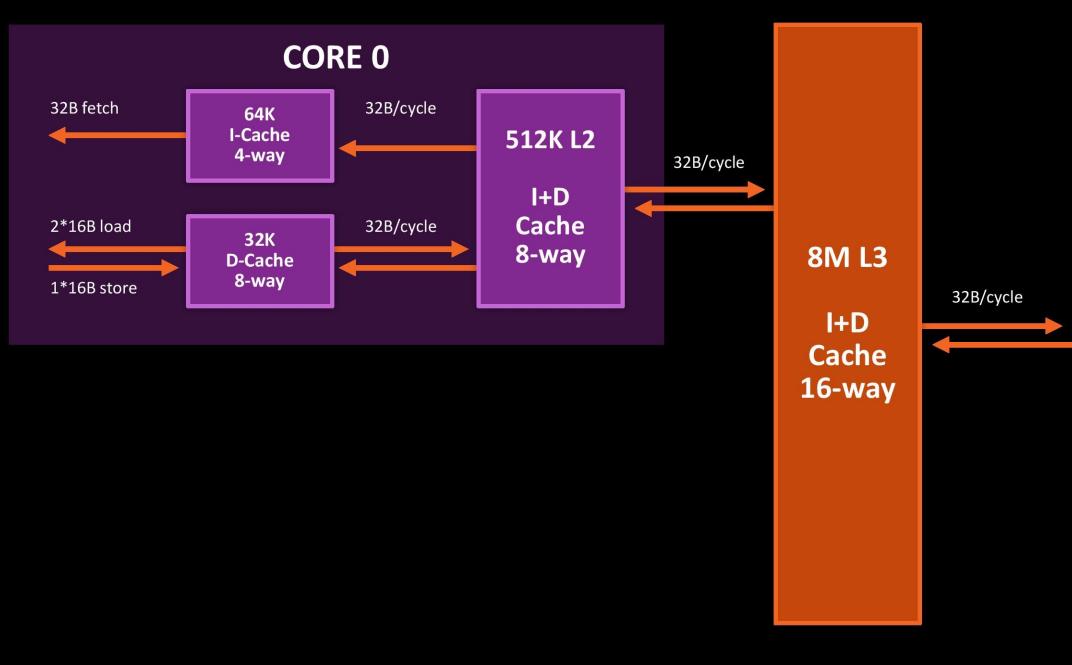
DECODE

- ▲ Inline Instruction-length Decoder
- ▲ Decode 4 x86 instructions
- ▲ Op cache
- ▲ Micro-op Queue
- ▲ Stack Engine
- ▲ Branch Fusion
- ▲ Memory File for Store to Load Forwarding



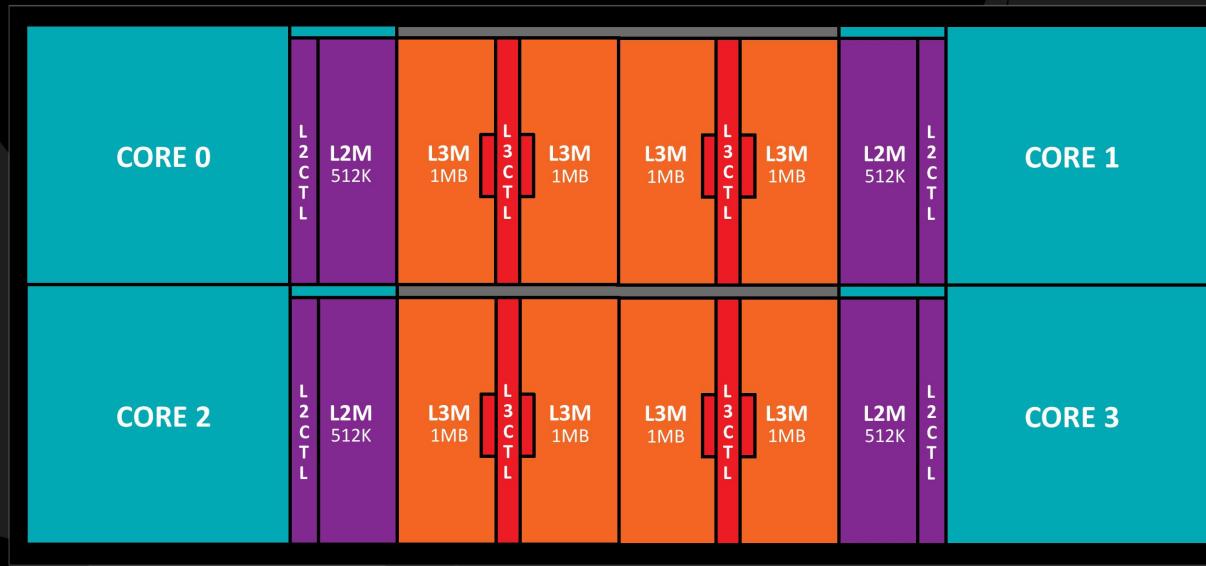
LOAD/STORE AND L2

- ▲ 72 Out of Order Loads
- ▲ 44 entry Store Queue
- ▲ Split TLB/Data Pipe, store pipe
- ▲ 64 entry L1 TLB, all page sizes
- ▲ 1.5K entry L2 TLB, no 1G pages
- ▲ 32K, 8 way Data Cache
 - Supports two 128-bit accesses
- ▲ Optimized L1 and L2 Prefetchers
- ▲ 512K, private (2 threads), inclusive L2



ZEN CACHE HIERARCHY

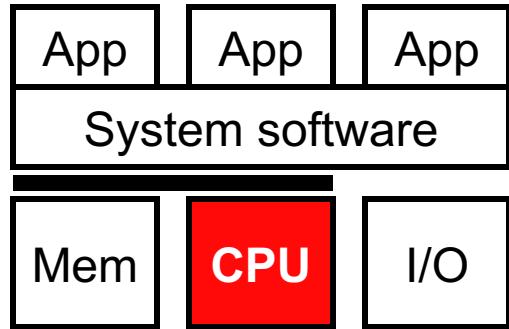
- ▲ Fast private 512K L2 cache
- ▲ Fast shared L3 cache
- ▲ High bandwidth enables prefetch improvements
- ▲ L3 is filled from L2 victims
- ▲ Fast cache-to-cache transfers
- ▲ Large Queues for Handling L1 and L2 misses



CPU COMPLEX

- ▲ A CPU complex (CCX) is four cores connected to an L3 Cache.
- ▲ The L3 Cache is 16-way associative, 8MB, mostly exclusive of L2.
- ▲ The L3 Cache is made of 4 slices, by low-order address interleave.
- ▲ Every core can access every cache with same average latency

Summary: Scheduling



- Code scheduling
 - To reduce pipeline stalls
 - To increase ILP (insn-level parallelism)
- Static scheduling by the compiler
 - Approach & limitations
- Dynamic scheduling in hardware
 - Register renaming
 - Instruction selection
 - Handling memory operations
- Up next: multicore