

# CIS 5710

## Computer Organization & Design

Meltdown & Spectre  
Hardware Speculative Execution Attacks

# Who discovered these bugs?

---

## Who reported Meltdown?

Meltdown was independently discovered and reported by three teams:

- [Jann Horn](#) ([Google Project Zero](#)),
- [Werner Haas](#), [Thomas Prescher](#) ([Cyberus Technology](#)),
- [Daniel Gruss](#), [Moritz Lipp](#), [Stefan Mangard](#), [Michael Schwarz](#) ([Graz University of Technology](#))

## Who reported Spectre?

Spectre was independently discovered and reported by two people:

- [Jann Horn](#) ([Google Project Zero](#)) and
- [Paul Kocher](#) in collaboration with, in alphabetical order, [Daniel Genkin](#) ([University of Pennsylvania](#) and [University of Maryland](#)), [Mike Hamburg](#) ([Rambus](#)), [Moritz Lipp](#) ([Graz University of Technology](#)), and [Yuval Yarom](#) ([University of Adelaide](#) and [Data61](#))

These bugs were discovered in 2017, and publicized in January 2018

# **Side and Covert Channels**

# Side Channels and Covert Channels

---

- A **side channel** is a mechanism that leads to inadvertent information transfer
  - e.g., going through your neighbor's recycling
- A **covert channel** is a mechanism used for illicit communication between two cooperating parties
  - e.g., a double-agent spy communicating with their handler
- Both channels result in undesired information transfer
- Computer programs can have such channels
  - even programs that utilize secrets, like encryption keys
  - We focus especially on side channels, as they're more dangerous

# Example side channel

---


```
boolean secret = true;
int array[1024] = {...};
if (secret) {
    array[0]++;
} else {
    array[500]++;
}
```

- Threat Model: what the attacker can do
  - can't change the hardware or OS, which work correctly
  - can run arbitrary programs at arbitrary times
- How can we learn the secret?
  - go through the program's recycling and see what we find

# Example side channel

```
boolean secret = true;
int array[1024] = {...};
if (secret) {
    array[0]++;
} else {
    array[500]++;
}
```

- We fill up the cache with known data
- Then, run the secret program
- Then, see what's left in the cache
  - access all the known data, see which accesses are fast and which are slow
- Caches can be a side channel
  - inadvertently revealing information about program execution



# Cache Side Channel Attacks

---

- Can be used for both:
  - **side channel** (unintentional communication)
  - covert channel (intentional communication)
- We can steal SSH private keys this way

## CACHE MISSING FOR FUN AND PROFIT

COLIN PERCIVAL

2005

ABSTRACT. Simultaneous multithreading — put simply, the sharing of the execution resources of a superscalar processor between multiple execution threads — has recently become widespread via its introduction (under the name “Hyper-Threading”) into Intel Pentium 4 processors. In this implementation, for reasons of efficiency and economy of processor area, the sharing of processor resources between threads extends beyond the execution units; of particular concern is that the threads share access to the memory caches.

We demonstrate that this shared access to memory caches pro-

# Side channels are delicate

---

- Might have to go through a lot of recycling to find what you're looking for
  - side channels often require many executions to recover a secret
- Other side channels:
  - power consumption
  - execution time
  - electromagnetic radiation
  - microarchitectural features like contention for execution ports
  - attackers always find new ones



# Abusing OoO Execution

# Meltdown code

---

Let's read kernel memory from userspace >:-)

r0, r1: temps

r2: kernel address we wish to read

r3: start of probe\_array

```
ldr r0 <- [r2]
sll r0 <- r0, #6
add r1 <- r3, r0
ldr r1 <- [r1]
```

What does this code do?

Based on the **data value** at [r2],  
we access a particular cache line  
of probe\_array.

# Meltdown code on an OoO processor

---

Let's read kernel memory from userspace >:-)

r0, r1: temps

r2: kernel address we wish to read

r3: start of probe\_array

```
ldr r0 <- [r2] <= this raises an exception...but not until Commit!
```

```
sll r0 <- r0, #6
```

```
add r1 <- r3, r0
```

```
ldr r1 <- [r1]
```

After the exn, we clear the  
ROB and other OoO structures  
**but not the caches**, leaving  
open a side channel!

# Speculative Execution Attacks

---

- An OoO processor speculates aggressively
  - on a mis-speculation, the incorrect instructions are wiped away and have no ISA-level visible impact
    - register writes, memory stores are all cleaned up
    - non-ISA level changes like cache contents **are not cleaned up**
    - insns that roll back can still leak information via a side channel!
- These attacks strike at core principles of OoO execution
  - undiscovered since 1990s?

Speculative execution can proceed far ahead of the main processor. For example, on an i7 Surface Pro 3 (i7-4650U) used for most of the testing, the code in Appendix A works with up to 188 simple instructions inserted in the source code between the 'if' statement and the line accessing array1/array2.

# Meltdown experiment

- Real Meltdown uses pages instead of cache lines, exploiting a TLB side channel




Figure 4: Even if a memory location is only accessed during out-of-order execution, it remains cached. Iterating over the 256 pages of `probe_array` shows one cache hit, exactly on the page that was accessed during the out-of-order execution.

# Meltdown Mitigations

---

- Meltdown affects Intel OoO and some ARM OoO processors, AMD is immune
  - Intel did page permission checks at Commit, not Execute
  - Intel cores have hardware fixes as of ~2019
- Meltdown can be patched in software
  - all major OSes released patches in 2018
  - performance impact for system calls
  - KPTI/KAISER Linux patch maps minimal kernel code/data

# Spectre Attacks

# What happens on a Context Switch

---

- Operating system responsible to handle context switching
  - Hardware support is just a timer interrupt
- Each process has an associated data structure which is used to record relevant state such as:
  - **Architected state**: PC, registers, Page table pointer
    - Save and restore them on context switches
    - Memory state?
  - **Non-architected state**: caches, predictor tables, etc.
    - **Ignore** or flush
    - ignoring is key aspect of Spectre/Meltdown vulnerabilities!
- OS swaps out values for old process, swaps in values for new process and restarts the system.



# Spectre Attack

---

- Goal: trick another process into leaking information
- Let's say my secure program contains this code:

```
if (x < array1_size) {  
    y = array2[array1[x] * 256];  
}
```

- If attacker controls  $x$ , value of  $\text{array1}[x]$  can be leaked
  - problem: how to get around the bounds check?

# Poison the branch predictor

---

- How do I poison the BHT/BTB?
  - If I know the BHT/BTB entries that will be used for the bounds check branch, I can train them in advance
    - e.g., branch with the same PC in the attacker process
  - Not hard when the attacker provides their own code
    - JavaScript, VM on EC2
- How do I get the victim to run the vulnerable code?
  - highly victim-specific
  - easiest in JIT environments
  - [Spectre paper](#) shows how to read Chrome-internal state from JavaScript

# Spectre Mitigations

---

- Spectre affects OoO chips from everyone
  - confirmed on Intel, ARM, AMD
- Software patches are incomplete, slow
  - Reduce use of indirect branches
  - Disable speculation through special instructions at “critical” code points
- Ultimately, Spectre attacks aren’t that easy to pull off